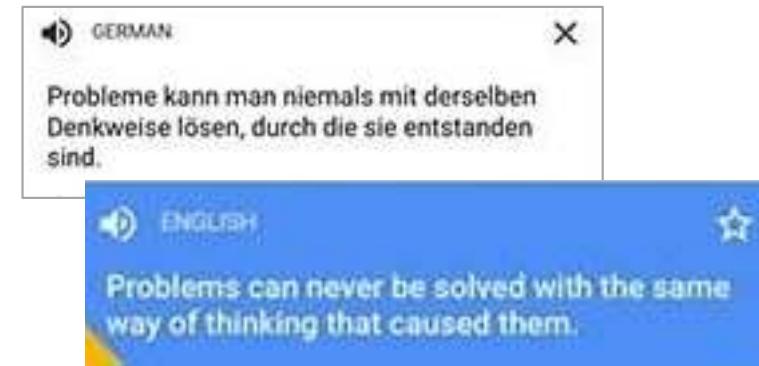
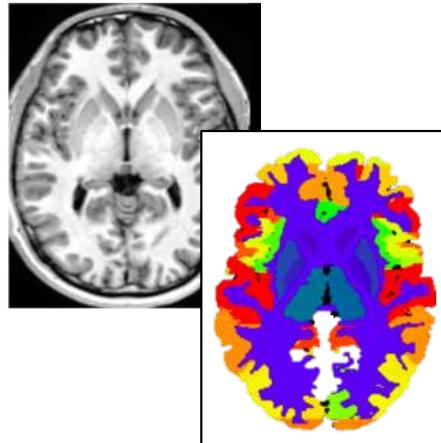
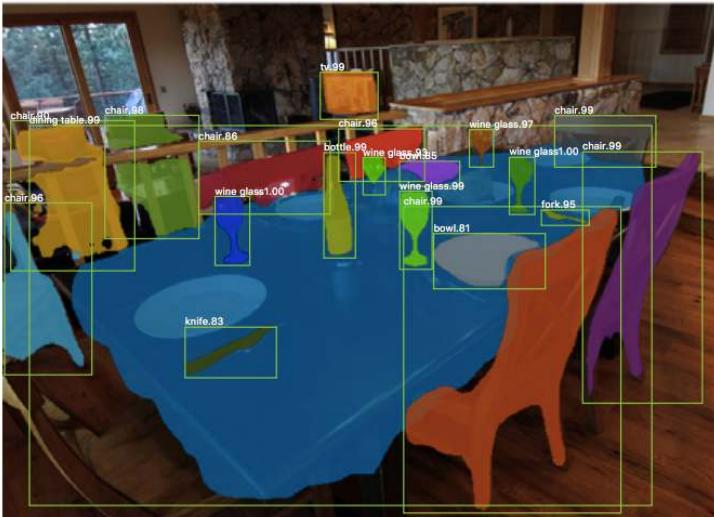
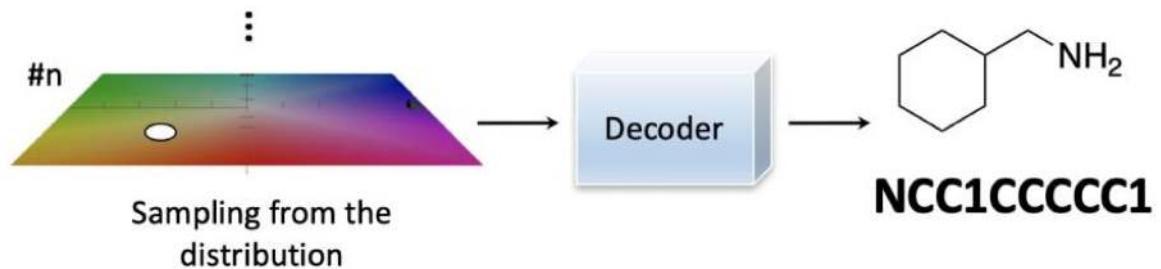
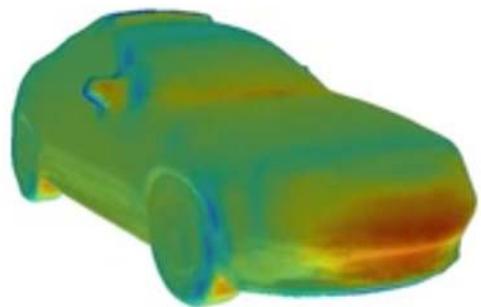


# Introduction to Deep Learning

Vincent Lepetit



Google Assistant



# Why Deep Learning is Currently so Popular?

- No need to engineer features;
- Very flexible framework. Originally developed for supervised learning, but can be extended to many other problems.
- *Why now?*
  - Faster computers (with GPUs); More training data; Better optimization algorithms; Easy to use and powerful libraries in Python; It took time to researchers to get convinced it actually works.

# Artificial Intelligence/Machine Learning/Deep Learning

## Artificial Intelligence

*Expert Systems*

$A^*$

*min-max*

## Machine Learning

*Nearest Neighbor classifier*

*Naive Bayes classifier*

*Support Vector Machines*

*Boosting*

*Random Forests*

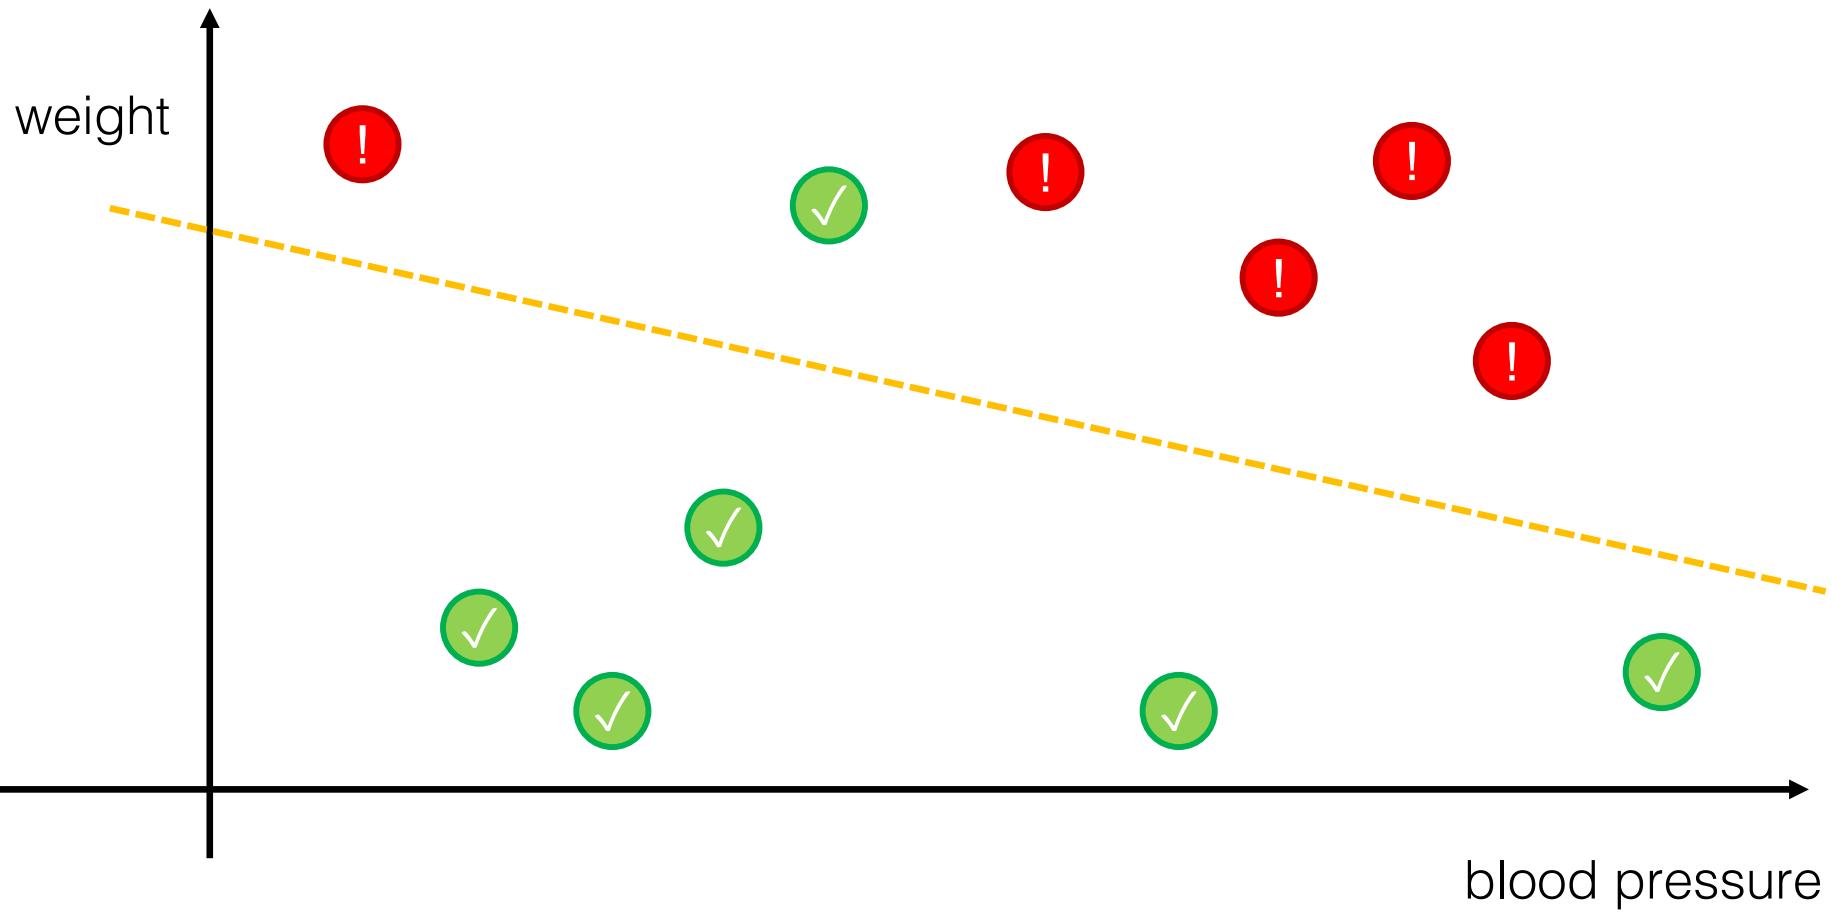
## Deep Learning

*Perceptron*

...

# From Early Approaches to Modern Deep Learning

# Linear Classifier / Perceptron



# formalization

The equation of the **separation** (boundary) is

$$w_1 \times (\text{blood pressure}) + w_2 \times \text{weight} + b = 0$$

Samples such that

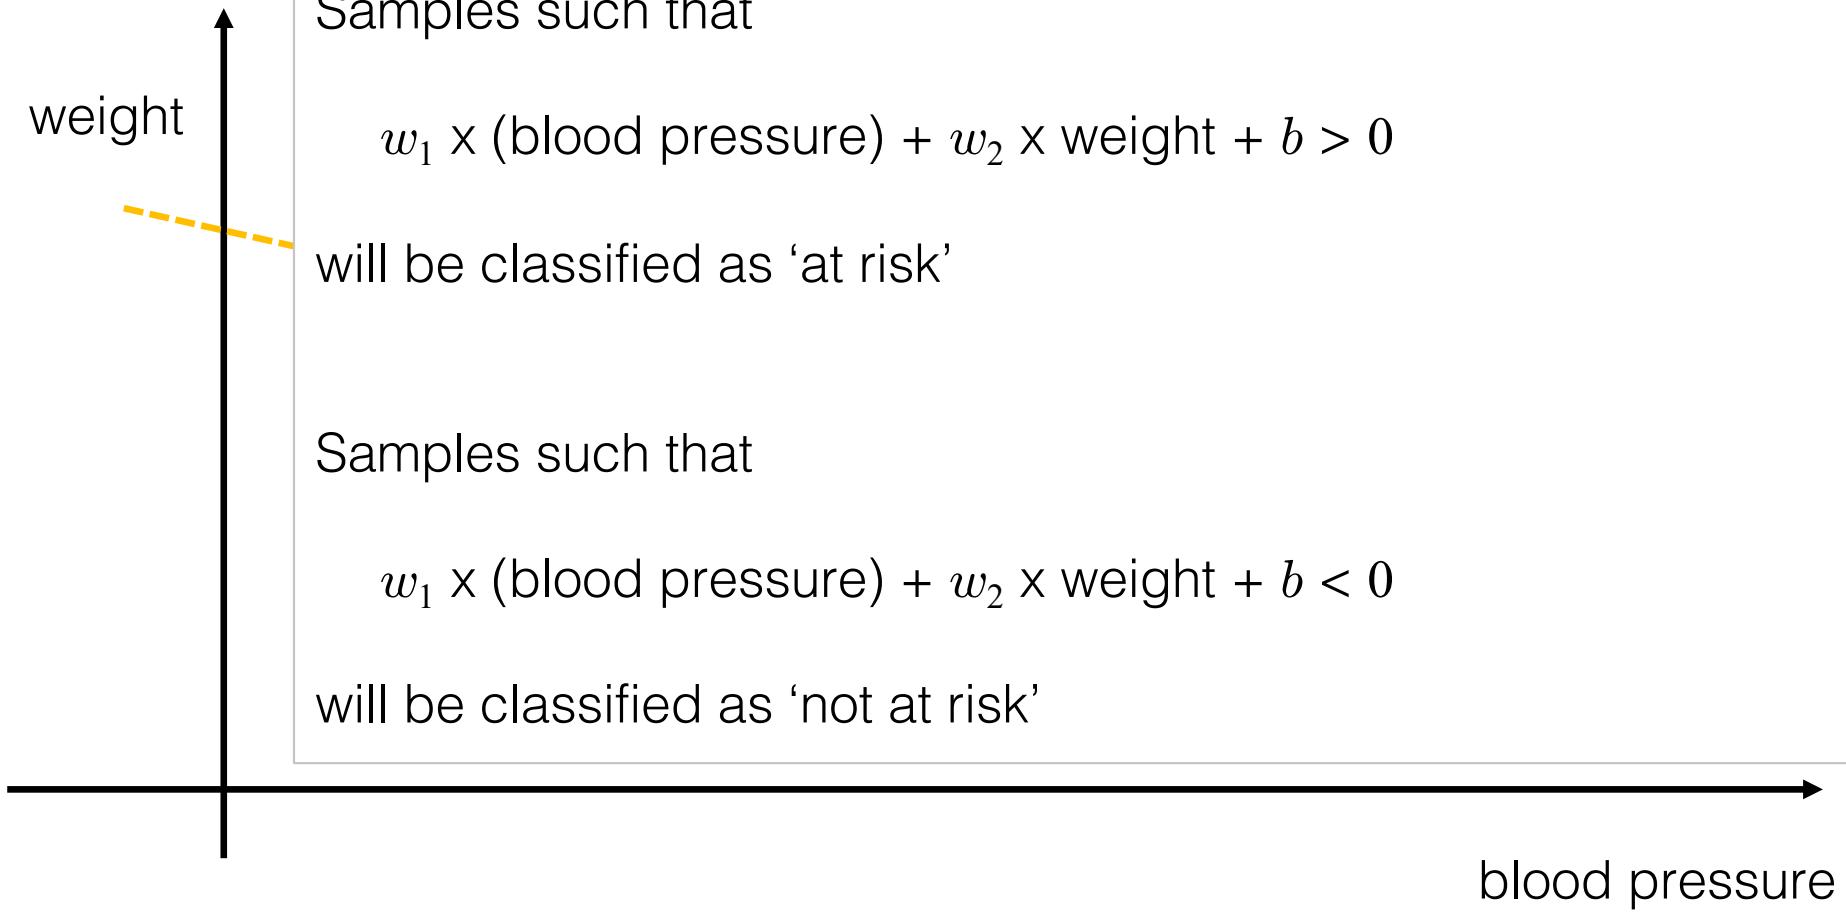
$$w_1 \times (\text{blood pressure}) + w_2 \times \text{weight} + b > 0$$

will be classified as 'at risk'

Samples such that

$$w_1 \times (\text{blood pressure}) + w_2 \times \text{weight} + b < 0$$

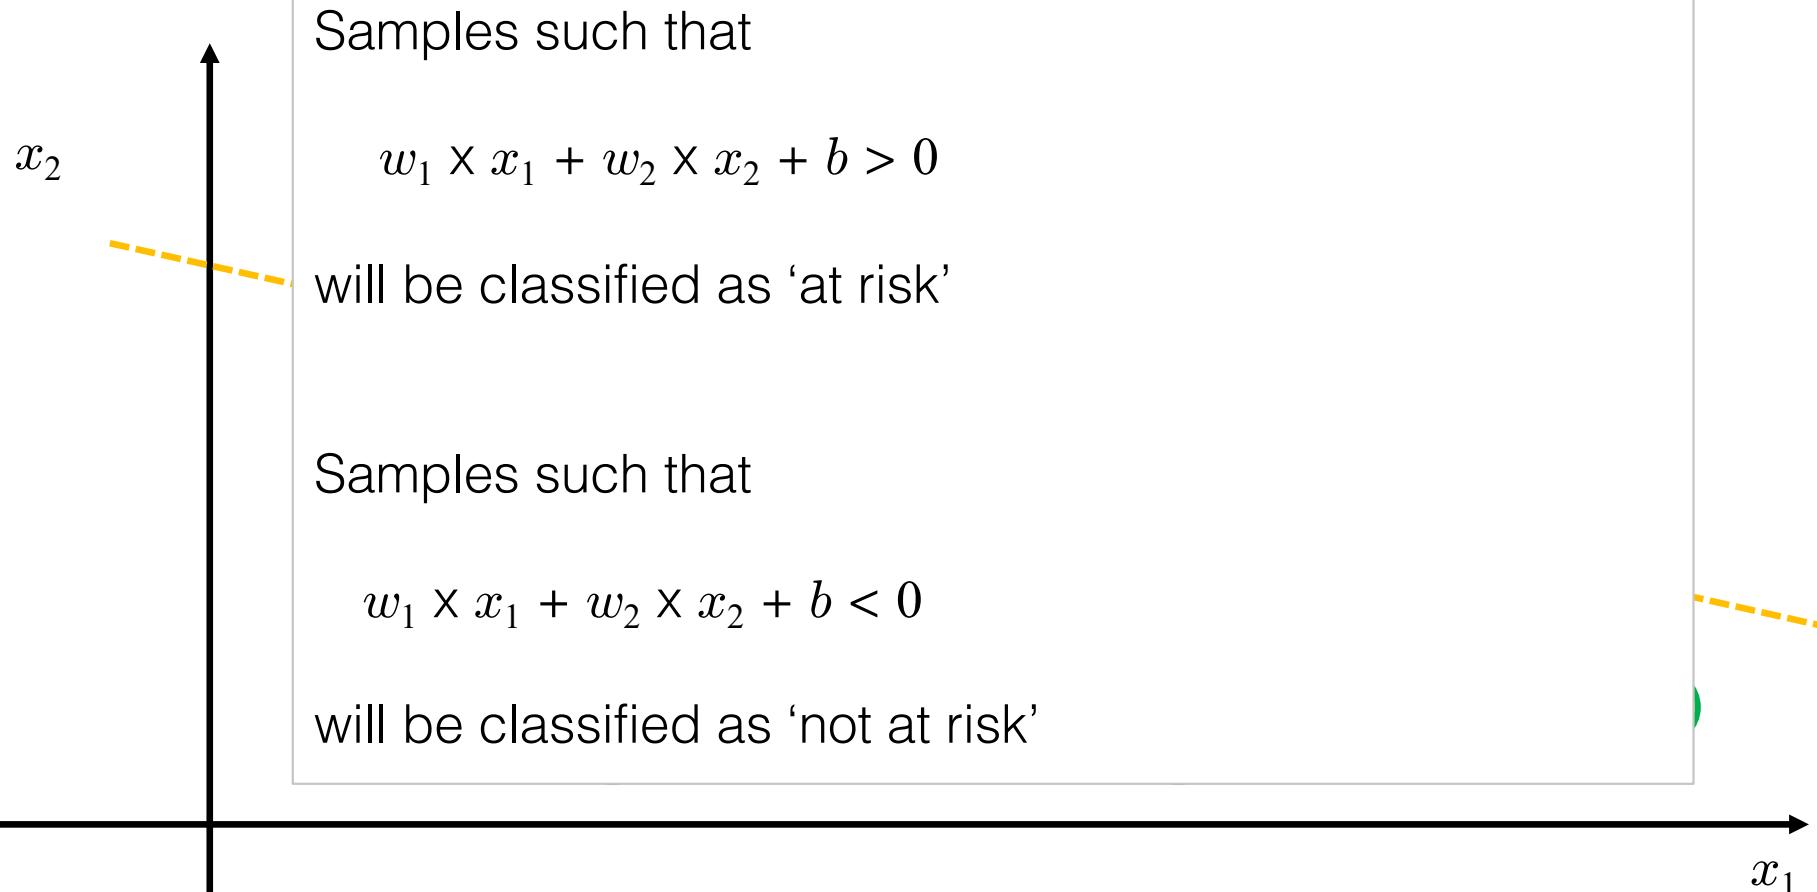
will be classified as 'not at risk'



# formalization (2)

The equation of the separation (boundary) is

$$w_1 \times x_1 + w_2 \times x_2 + b = 0$$



$$\begin{matrix} x_1 \\ x_2 \end{matrix} \quad \begin{matrix} w_1 \\ w_2 \end{matrix}$$

$$o = w_1x_1 + w_2x_2$$

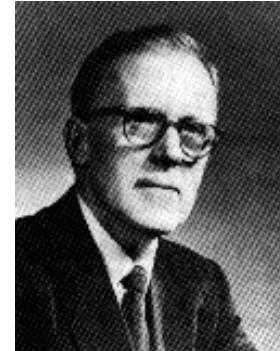
Prediction based on the sign of  $o$

$$\begin{array}{c}
 x_1 \\
 x_2 \\
 \vdots \\
 \vdots \\
 x_n
 \end{array}
 \quad
 \begin{array}{c}
 w_1 \\
 w_2 \\
 \vdots \\
 \vdots \\
 w_n
 \end{array}$$

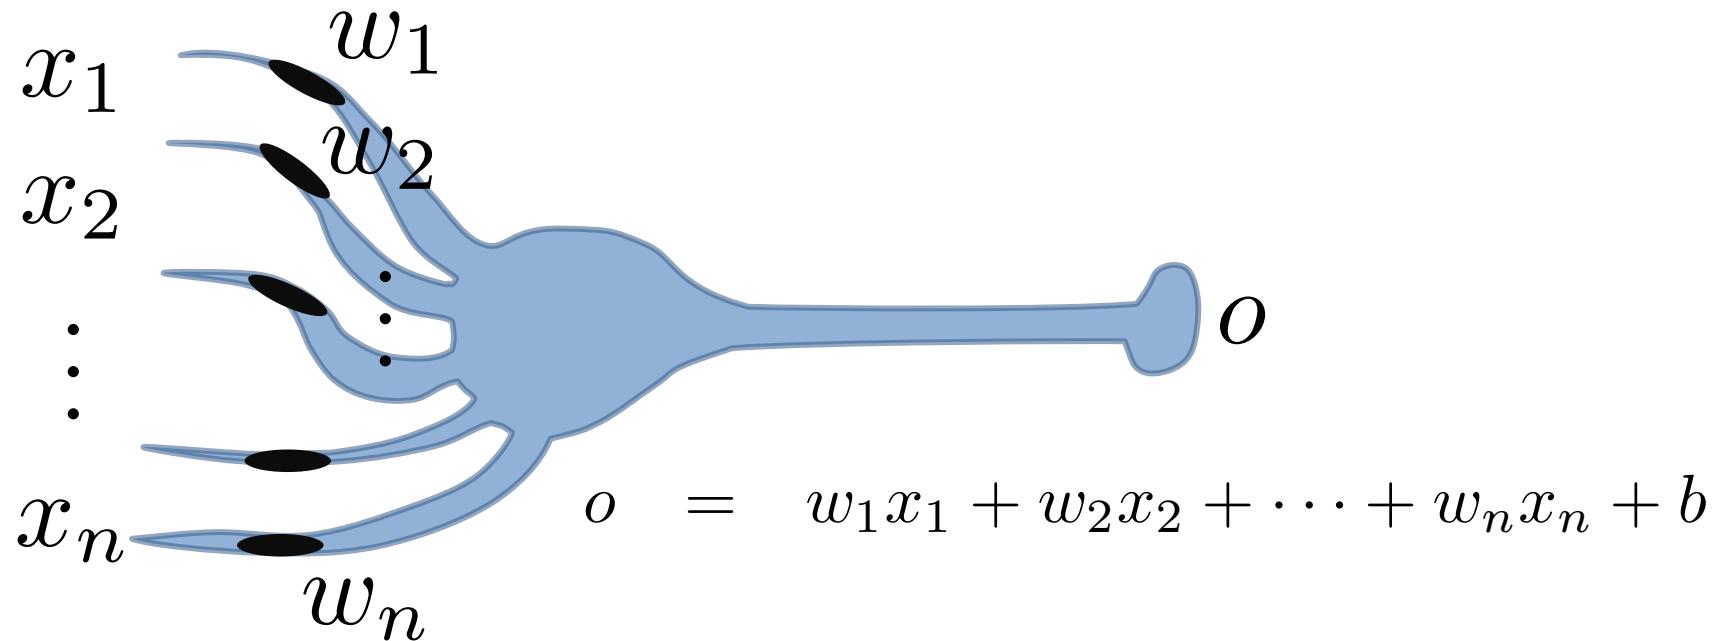
$$o = w_1x_1 + w_2x_2 + \cdots + w_nx_n + b$$

Prediction based on the sign of  $o$

# PERCEPTRON



Inspired by the work of  
Donald Hebb (1949)



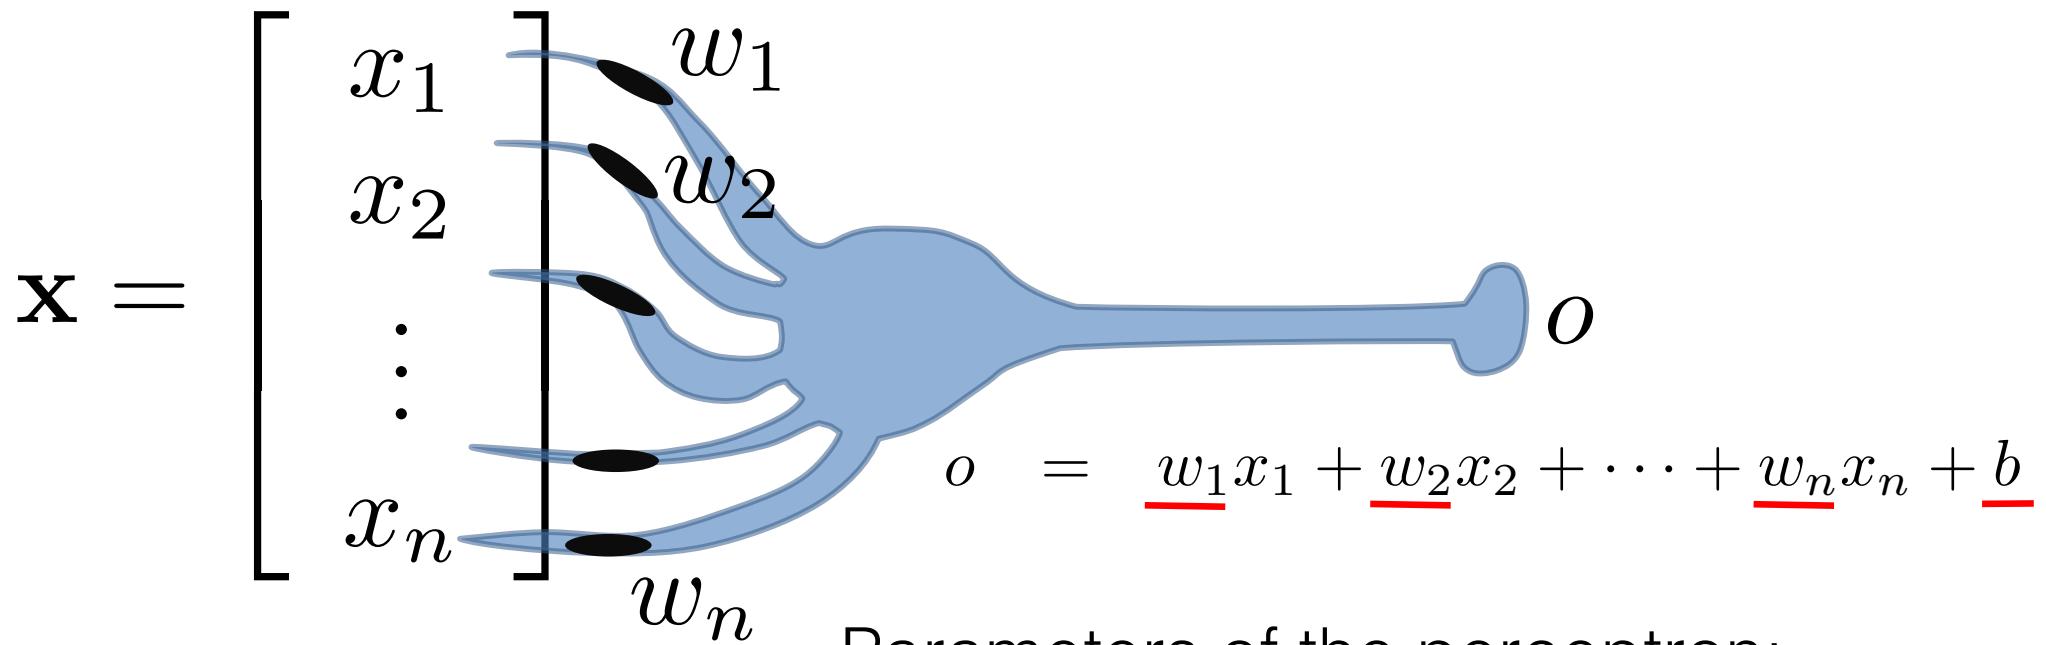
Prediction based on the sign of  $O$

# PERCEPTRON

$$\mathbf{x} = \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{bmatrix} \quad \begin{aligned} o &= w_1x_1 + w_2x_2 + \cdots + w_nx_n + b \\ &= \sum_i w_i x_i + b \\ &= \mathbf{w}^\top \mathbf{x} + b \end{aligned}$$

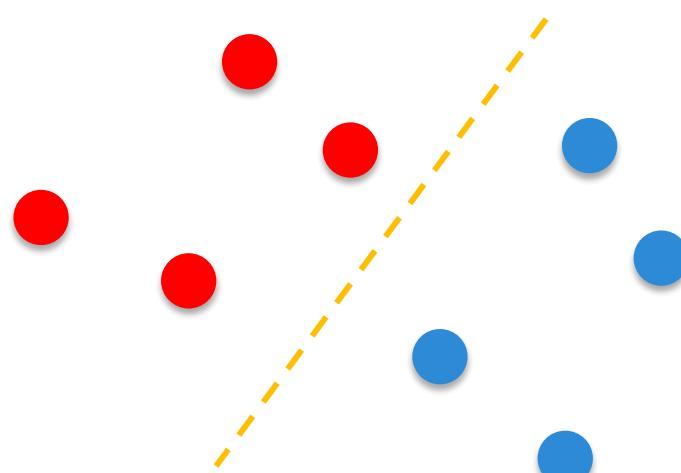
Prediction based on the sign of  $o$

# Perceptron

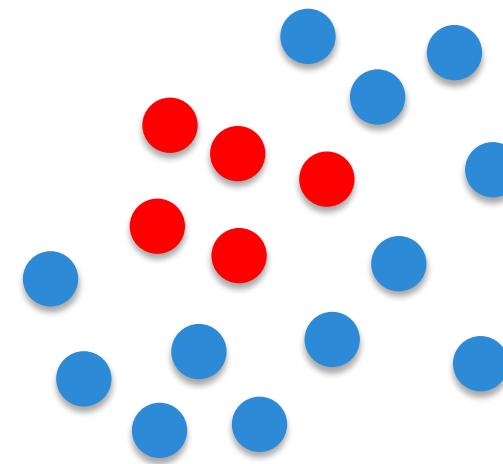


Parameters of the perceptron:  
we need to find good values for them  
(we will see that later)

A perceptron can only correctly classify data points that are linearly separable:

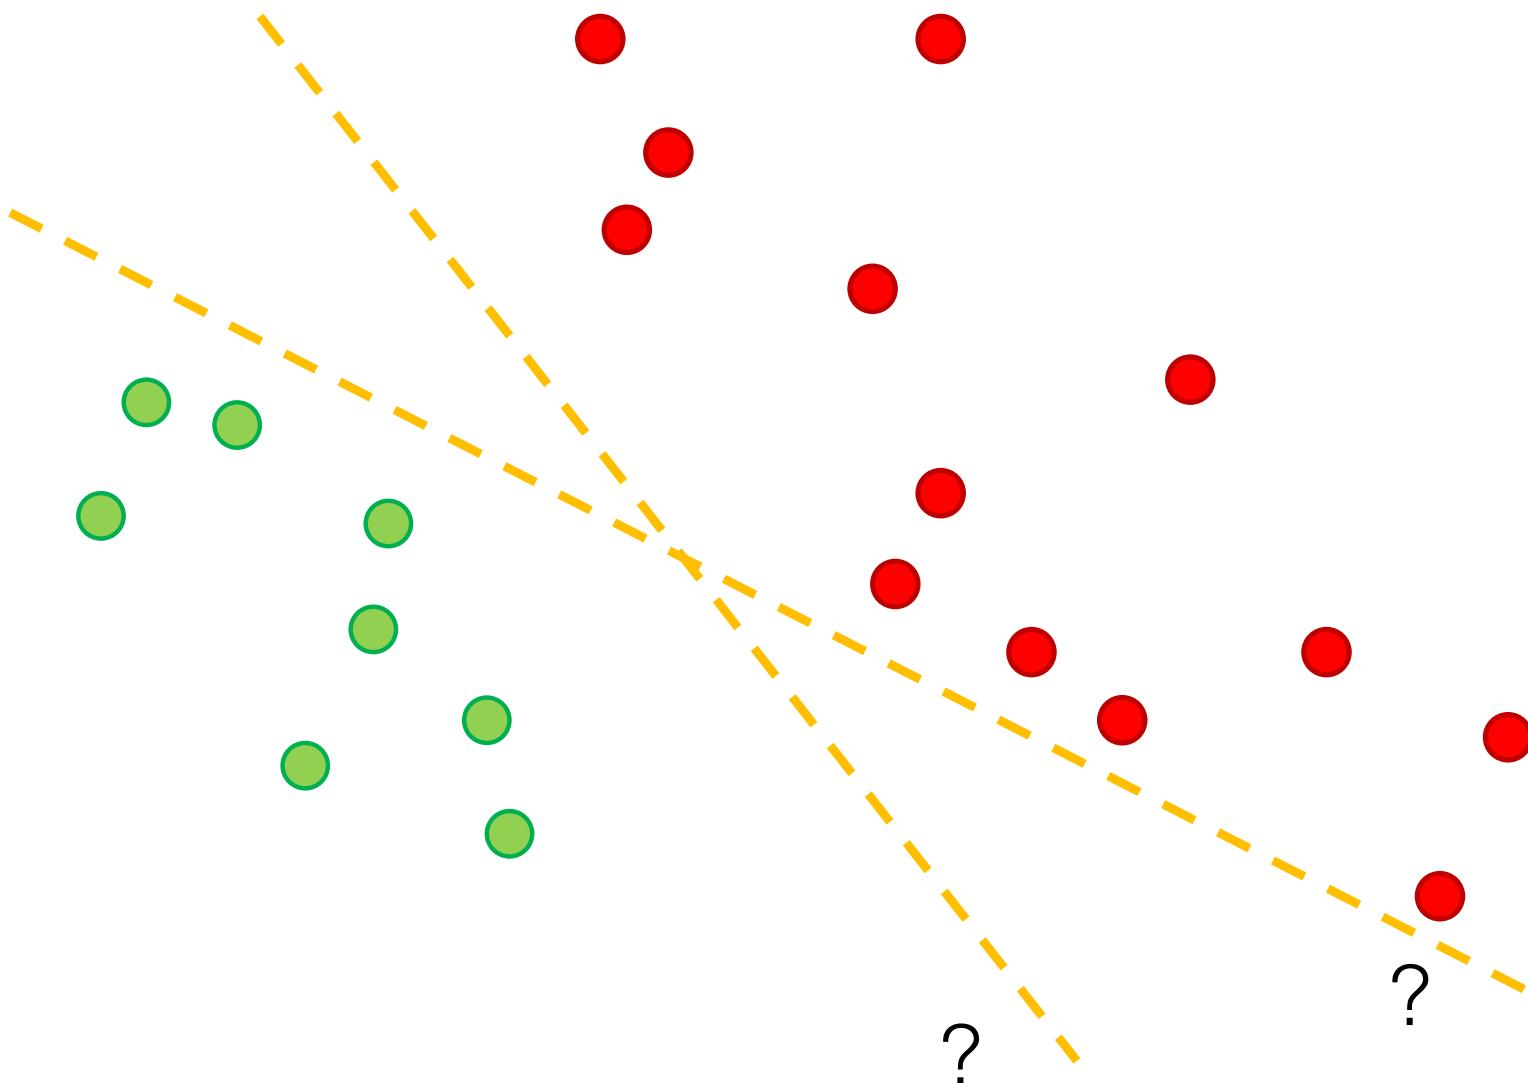


linearly separable

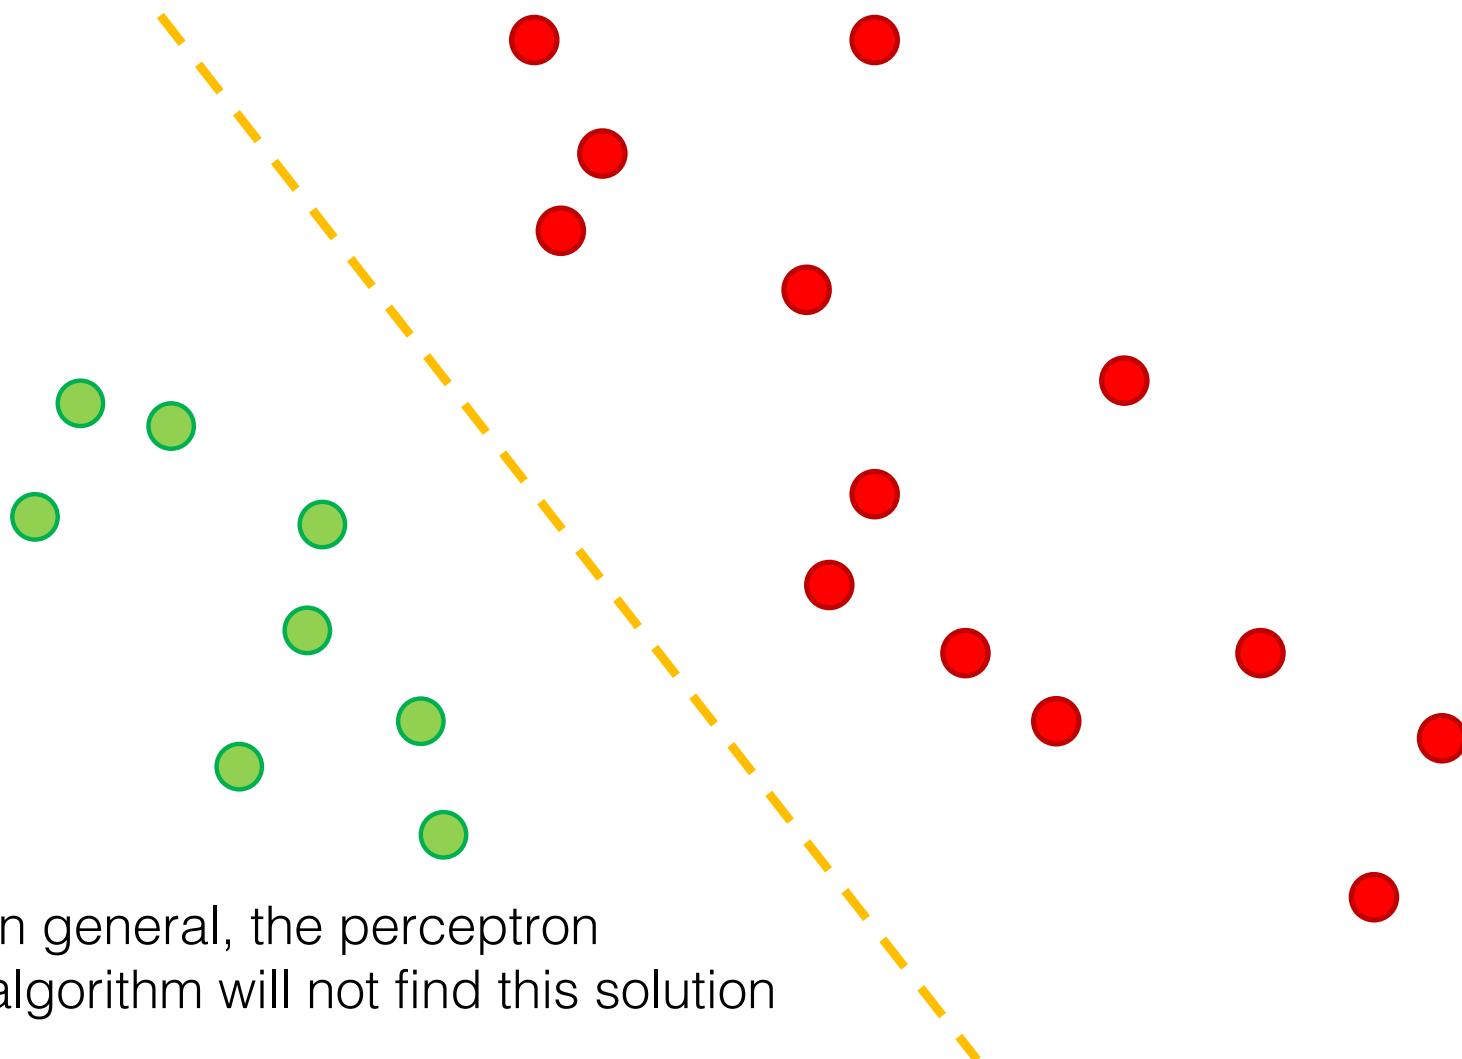


nonlinearly separable

# What is the Best Linear Classifier?



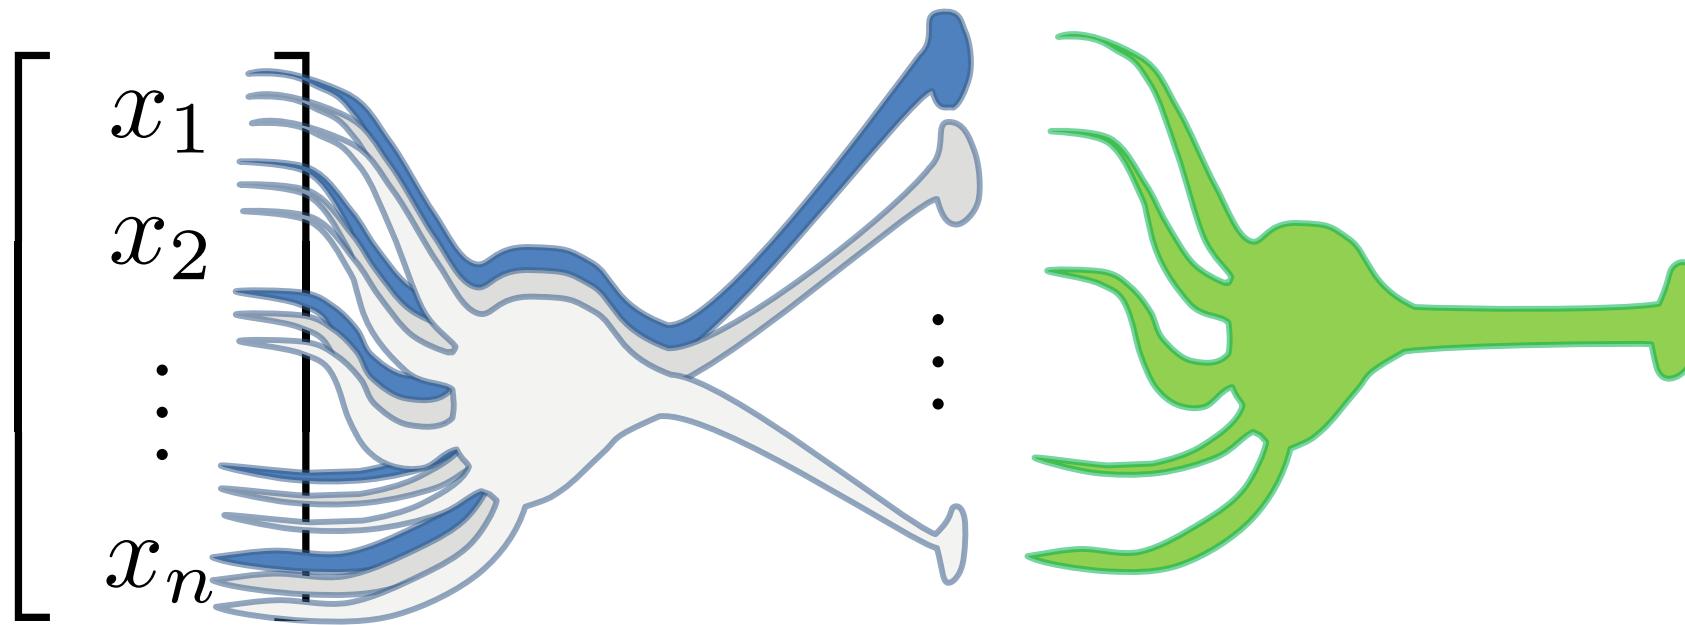
# What is the Best Linear Classifier?



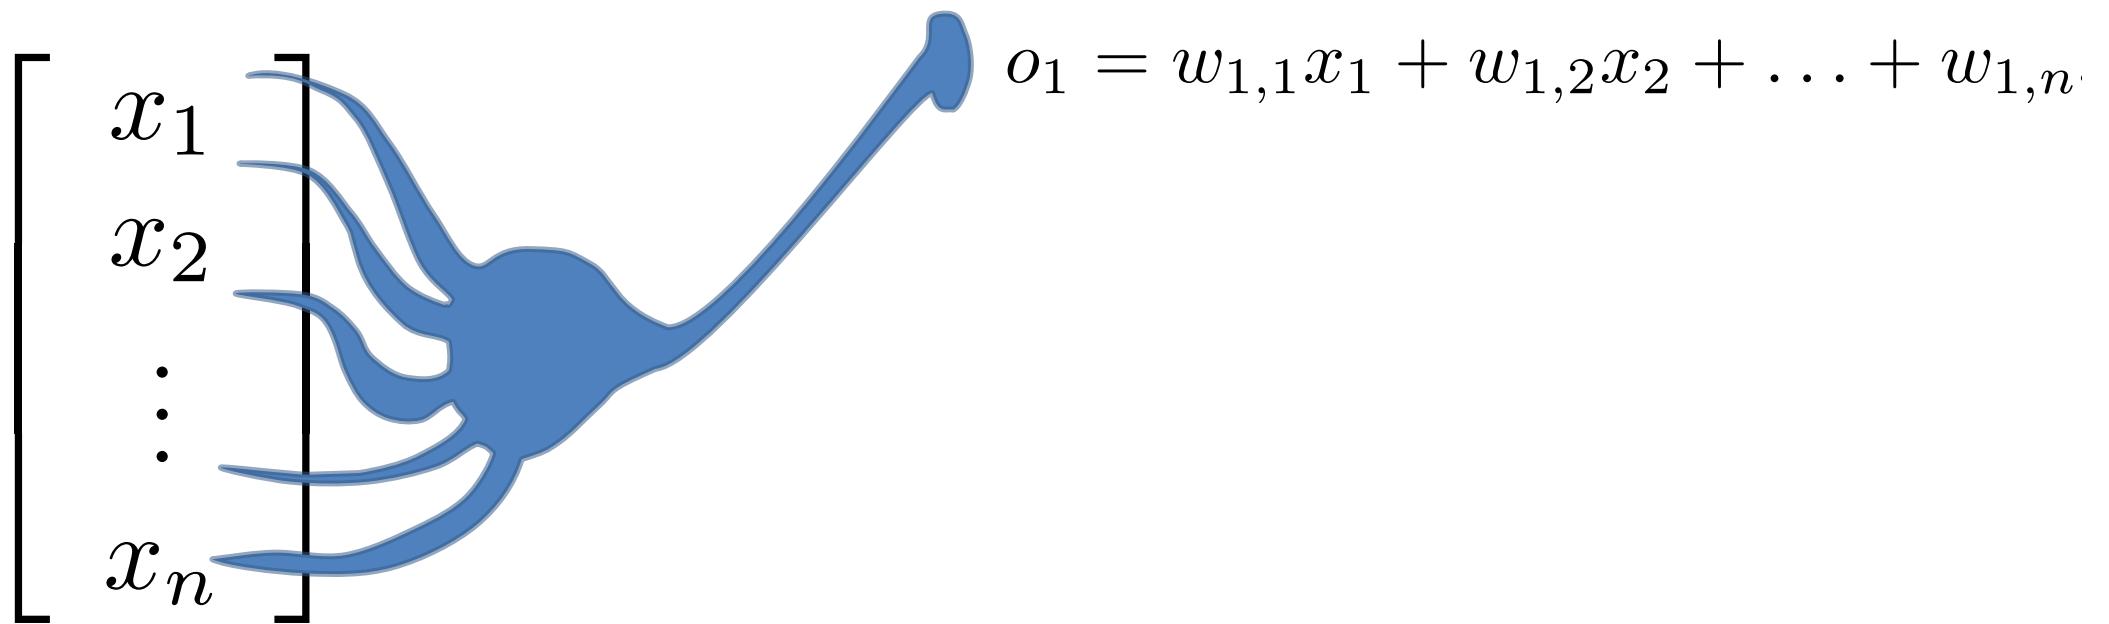
TWO-LAYER NETWORKS,

MULTI-LAYER PERCEPTRONS (MLP)

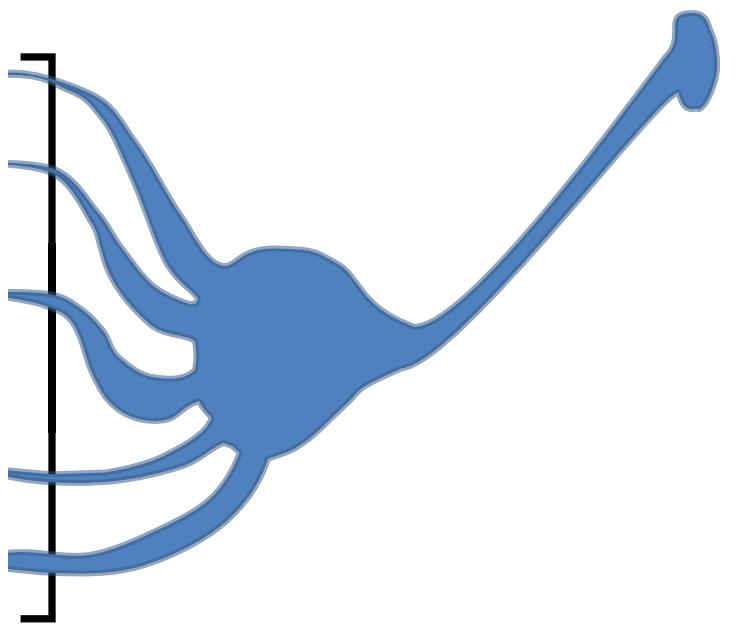
## Two-Layer Network (~1980)



## Two-Layer Network (~1980)

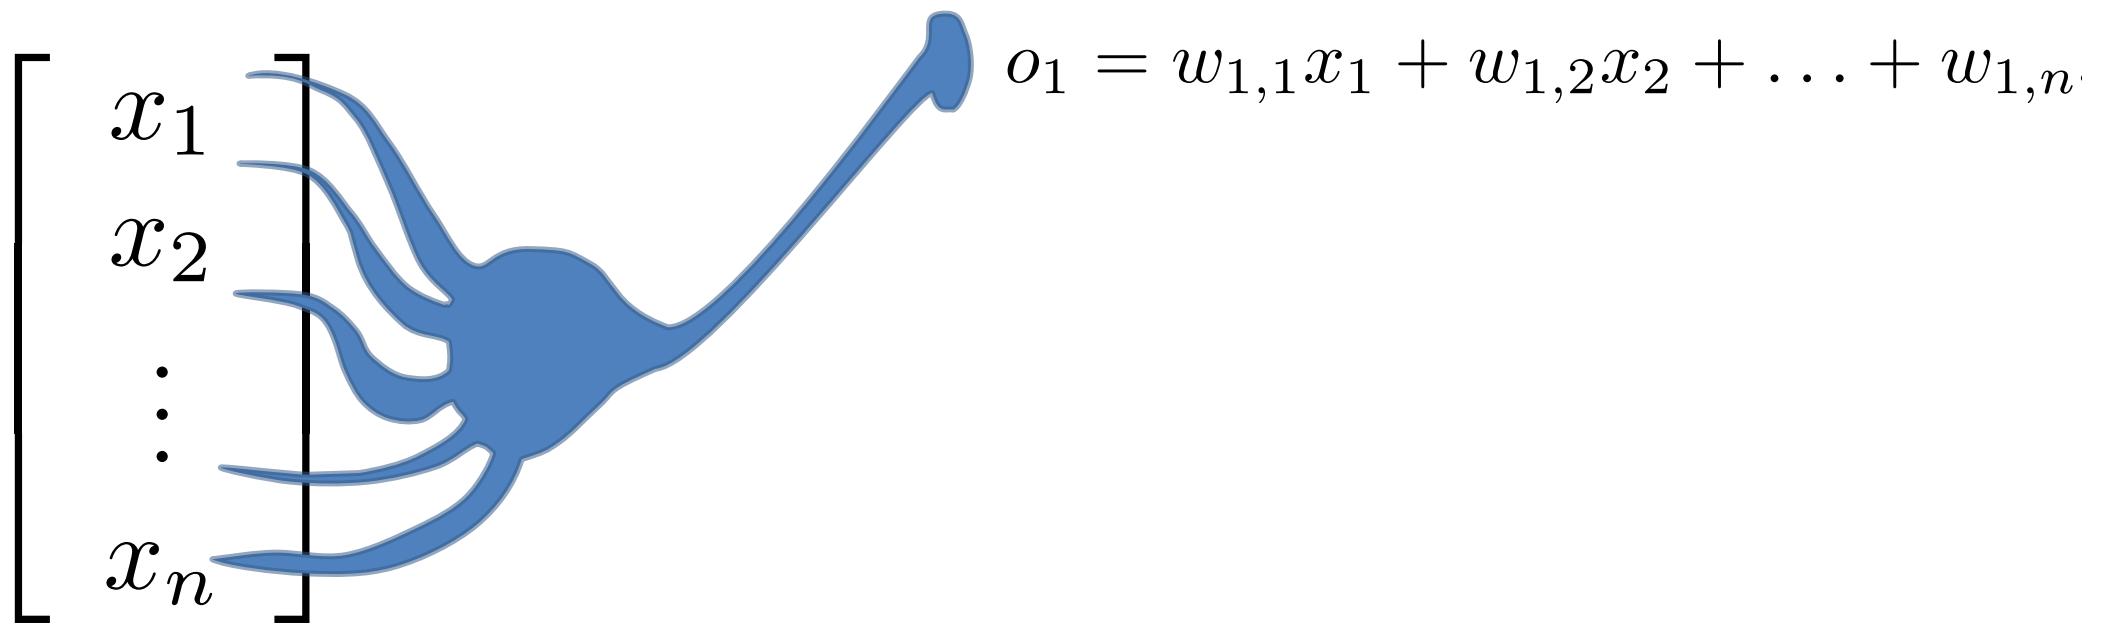


## Two-Layer Network (~1980)

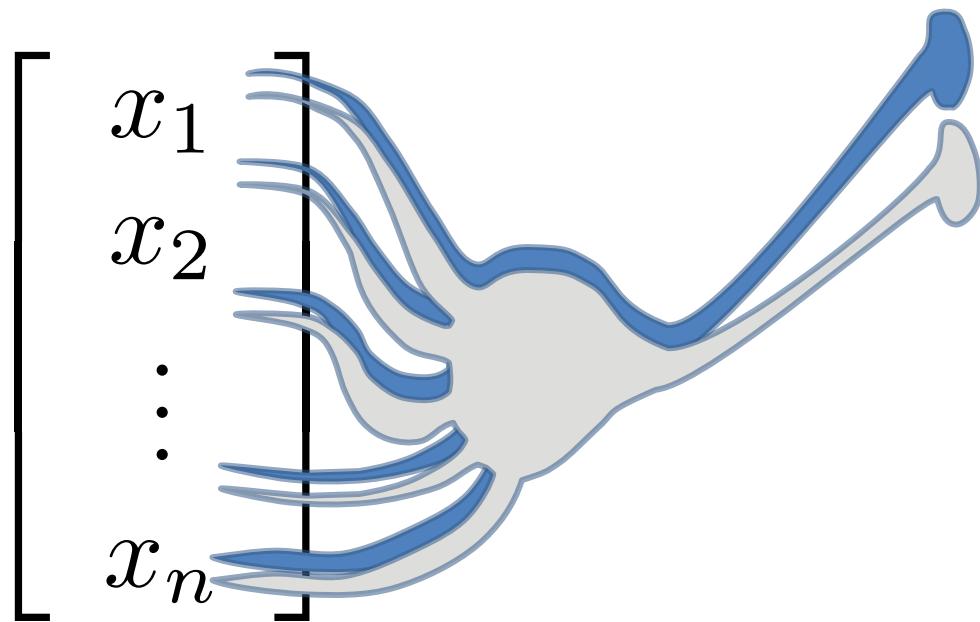


$$o_1 = w_{1,1}x_1 + w_{1,2}x_2 + \dots + w_{1,n}x_n + b_1$$

## Two-Layer Network (~1980)



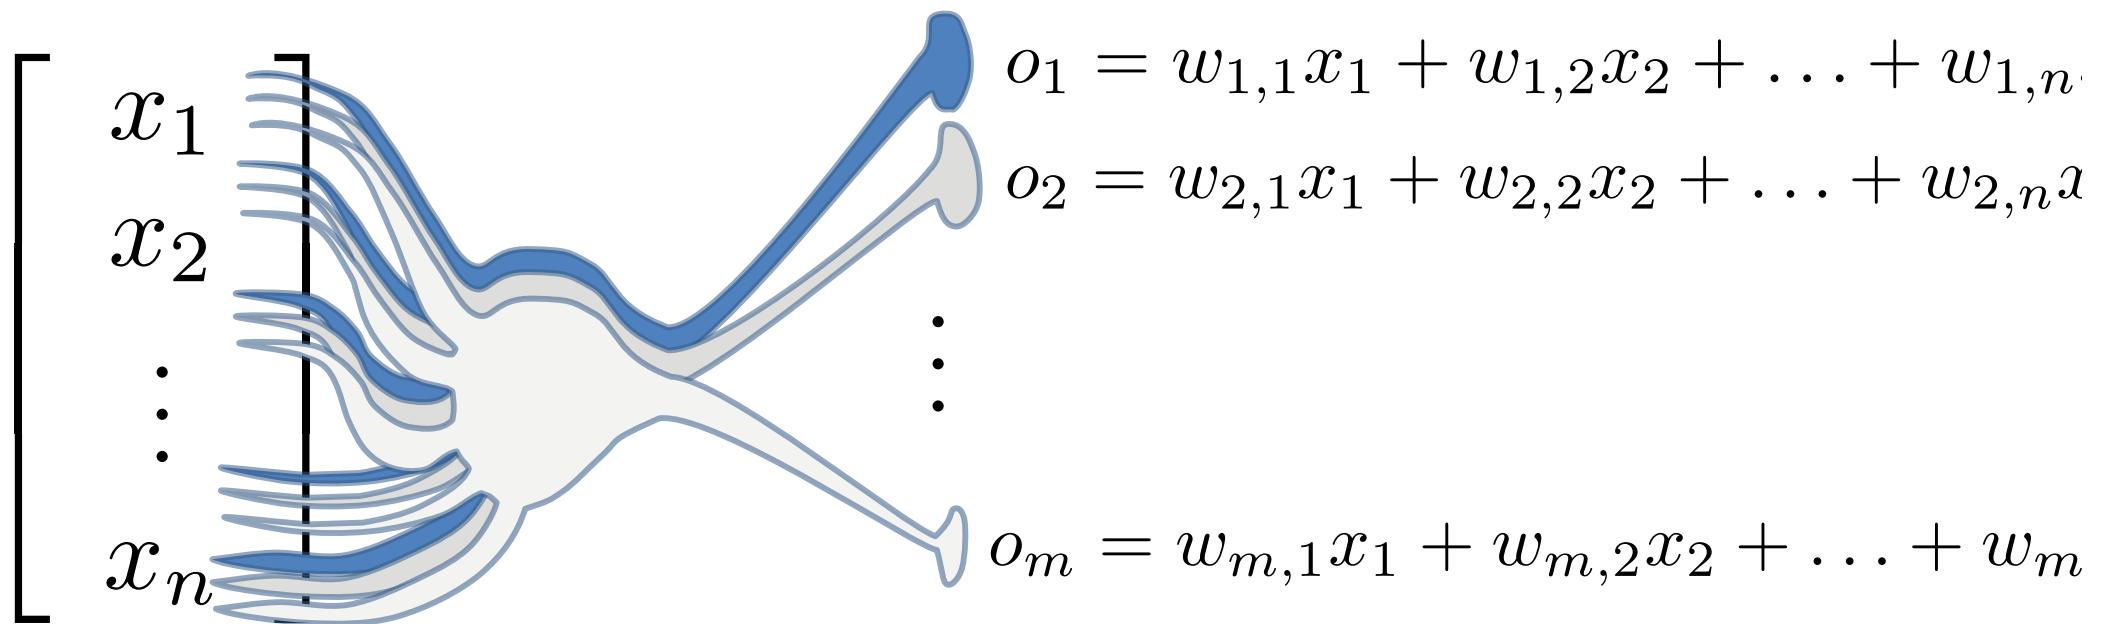
## Two-Layer Network (~1980)



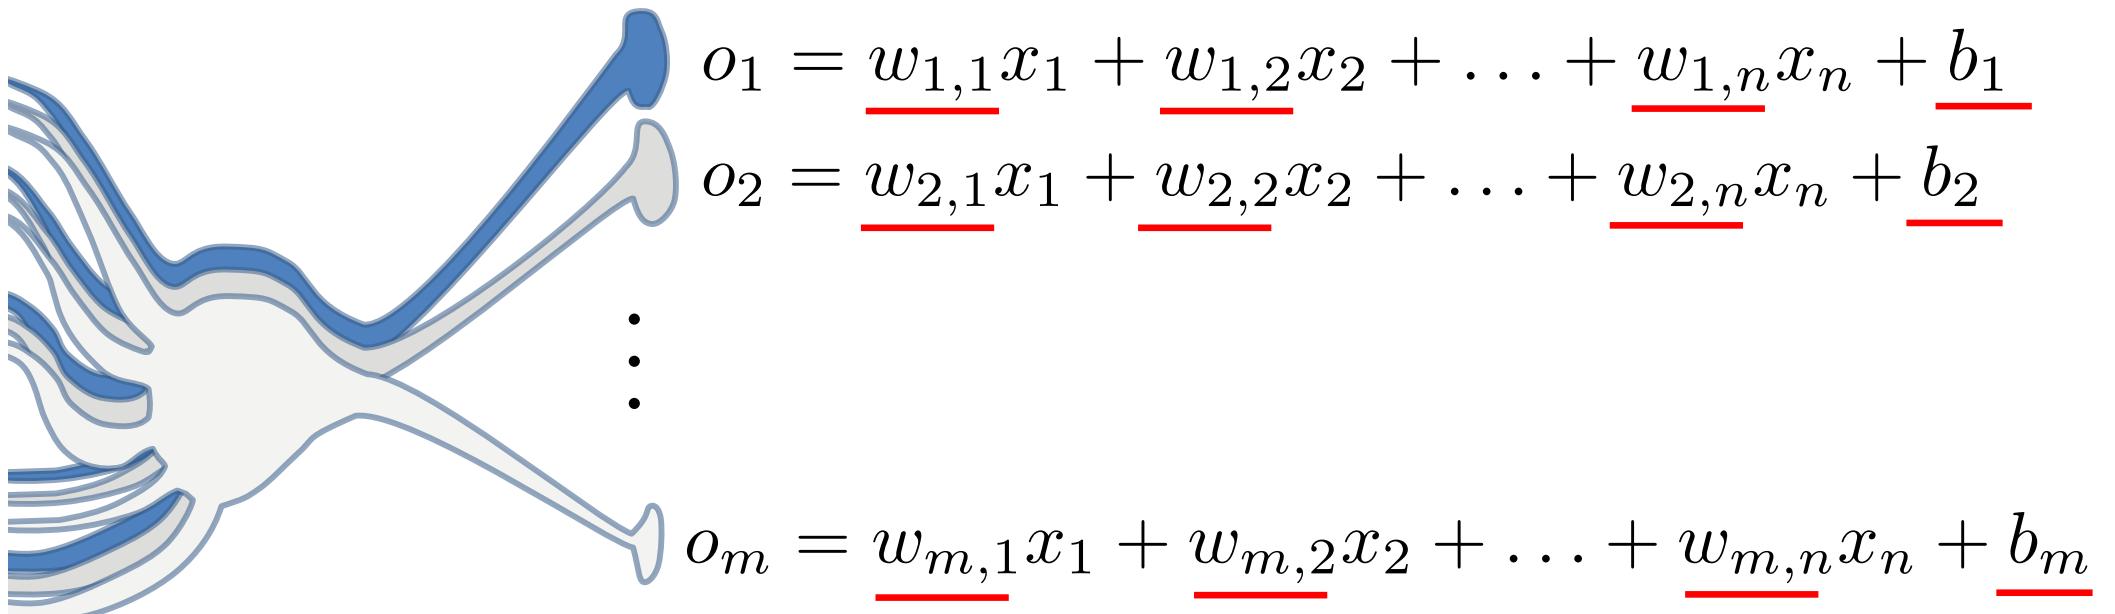
$$o_1 = w_{1,1}x_1 + w_{1,2}x_2 + \dots + w_{1,n}x_n$$

$$o_2 = w_{2,1}x_1 + w_{2,2}x_2 + \dots + w_{2,n}x_n$$

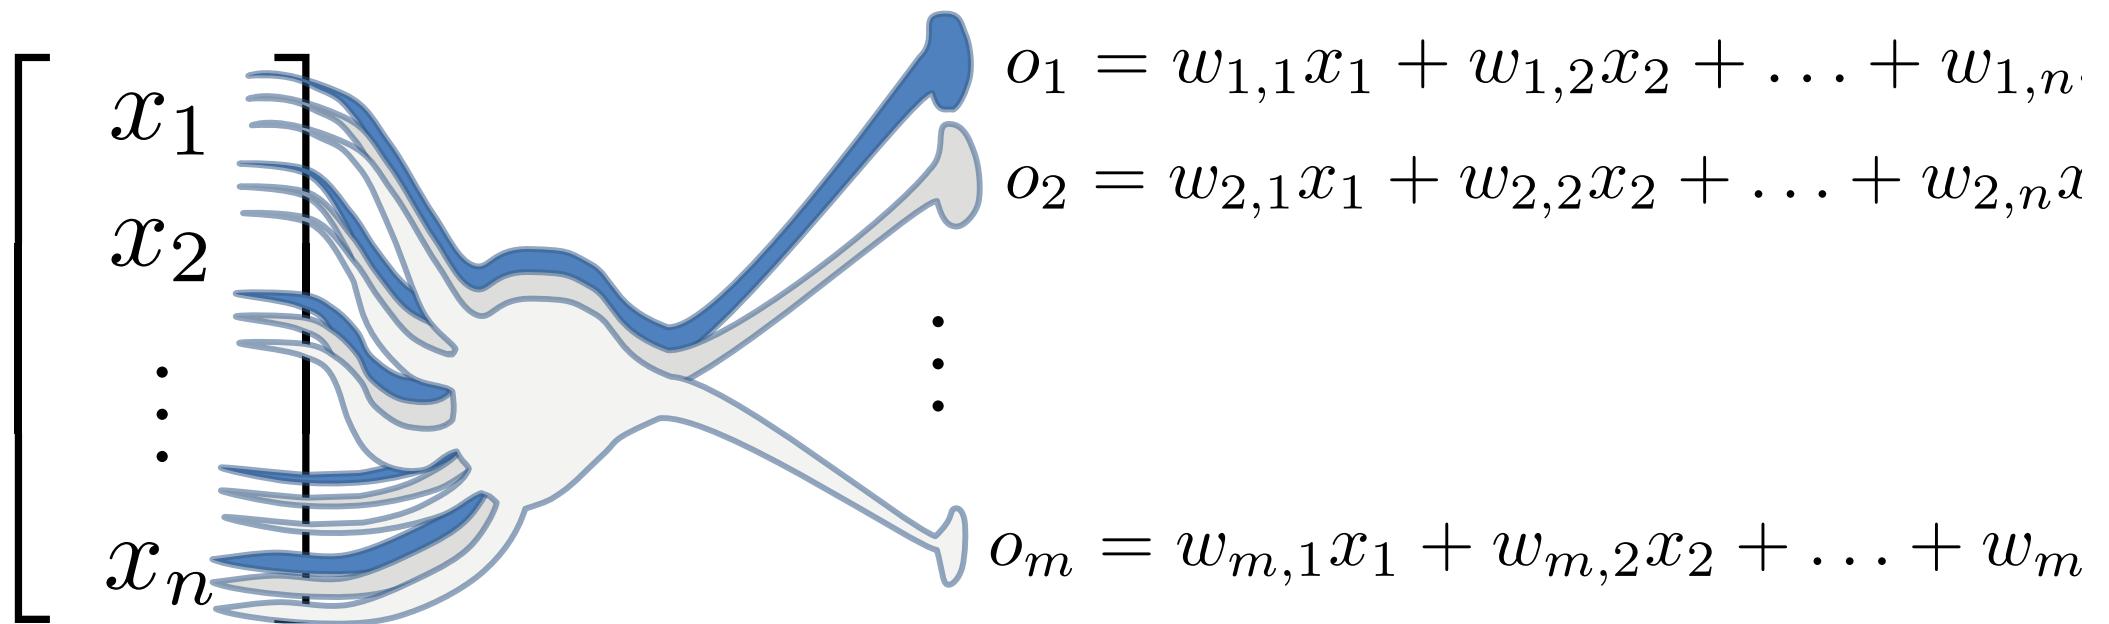
## Two-Layer Network (~1980)



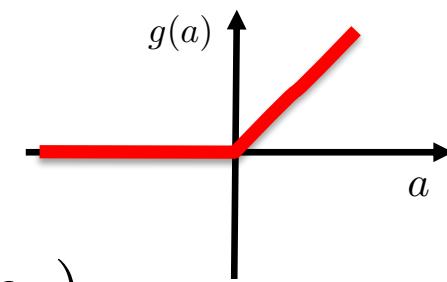
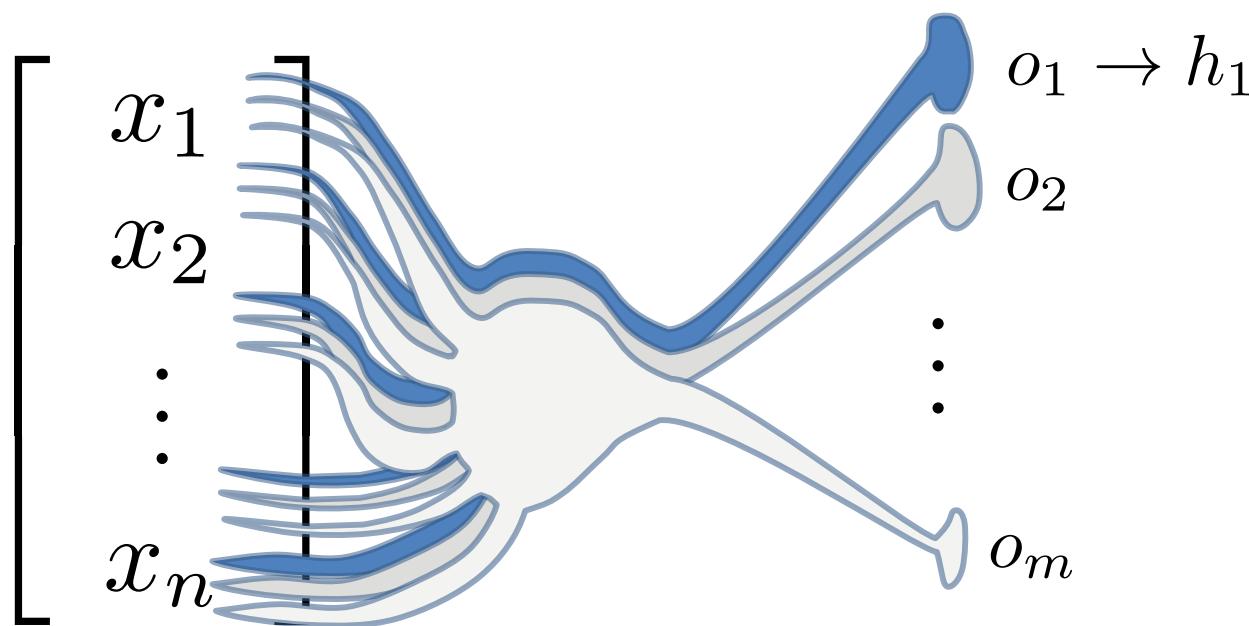
## the first network's parameters



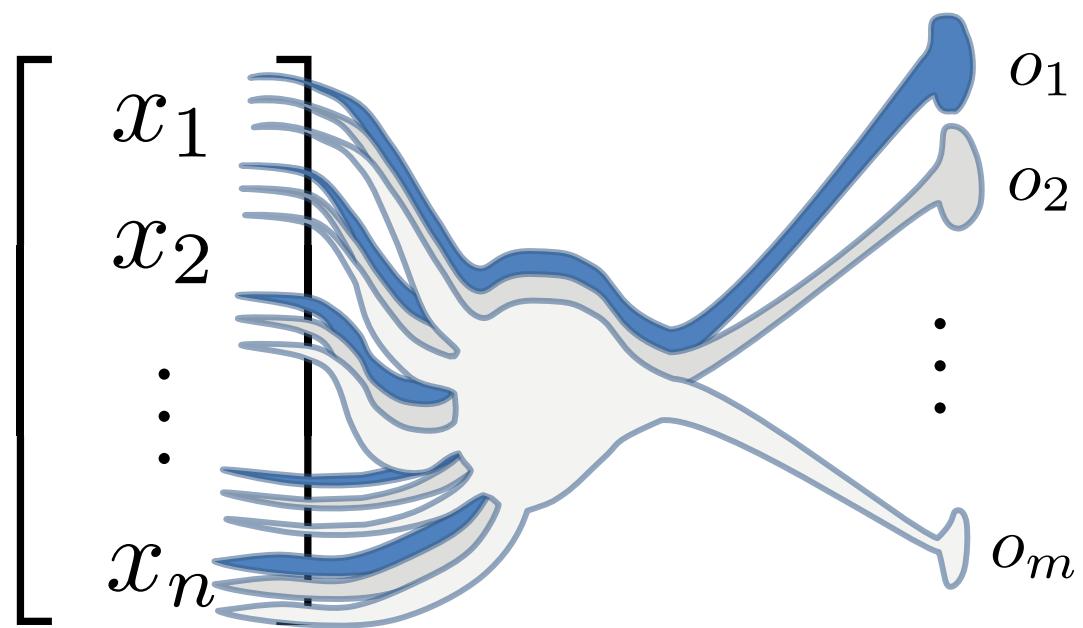
## Two-Layer Network (~1980)



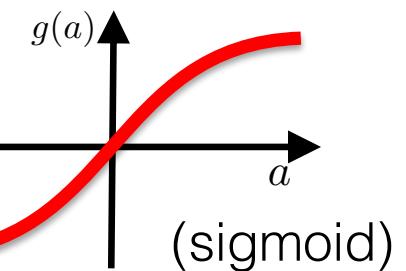
## Two-Layer Network (~1980)



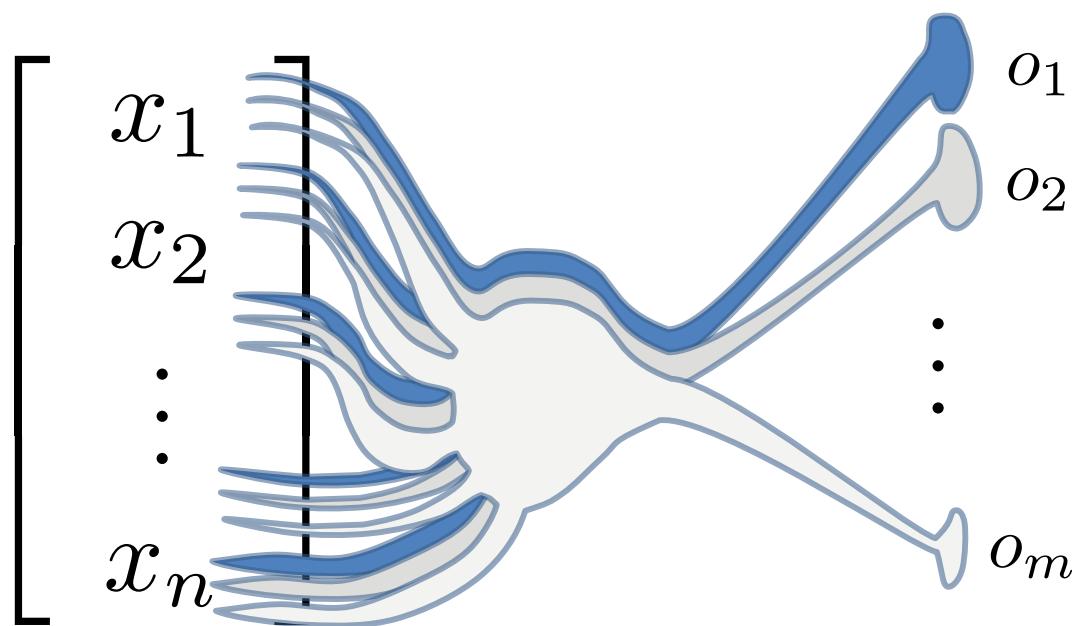
## Two-Layer Network (~1980)



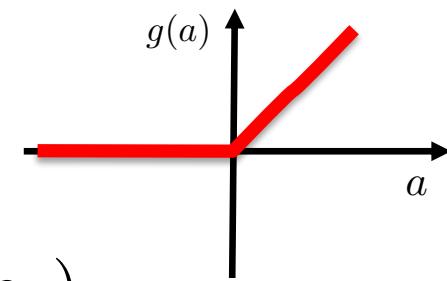
$$o_1 \rightarrow h_1 = g(o_1)$$



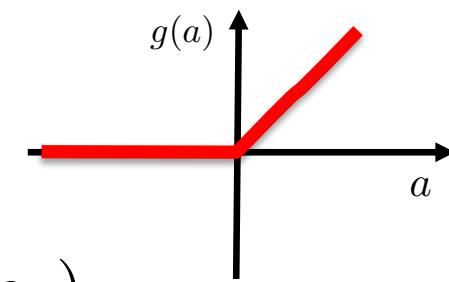
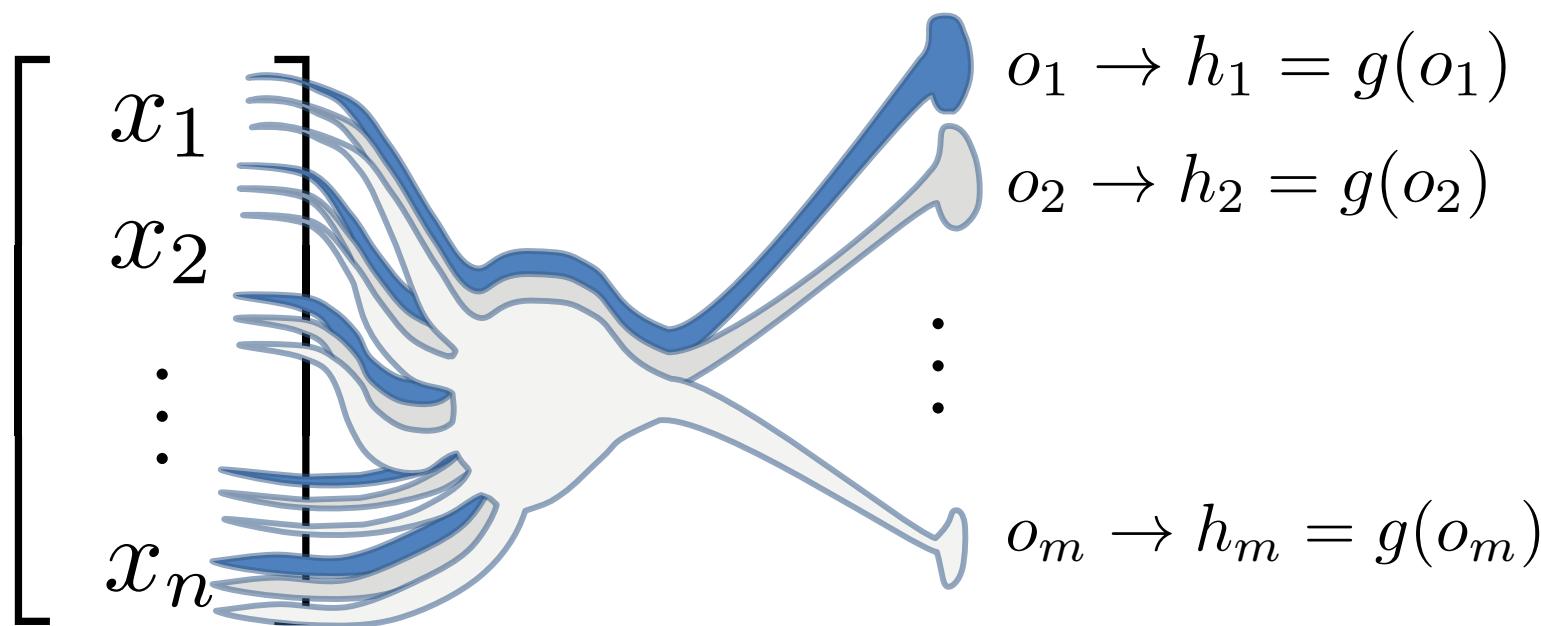
## Two-Layer Network (~1980)



$$o_1 \rightarrow h_1 = g(o_1)$$



## Two-Layer Network (~1980)

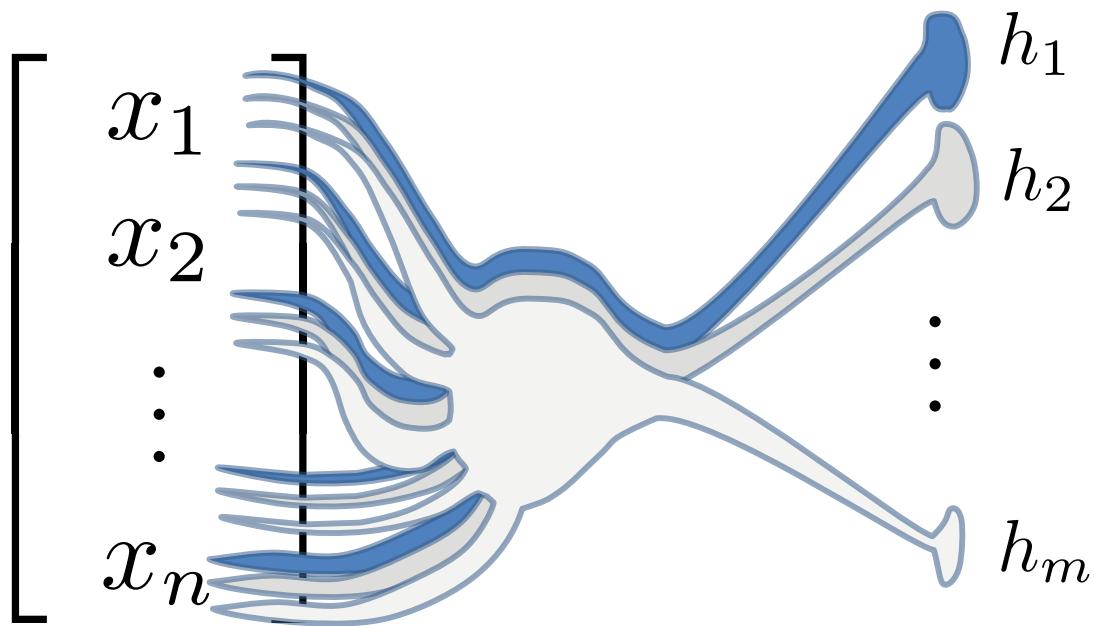


$$o_1 \rightarrow h_1 = g(o_1)$$

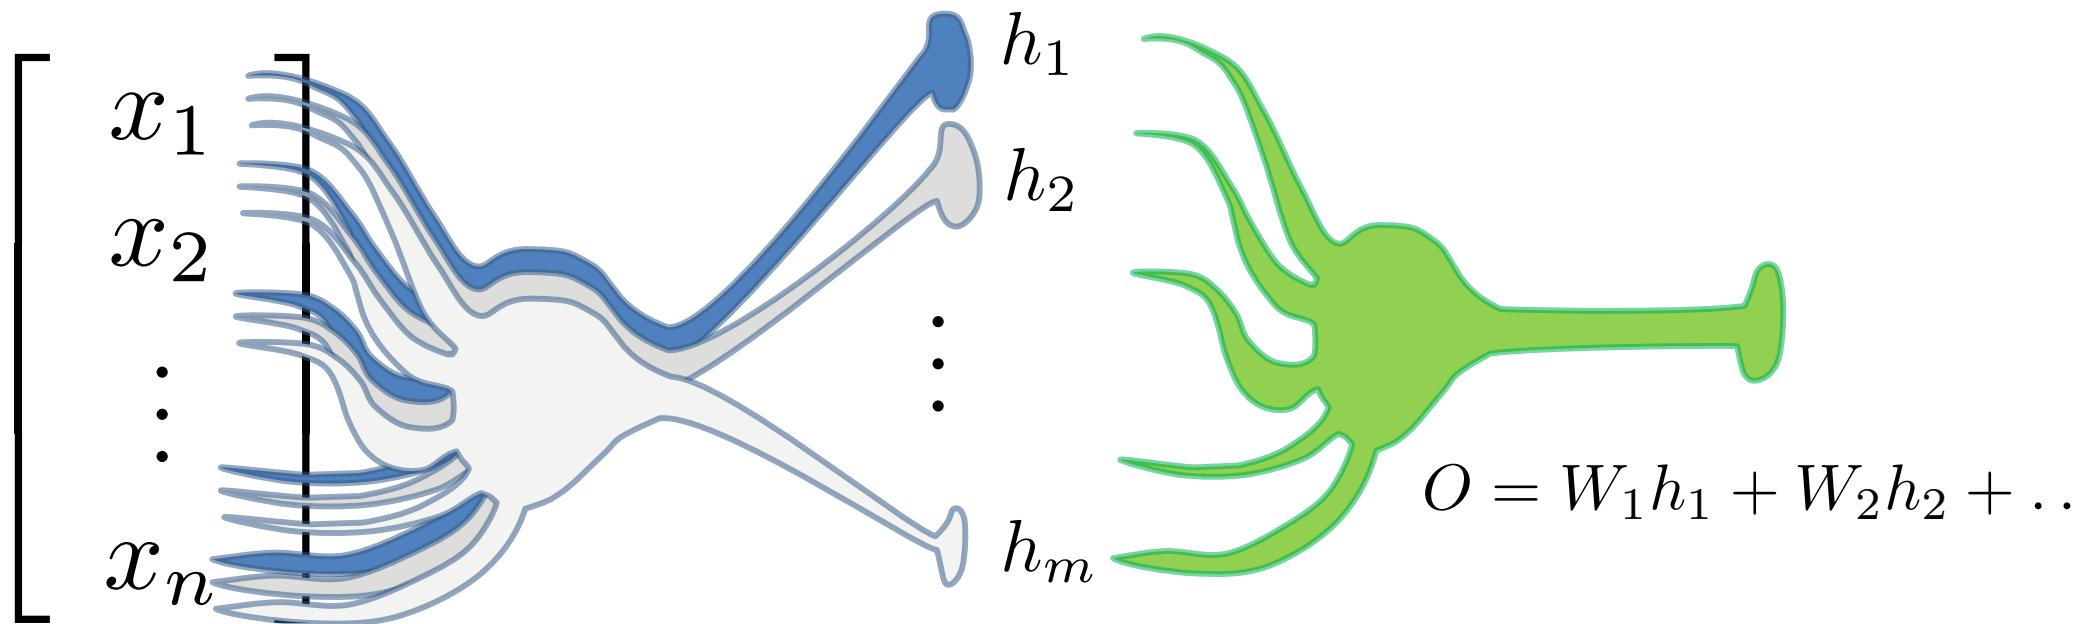
$$o_2 \rightarrow h_2 = g(o_2)$$

$$o_m \rightarrow h_m = g(o_m)$$

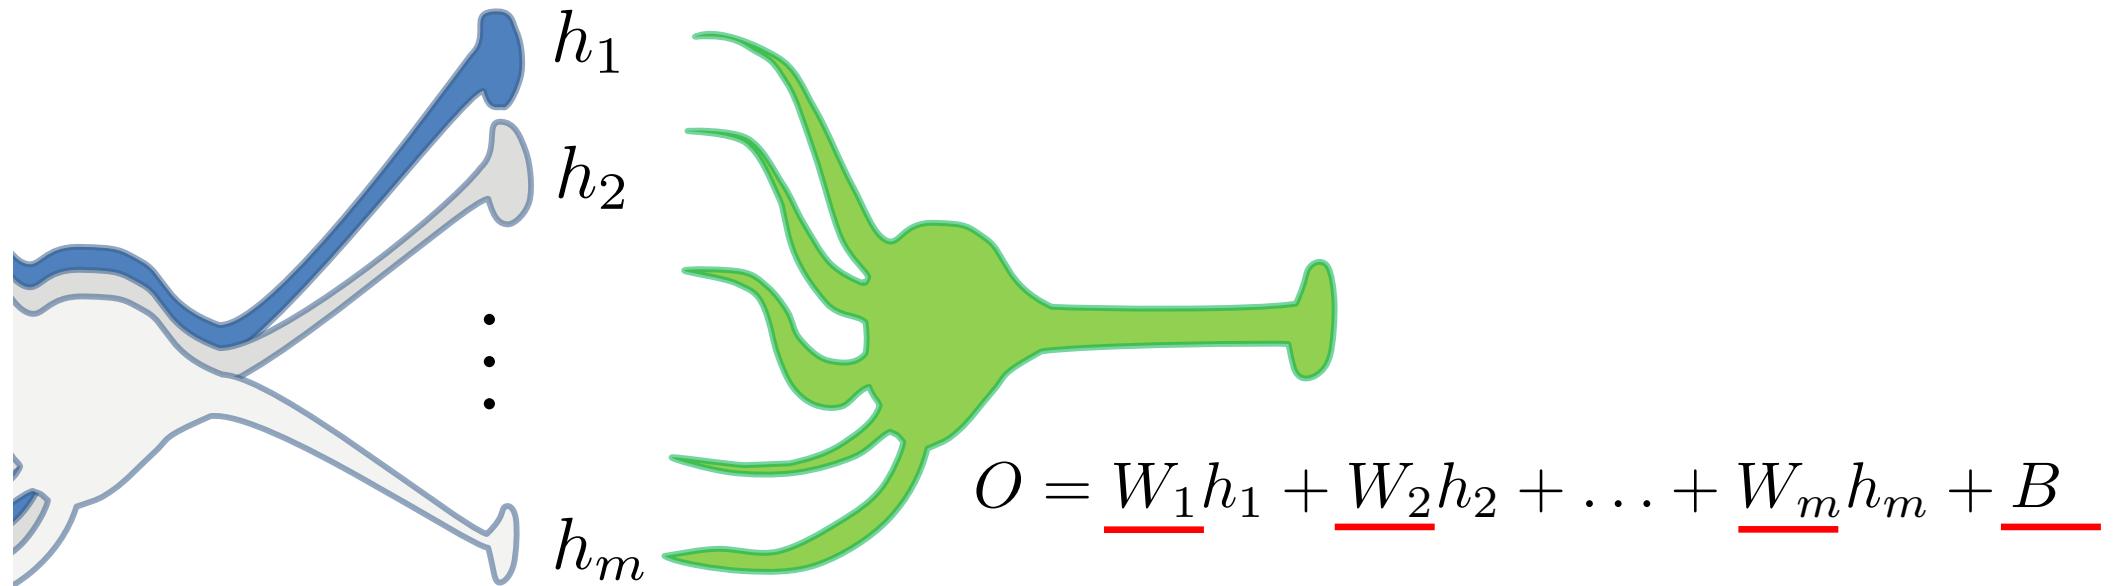
## Two-Layer Network (~1980)



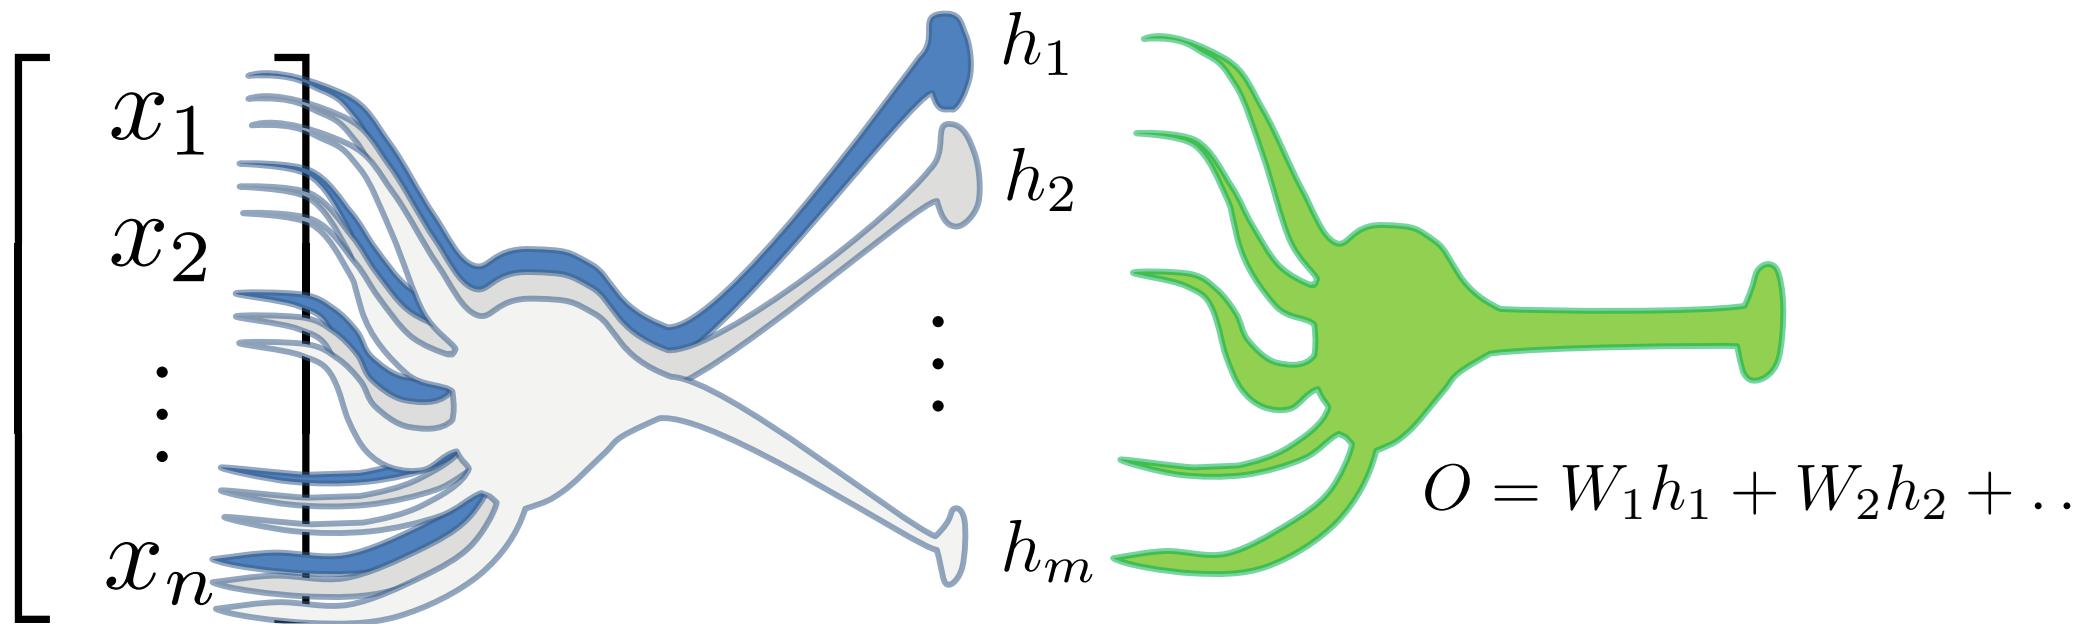
## Two-Layer Network (~1980)



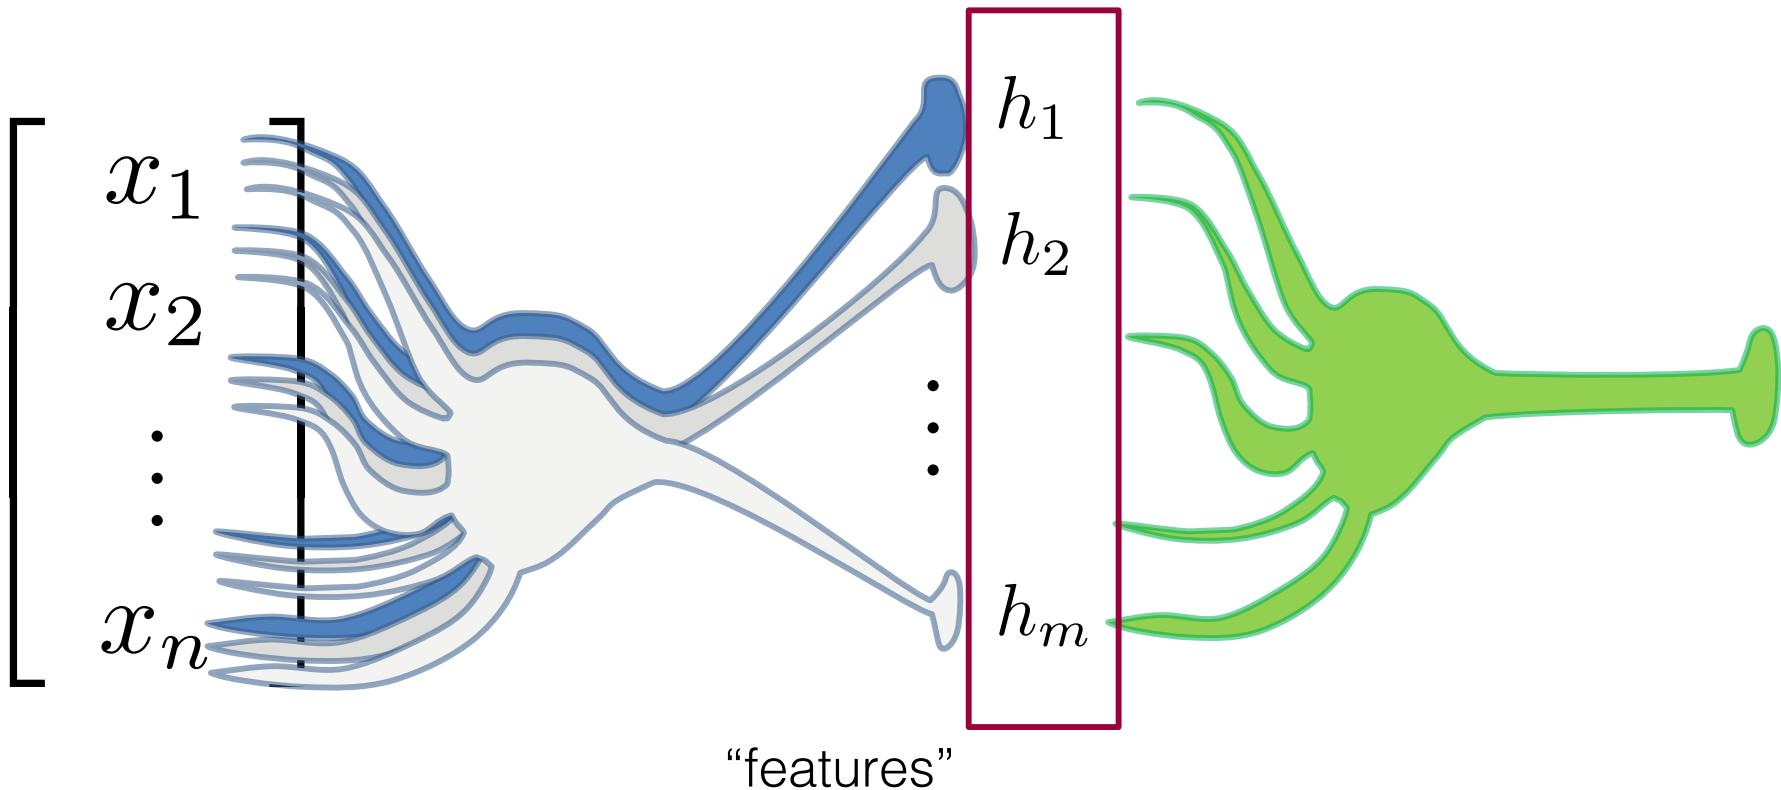
the rest of the network's parameters



## Two-Layer Network (~1980)



## intermediate layer

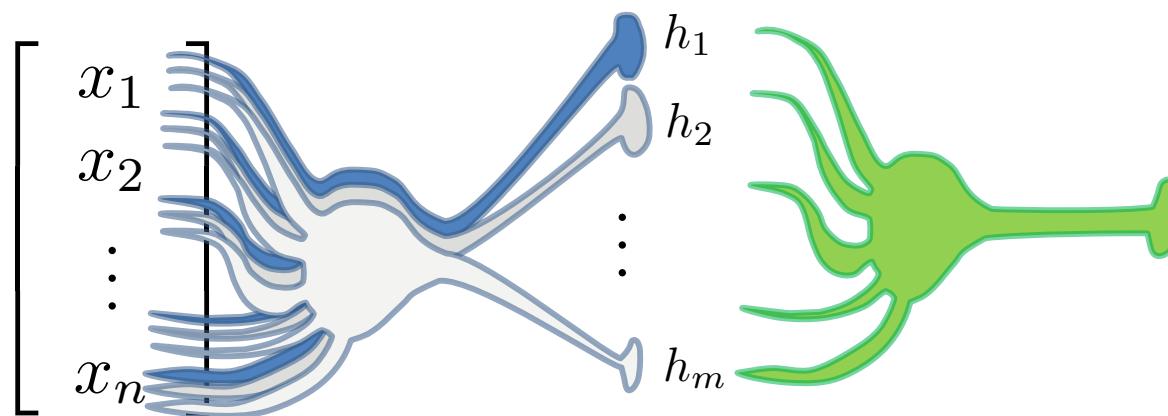


- the number  $m$  of “features” is a hyperparameter: it needs to be chosen;
- the values of the features are difficult to interpret (individually);

<https://playground.tensorflow.org/>

# Universal Approximation Theorem

(almost) any classification problem can be solved by a two-layer network

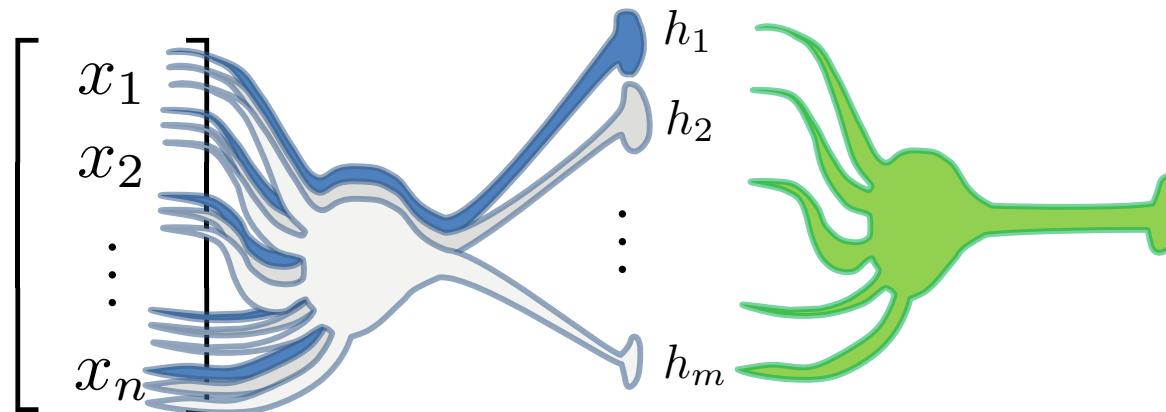


# Universal Approximation Theorem

(almost) any classification problem can be solved by a two-layer network

BUT:

1. This assumes we have enough training data.
2. For difficult problems, using a two-layer network is intractable.

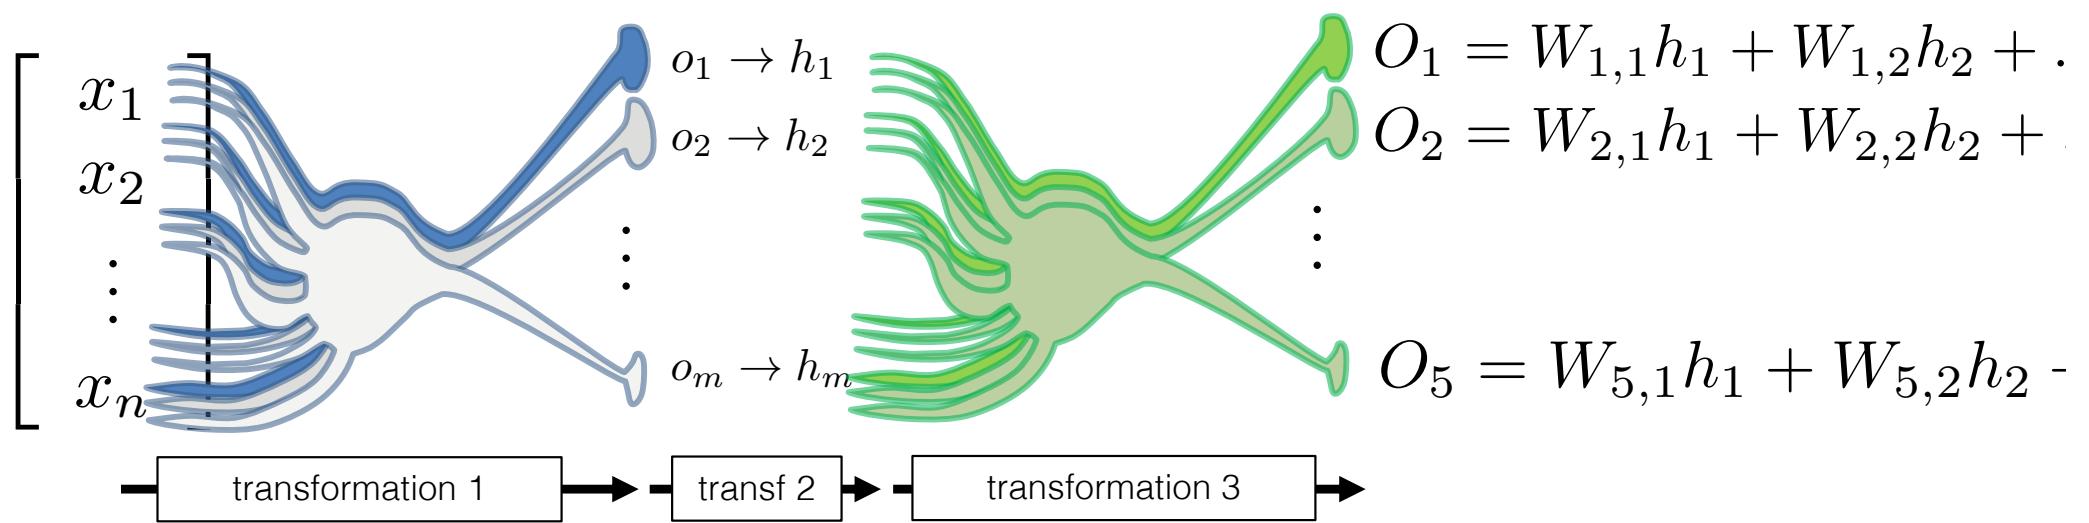


Fortunately, we can use more layers ie a deep network instead.

# MULTI-CLASS PROBLEMS

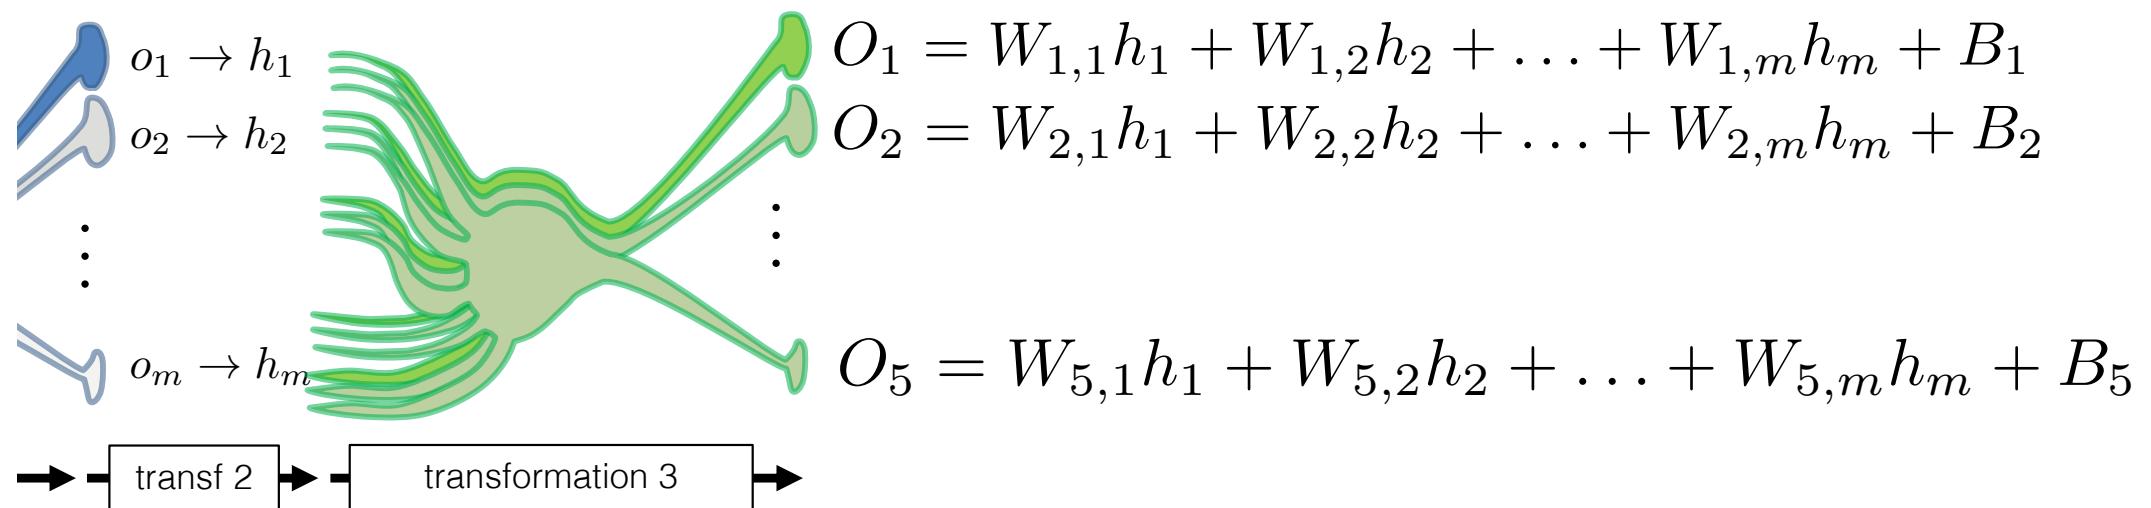
# Multi-class problems

With a two-layer network, and 5 classes:



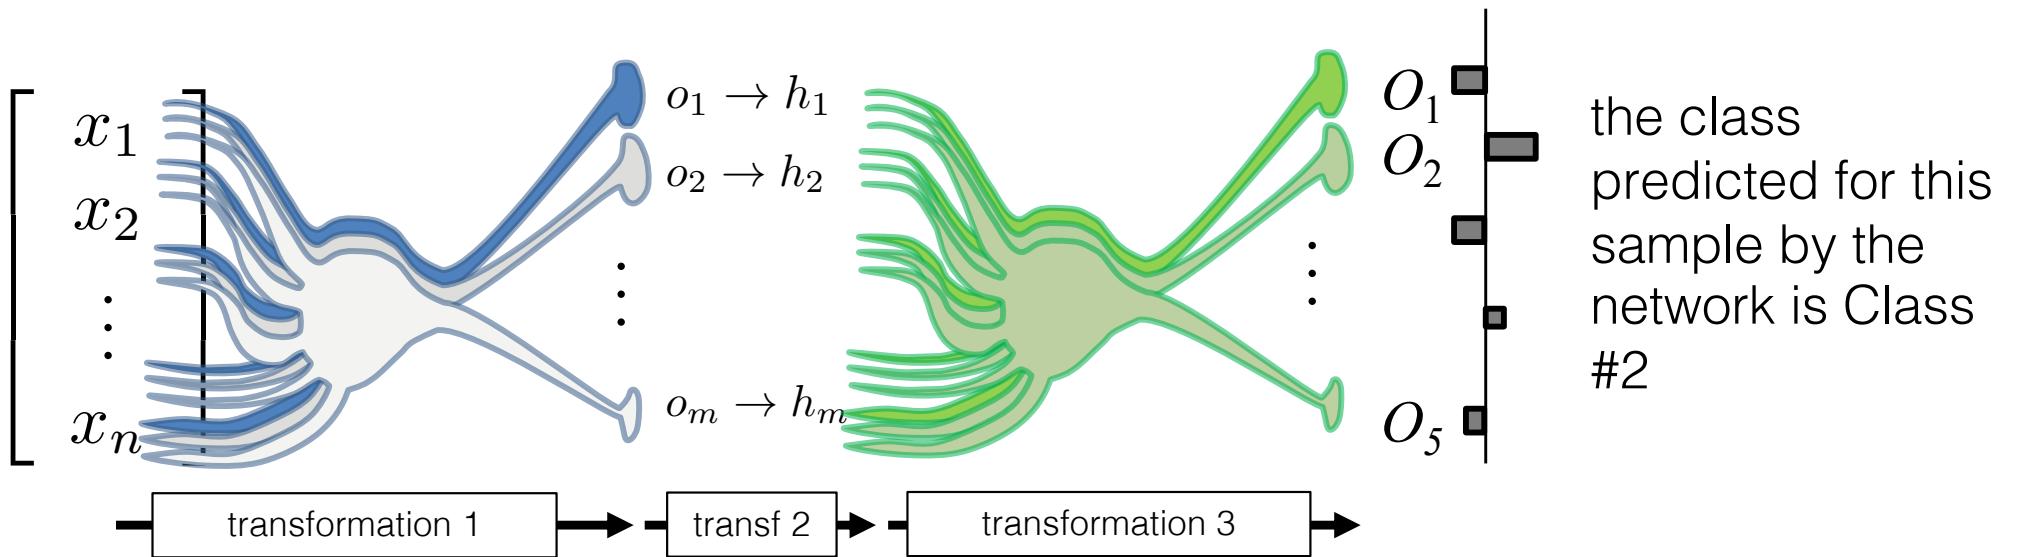
# Multi-class problems

With a two-layer network, and 5 classes:



# Multi-class problems

With a two-layer network, and 5 classes:

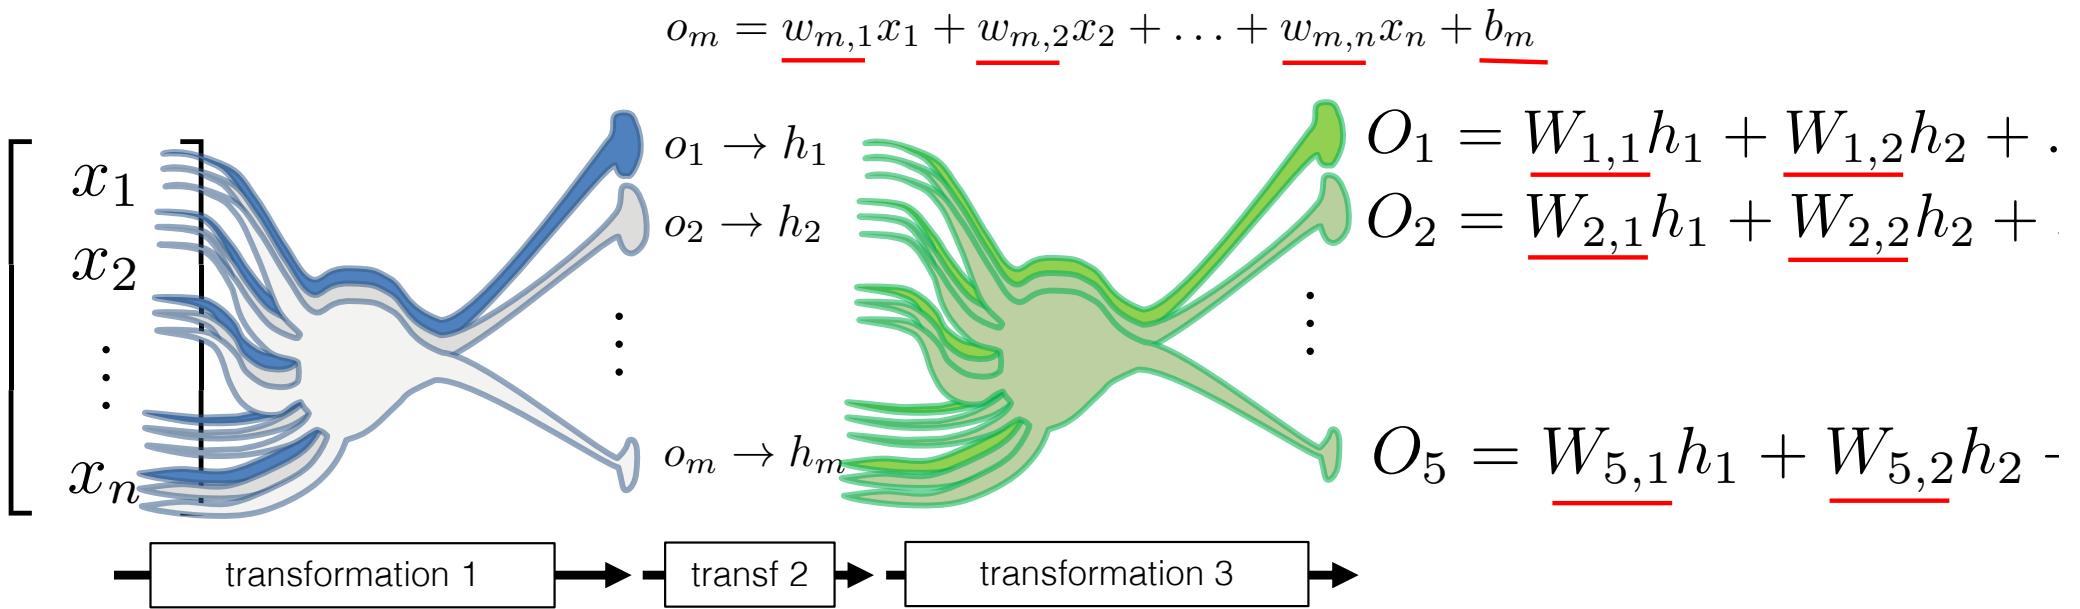


# FINDING THE NETWORK'S PARAMETERS

## the network's parameters

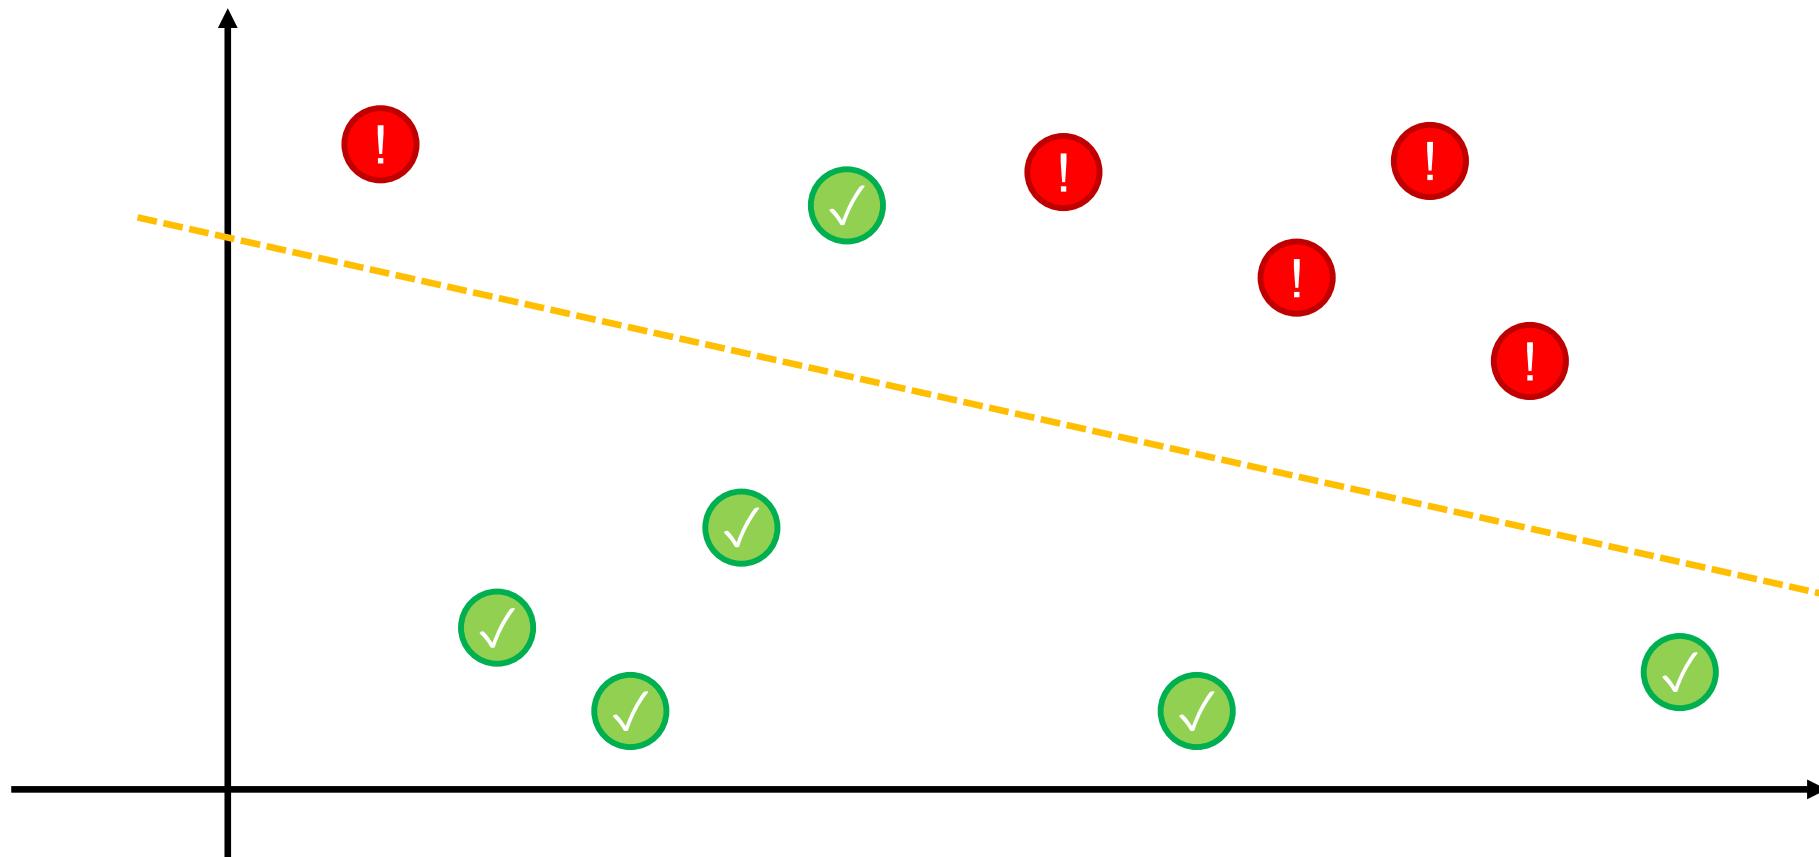
$$o_1 = \underline{w_{1,1}}x_1 + \underline{w_{1,2}}x_2 + \dots + \underline{w_{1,n}}x_n + \underline{b_1}$$

$$o_2 = \underline{w_{2,1}}x_1 + \underline{w_{2,2}}x_2 + \dots + \underline{w_{2,n}}x_n + \underline{b_2}$$

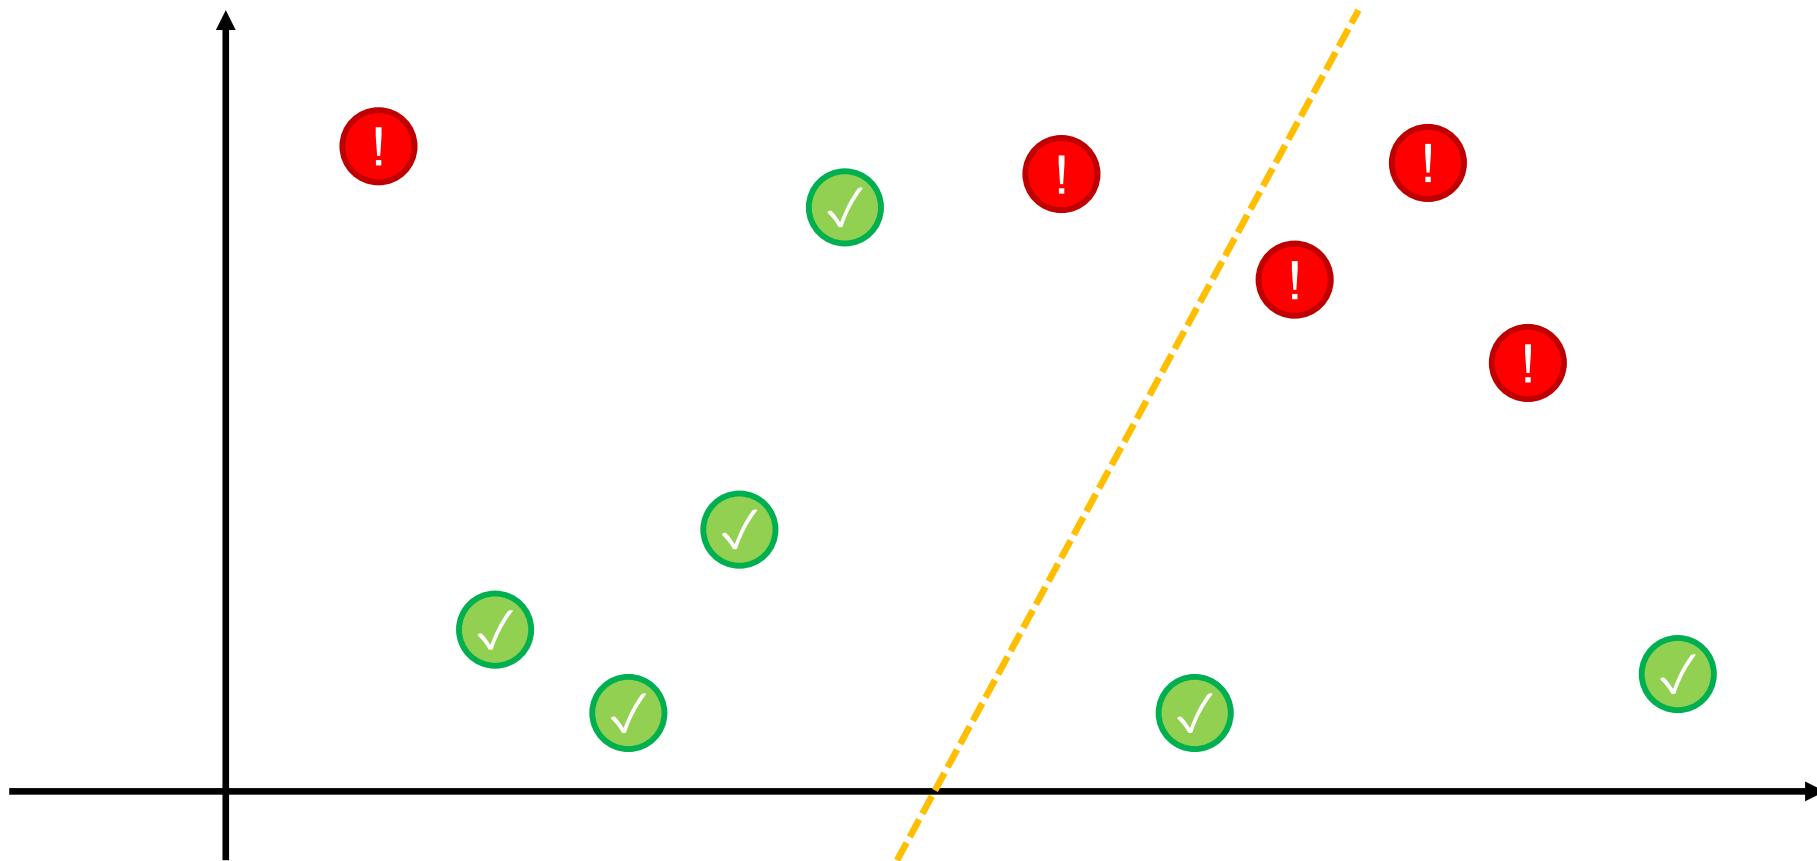


Given a training set, how can we find good values for these parameters?

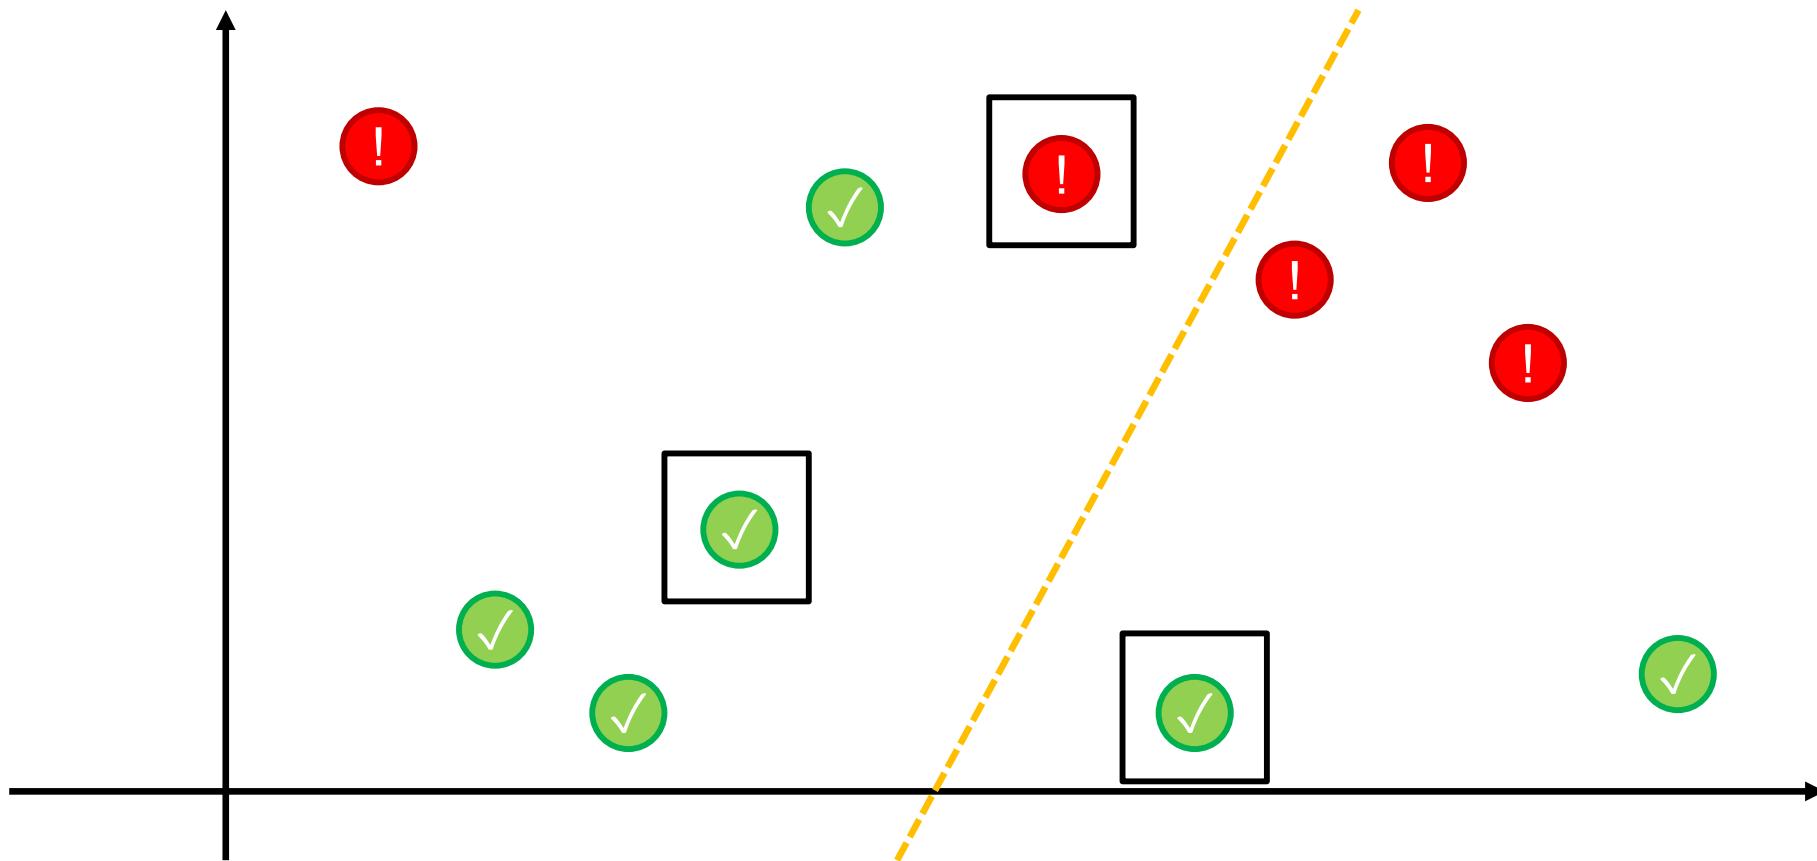
# finding the parameters of a perceptron



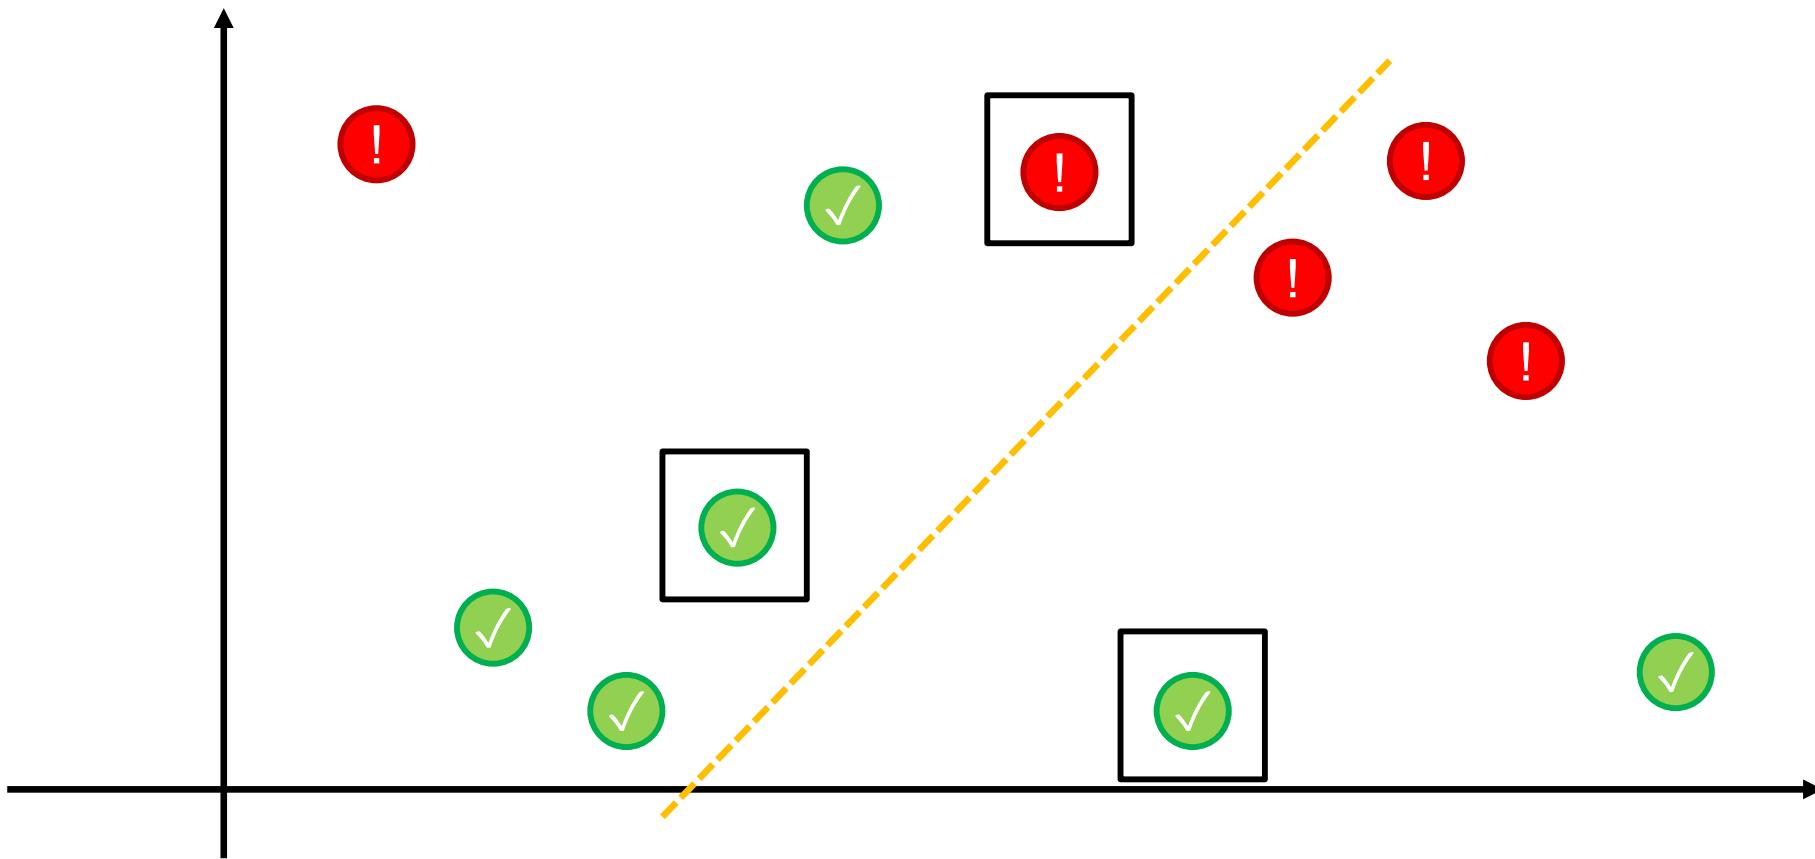
# random initialization



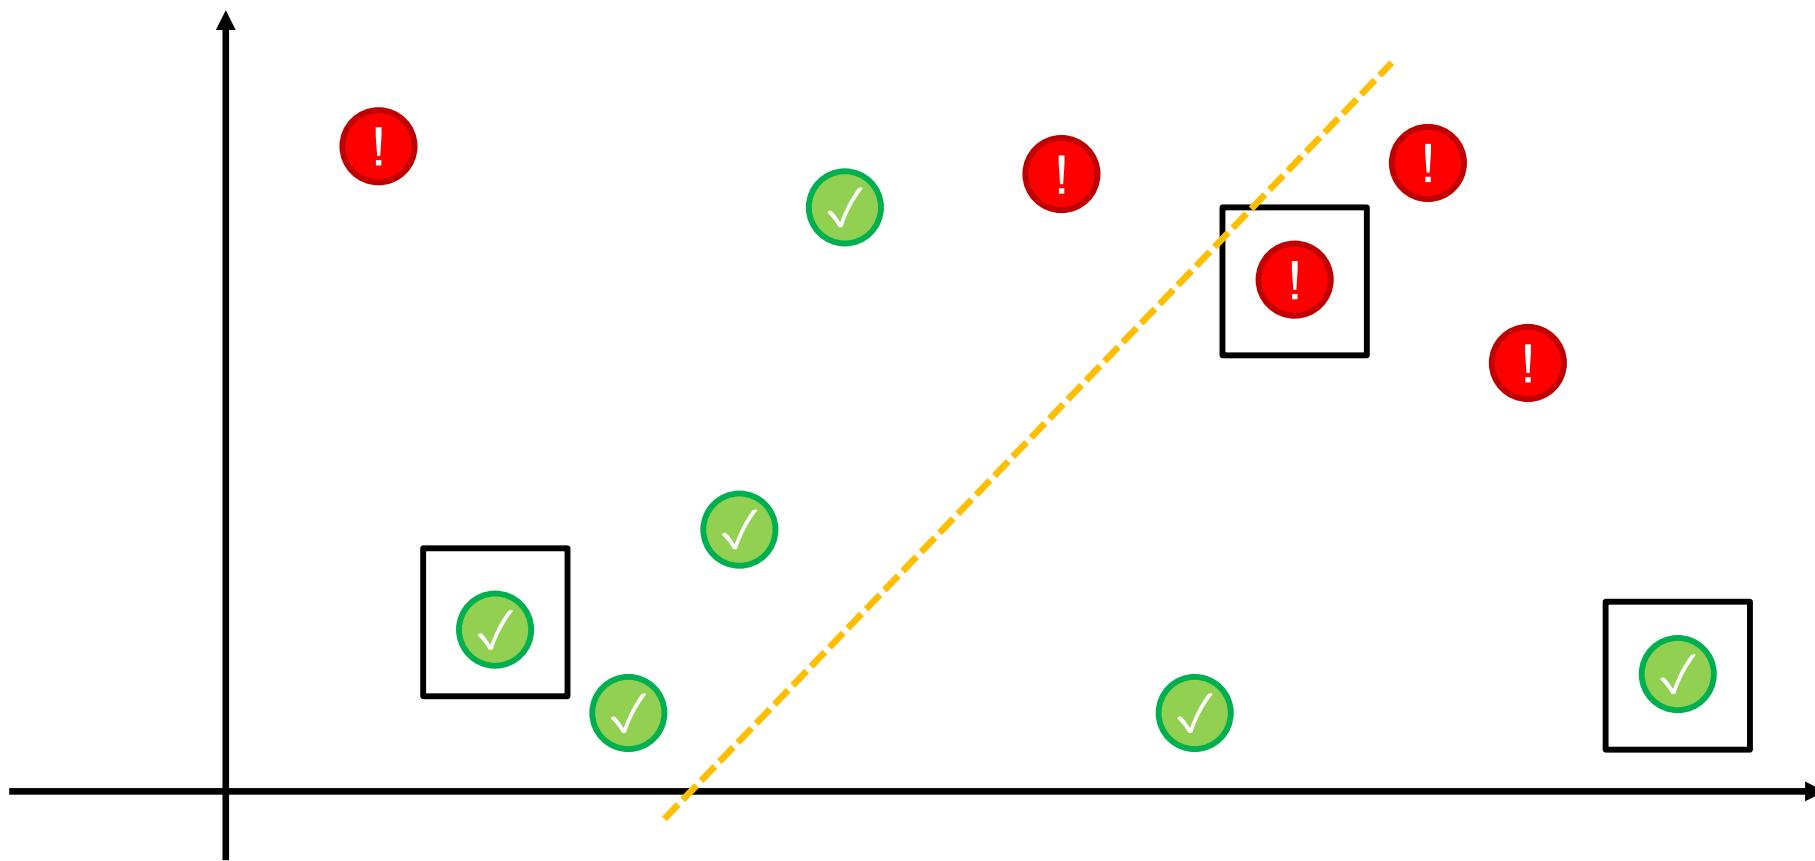
pick training samples  
randomly



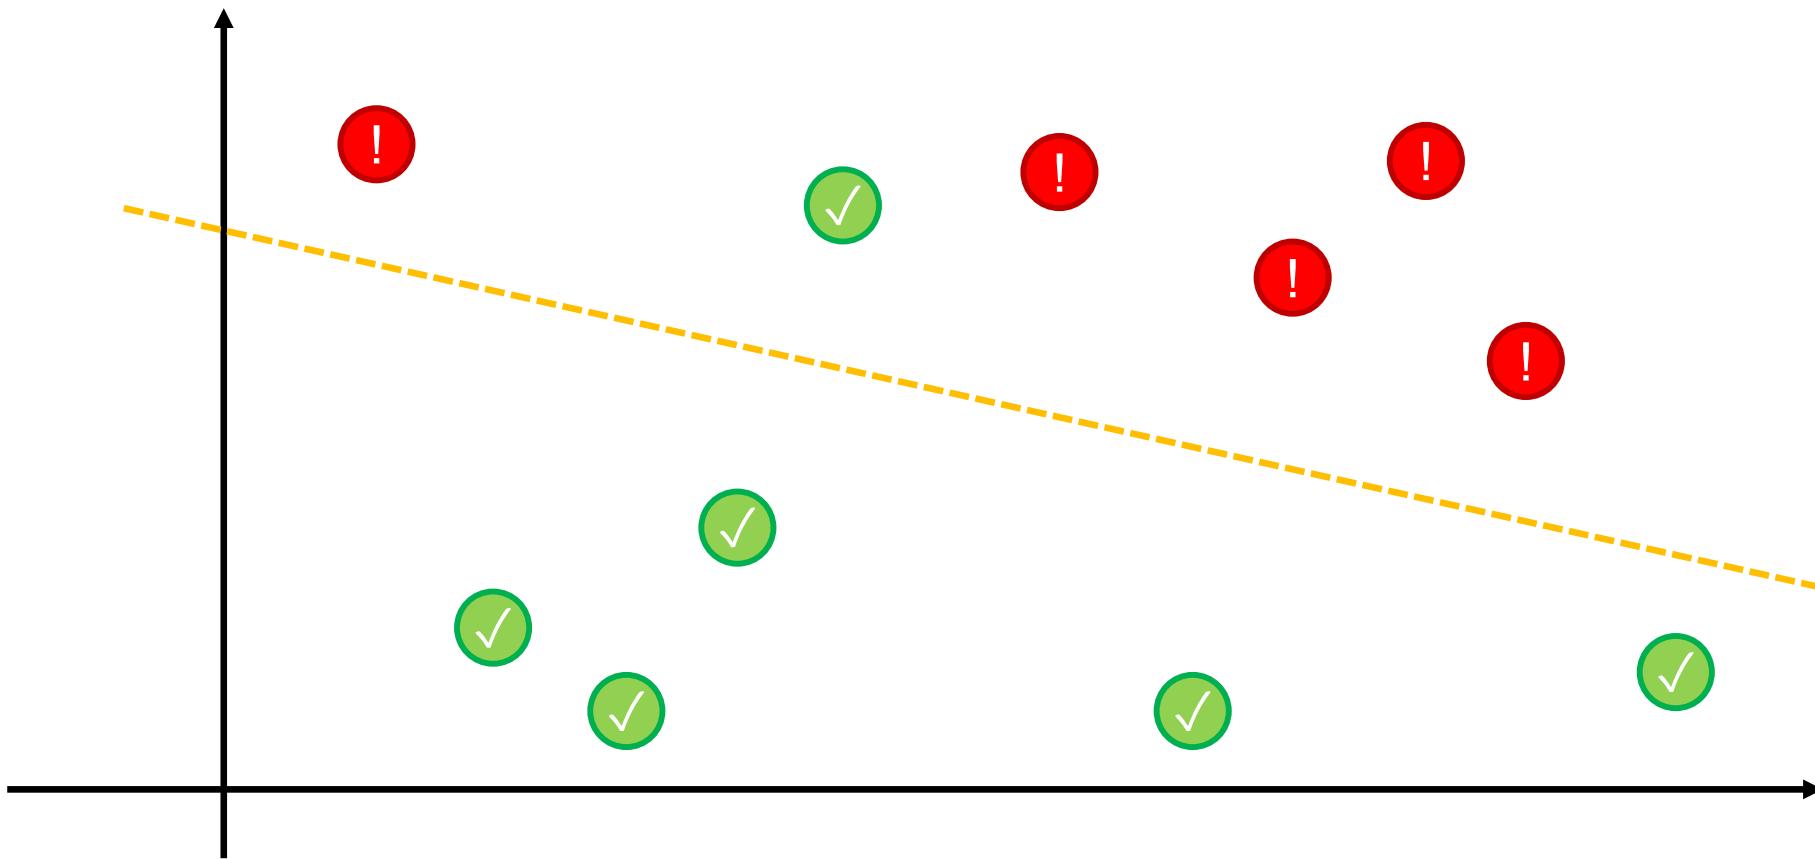
improve the parameters  
for these samples



# iterate



until we get a good  
solution



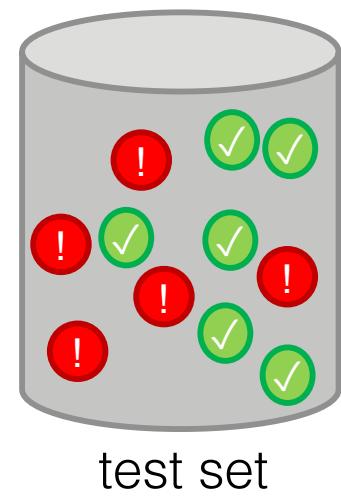
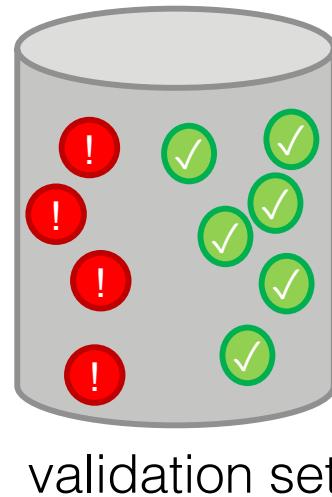
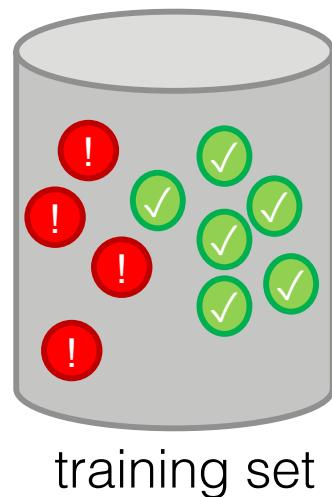
It can happen that we have

- good performance on the training set;
- poor performance on the validation set;
- poor performance on the test set.

This is called ‘**overfitting**’ (*surapprentissage*). “The method does not generalize well”

This may be due to many reasons *e.g.*

- the training and validation sets are not representative of the test set, or
- when the method has too many parameters.



to estimate the parameters  
ie “find the boundary”

to estimate the  
hyperparameters

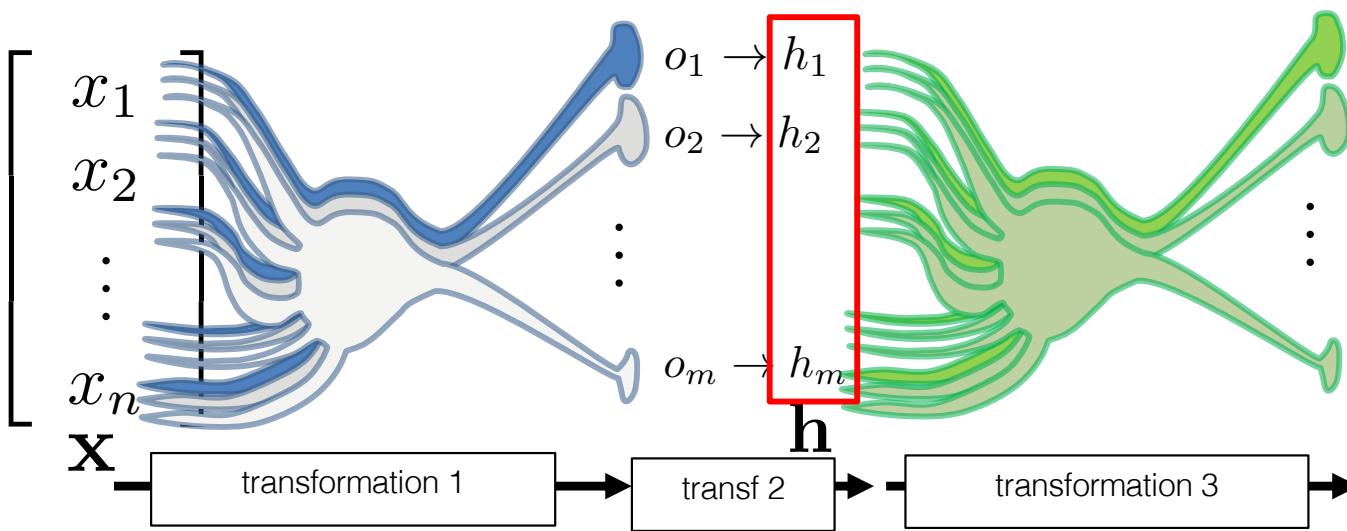
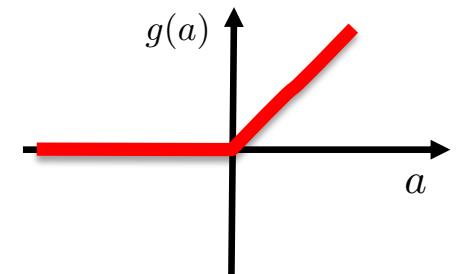
use it to evaluate the  
classifier

# GOING BACK TO THE TWO-LAYER NETWORK

# transformations

transformation 1:  $\mathbf{o} = \underline{\mathbf{W}_1 \mathbf{x}} + \underline{\mathbf{b}_1}$

transformation 2:  $\mathbf{h} = g(\mathbf{o}) = g(\mathbf{W}_1 \mathbf{x} + \mathbf{b}_1)$



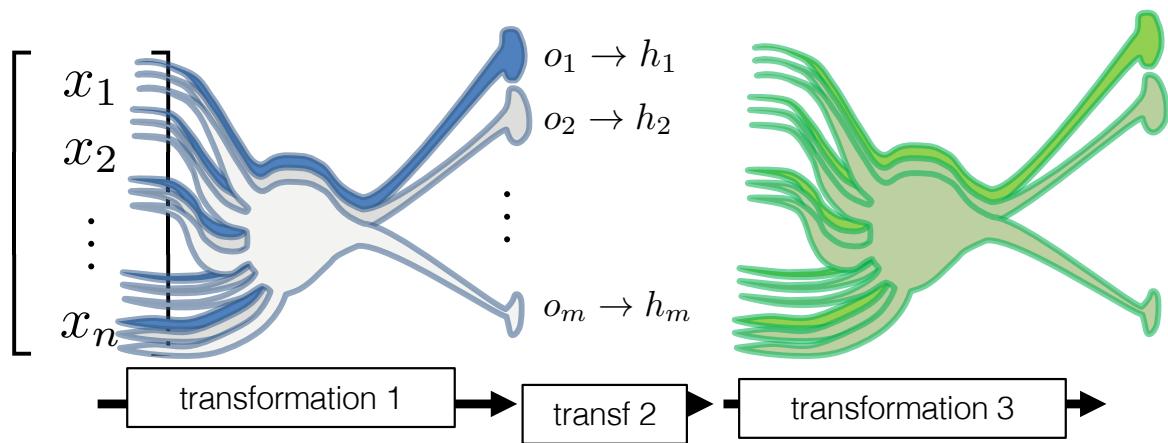
transformation 3:  
 $\mathbf{O} = \underline{\mathbf{W}_2 \mathbf{h}} + \underline{\mathbf{b}_2}$   
(linear/affine)  
transformation

$$\Theta = (w_{1,1}, w_{1,2}, \dots, w_{1,n}, b_1, \dots, b_m, W_{1,1}, W_{1,2}, \dots)$$

# A FIRST "DEEP NETWORK"

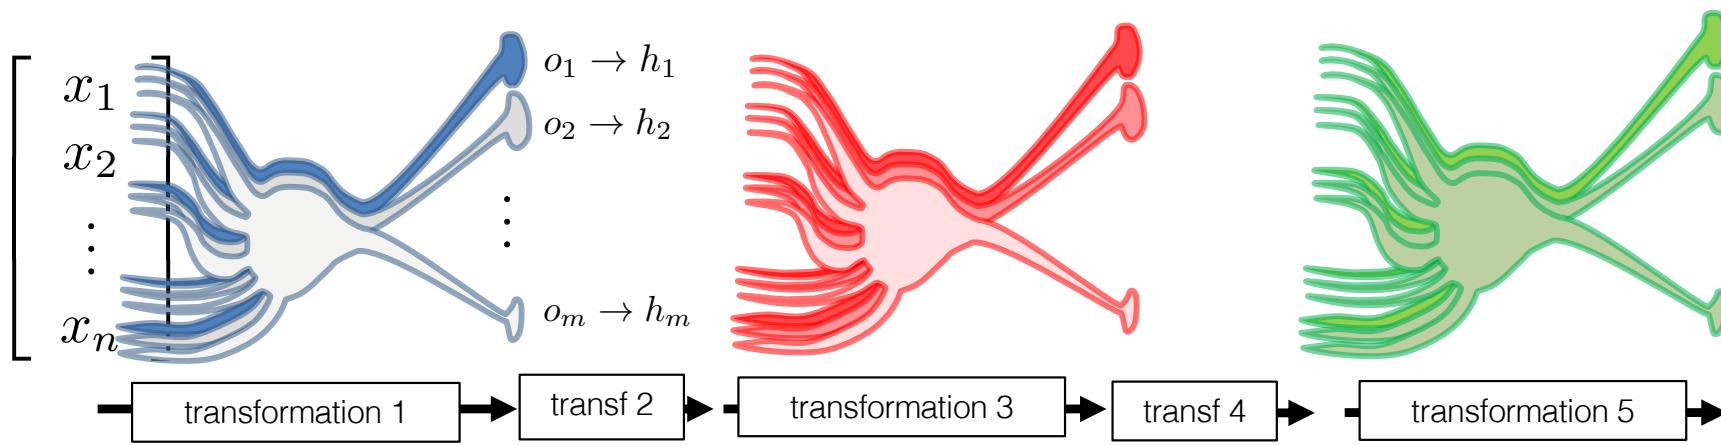
# Multi-Layer Networks/ Multi-Layer Perceptrons

Two layers:



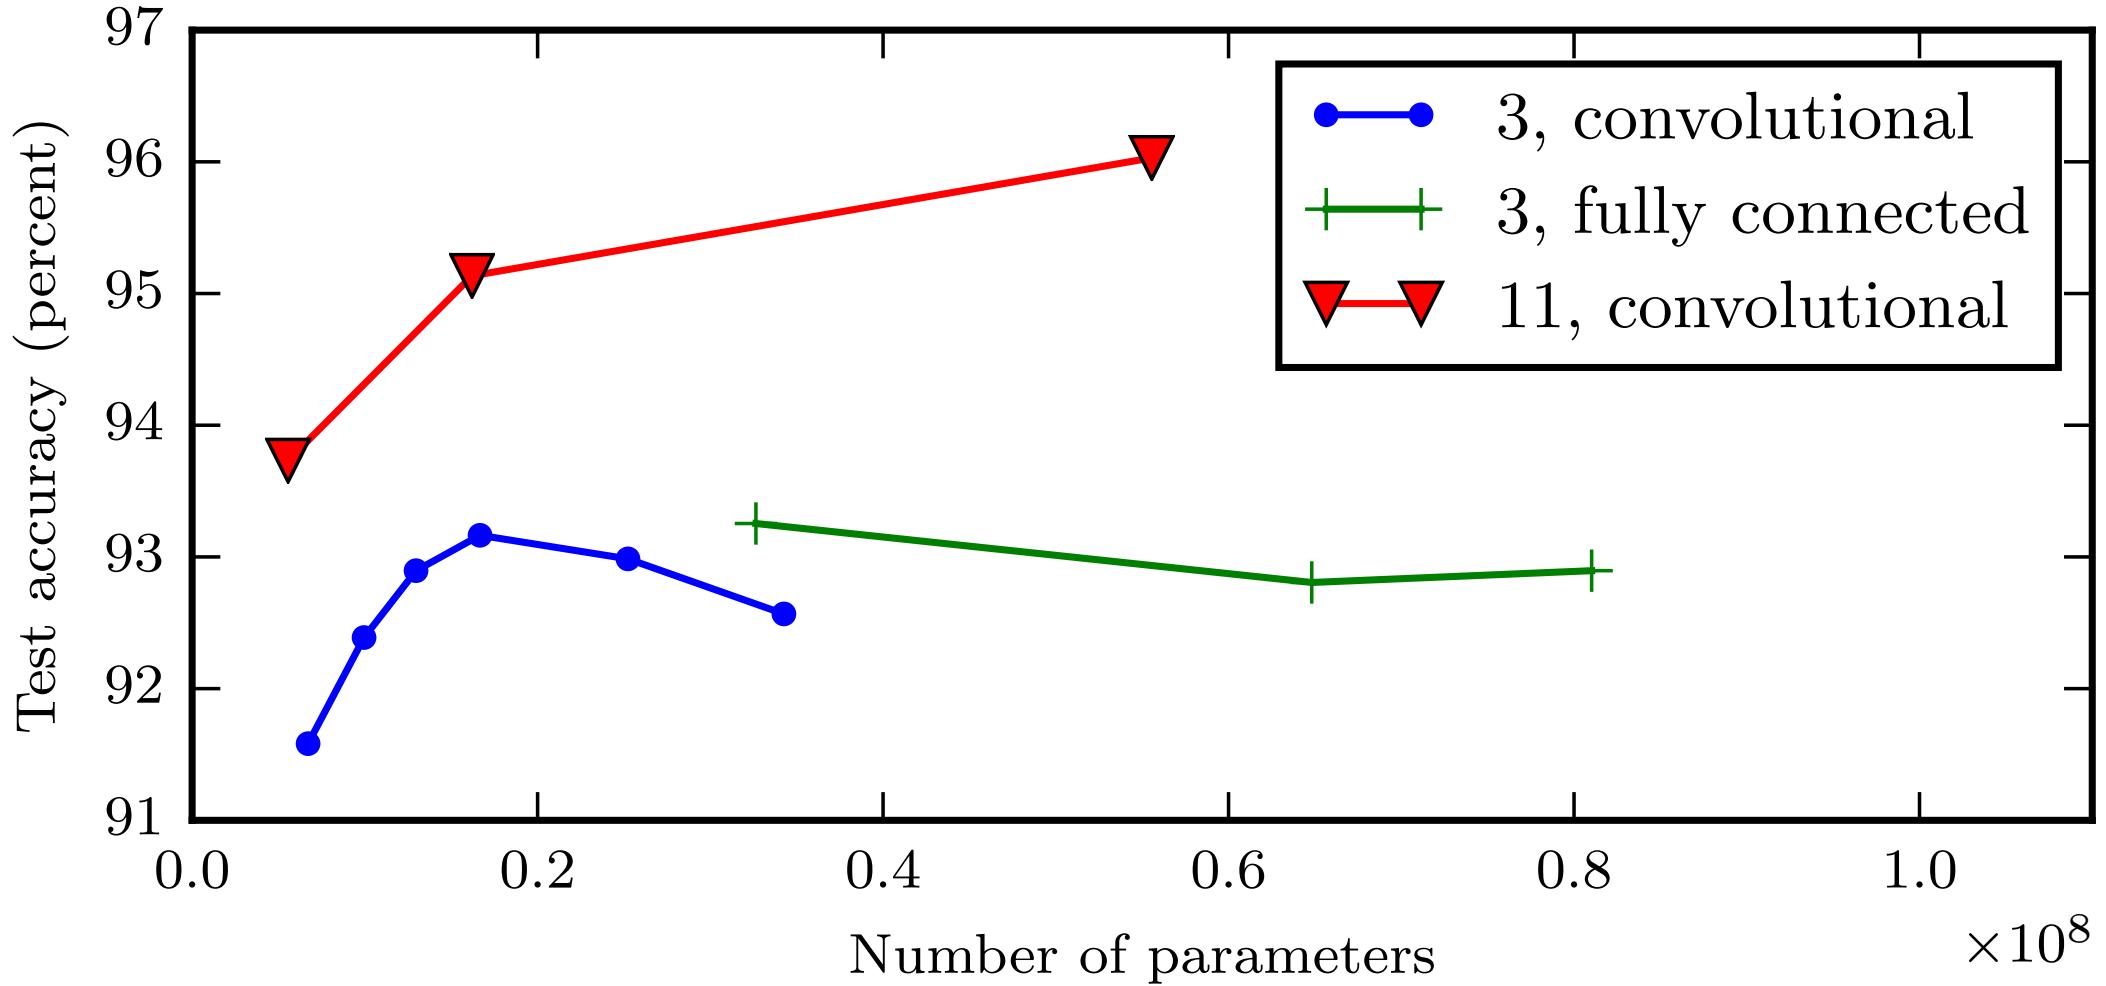
# Multi-Layer Networks/ Multi-Layer Perceptrons: a first „deep network“

Three layers:



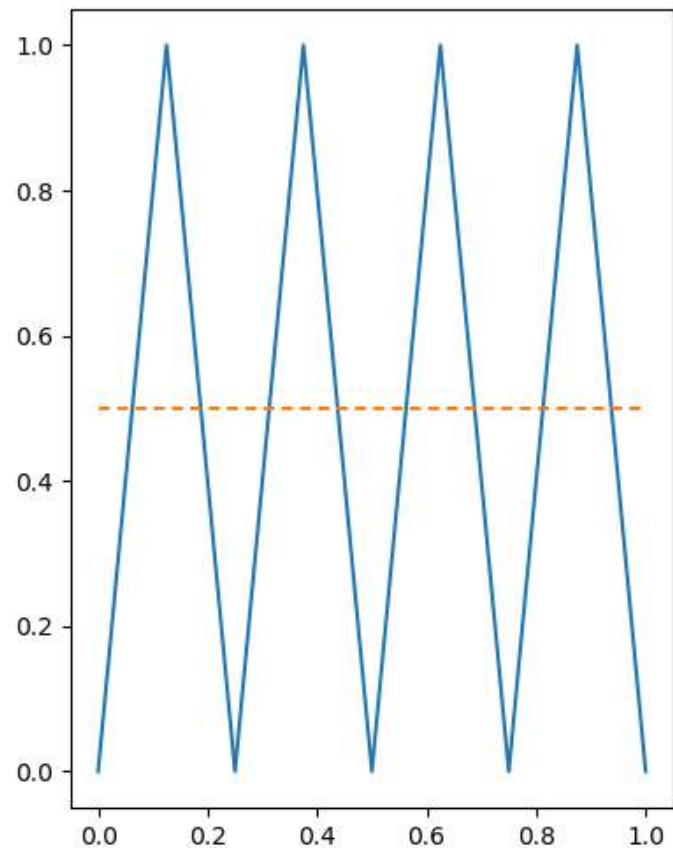
Can be used exactly the same way as a two-layer network;  
Advantage: Can work better than a two-layer network.

# Shallow Networks Overfit More

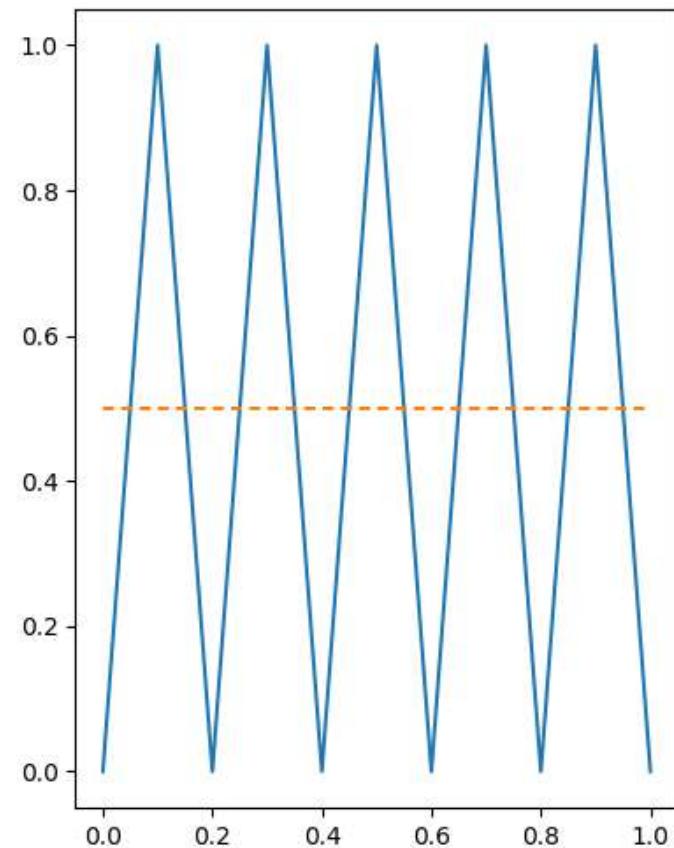


# the power of compositions

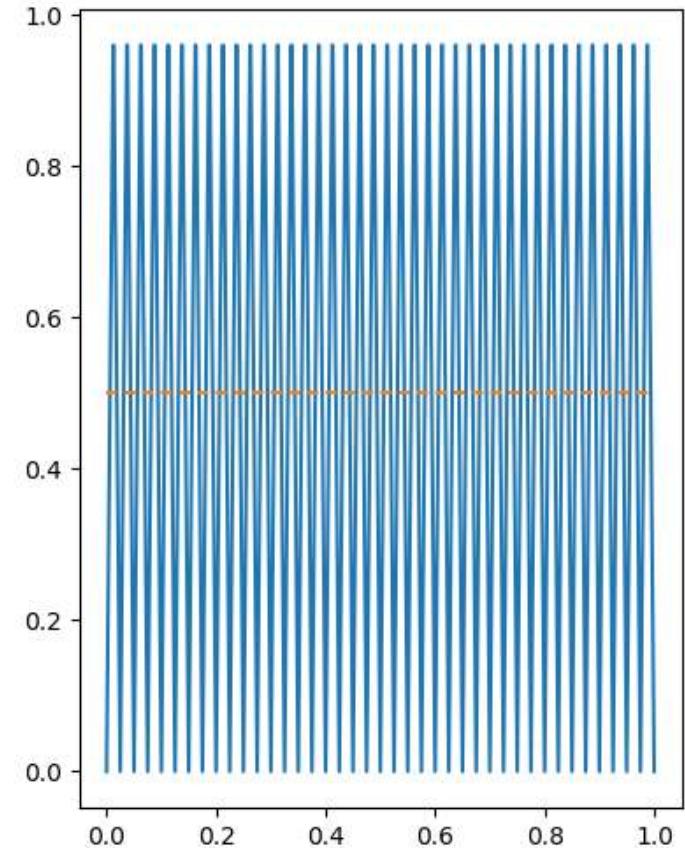
$f_1$



$f_2$

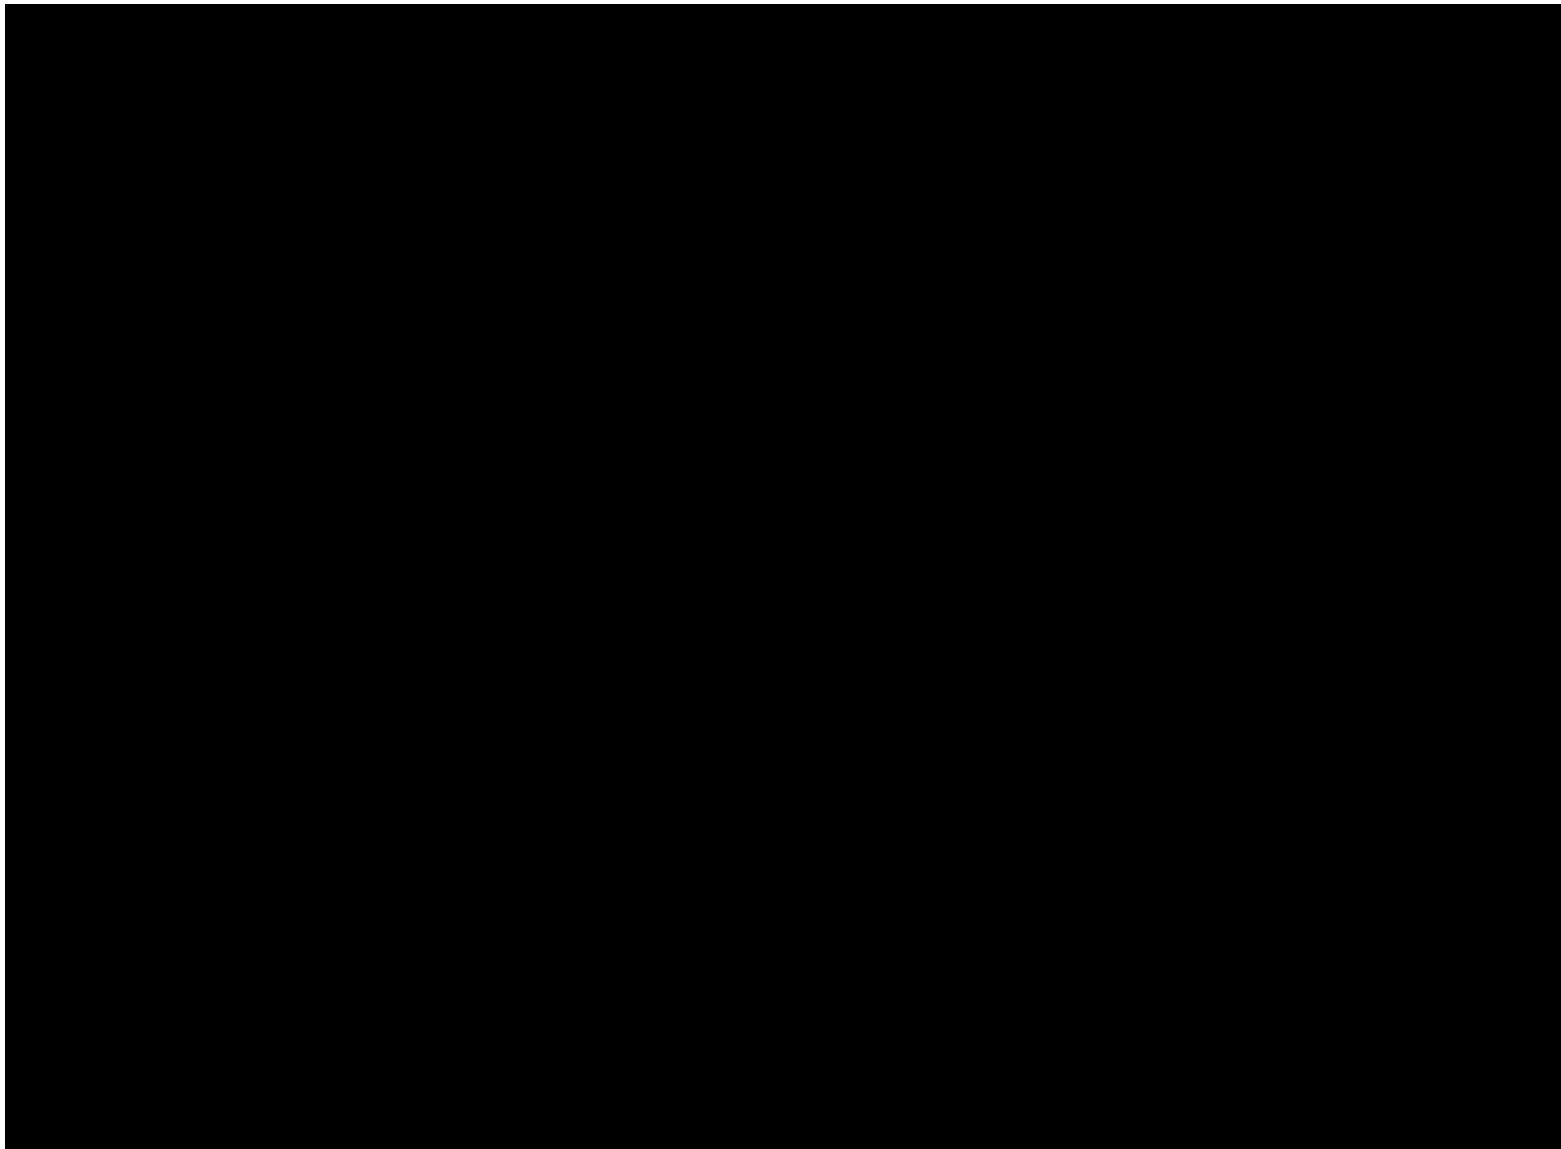


$f_2(f_1)$

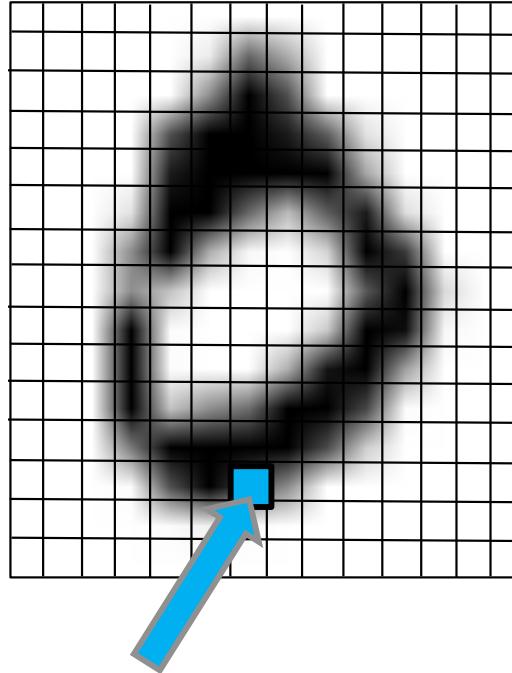
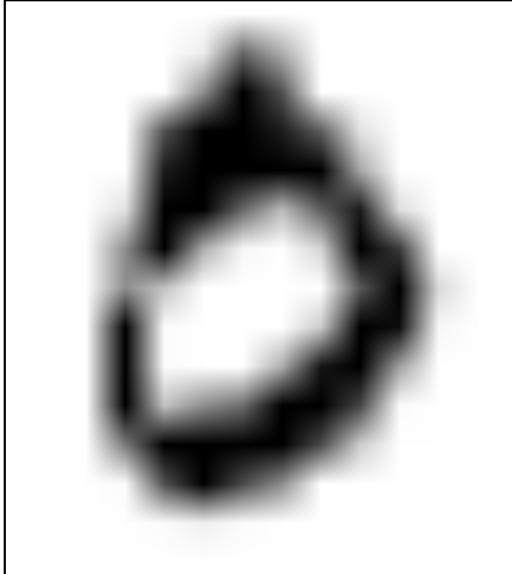


# A DEEPER NETWORK

# LeNet5 (LeCun, 1992)



# An Image is a Set of Values



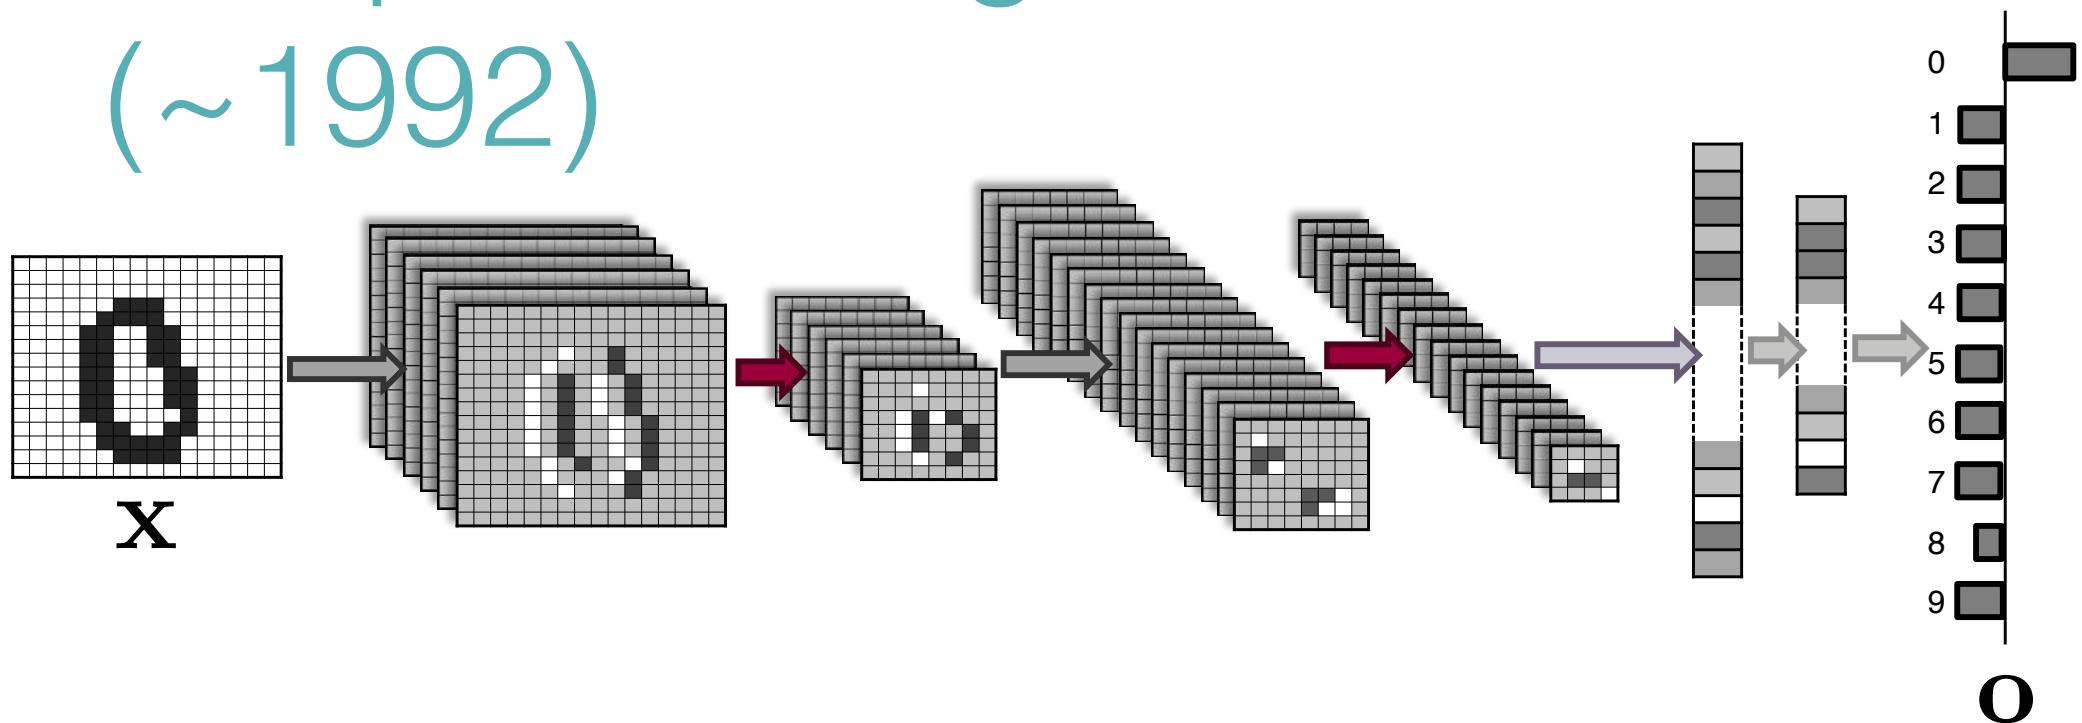
To each pixel correspond 1 value (3 for color images).

(for grayscale images: 0 for black, 255 for white)

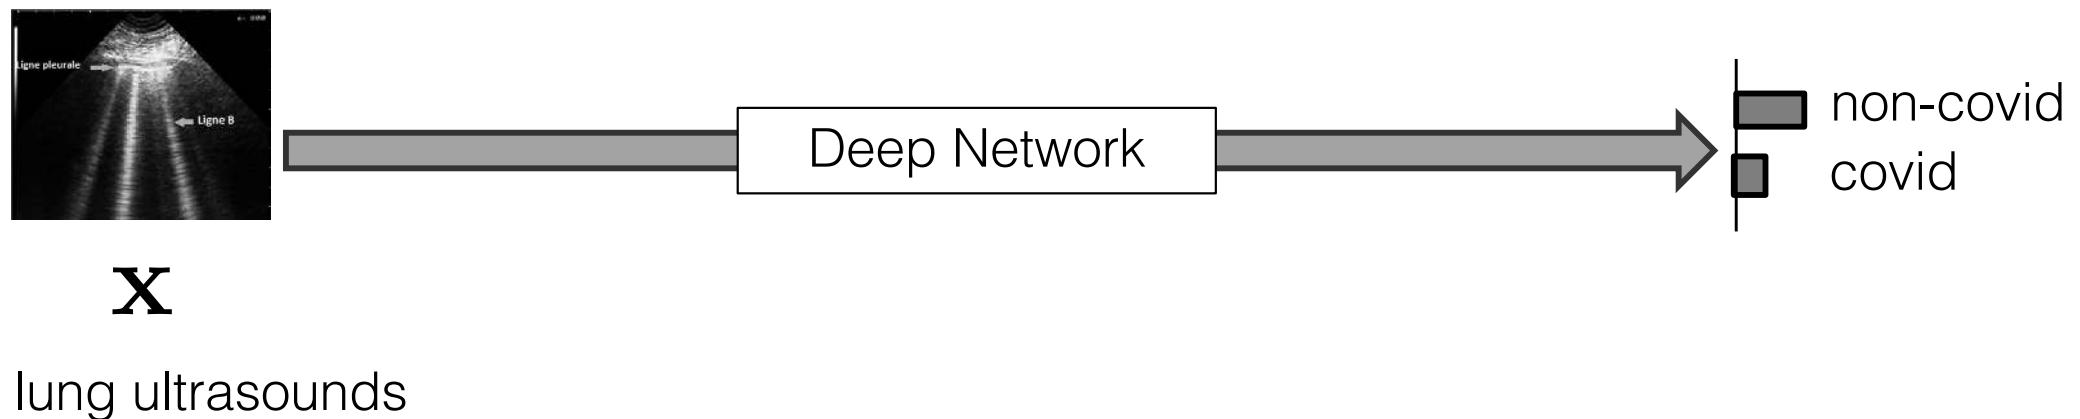
210 values for an image of size  $14 \times 15$  pixels.

262'000+ values for an image of size  $512 \times 512$  pixels.

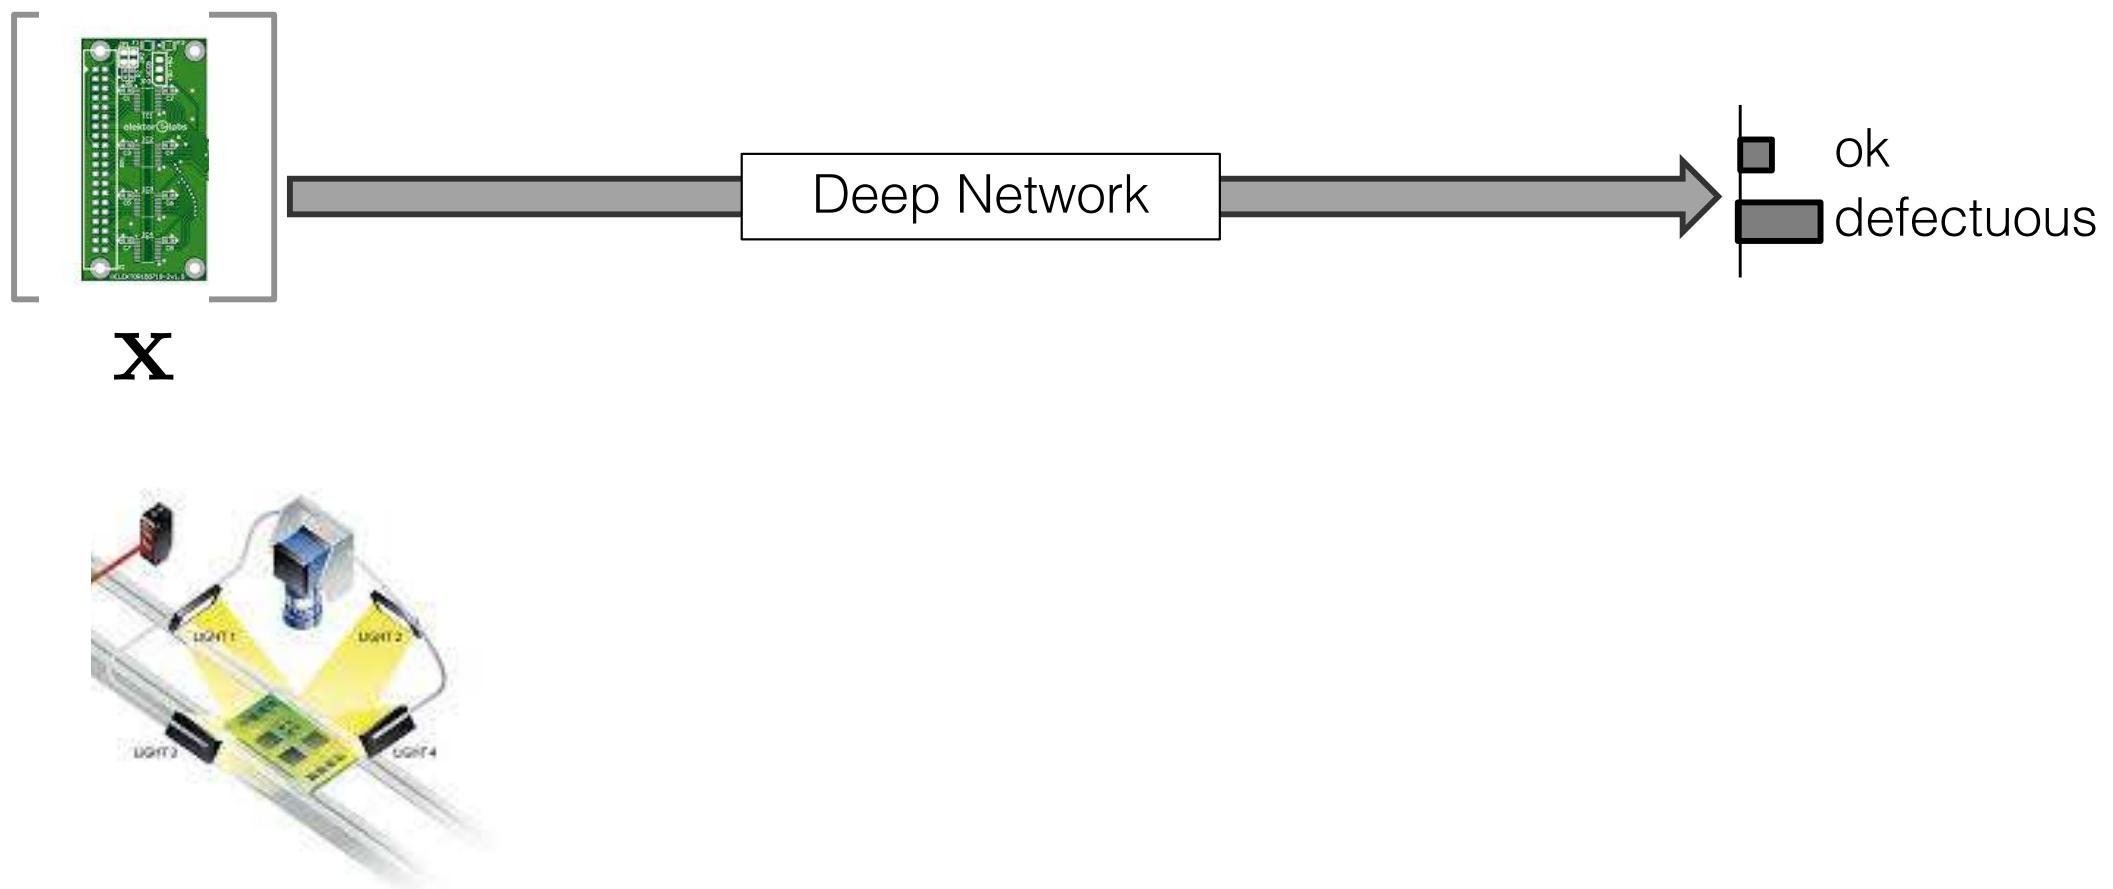
# Deep Learning *a la* LeCun (~1992)

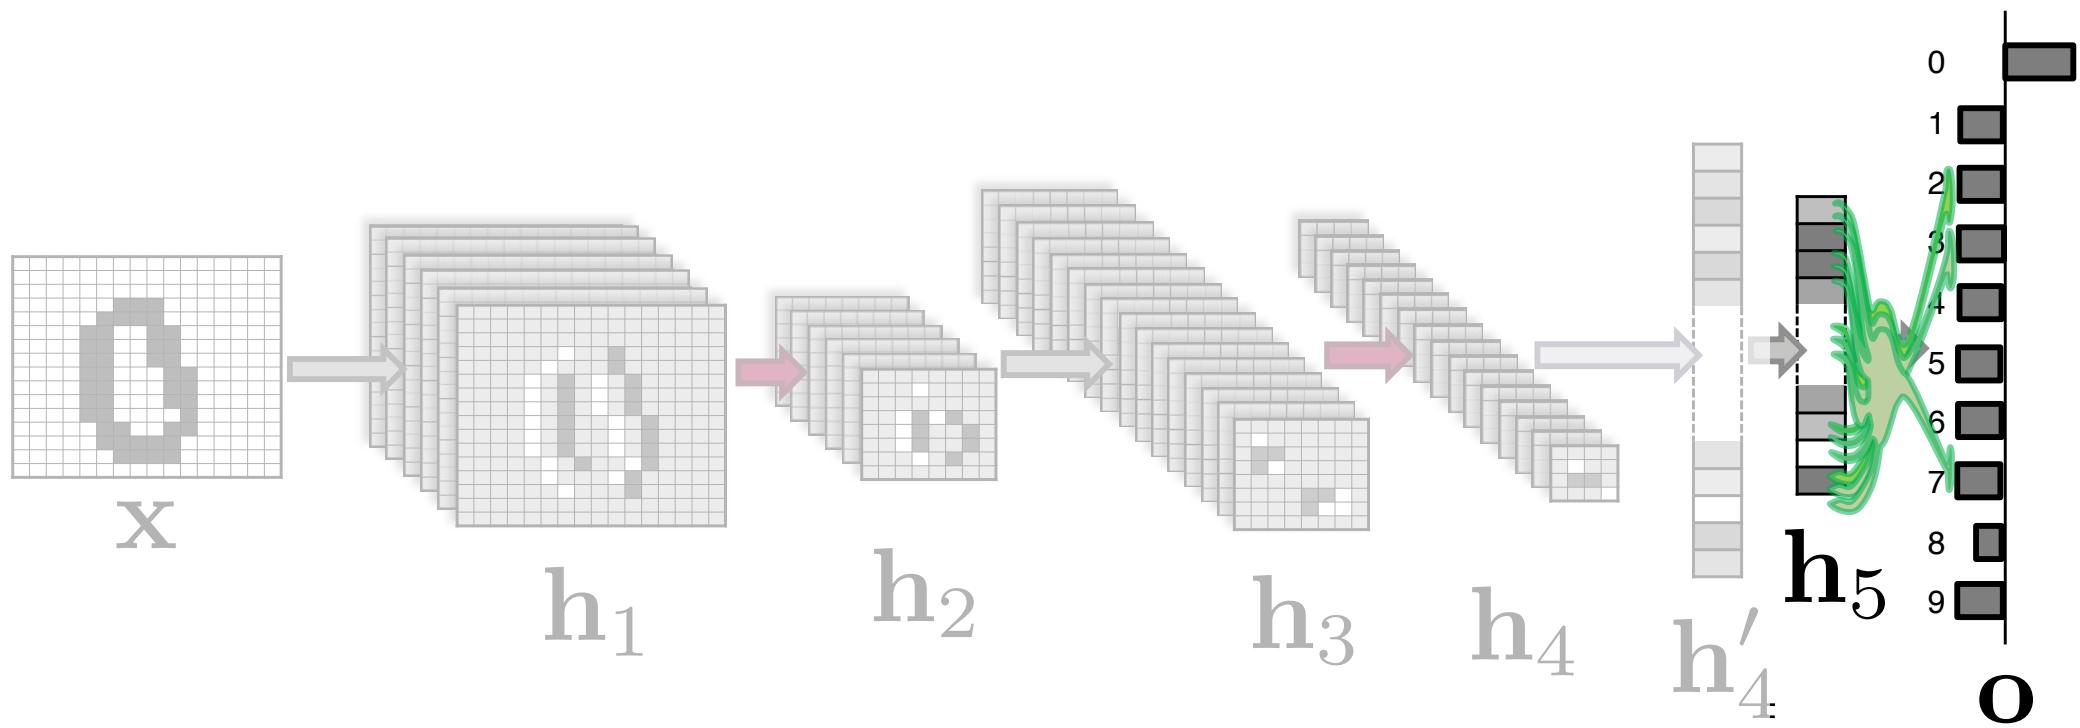


# application example

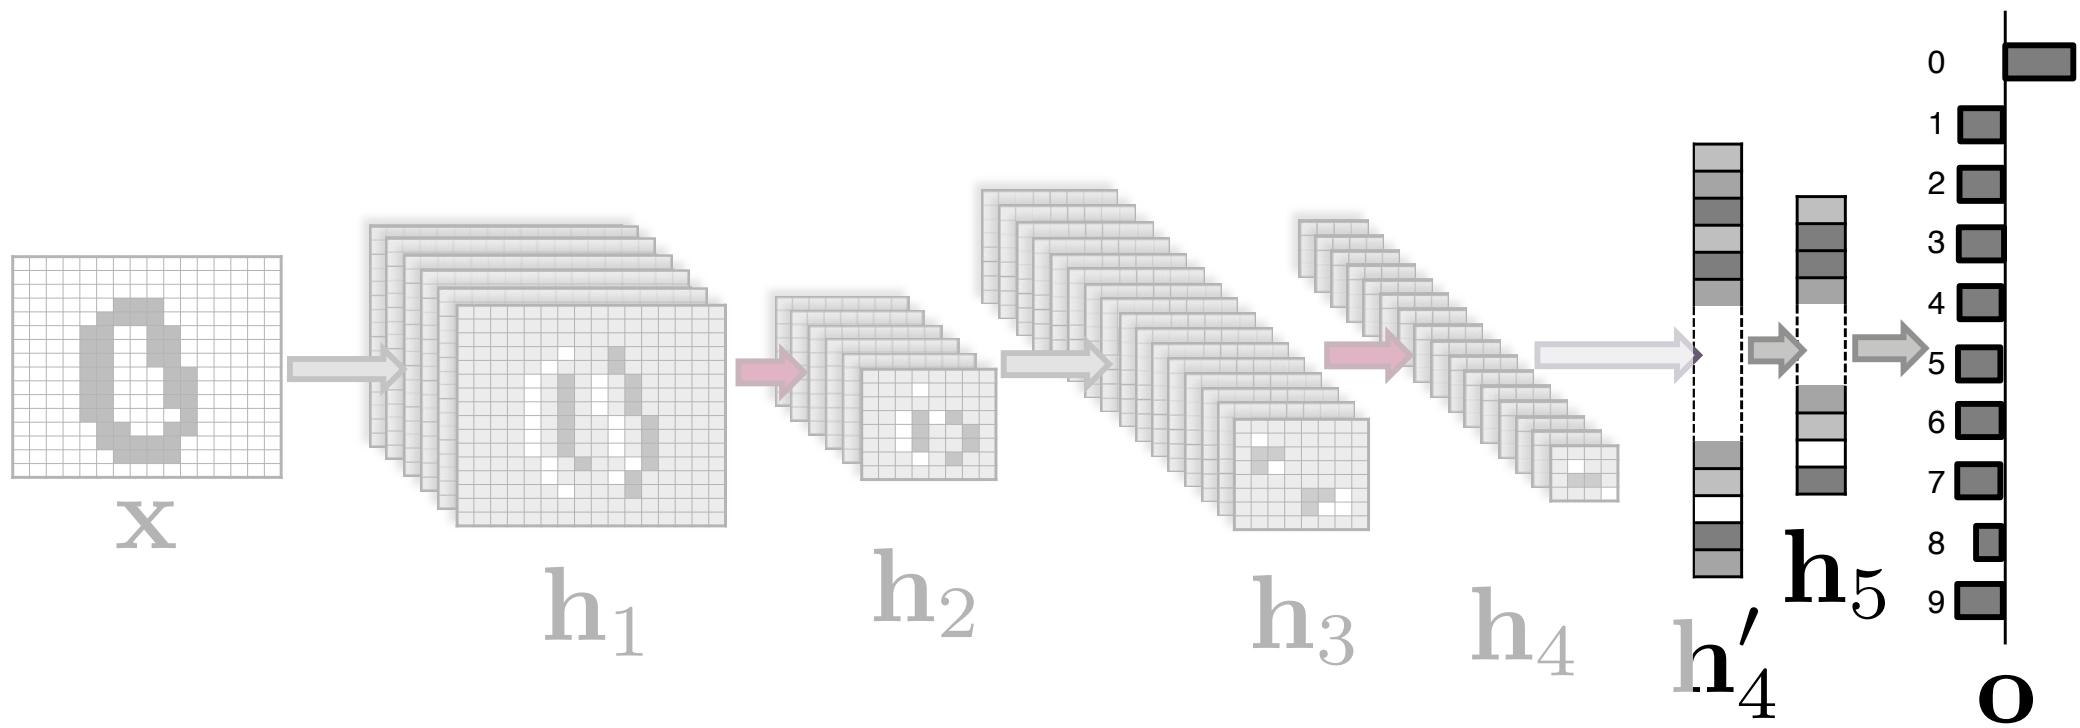


# application example (2)



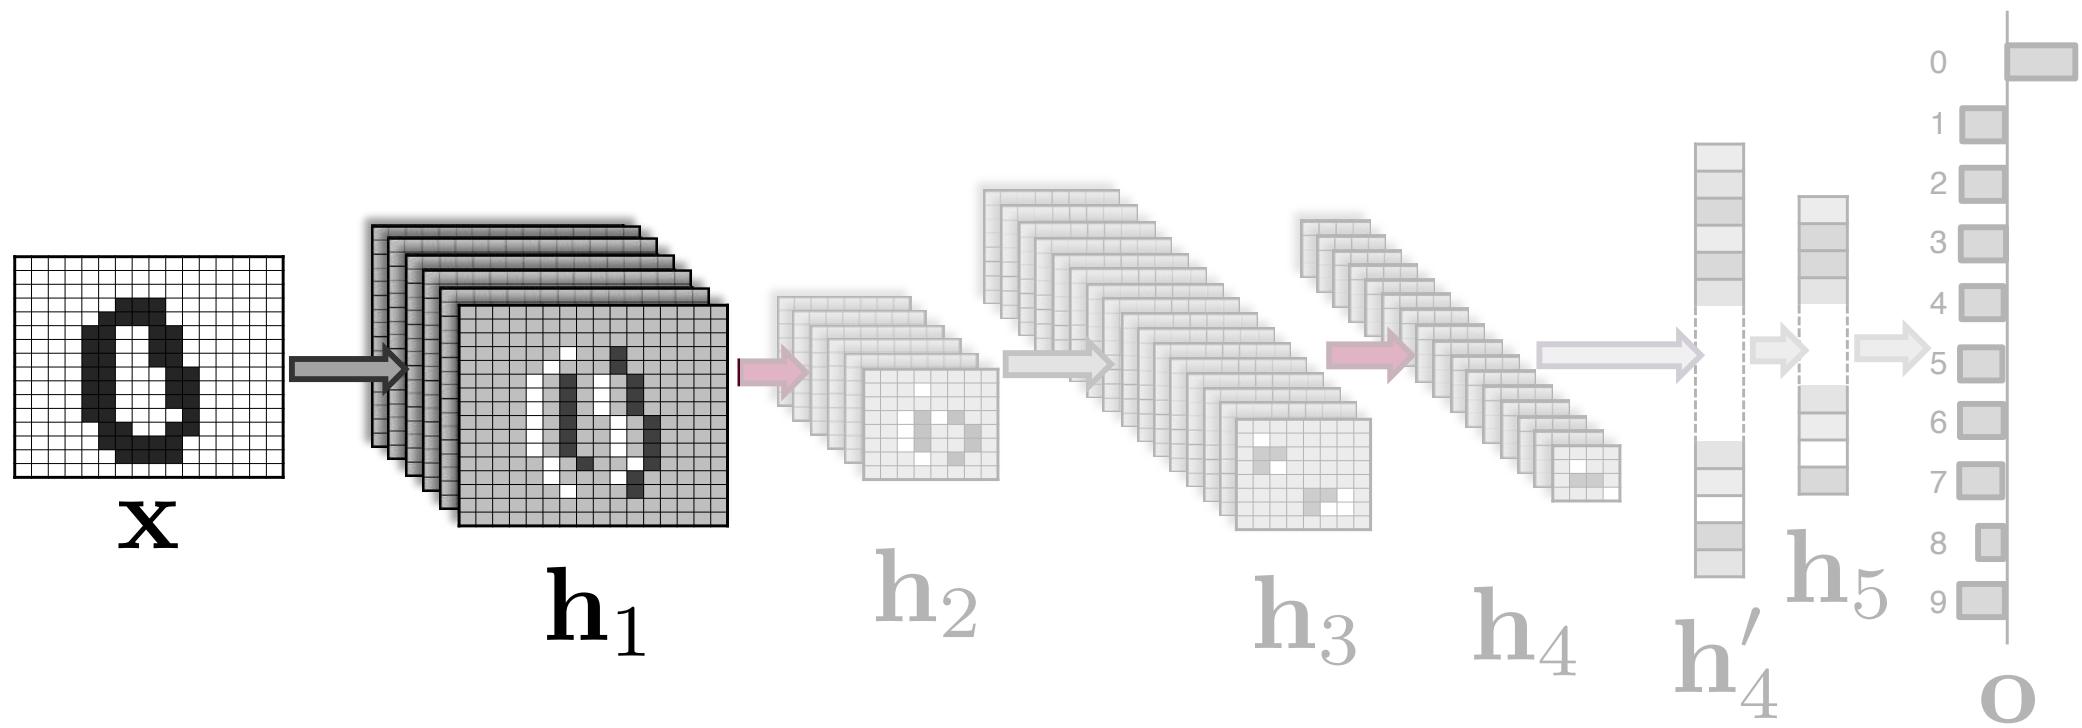


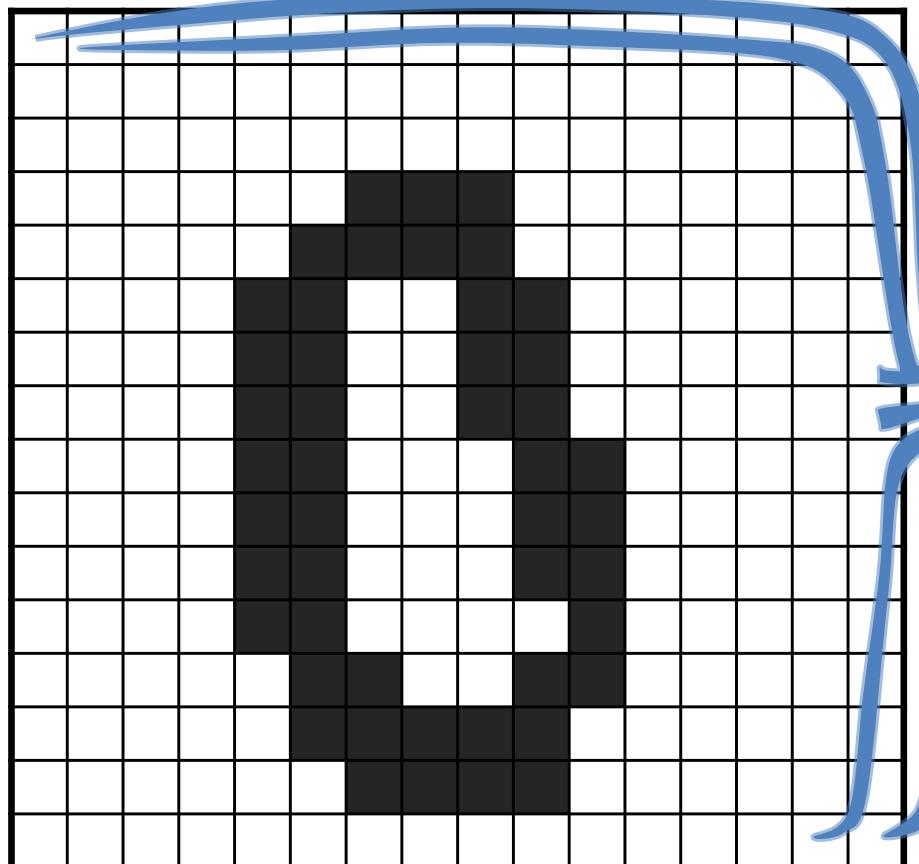
$$\mathbf{o} = \mathbf{W}_6 \mathbf{h}_5 + \mathbf{b}_6$$



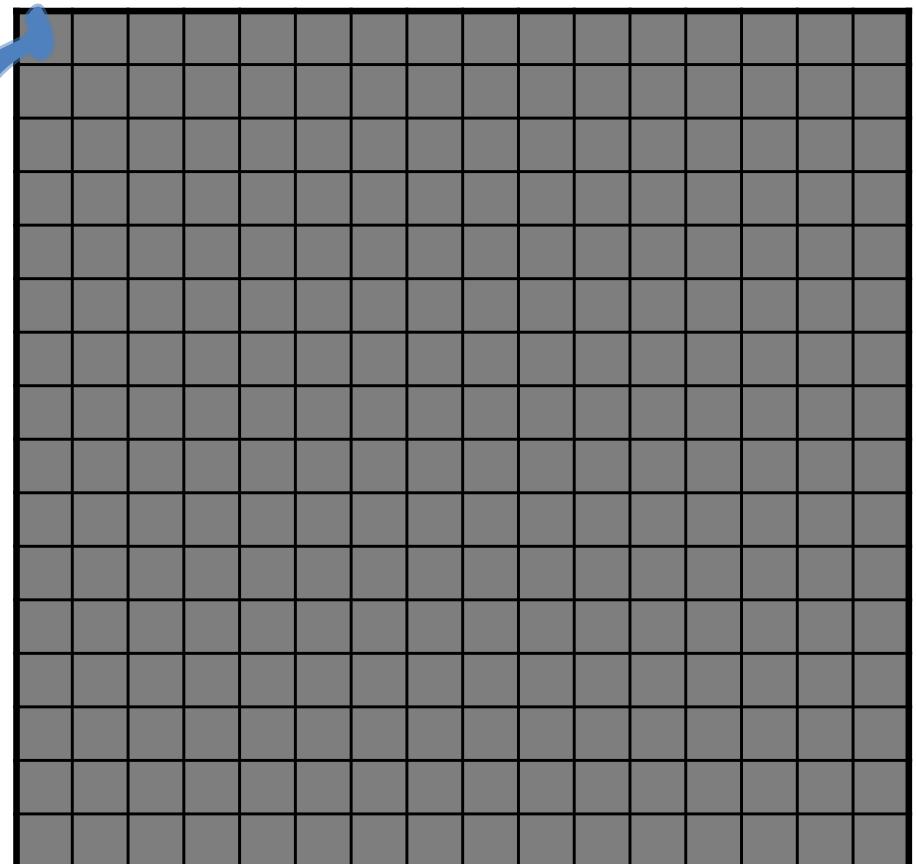
$$h_5 = g(\mathbf{W}_5 h'_4 + \mathbf{b}_5)$$

$$\mathbf{o} = \mathbf{W}_6 h_5 + \mathbf{b}_6$$



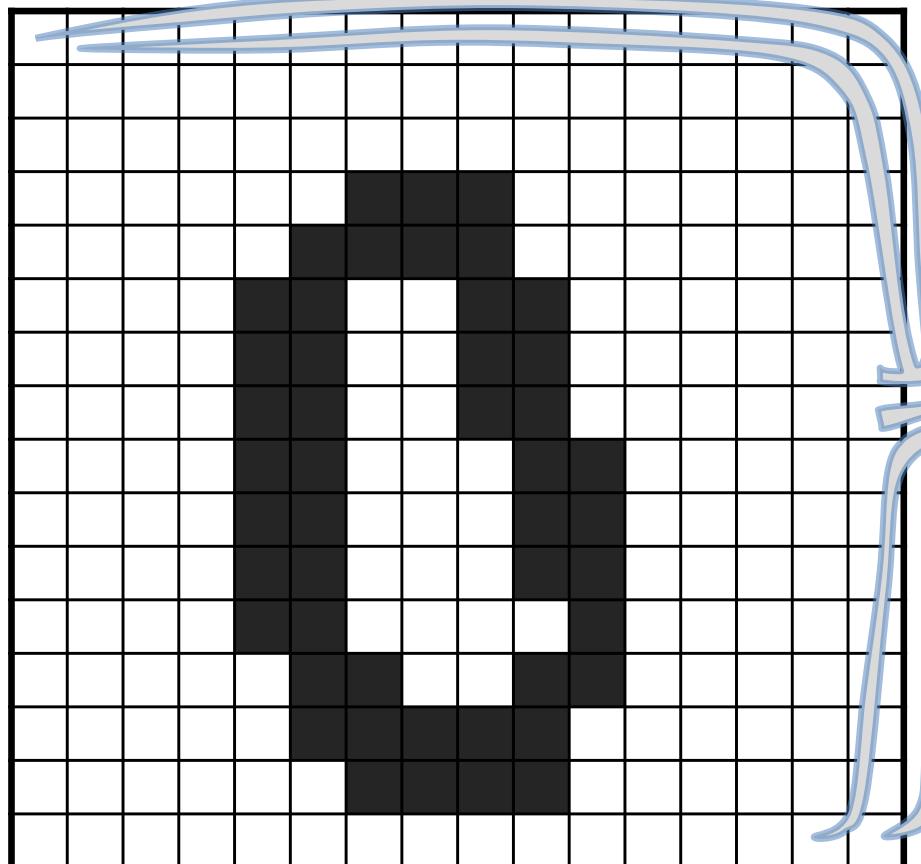


$\mathbf{x}$

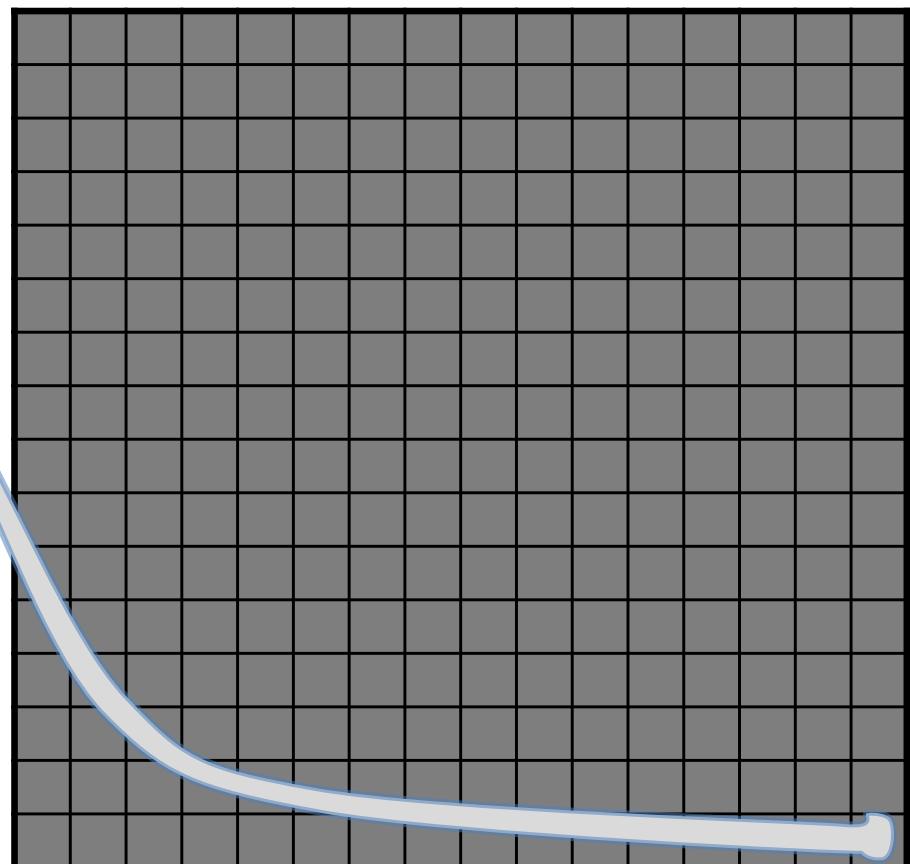


$\mathbf{h}_1$

$$\mathbf{h}_1 = g(\mathbf{W}\mathbf{x} + \mathbf{b}) ?$$

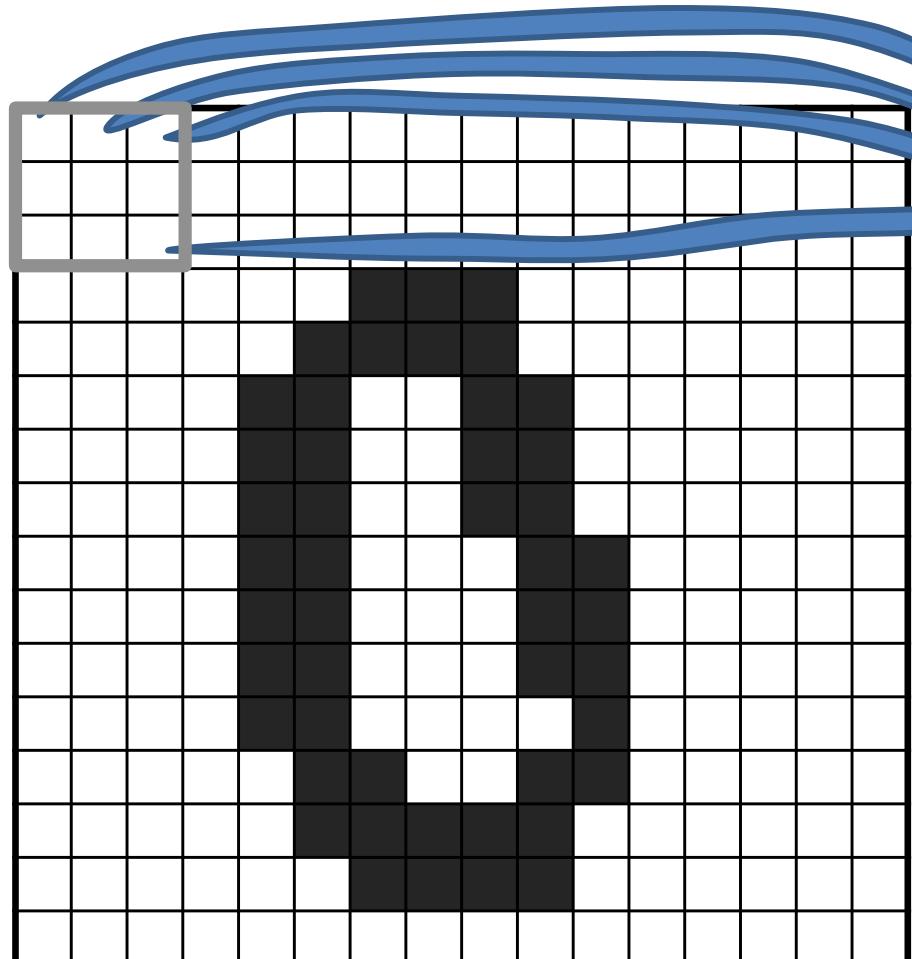


$\mathbf{x}$

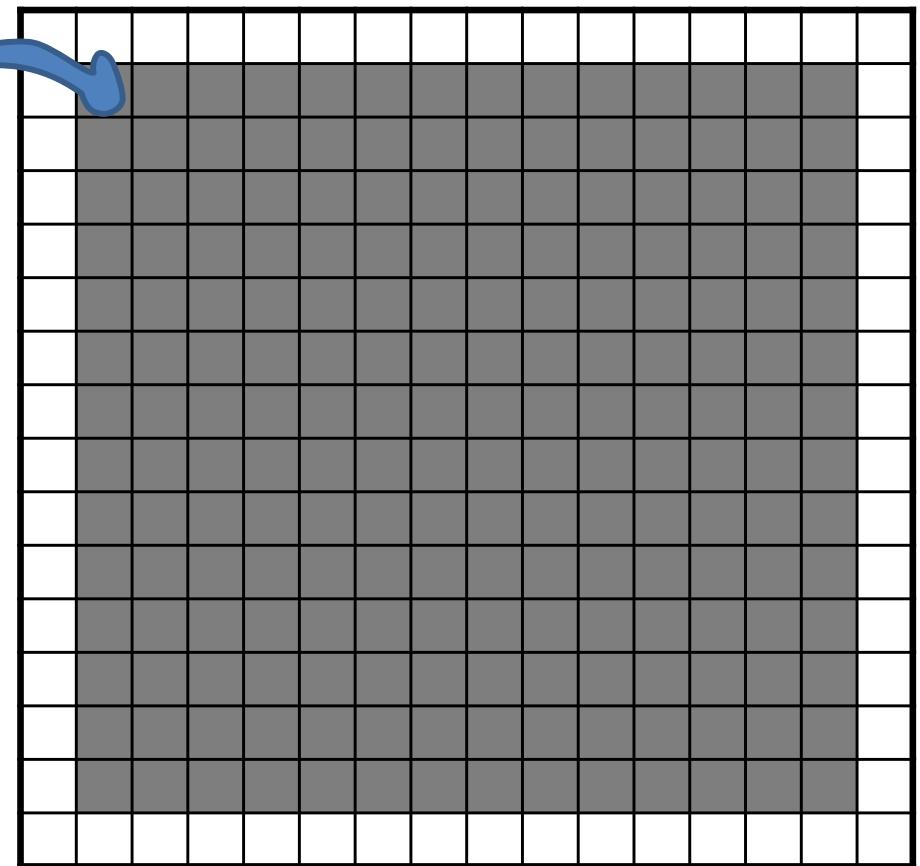


$\mathbf{h}_1$

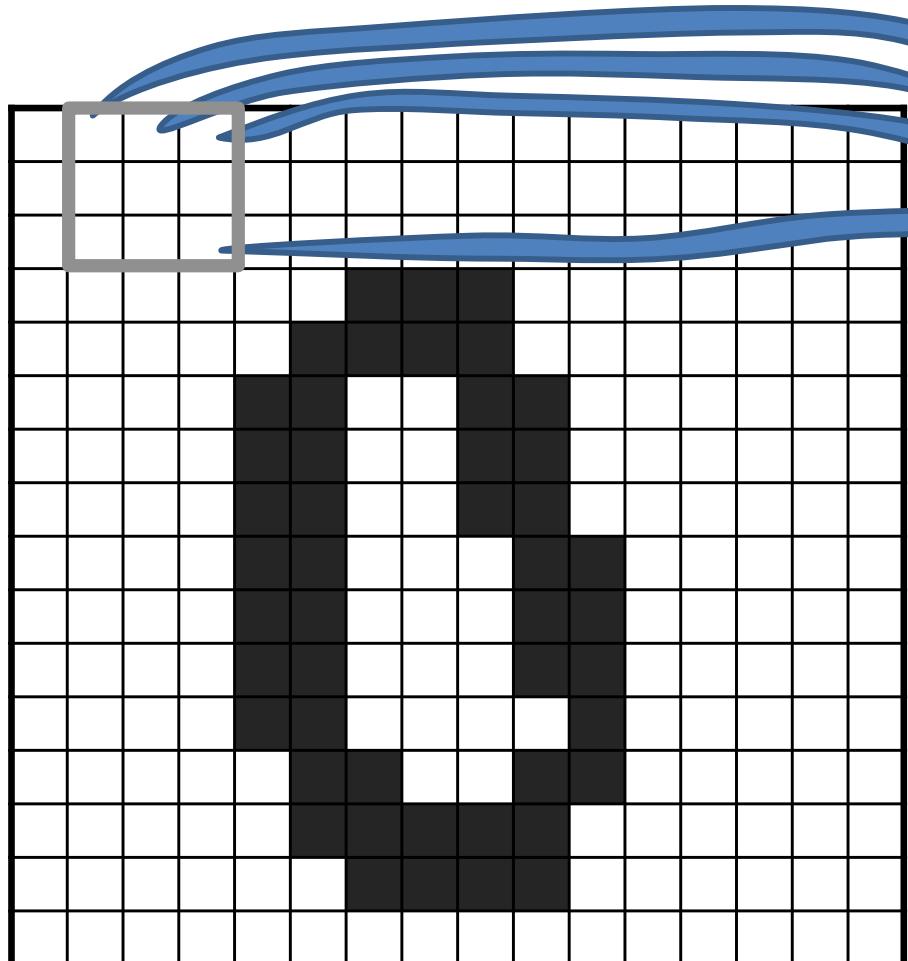
$$\mathbf{h}_1 = g(\mathbf{W}\mathbf{x} + \mathbf{b}) ?$$



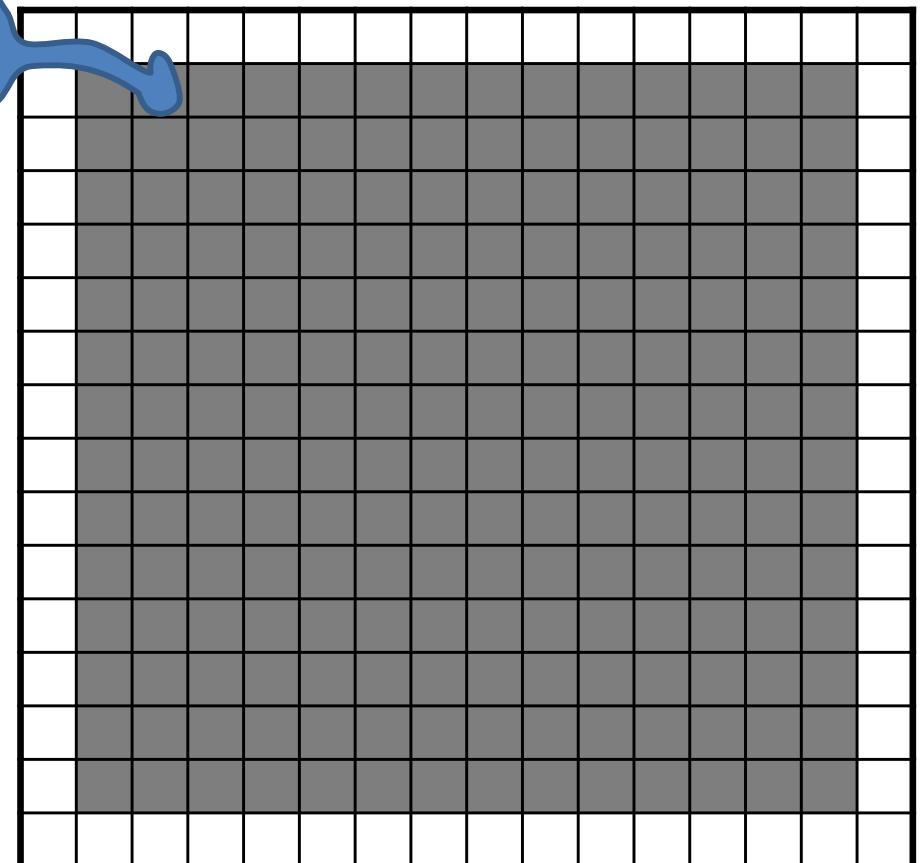
**X**



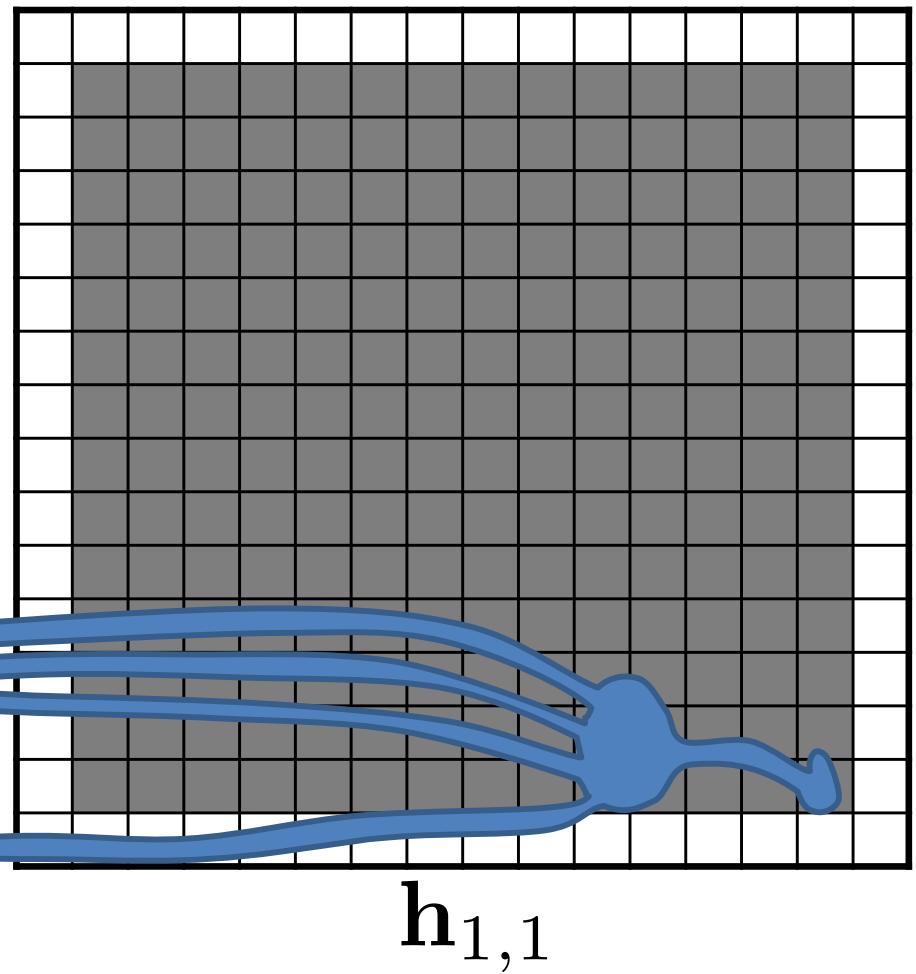
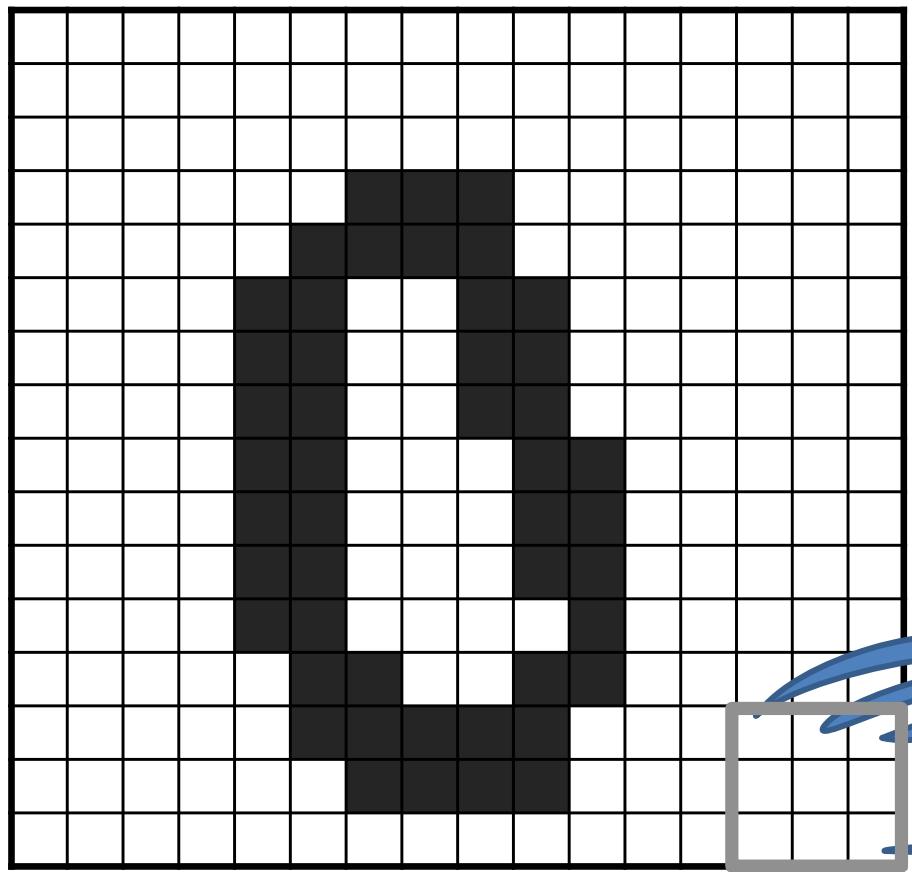
**$h_{1,1}$**



**X**

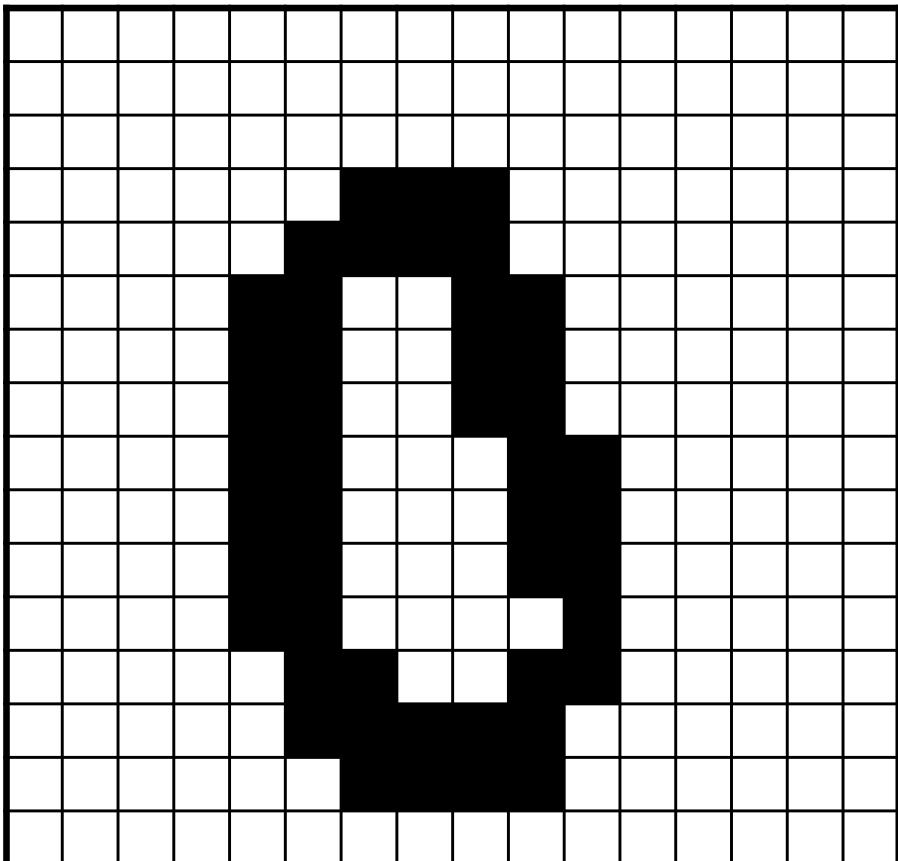


**$h_{1,1}$**



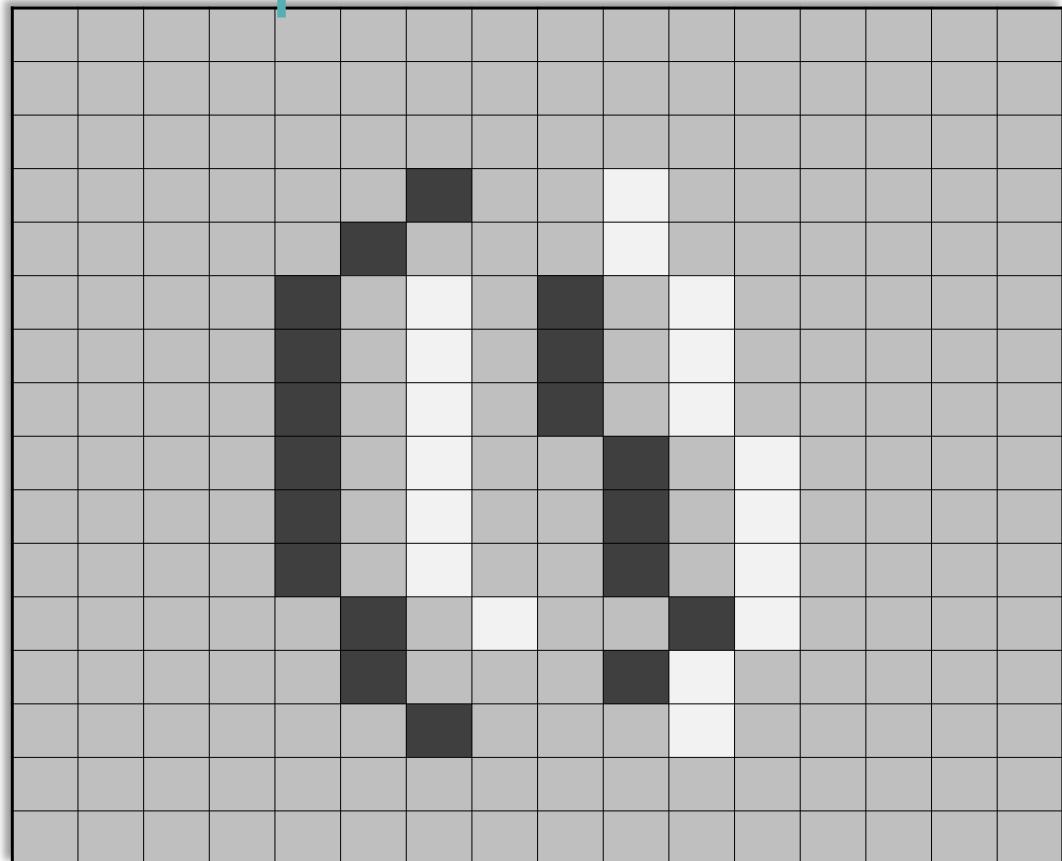
Product of convolution:  $\mathbf{h}_{1,1} = g(\mathbf{f}_{1,1} * \mathbf{x} + \mathbf{b}_{1,1})$

# Convolution: Example



**x**

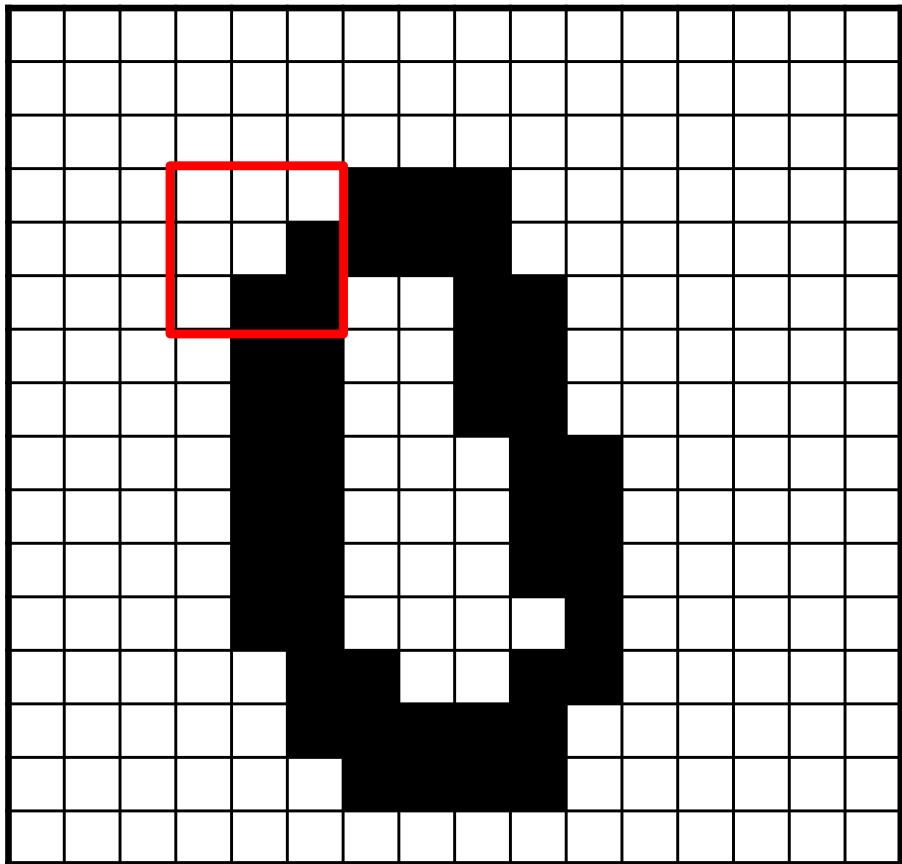
$$f_{1,1} = \begin{array}{|c|c|c|}\hline -1 & 0 & +1 \\ \hline -1 & 0 & +1 \\ \hline -1 & 0 & +1 \\ \hline\end{array}$$



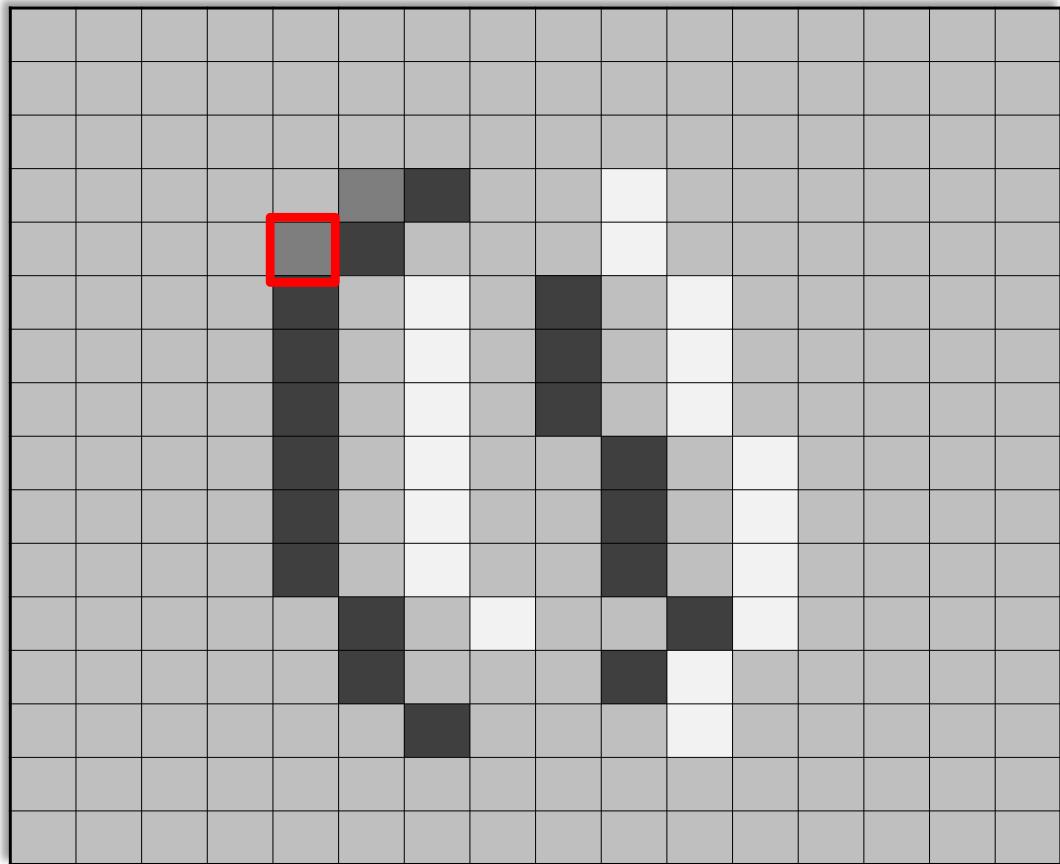
**$h_{1,1}$**

$$h_{1,1} = g(f_{1,1} * x + b_{1,1})$$

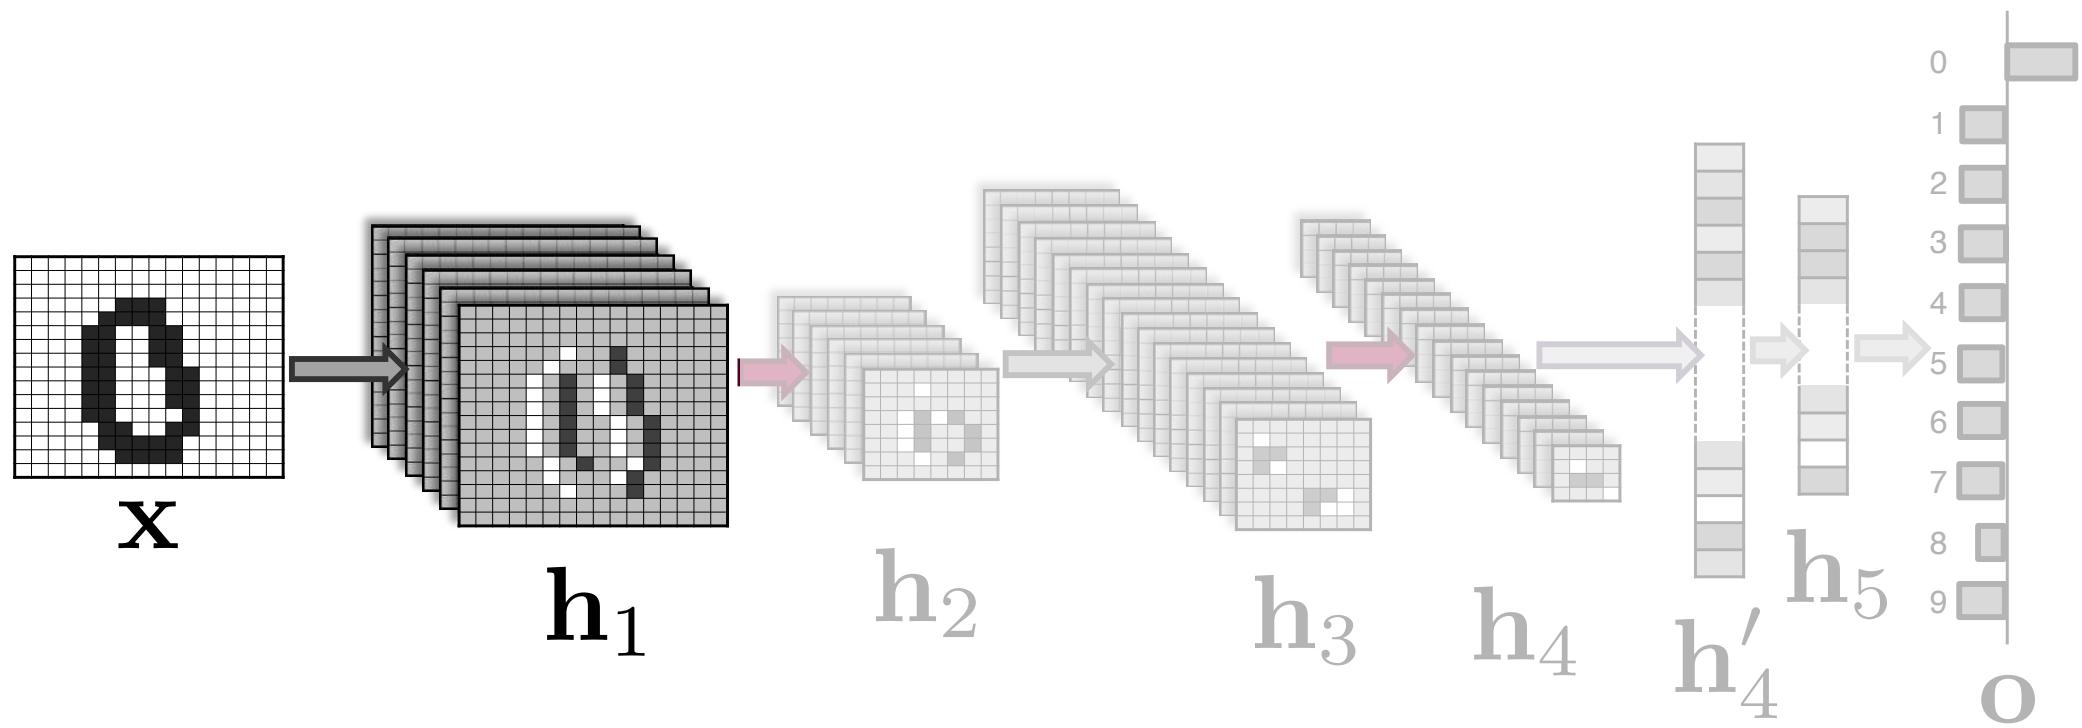
# Numerical Example



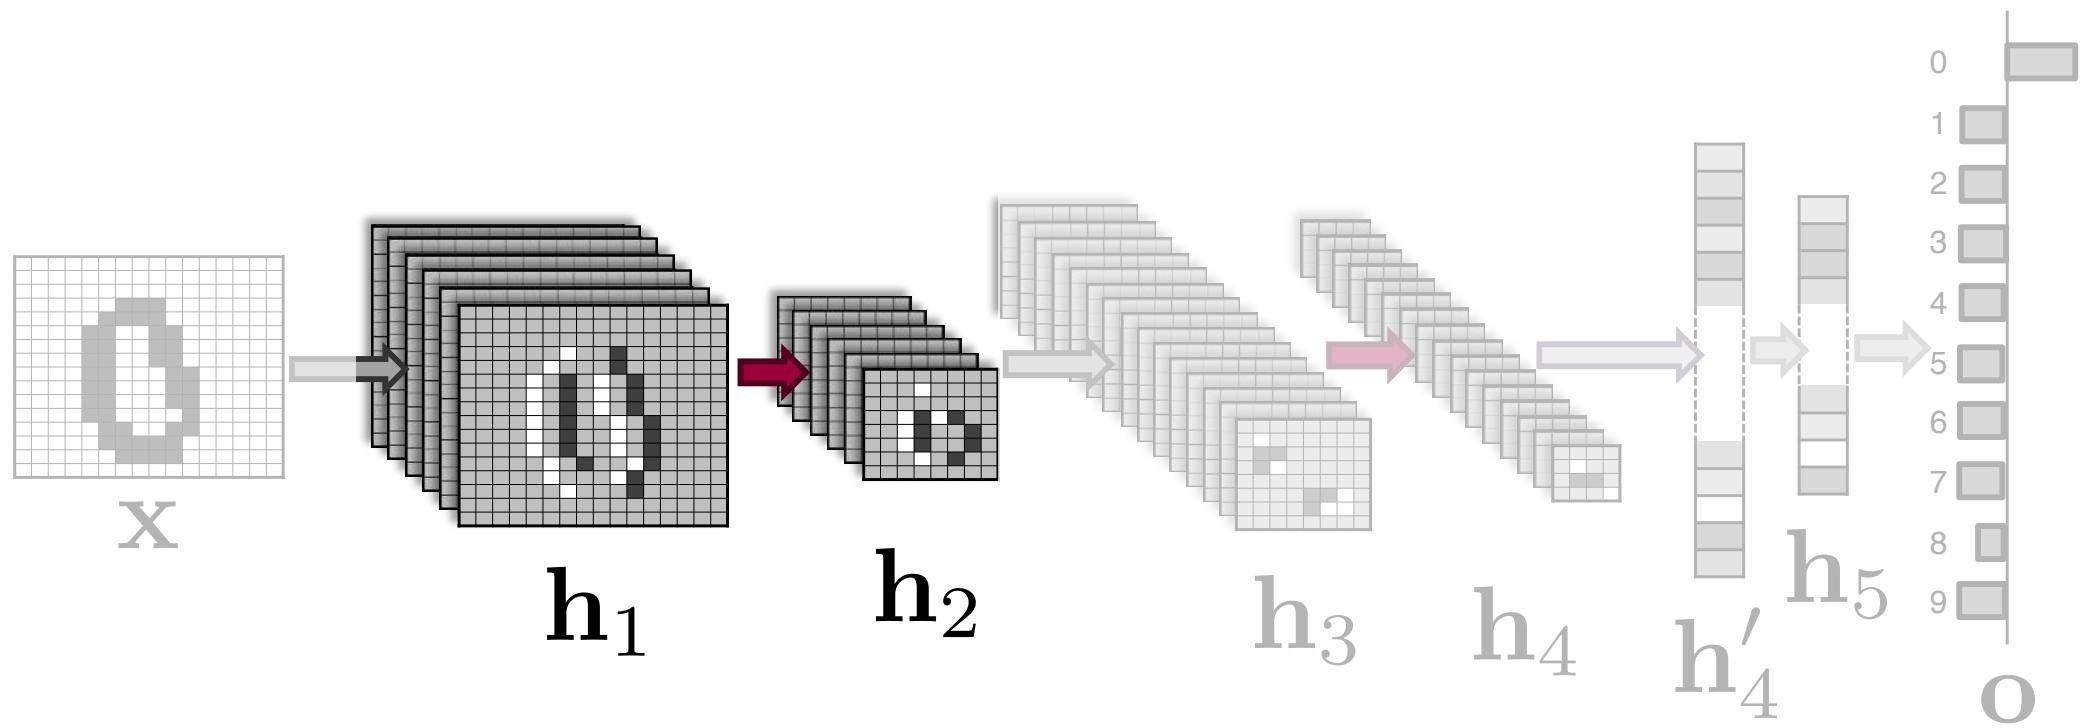
$$f_{1,1} = \begin{array}{|c|c|c|}\hline -1 & 0 & +1 \\ \hline -1 & 0 & +1 \\ \hline -1 & 0 & +1 \\ \hline\end{array}$$



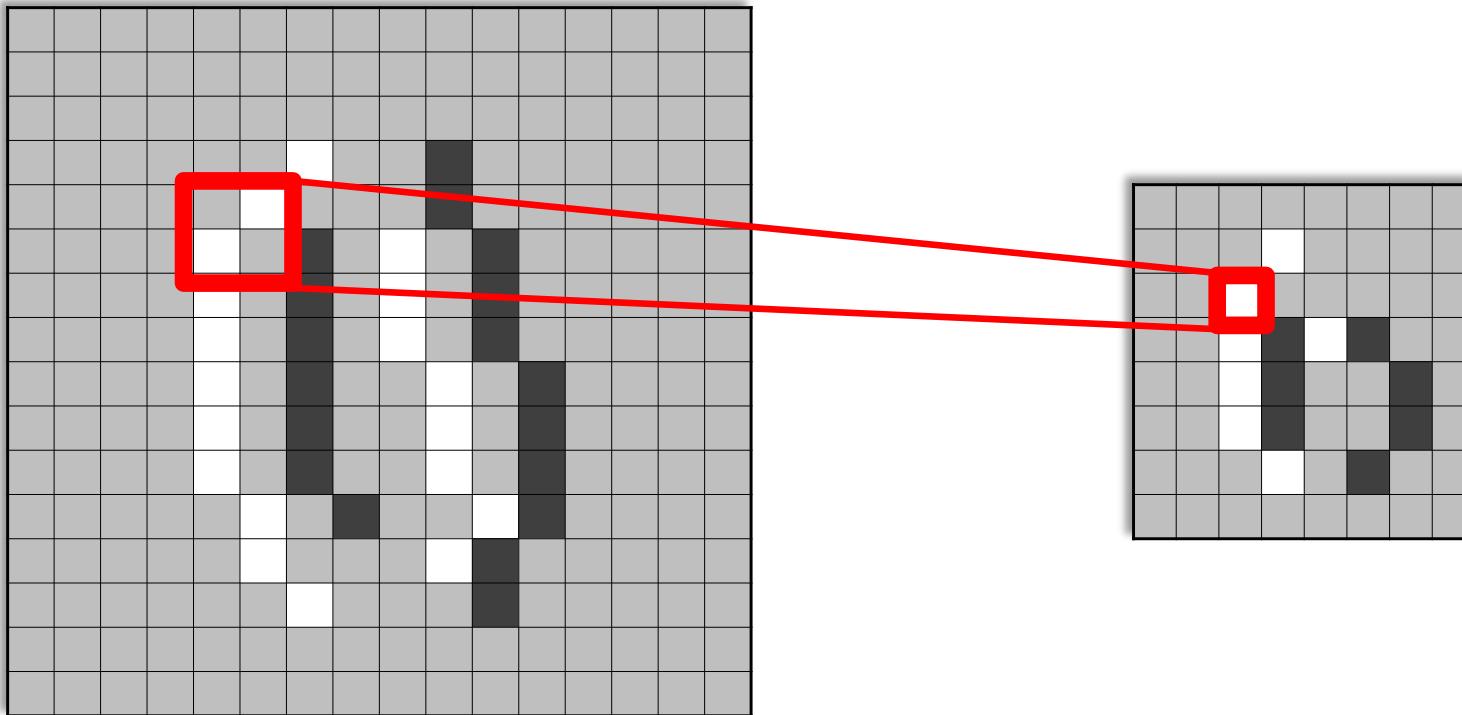
$$\begin{aligned} h &= (-1) \times 255 + 0 \times 255 + (+1) \times 255 + \\ &\quad (-1) \times 255 + 0 \times 255 + (+1) \times 0 + \\ &\quad (-1) \times 255 + 0 \times 0 + (+1) \times 0 \\ &= -255 + 0 + 255 \\ &\quad -255 + 0 + 0 \\ &\quad -255 + 0 + 0 \\ &= -510 \end{aligned}$$



$$\mathbf{h}_1 = [g(\mathbf{f}_{1,1} * \mathbf{x}), \dots, g(\mathbf{f}_{1,m} * \mathbf{x})]$$

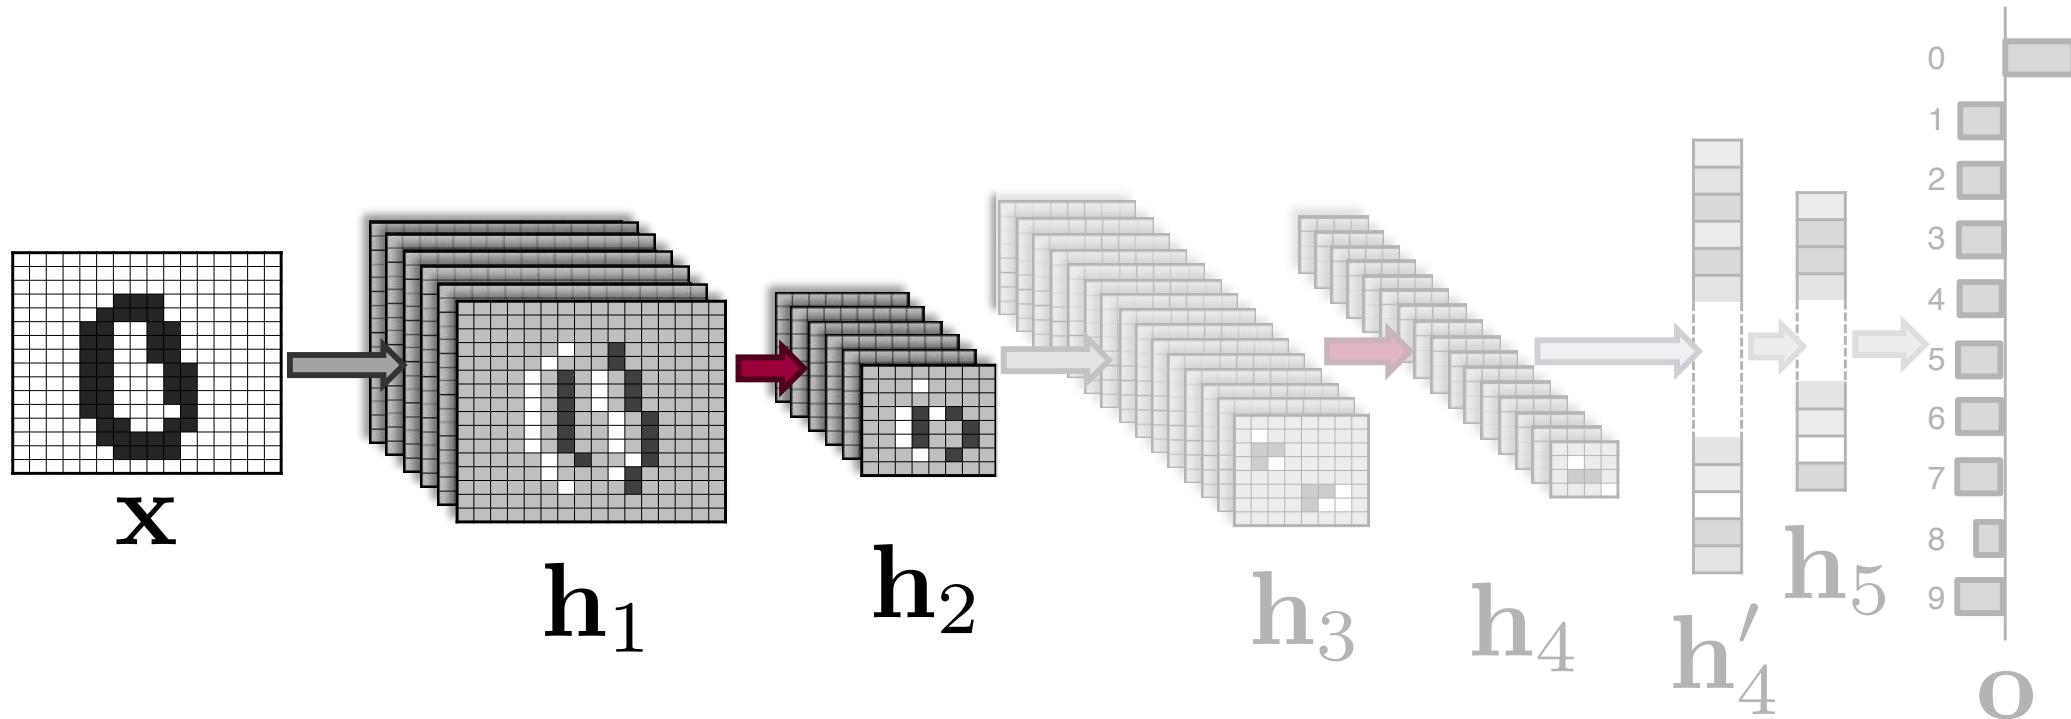


# Subsampling / Pooling



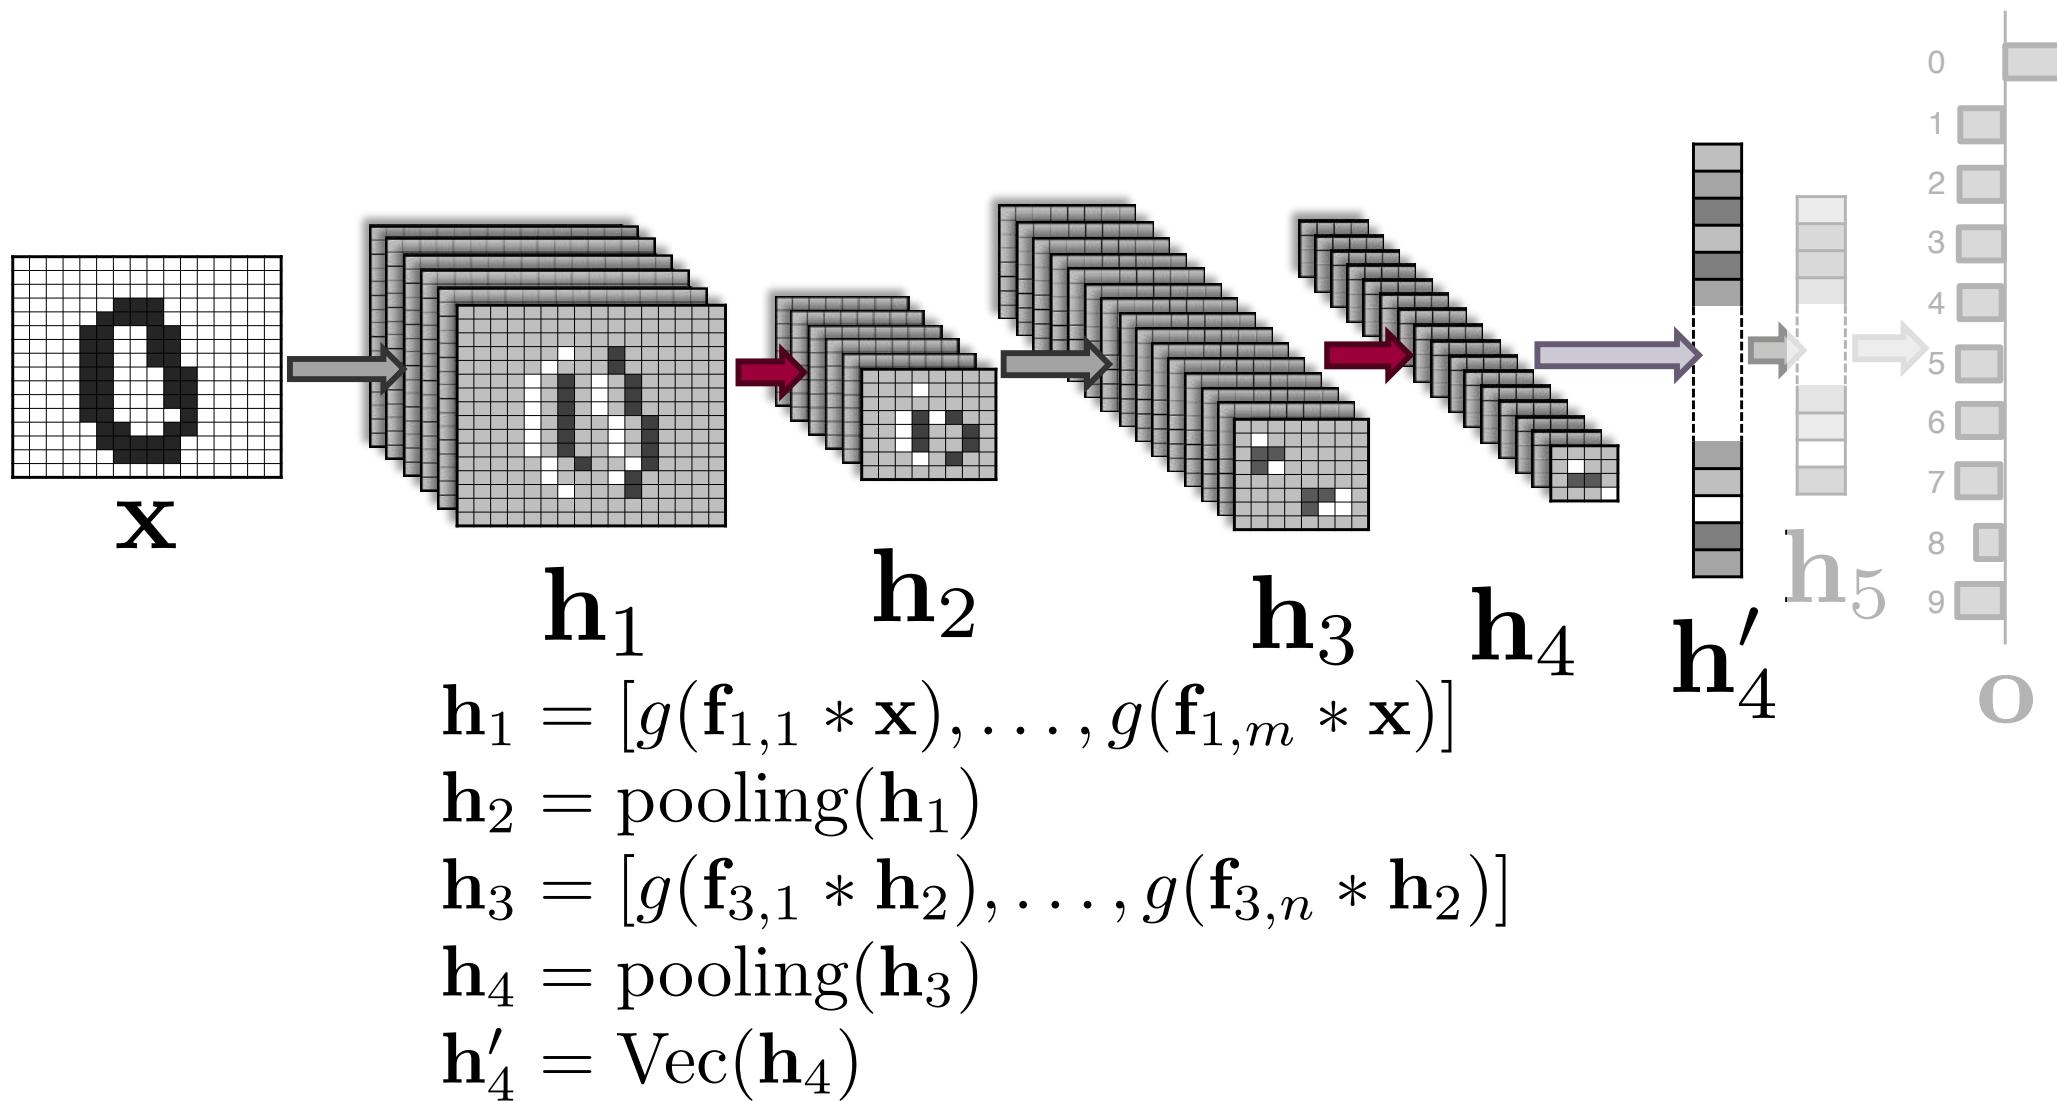
For example, max-pooling:

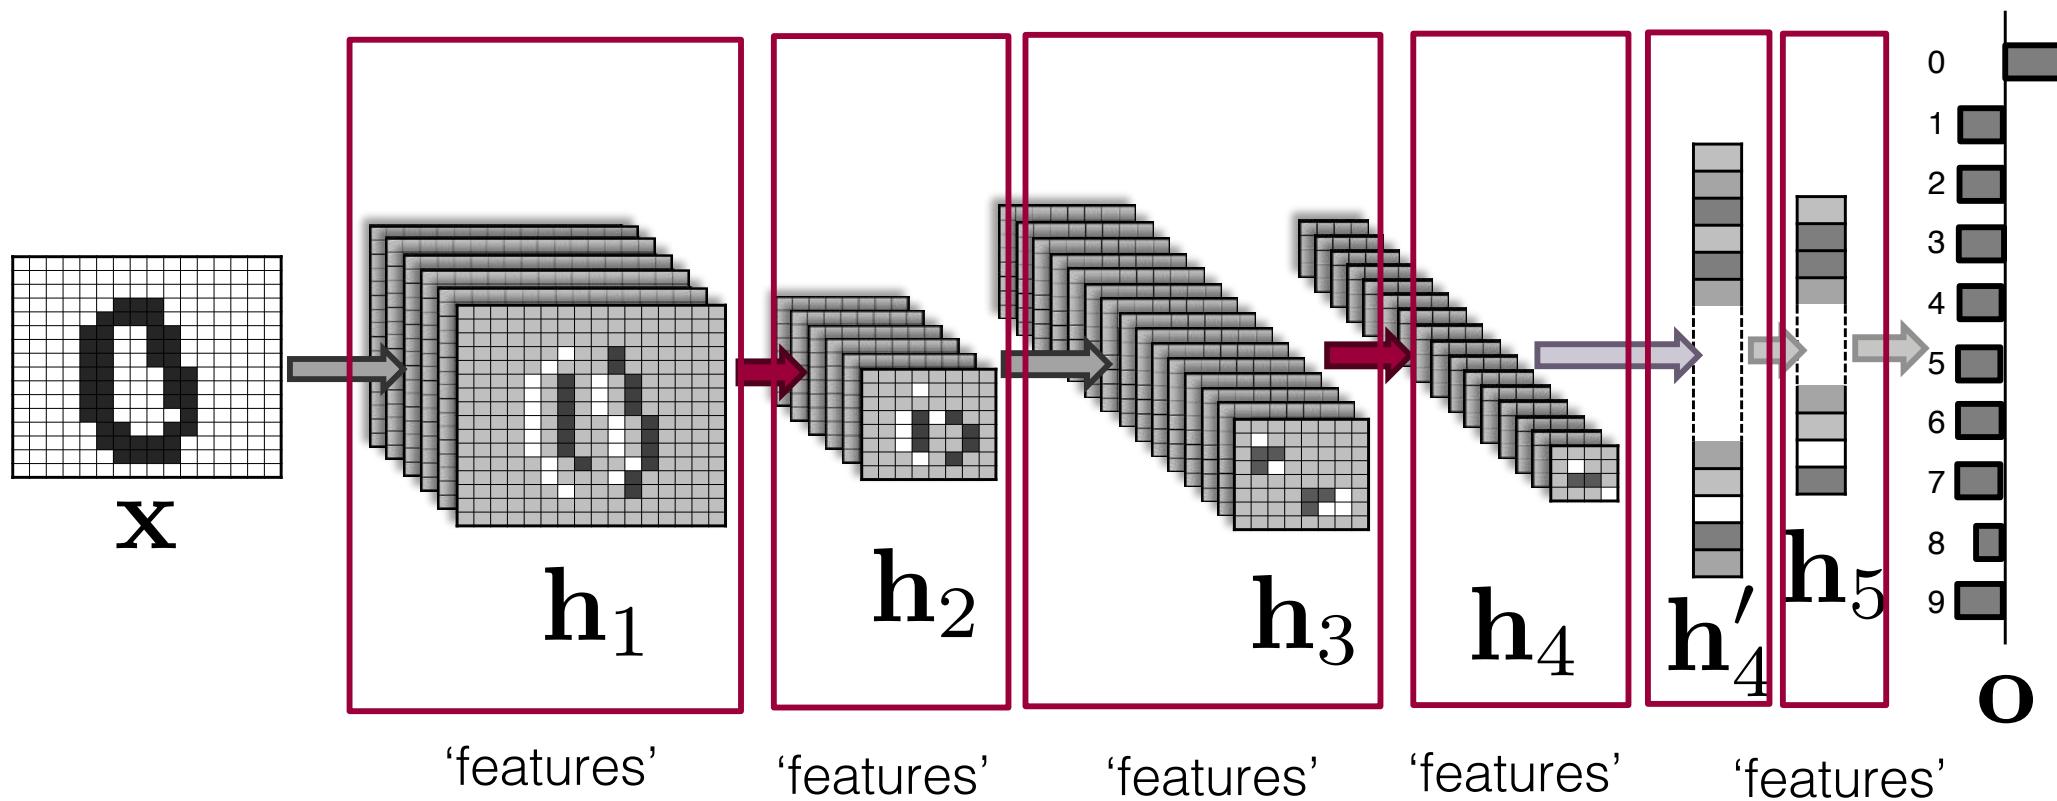
$$\mathbf{h}_i[u, v] = \max\{ \begin{array}{ll} \mathbf{h}_{i-1}[2u, & 2v], \\ \mathbf{h}_{i-1}[2u, & 2v + 1], \\ \mathbf{h}_{i-1}[2u + 1, & 2v], \\ \mathbf{h}_{i-1}[2u + 1, & 2v + 1] \end{array} \}$$



$$\begin{aligned}\mathbf{h}_1 &= [g(\mathbf{f}_{1,1} * \mathbf{x}), \dots, g(\mathbf{f}_{1,m} * \mathbf{x})] \\ \mathbf{h}_2 &= \text{pooling}(\mathbf{h}_1)\end{aligned}$$

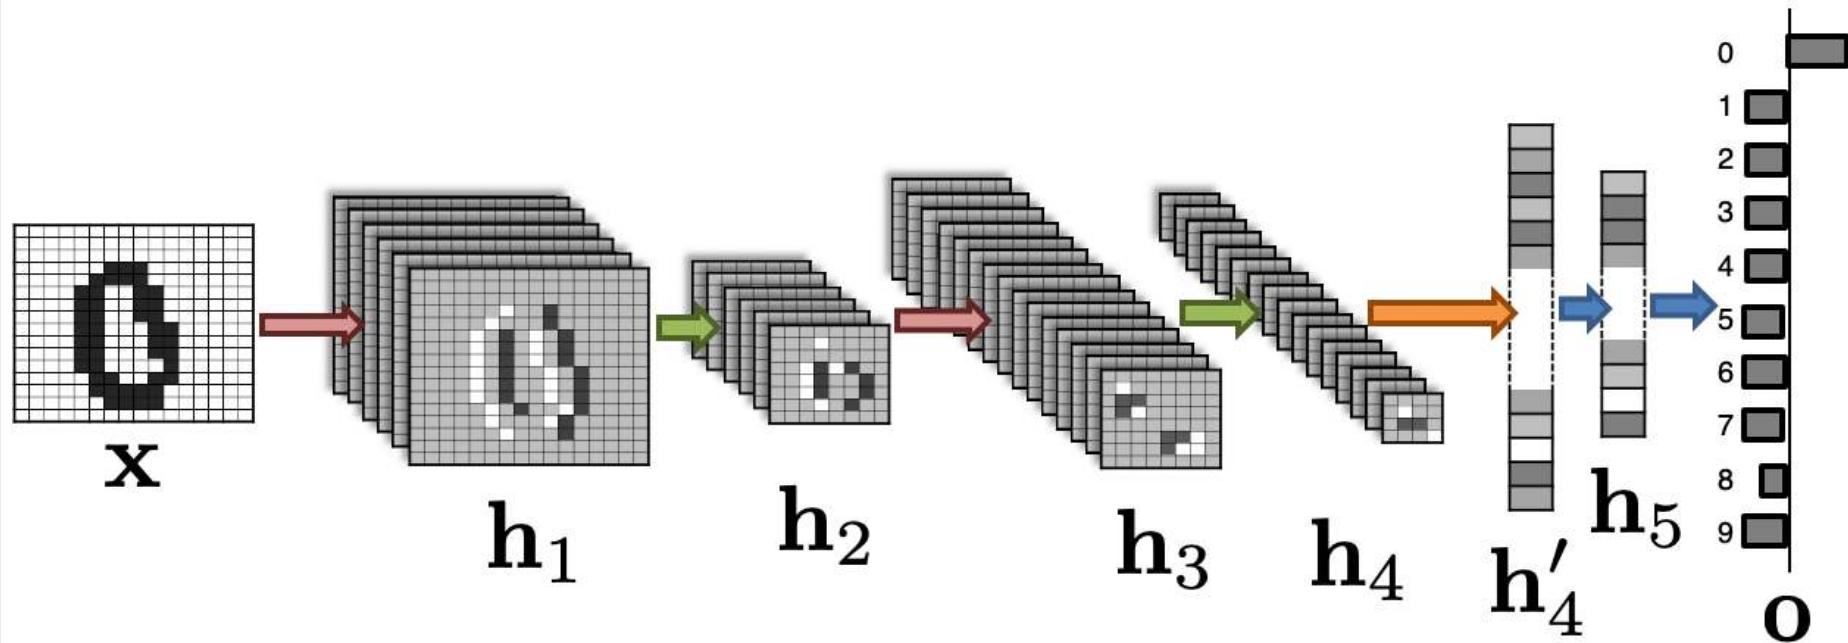
Inspired by the theory of Hubel and Wiesel on V1  
(Nobel prize in 1981)





First features may be ‘interpreted’, but the rest is usually more difficult.

# How Can We Find the Network's Parameters?



$$\mathbf{h}_1 = [g(\mathbf{f}_{1,1} * \mathbf{x}), \dots, g(\mathbf{f}_{1,m} * \mathbf{x})]$$

$$\mathbf{h}_2 = \text{pooling}(\mathbf{h}_1)$$

$$\mathbf{h}_3 = [g(\mathbf{f}_{3,1} * \mathbf{h}_2), \dots, g(\mathbf{f}_{3,n} * \mathbf{h}_2)]$$

$$\mathbf{h}_4 = \text{pooling}(\mathbf{h}_3)$$

$$\mathbf{h}'_4 = \text{Vec}(\mathbf{h}_4)$$

$$\mathbf{h}_5 = g(\mathbf{W}_5 \mathbf{h}'_4 + \mathbf{b}_5)$$

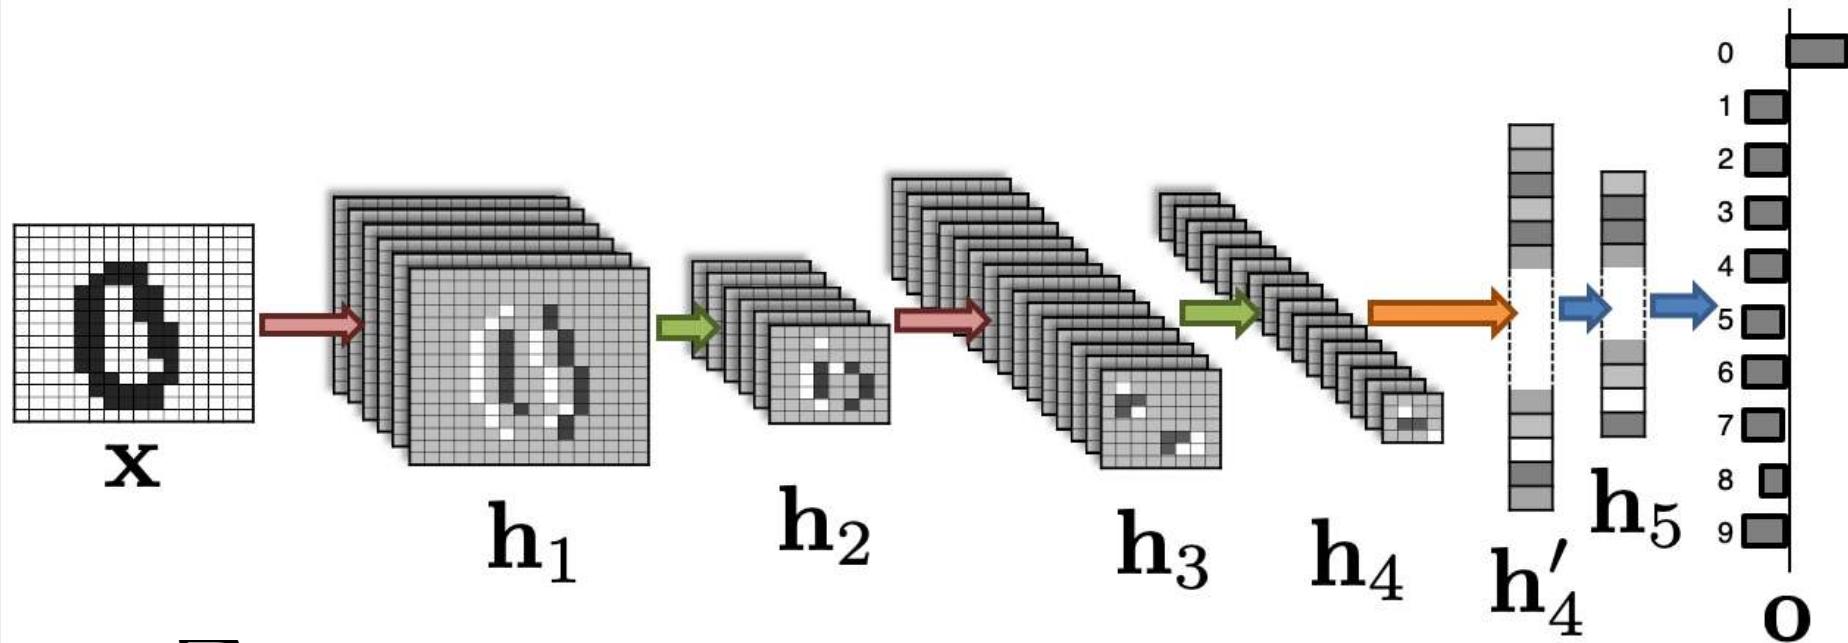
$$\mathbf{o} = \mathbf{W}_6 \mathbf{h}_5 + \mathbf{b}_6$$

$$\mathcal{L}(\Theta) = \sum_i L(c_i, f(\mathbf{x}_i; \Theta))$$

$$\Theta = (\mathbf{f}_{1,1}, \dots, \mathbf{f}_{1,m}, \dots, \mathbf{W}_5, \mathbf{b}_5, \mathbf{W}_6, \mathbf{b}_6)$$

# OPTIMIZING / TRAINING A DEEP NETWORK

# How Can We Find the Network's Parameters?



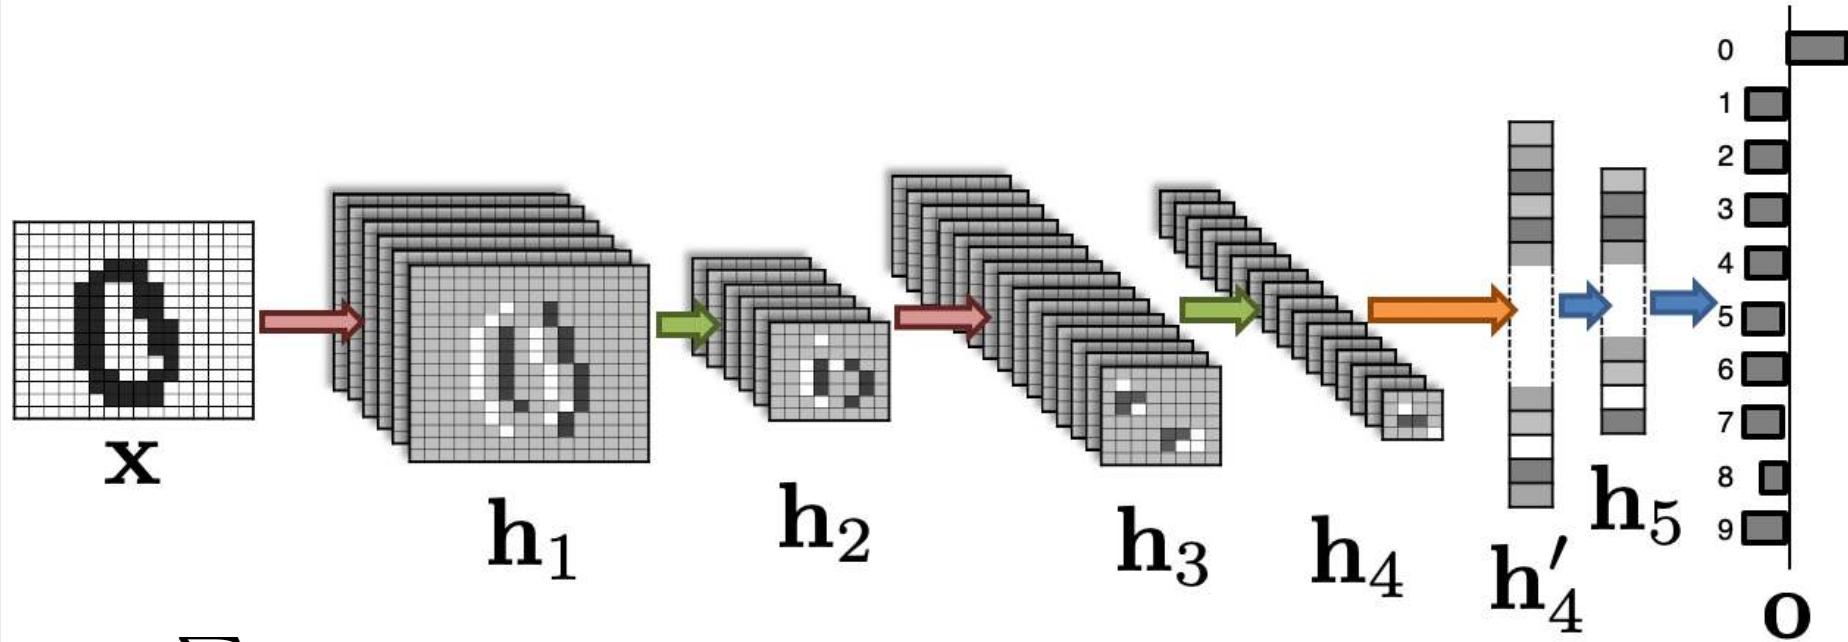
$$\mathcal{L}(\Theta) = \sum_i L(c_i, f(\mathbf{x}_i; \Theta))$$

$$\Theta = (\mathbf{f}_{1,1}, \dots, \mathbf{f}_{1,m}, \dots, \mathbf{W}_5, \mathbf{b}_5, \mathbf{W}_6, \mathbf{b}_6)$$

As before:

- Initialize  $\Theta$  randomly;
- Optimize  $\Theta$  using gradient descent.

# backpropagation

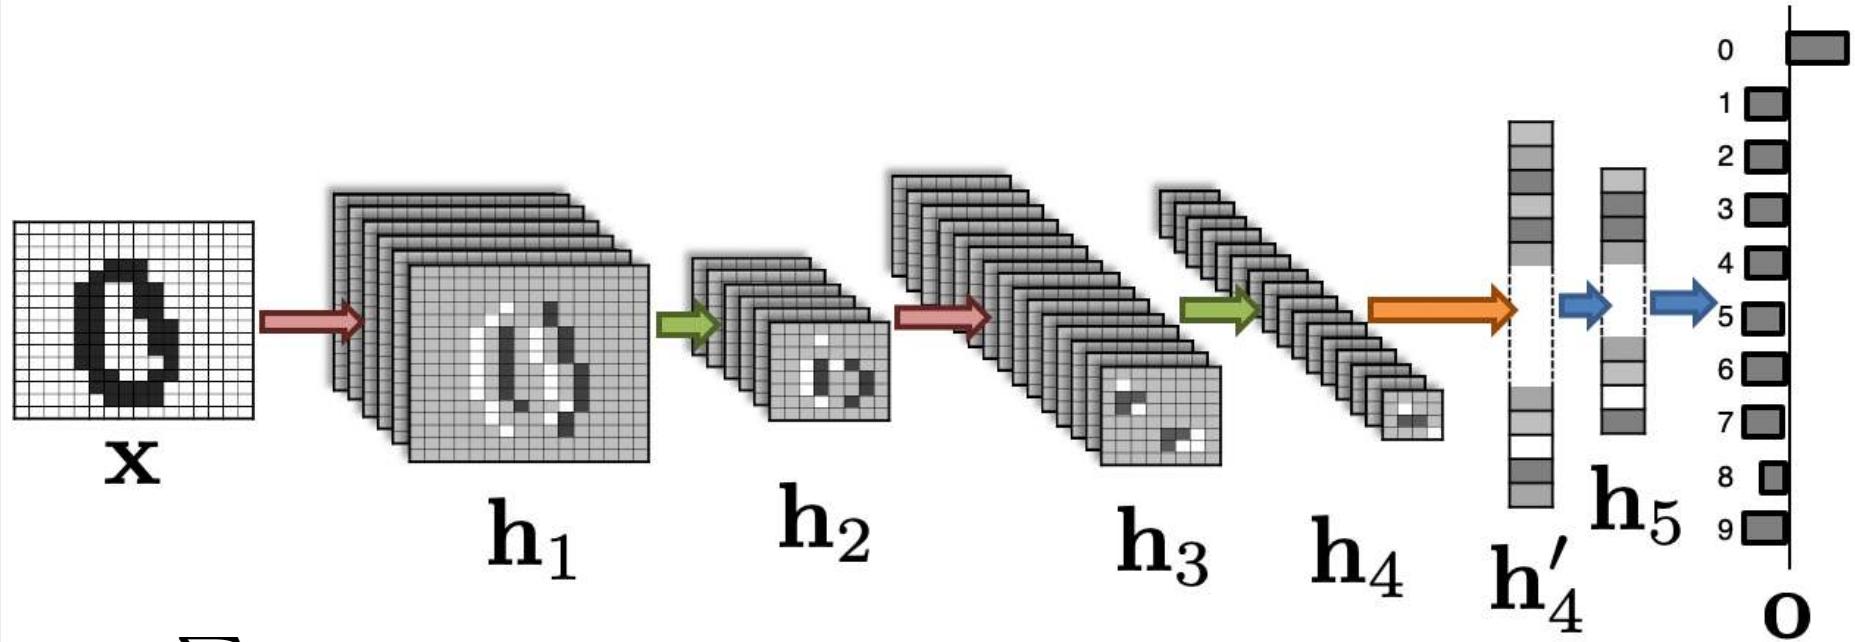


$$\mathcal{L}(\Theta) = \sum_i L(c_i, f(\mathbf{x}_i; \Theta))$$

$$\Theta = (\mathbf{f}_{1,1}, \dots, \mathbf{f}_{1,m}, \dots, \mathbf{W}_5, \mathbf{b}_5, \mathbf{W}_6, \mathbf{b}_6)$$

**Back-propagation:** An efficient method to compute the gradient of the objective function;

# backpropagation



$$\mathcal{L}(\Theta) = \sum_i L(c_i, f(\mathbf{x}_i; \Theta))$$

$$\Theta = (\mathbf{f}_{1,1}, \dots, \mathbf{f}_{1,m}, \dots, \mathbf{W}_5, \mathbf{b}_5, \mathbf{W}_6, \mathbf{b}_6)$$

Several variants of gradient descent have been developed recently (Adam, RMSProp).

# Python Libraries

TensorFlow (Google), Keras (Google – higher level), PyTorch (Facebook), ..

```
from keras.models import Sequential
from keras.layers import Conv2D, MaxPooling2D

model = Sequential()
model.add(Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28, 1)))
from keras.layers import Flatten
model.add(Flatten())

from keras.layers import Dense
model.add(Dense(128, activation='relu'))

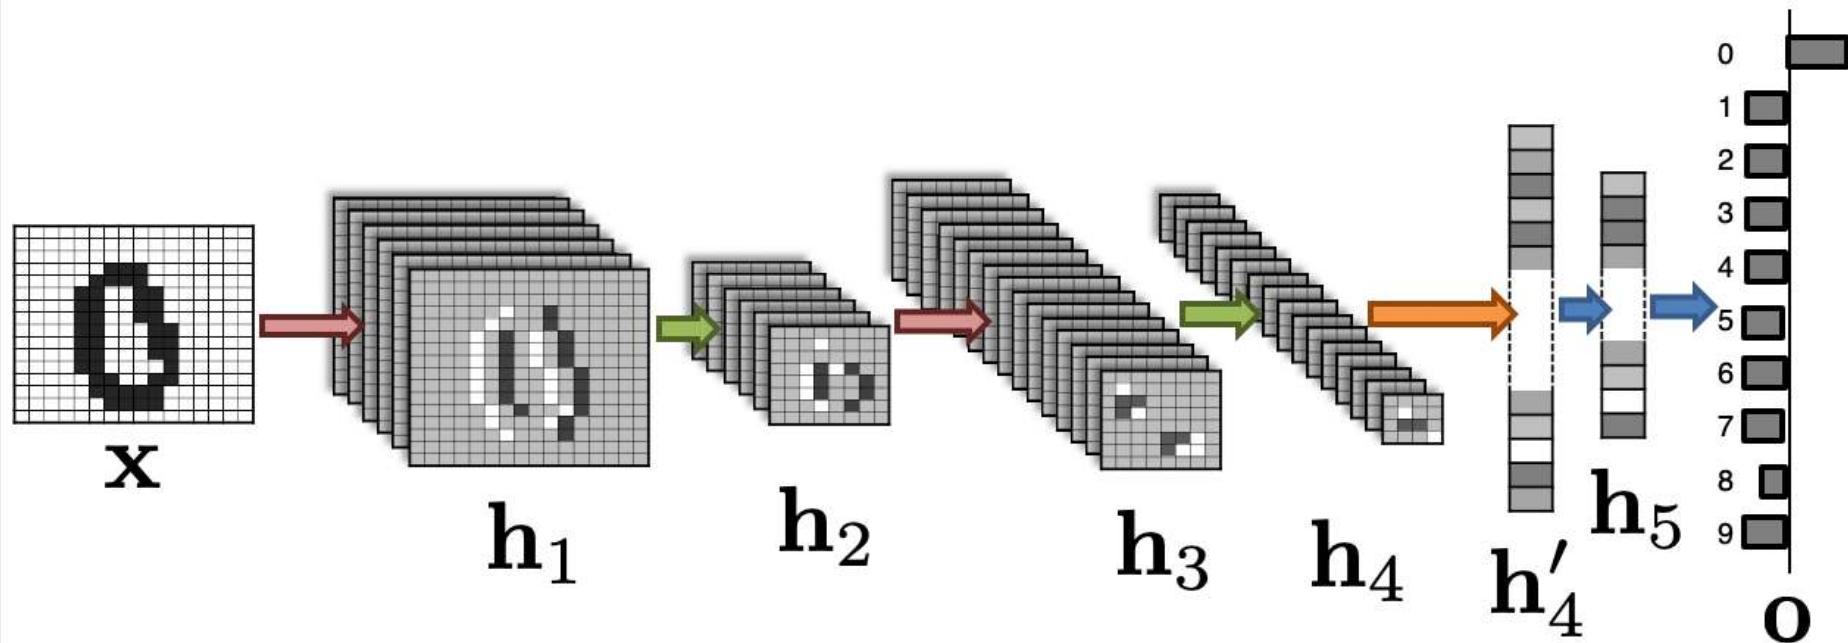
model.add(Dense(10, activation='softmax'))
model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])

model.fit(X_train, Y_train, batch_size=32, epochs=10, verbose=1)
```

# How to Choose the Number of Layers?

The Number of Filters per Layer?

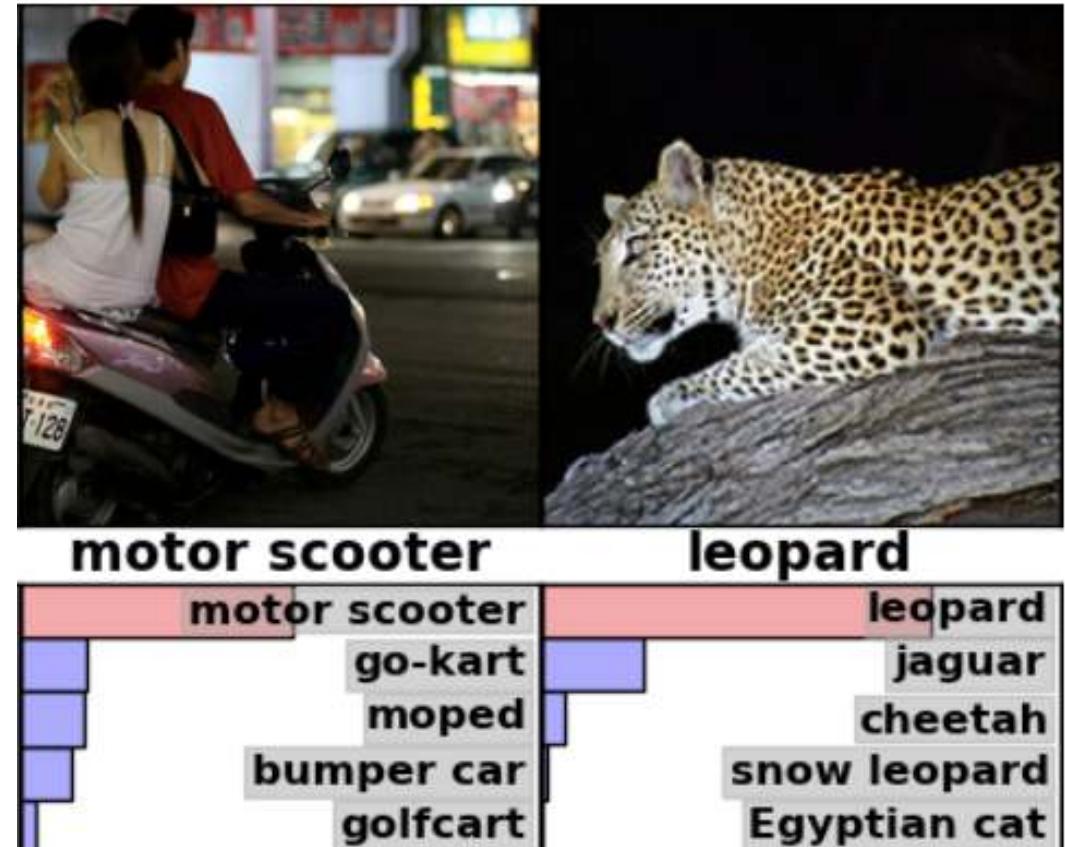
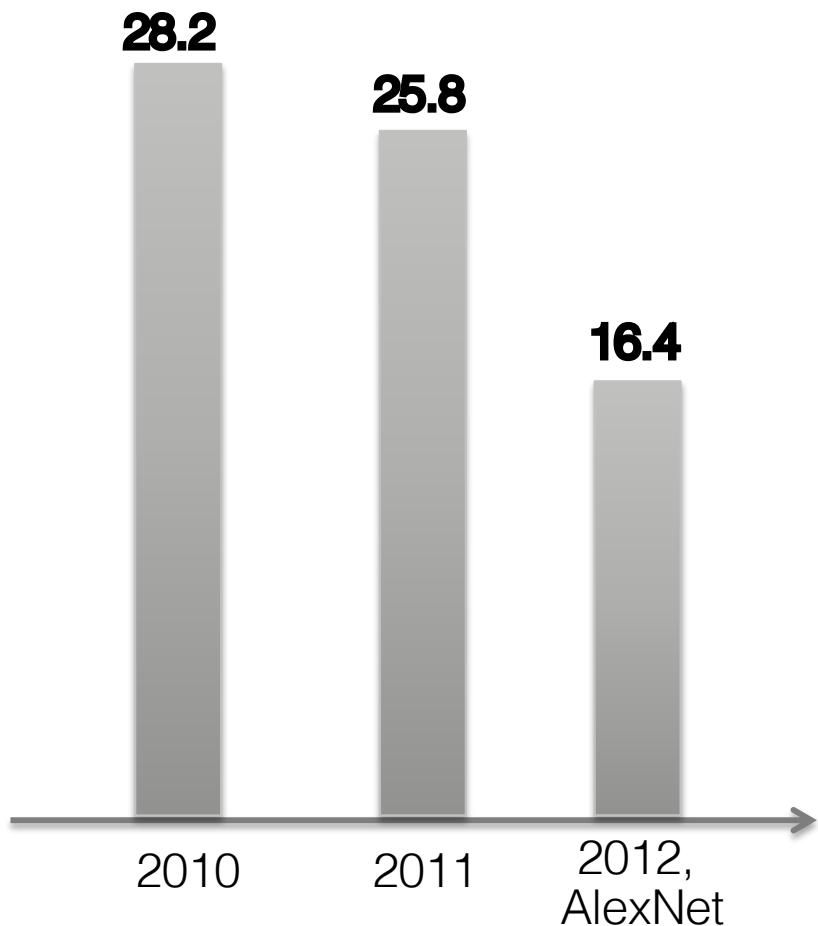
The Sizes of the Filters?



- experience;
- trial-and-error;
- Auto-ML: automated methods to find a good **architecture**.

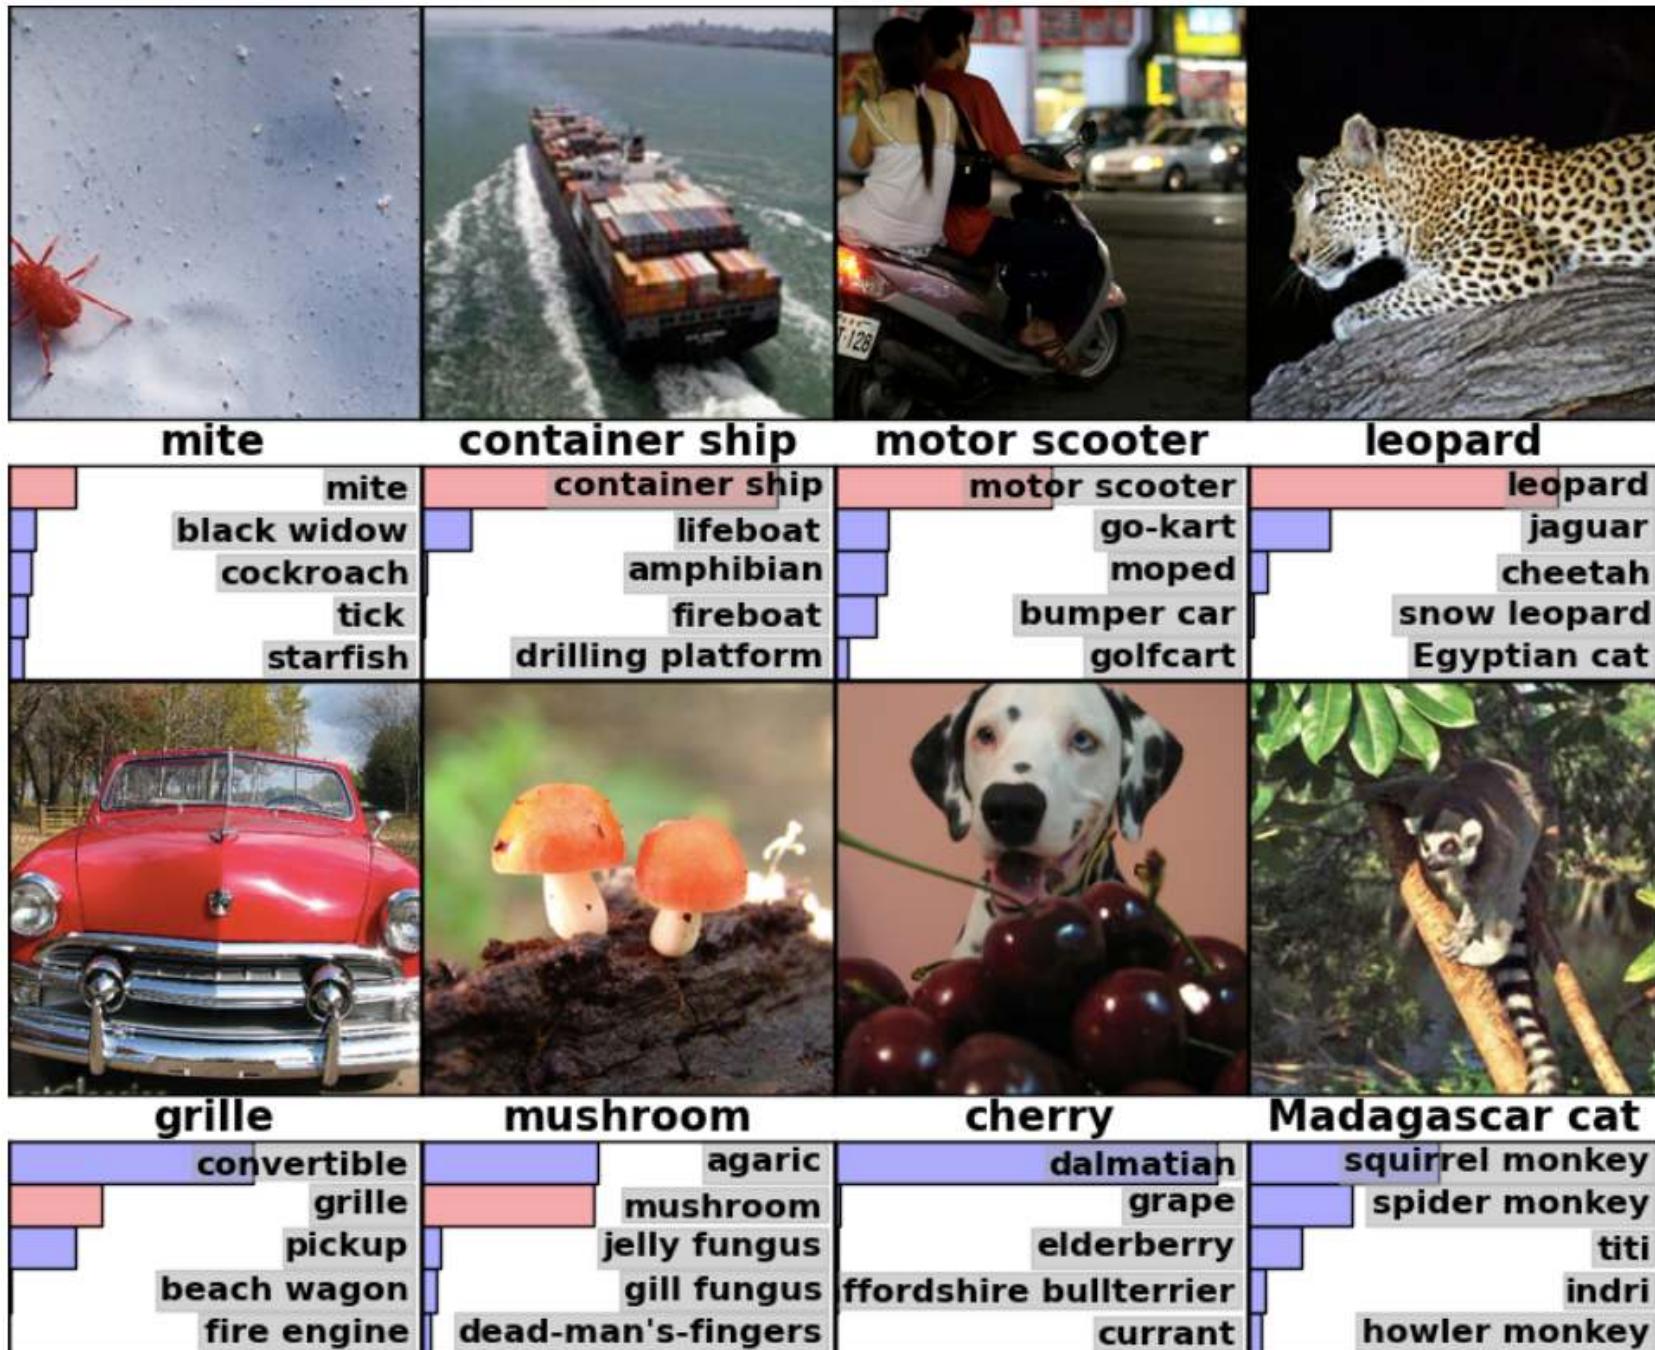
## MORE RECENT RESULTS

# AlexNet (2010)

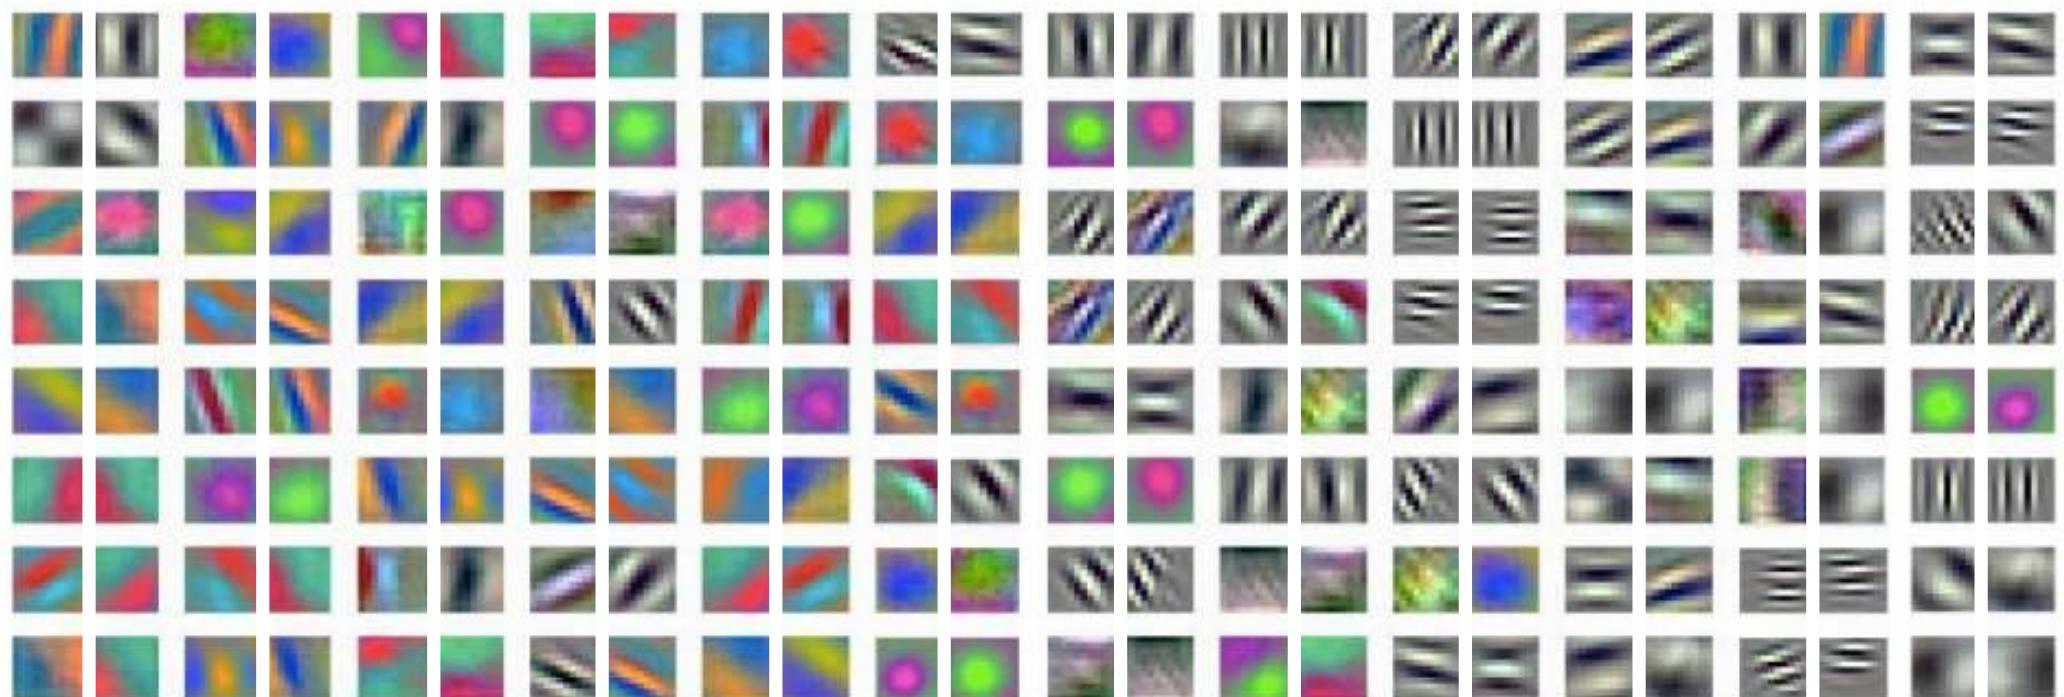


Classification task on ImageNet  
challenge top-5 error.

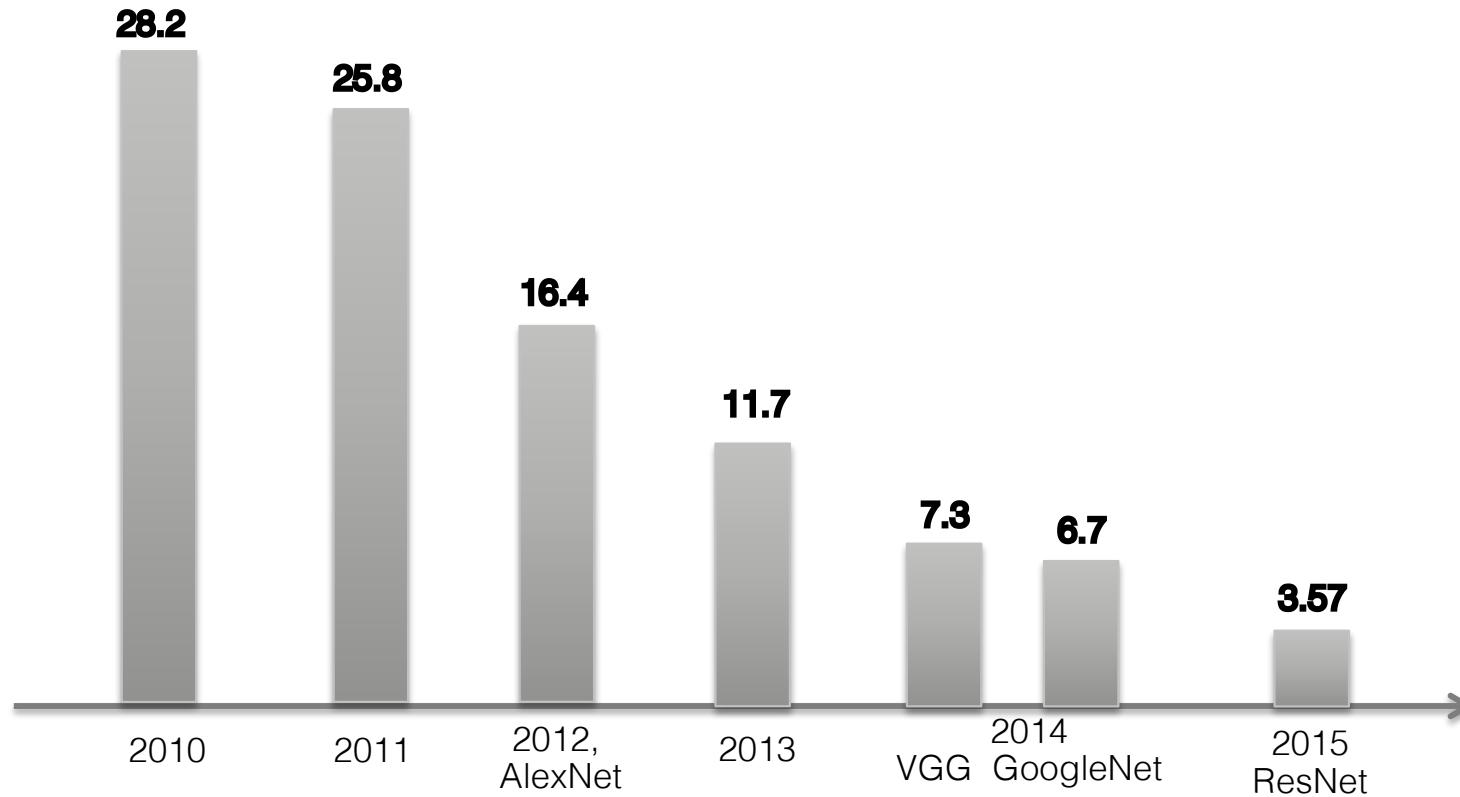
# AlexNet: Some results on Imagenet



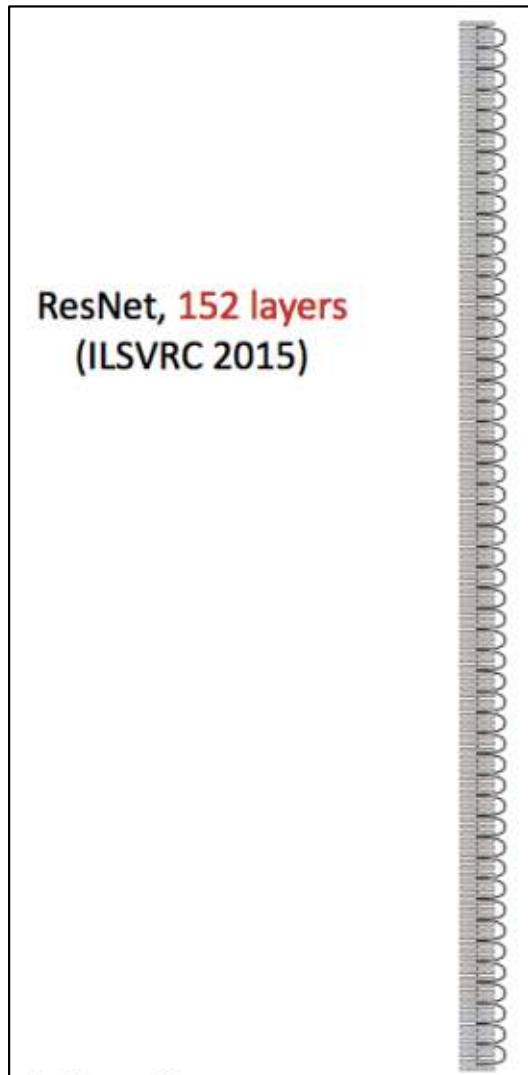
# AlexNet: Learned Filters for the First Layer



# Classification task on ImageNet challenge top-5 error

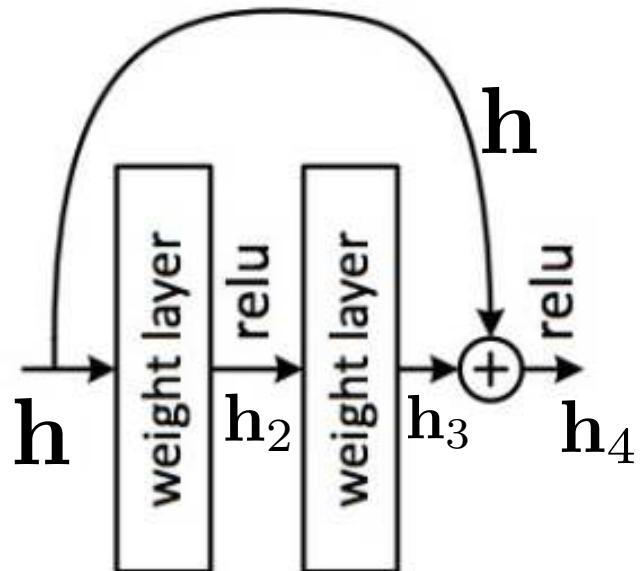


# ResNet [He et al, CVPR 2016] ILSVRC 2015 Winner



# The Residual Module (resnet block)

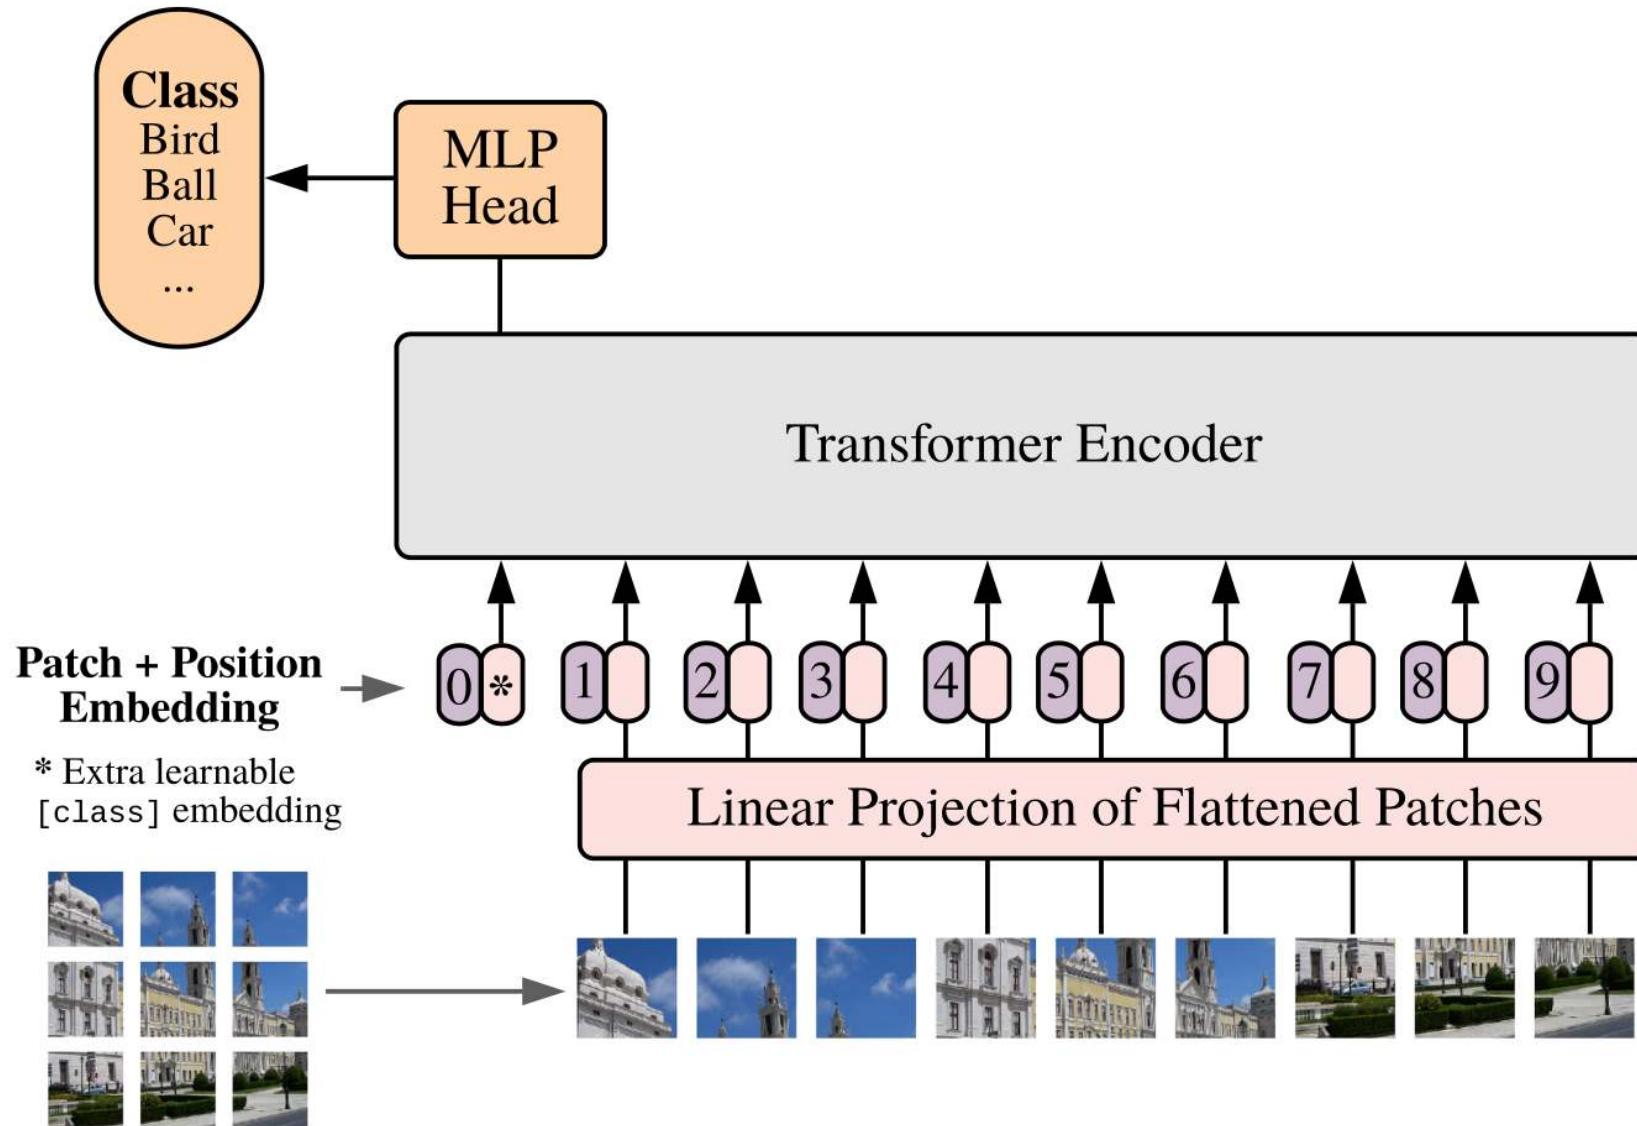
Uses 'skip' or 'shortcut' connections:



$$\begin{aligned}\mathbf{h}_2 &= g(\mathbf{W}_2 \mathbf{h} + \mathbf{b}_2) \\ \mathbf{h}_3 &= \mathbf{W}_3 \mathbf{h}_2 + \mathbf{b}_3 \\ \mathbf{h}_4 &= g(\mathbf{h}_3 + \mathbf{h})\end{aligned}$$

- Makes it easy for network layers to represent the identity mapping;
- Limits vanishing and exploding gradients.

# Transformers



BEING CAREFUL ABOUT  
WHAT A DEEP NETWORK  
REALLY LEARNS

# recognizing objects

traffic light (99) leaf beetle (99)



racket (51)



tree frog (99)



cash machine (97)



beacon (99)



padlock (99)



ice lolly (99)



# recognizing objects



(a) Output prediction on original images.



(b) Prediction when foreground is whitened.

# recognizing objects

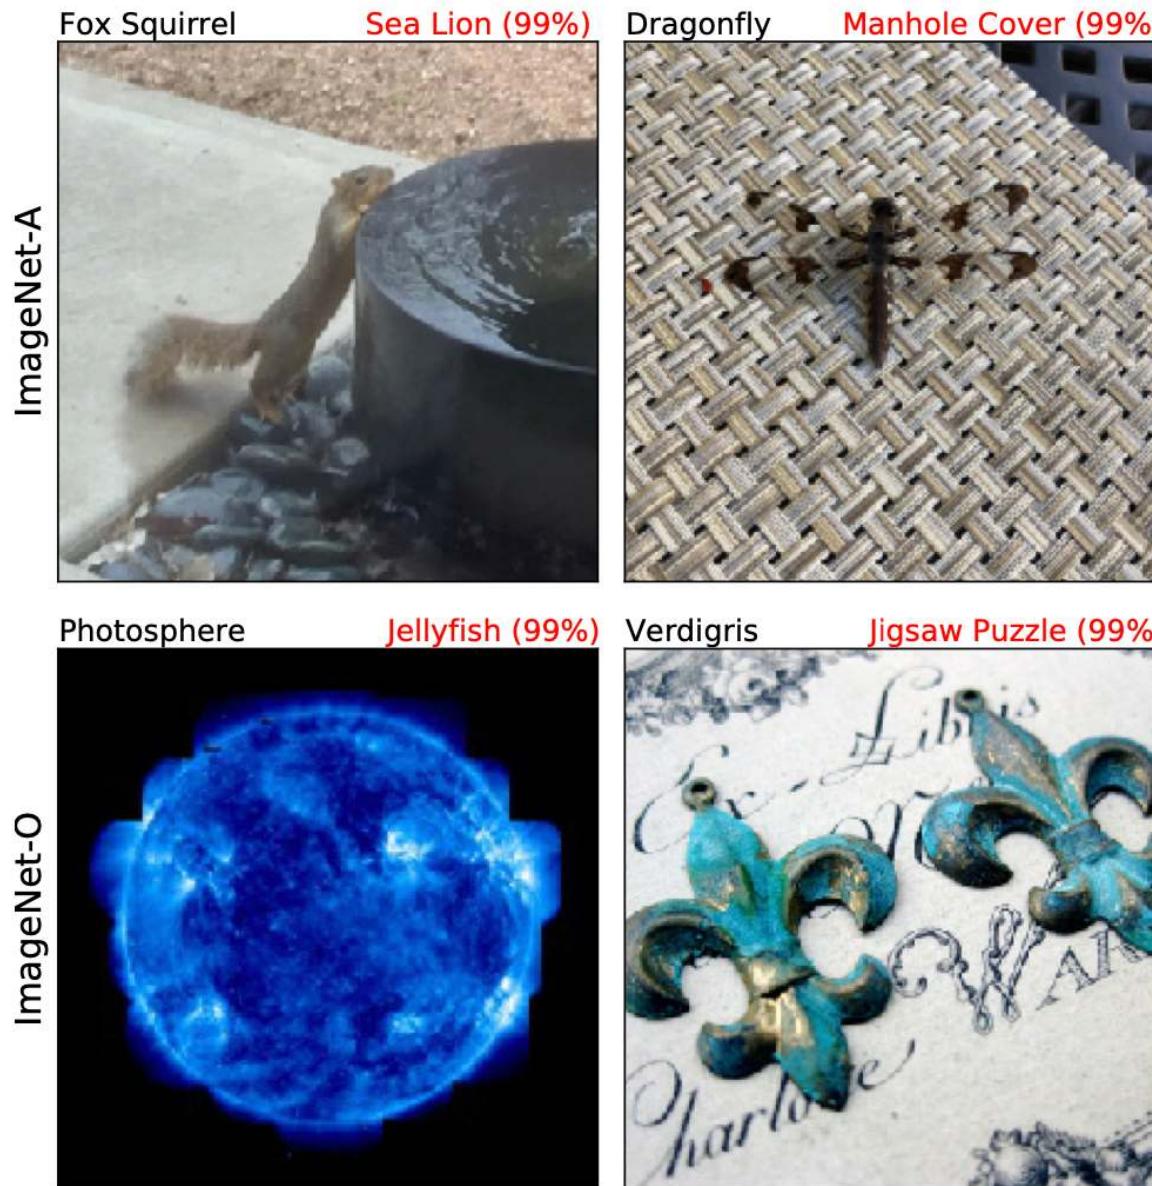


(a) Output prediction on original images.



(b) Prediction when foreground is whitened.

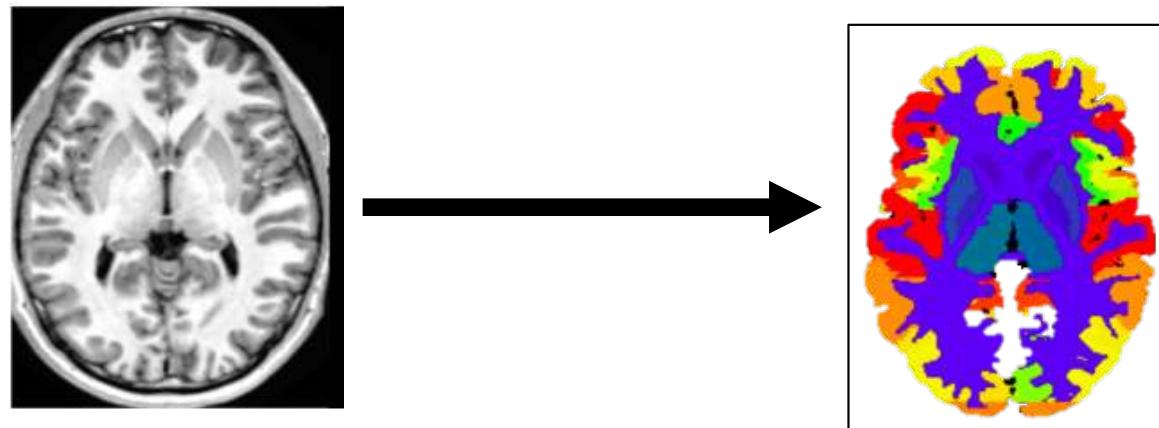
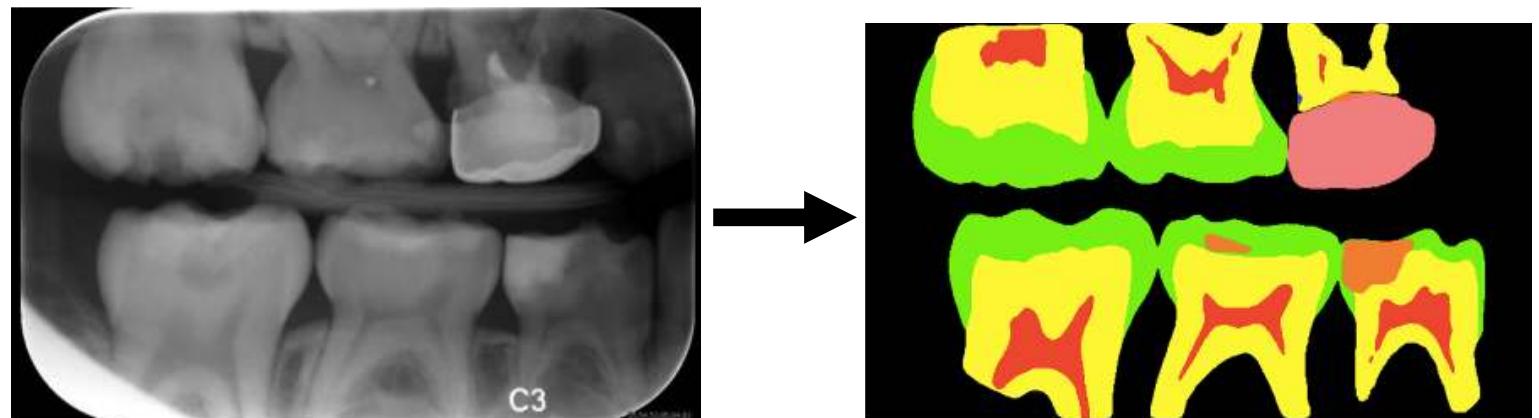
# some failures



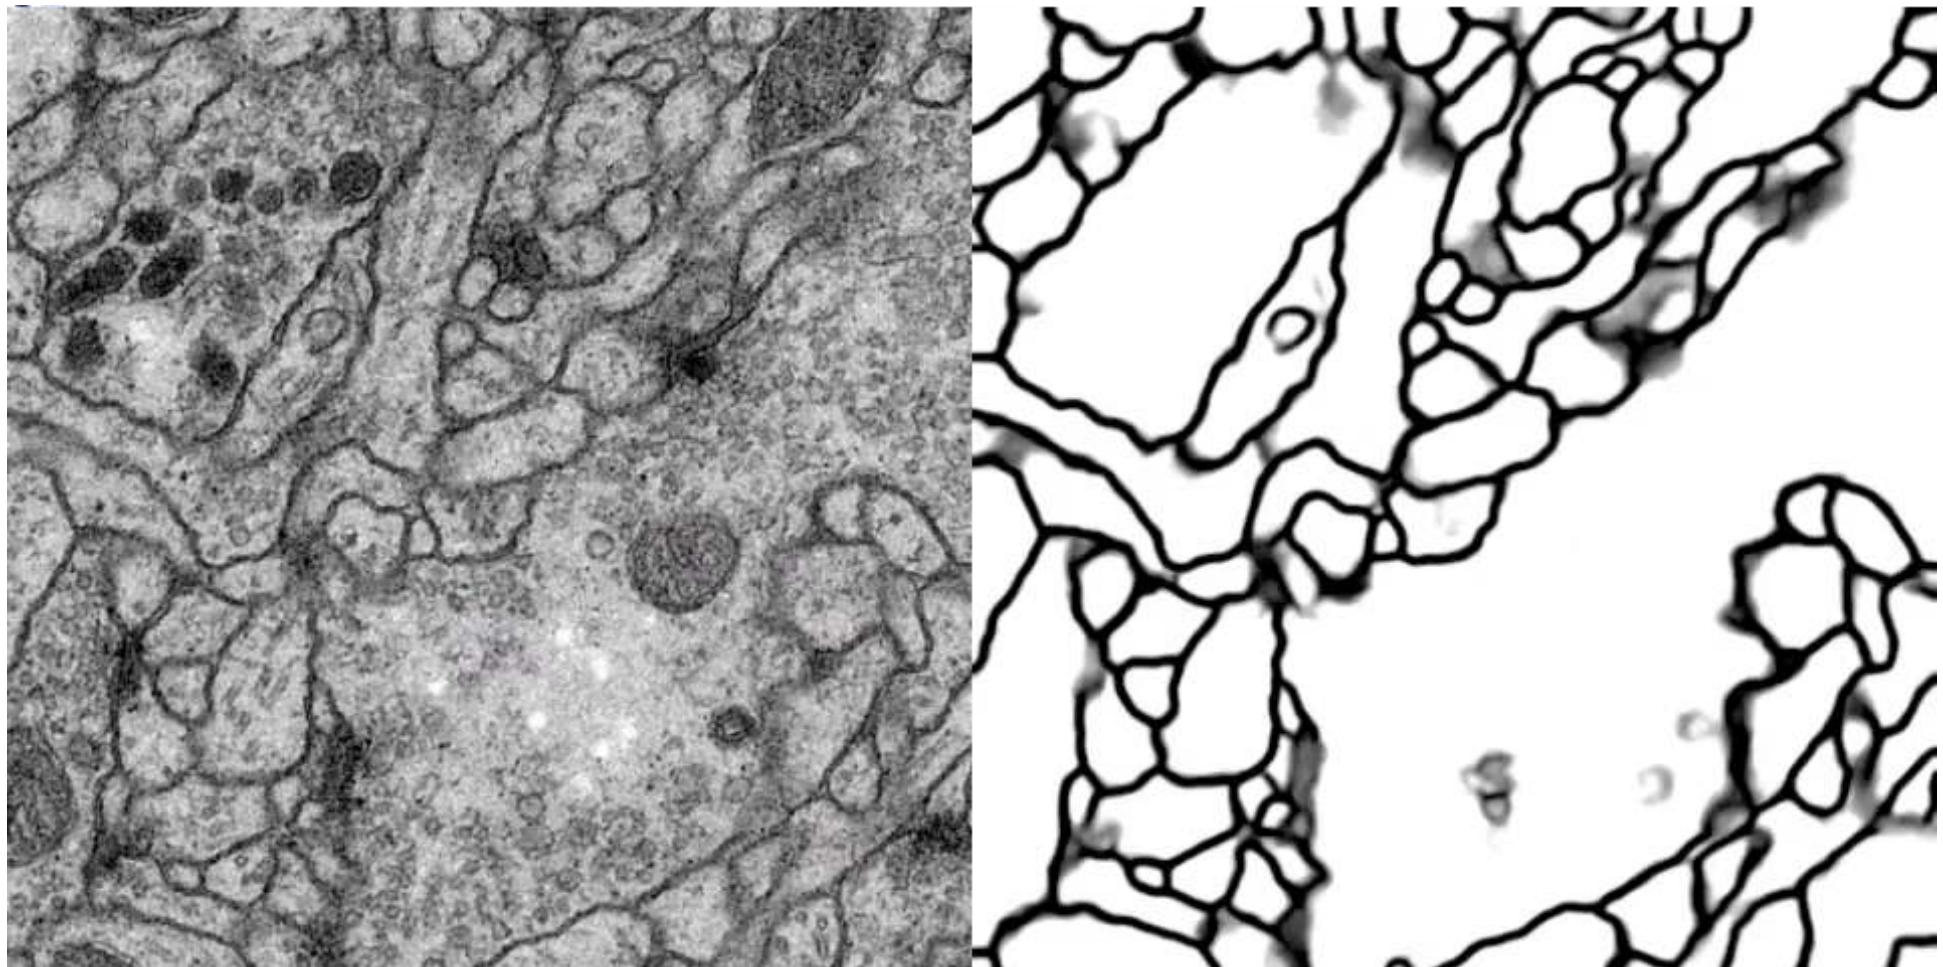
# IMAGE SEGMENTATION

# image segmentation

these can be  
seen as 1  
classification  
problem for  
each pixel



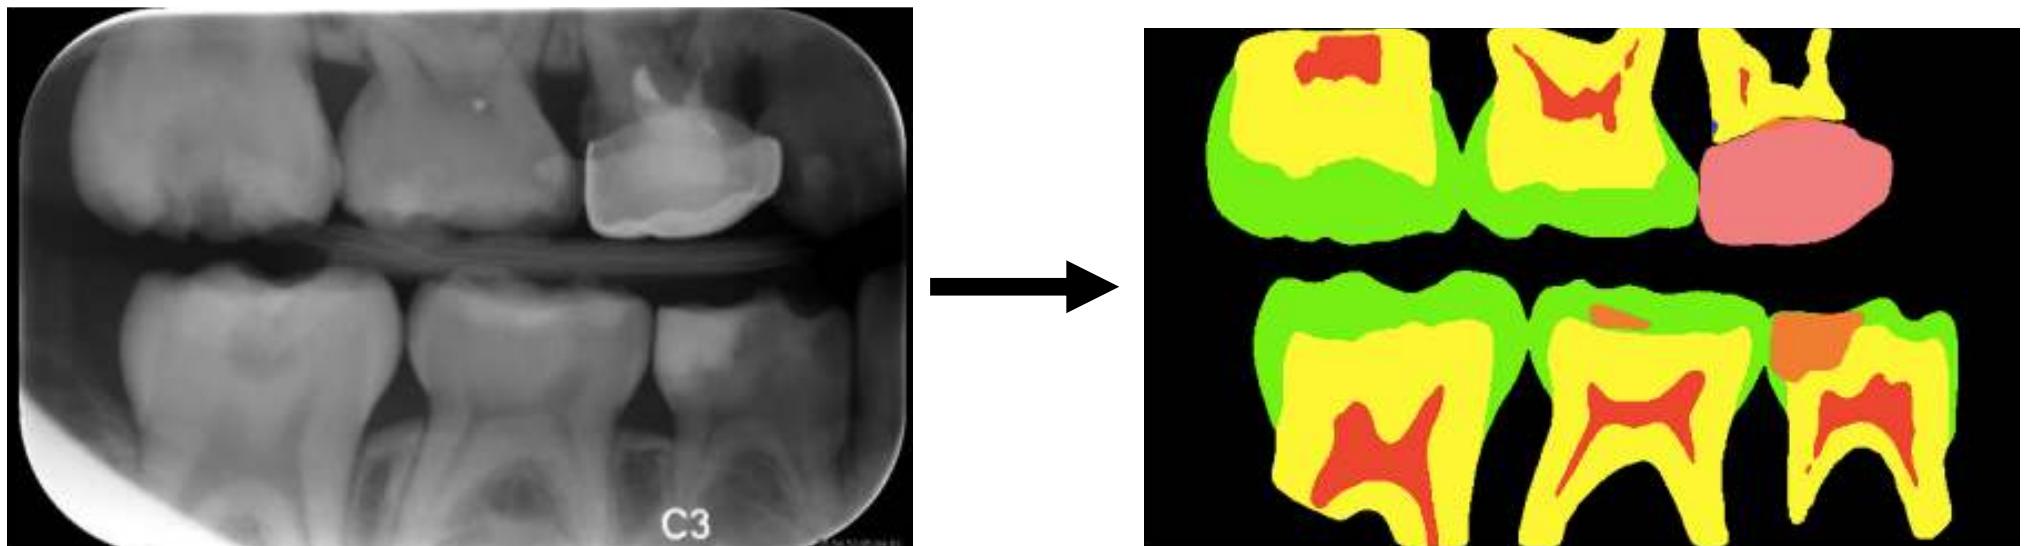
# U-Net: Results



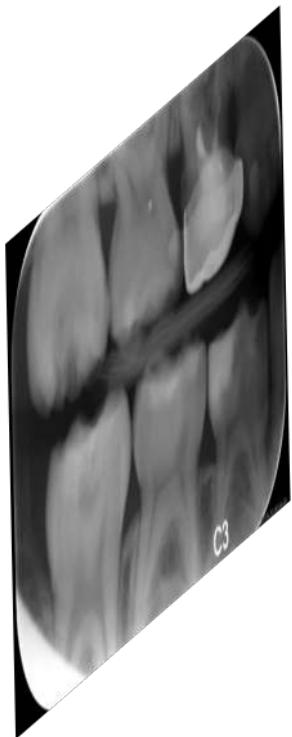
Input image

Our result: 0.000353 warping error  
(**New best score** at submission march 6th, 2015)  
Sliding-window CNN: 0.000420  
Training time: 10h, Application: 1s per image

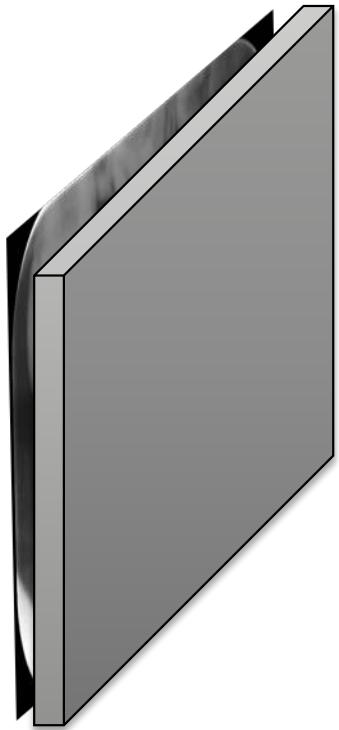
# image segmentation



# U-Net: Architecture

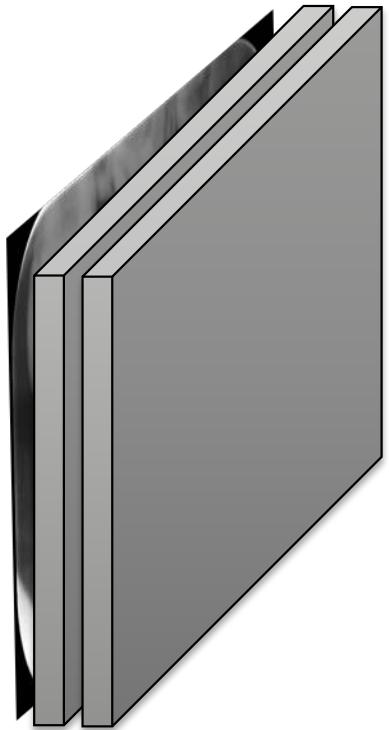


# U-Net: Architecture



$$\mathbf{h}_1 = [g(\mathbf{f}_{1,1} * \mathbf{x}), \dots, g(\mathbf{f}_{1,m} * \mathbf{x})]$$

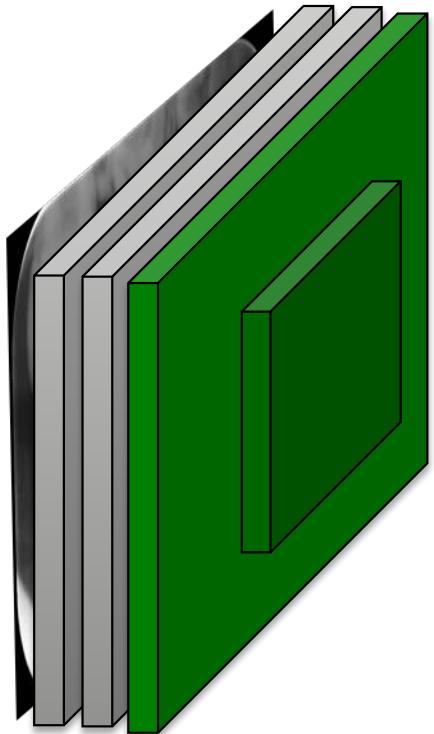
# U-Net: Architecture



$$\mathbf{h}_1 = [g(\mathbf{f}_{1,1} * \mathbf{x}), \dots, g(\mathbf{f}_{1,m} * \mathbf{x})]$$

$$\mathbf{h}_2 = [g(\mathbf{f}_{2,1} * \mathbf{h}_1), \dots, g(\mathbf{f}_{2,m_2} * \mathbf{h}_1)]$$

# U-Net: Architecture

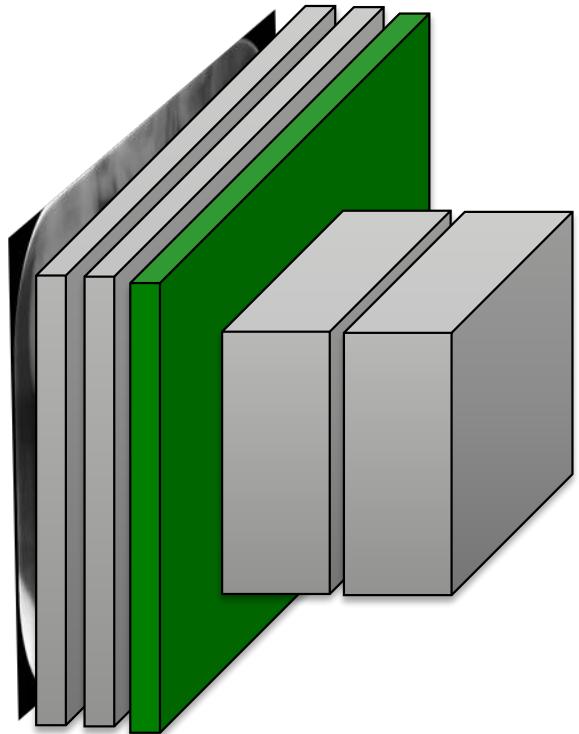


$$\mathbf{h}_1 = [g(\mathbf{f}_{1,1} * \mathbf{x}), \dots, g(\mathbf{f}_{1,m} * \mathbf{x})]$$

$$\mathbf{h}_2 = [g(\mathbf{f}_{2,1} * \mathbf{h}_1), \dots, g(\mathbf{f}_{2,m_2} * \mathbf{h}_1)]$$

$$\mathbf{h}_3 = \text{pooling}(\mathbf{h}_2)$$

# U-Net: Architecture



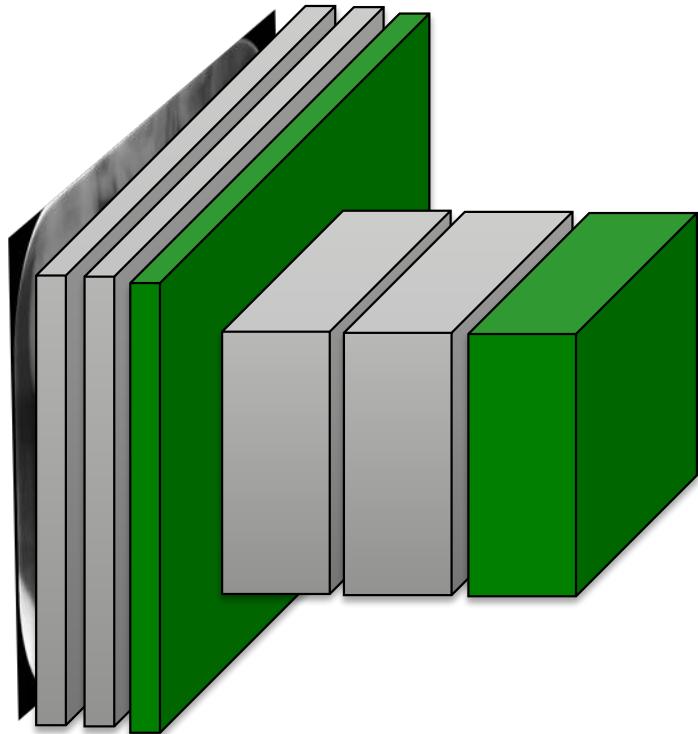
$$\mathbf{h}_1 = [g(\mathbf{f}_{1,1} * \mathbf{x}), \dots, g(\mathbf{f}_{1,m} * \mathbf{x})]$$

$$\mathbf{h}_2 = [g(\mathbf{f}_{2,1} * \mathbf{h}_1), \dots, g(\mathbf{f}_{2,m_2} * \mathbf{h}_1)]$$

$$\mathbf{h}_3 = \text{pooling}(\mathbf{h}_2)$$

$$\mathbf{h}_4 = [g(\mathbf{f}_{4,1} * \mathbf{h}_3), \dots, g(\mathbf{f}_{4,m_4} * \mathbf{h}_3)]$$

# U-Net: Architecture



$$\mathbf{h}_1 = [g(\mathbf{f}_{1,1} * \mathbf{x}), \dots, g(\mathbf{f}_{1,m} * \mathbf{x})]$$

$$\mathbf{h}_2 = [g(\mathbf{f}_{2,1} * \mathbf{h}_1), \dots, g(\mathbf{f}_{2,m_2} * \mathbf{h}_1)]$$

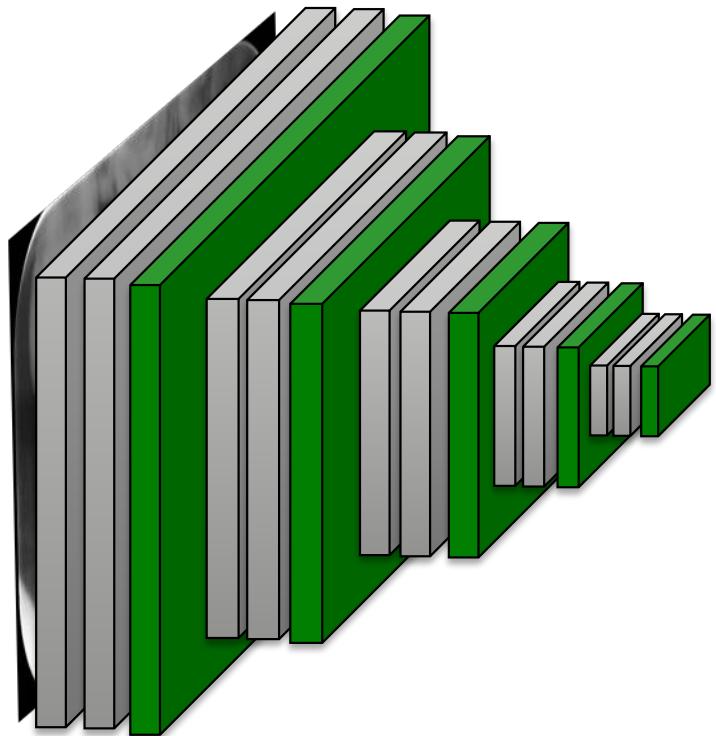
$\mathbf{h}_3 = \text{pooling}(\mathbf{h}_2)$

$$\mathbf{h}_4 = [g(\mathbf{f}_{4,1} * \mathbf{h}_3), \dots, g(\mathbf{f}_{4,m_4} * \mathbf{h}_3)]$$

$$\mathbf{h}_5 = [g(\mathbf{f}_{5,1} * \mathbf{h}_4), \dots, g(\mathbf{f}_{5,m_5} * \mathbf{h}_4)]$$

$\mathbf{h}_6 = \text{pooling}(\mathbf{h}_5)$

# U-Net: Architecture

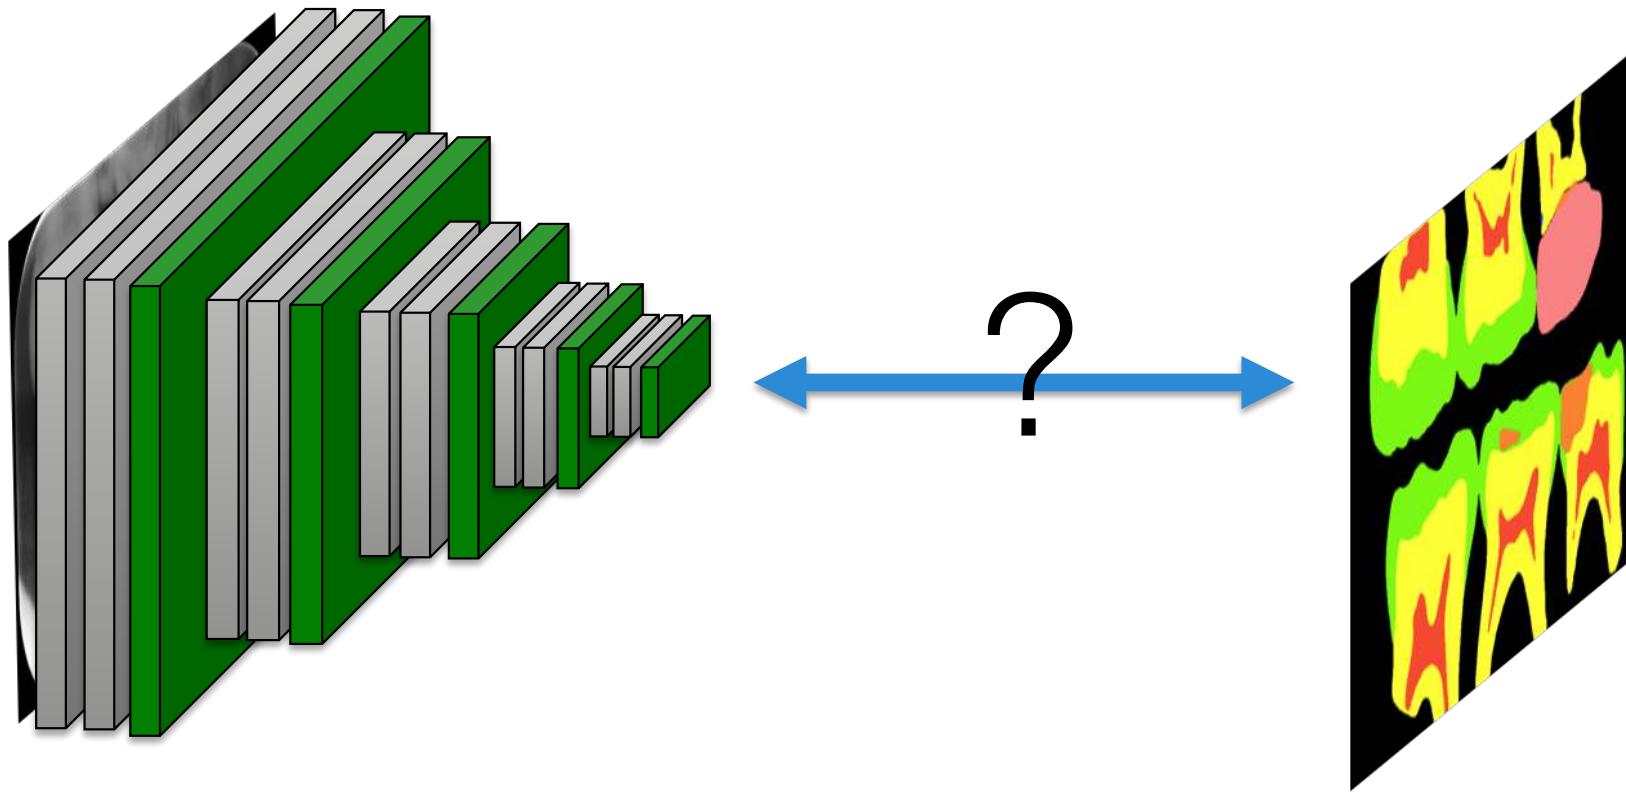


$$\mathbf{h}_{13} = [g(\mathbf{f}_{13,1} * \mathbf{h}_{12}), \dots, g(\mathbf{f}_{13,m_{13}} * \mathbf{h}_{12})]$$

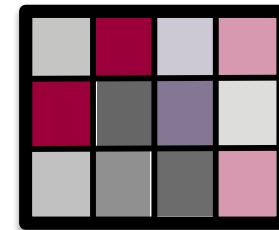
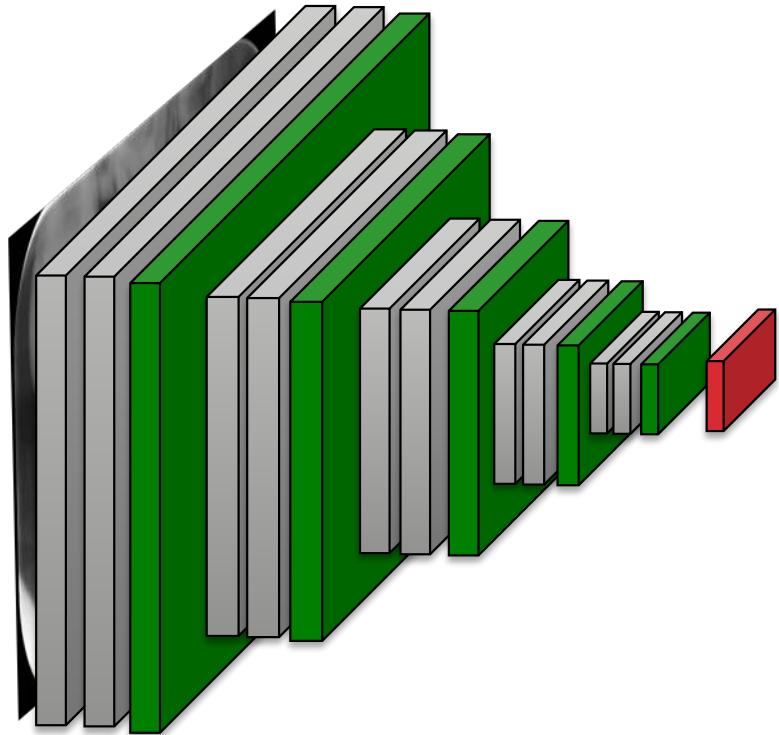
$$\mathbf{h}_{14} = [g(\mathbf{f}_{14,1} * \mathbf{h}_{13}), \dots, g(\mathbf{f}_{14,m_{14}} * \mathbf{h}_{13})]$$

$$\mathbf{h}_{15} = \text{pooling}(\mathbf{h}_{14})$$

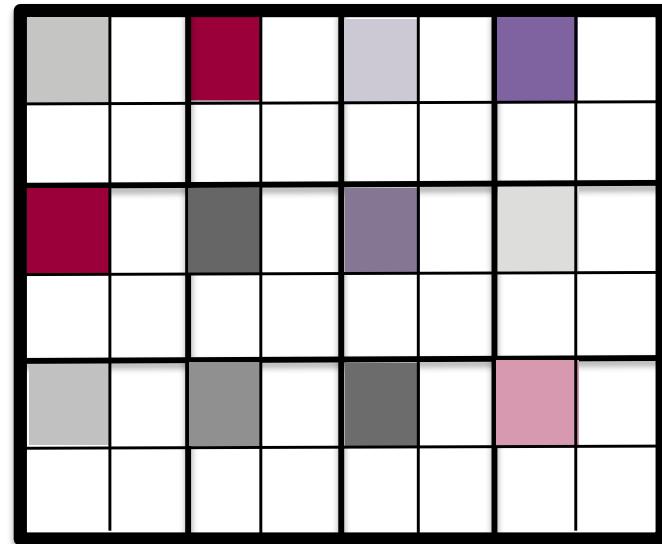
# U-Net: Architecture



# U-Net: Architecture

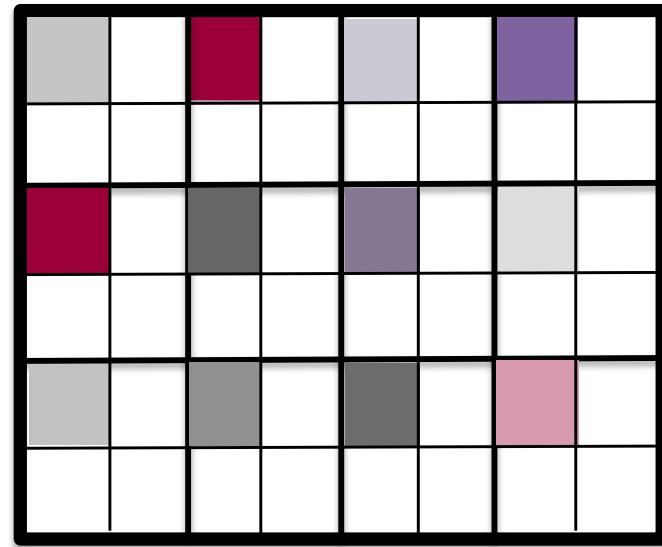
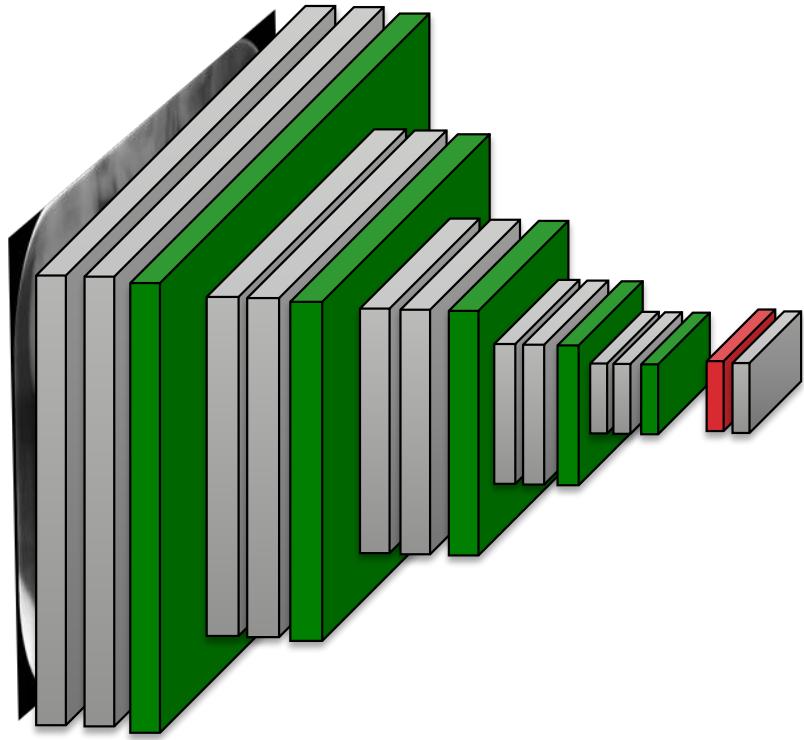


UpSampling



$$\mathbf{h}_{16} = \text{UpSampling}(\mathbf{h}_{15})$$

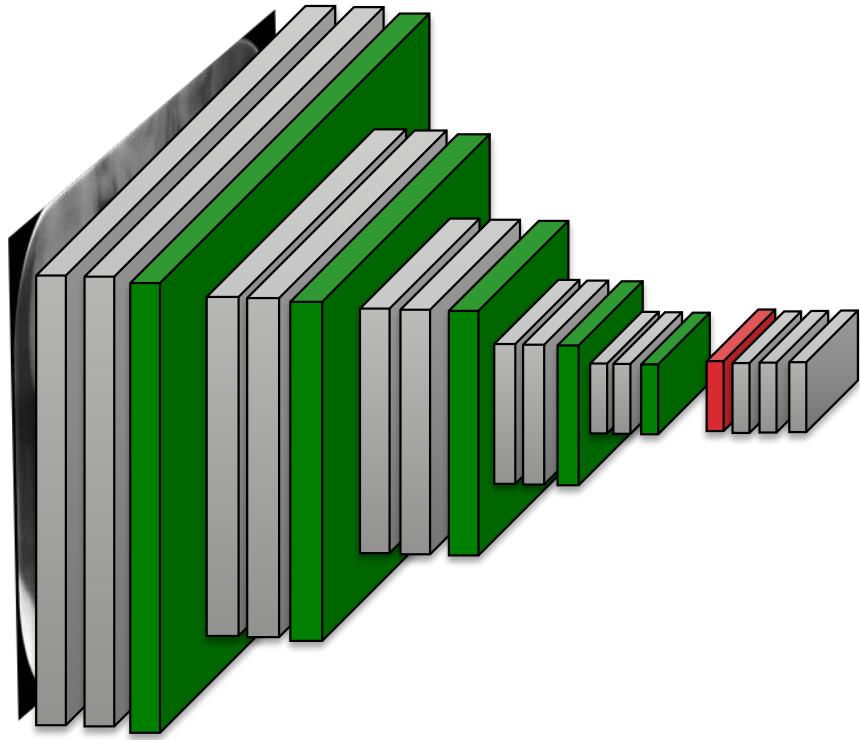
# U-Net: Architecture



$$\mathbf{h}_{16} = \text{UpSampling}(\mathbf{h}_{15})$$

$$\mathbf{h}_{17} = [g(\mathbf{f}_{17,1} * \mathbf{h}_{16}), \dots, g(\mathbf{f}_{17,m_{17}} * \mathbf{h}_{16})]$$

# U-Net: Architecture

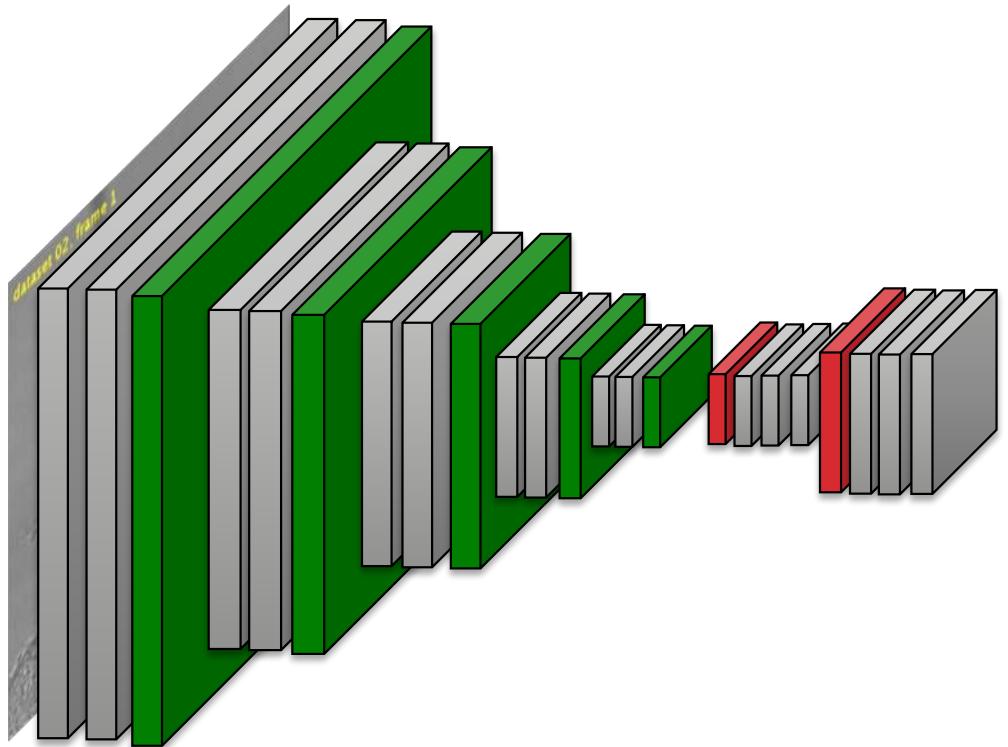


$$\mathbf{h}_{16} = \text{UpSampling}(\mathbf{h}_{15})$$

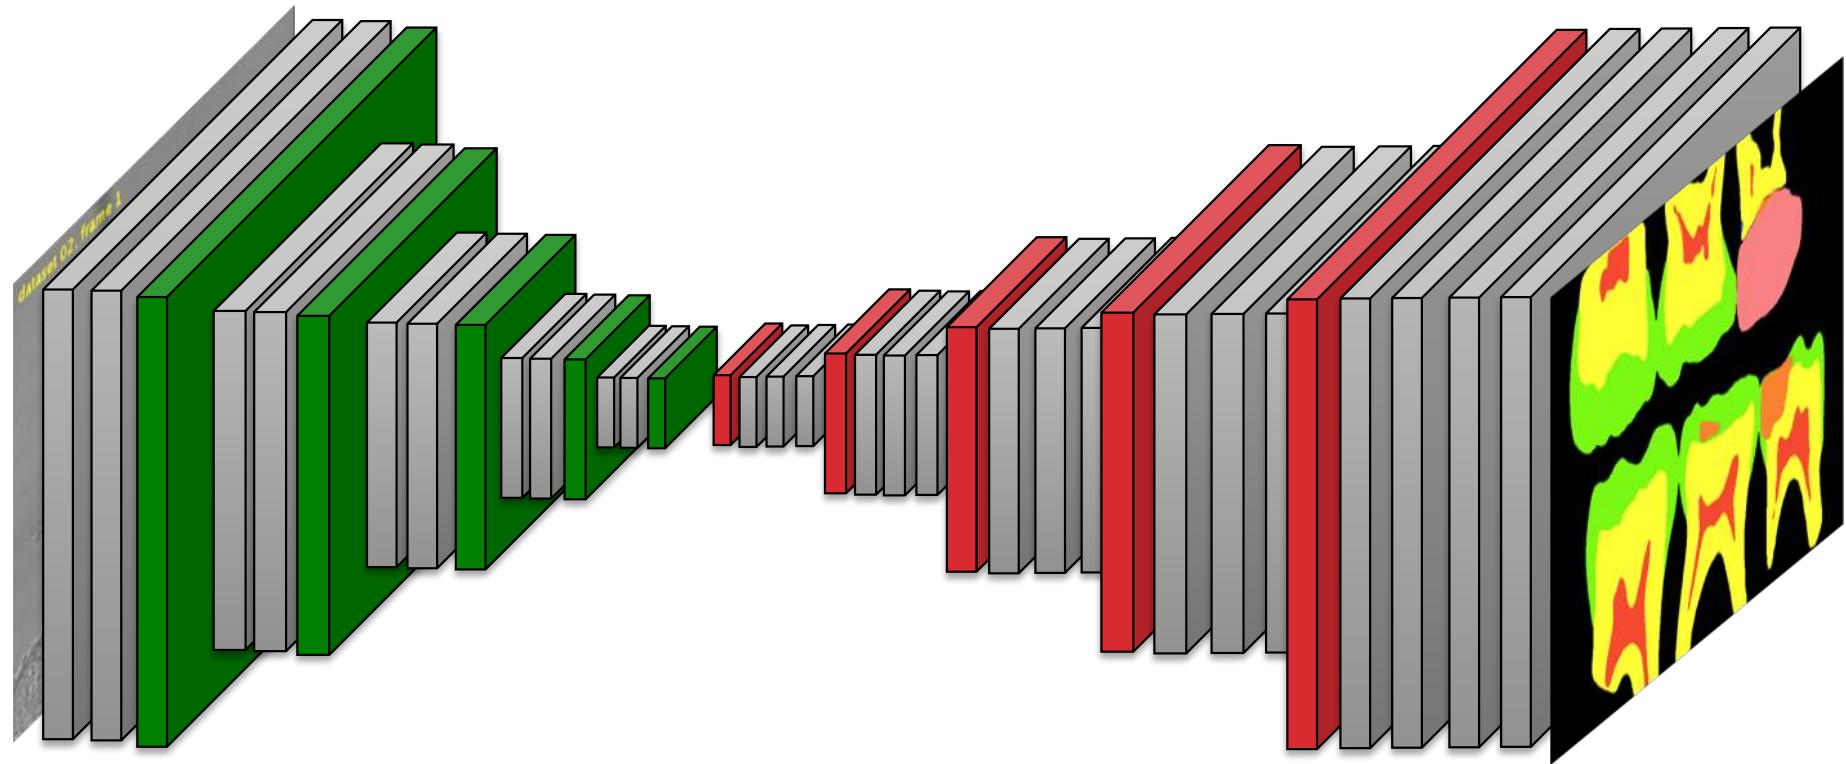
$$\mathbf{h}_{17} = [g(\mathbf{f}_{17,1} * \mathbf{h}_{16}), \dots, g(\mathbf{f}_{17,m_{17}} * \mathbf{h}_{16})]$$

$$\mathbf{h}_{18} = [g(\mathbf{f}_{18,1} * \mathbf{h}_{17}), \dots, g(\mathbf{f}_{18,m_{18}} * \mathbf{h}_{17})]$$

# U-Net: Architecture



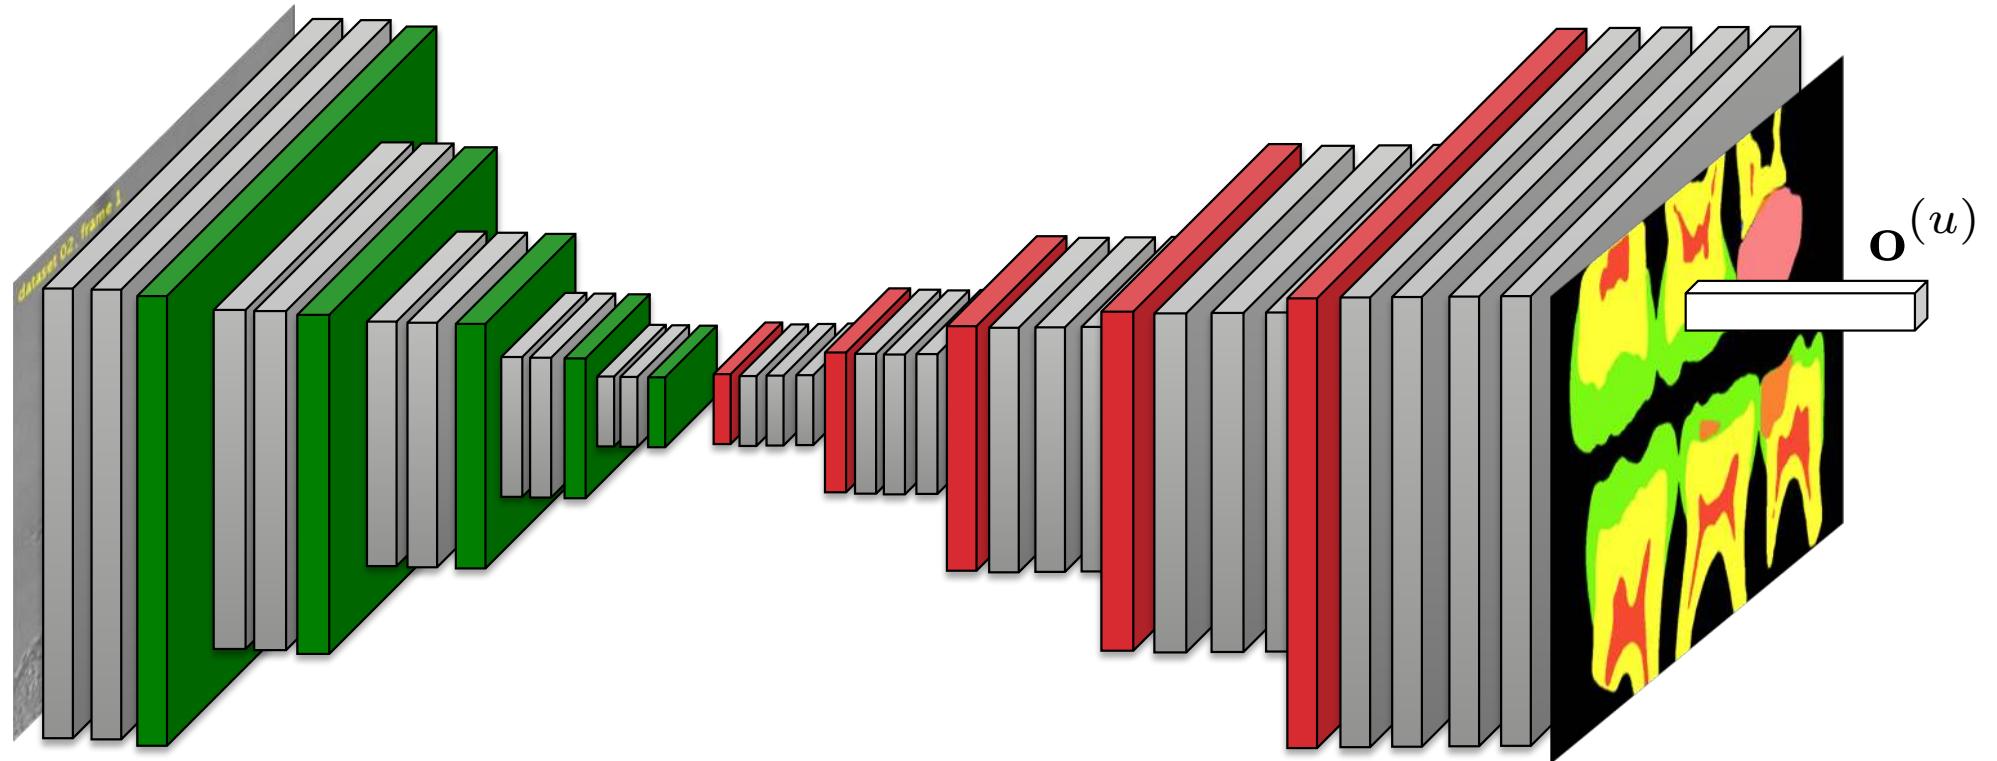
# U-Net: Architecture



What is the output of the network exactly?



# U-Net: Architecture



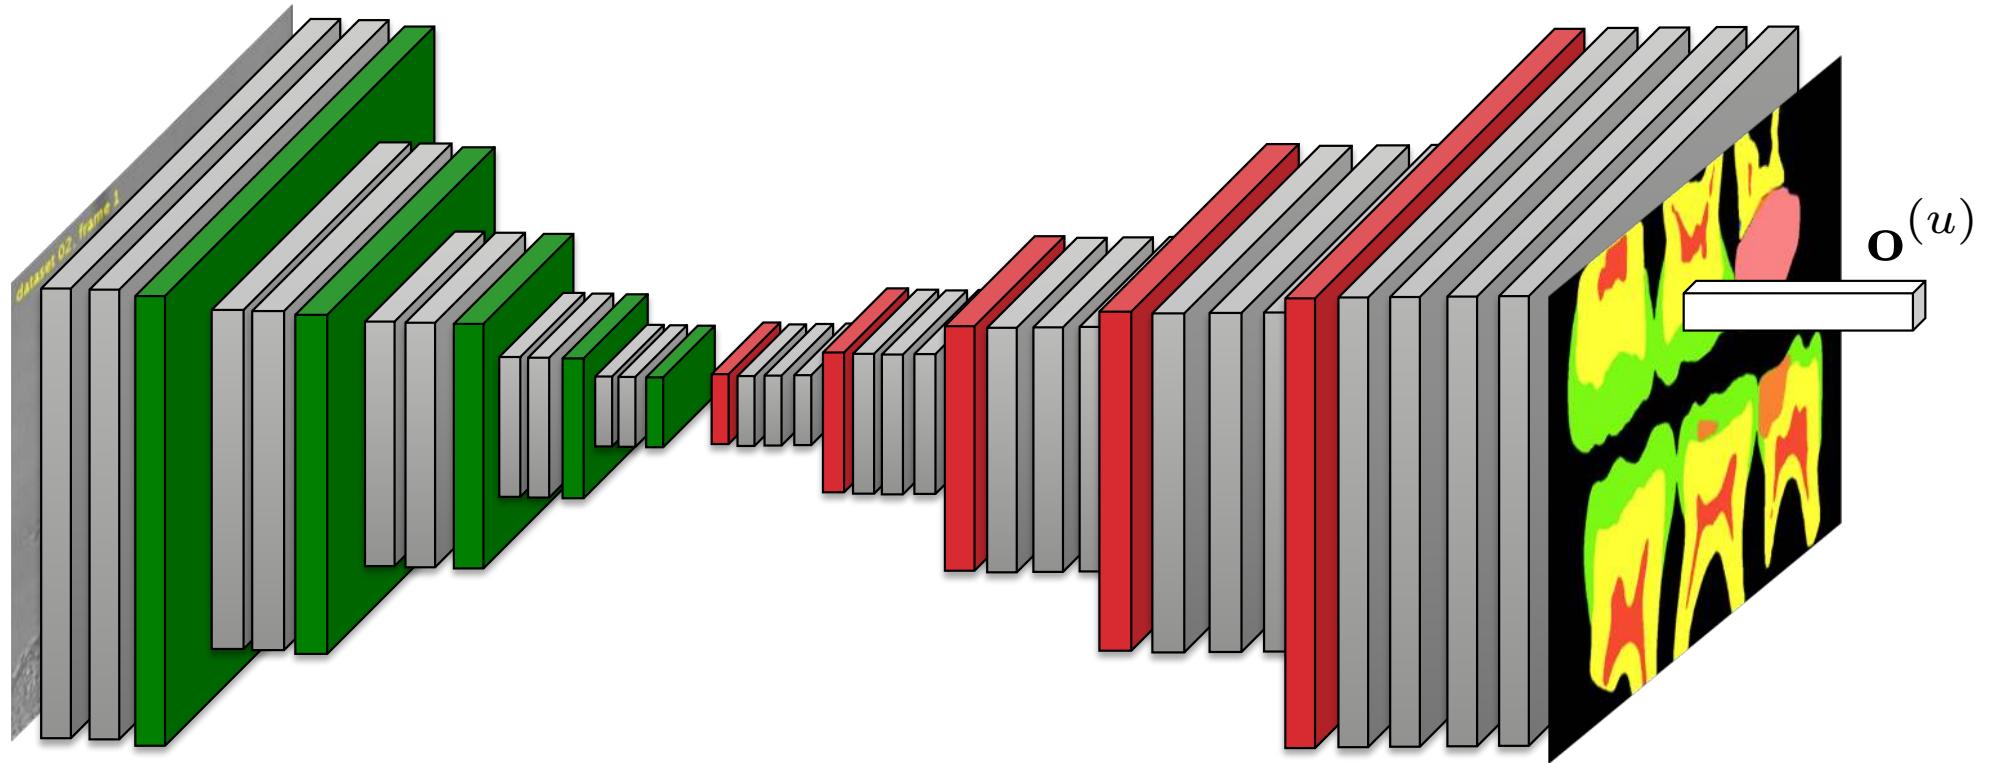
What is the output of the network exactly?

For each pixel, we have 6 possible classes.

$\mathbf{o}^{(u)}$ , where  $u$  is a pixel, is a vector of 6 values.



# U-Net: Loss Function



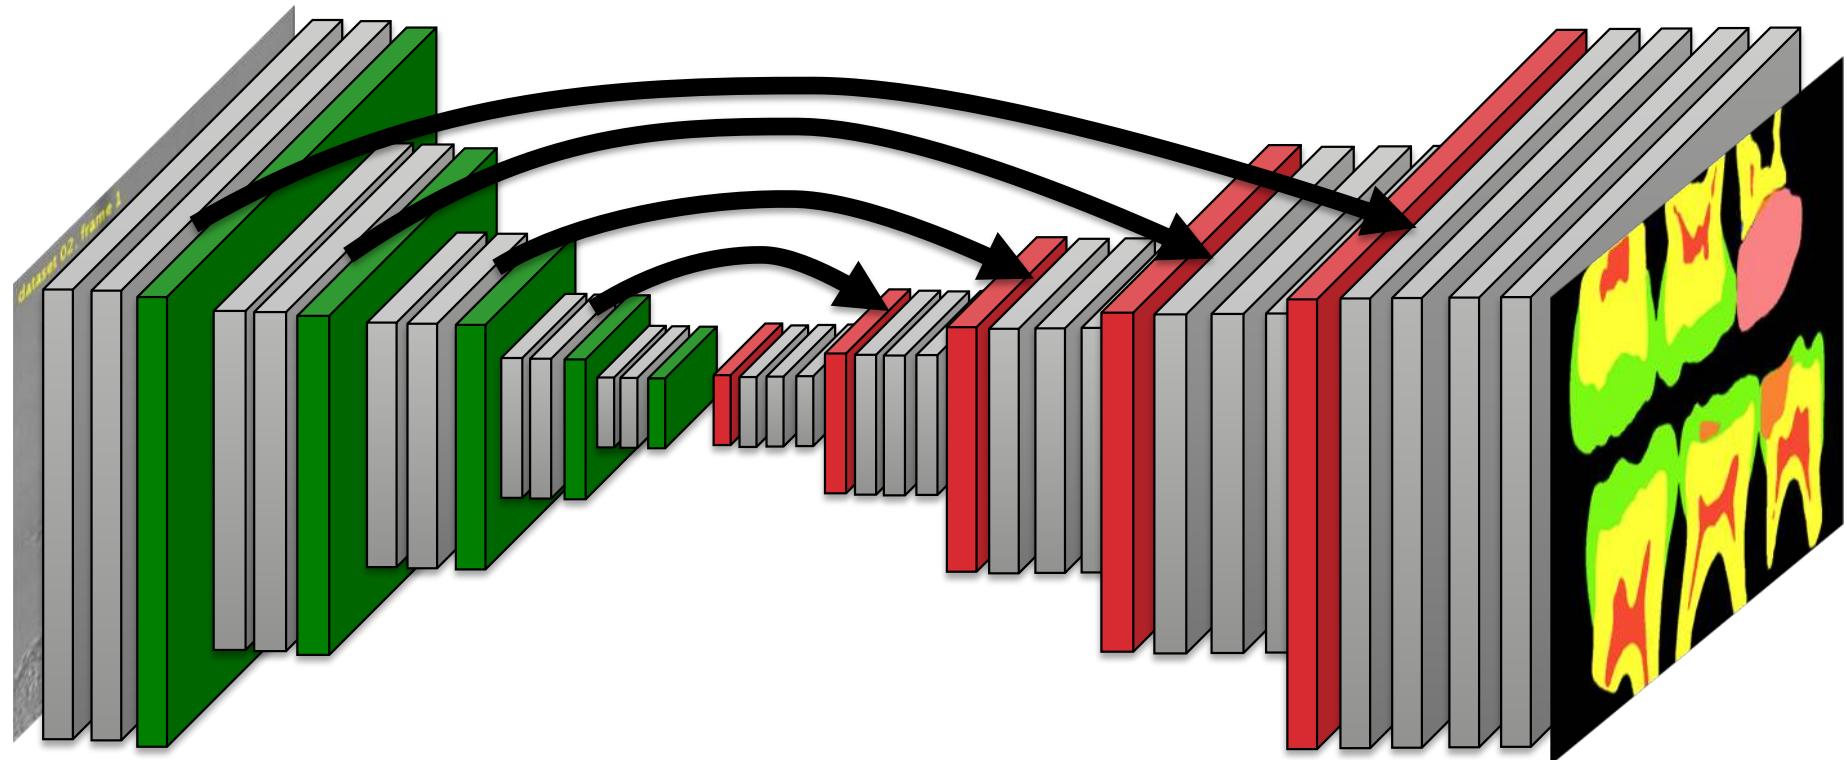
$$\mathcal{L} = - \sum_{(\mathbf{x}, \mathbf{d}) \in \mathcal{T}} \sum_u \log y(\mathbf{x}, \mathbf{d}, u)$$

with  $y(\mathbf{x}, \mathbf{d}, u) = \text{softmax}(\mathbf{o}^{(u)})_{\mathbf{d}^{(u)}}$

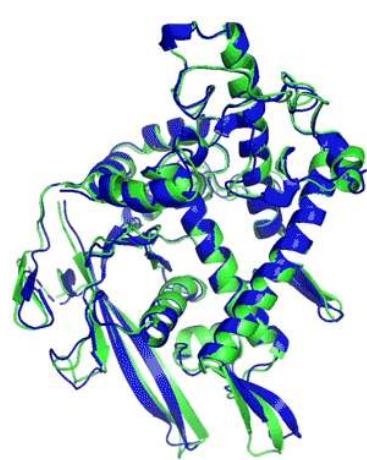
where

- ▷  $\mathbf{o}^{(u)}$  is the network output for image  $\mathbf{x}$  for pixel  $u$ , and
- ▷  $\mathbf{d}^{(u)}$  is the desired class for pixel  $u$  in image  $\mathbf{x}$  for pixel  $u$

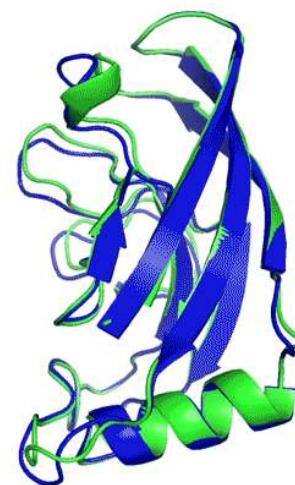
# U-Net: Skip Connections



# ALPHA FOLD



T1O37 / 6vr4  
90.7 GDT  
(RNA polymerase domain)



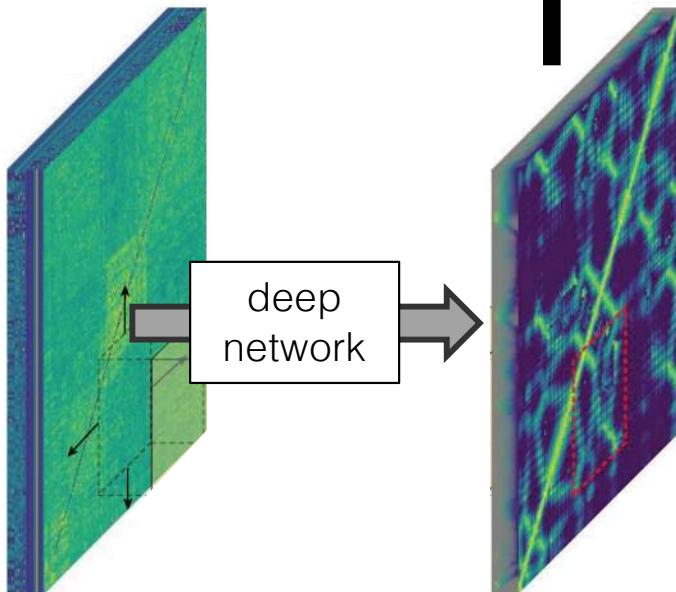
T1O49 / 6y4f  
93.3 GDT  
(adhesin tip)

- Experimental result
- Computational prediction

Protein  
sequence and  
MSA features



covariance  
matrix of the  
MSA features



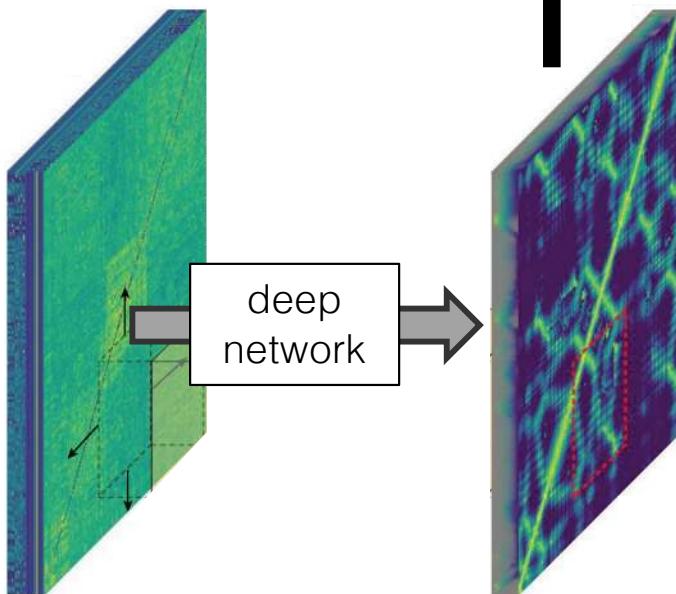
the predicted  
distagram is  
used to define a  
potential: a  
protein structure  
with a similar  
distagram has a  
low potential



'distagram': distribution over  
the distances and torsions  
between every pair of  
residues

Protein  
sequence and  
MSA features

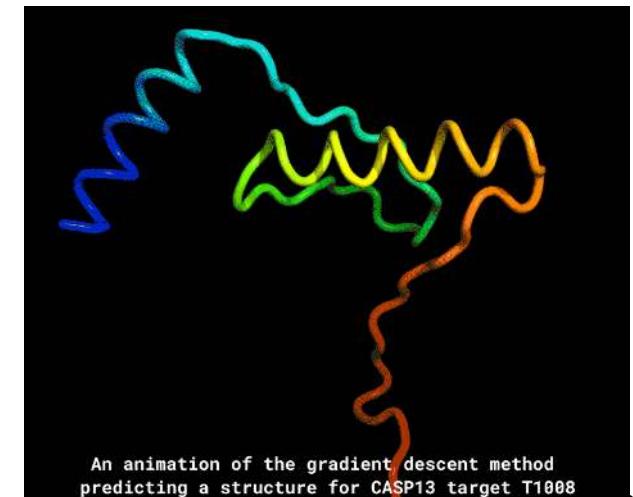
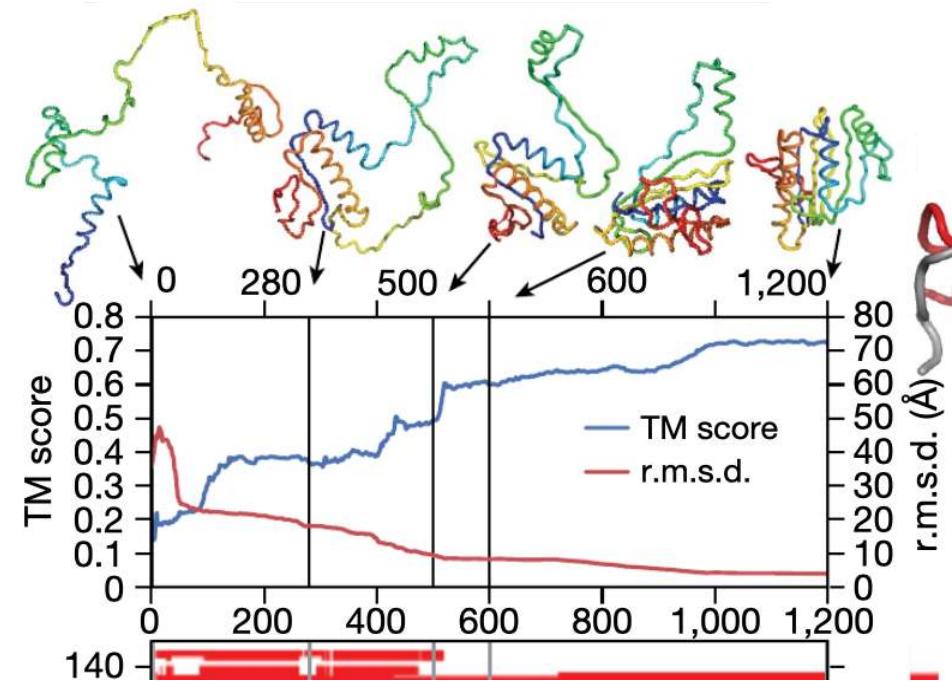
↓  
covariance  
matrix of the  
MSA features



'distogram': distribution over  
the distances and torsions  
between every pair of  
residues

the predicted  
distagram is  
used to define a  
potential: a  
protein structure  
with a similar  
distagram has a  
low potential

→  
start with a random 3d structure for the  
protein sequence and deform it to  
minimize its potential using gradient  
descent

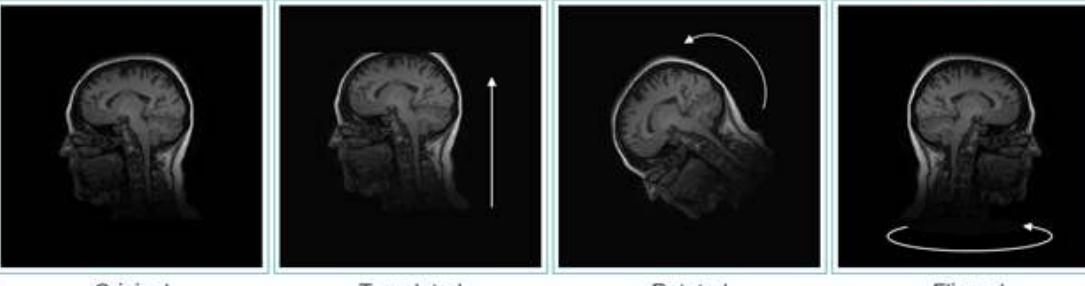


# WHAT TO DO WHEN ONLY FEW TRAINING DATA IS AVAILABLE

# Data Augmentation

**AUGMENTATION**  
4 rigid examples

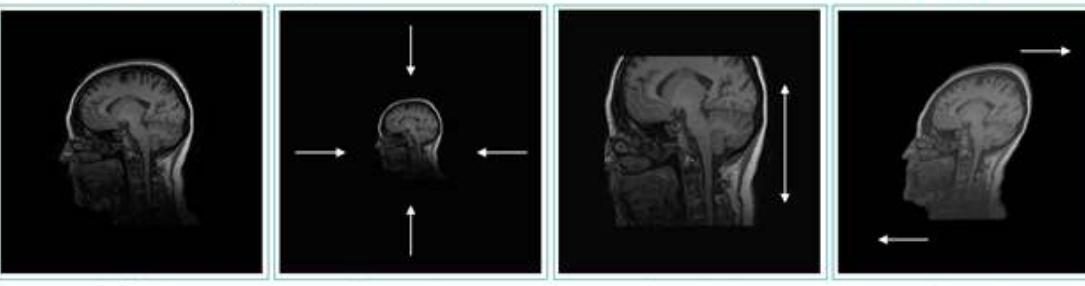
 Quantib



Original      Translated      Rotated      Flipped

**AUGMENTATION**  
3 stretch and sheering examples

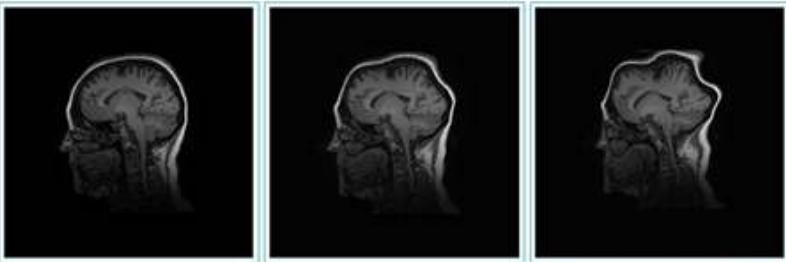
 Quantib



Original      Zoomed out      Stretched      Sheared

**AUGMENTATION**  
2 elastic deformation examples

 Quantib



Original      Elastic deformation      Too much elastic deformation

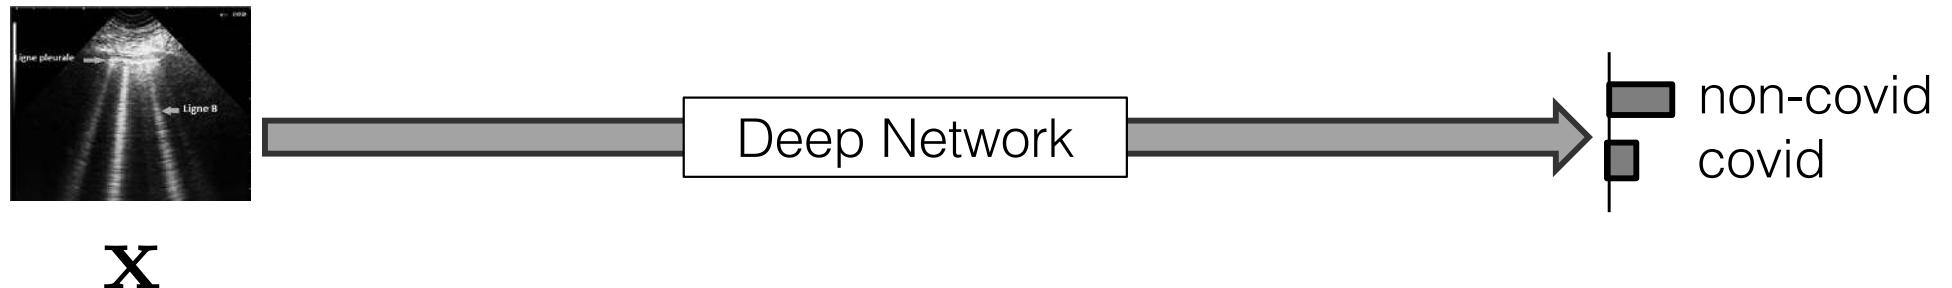
# transfer learning / domain transfer

Transfer learning:

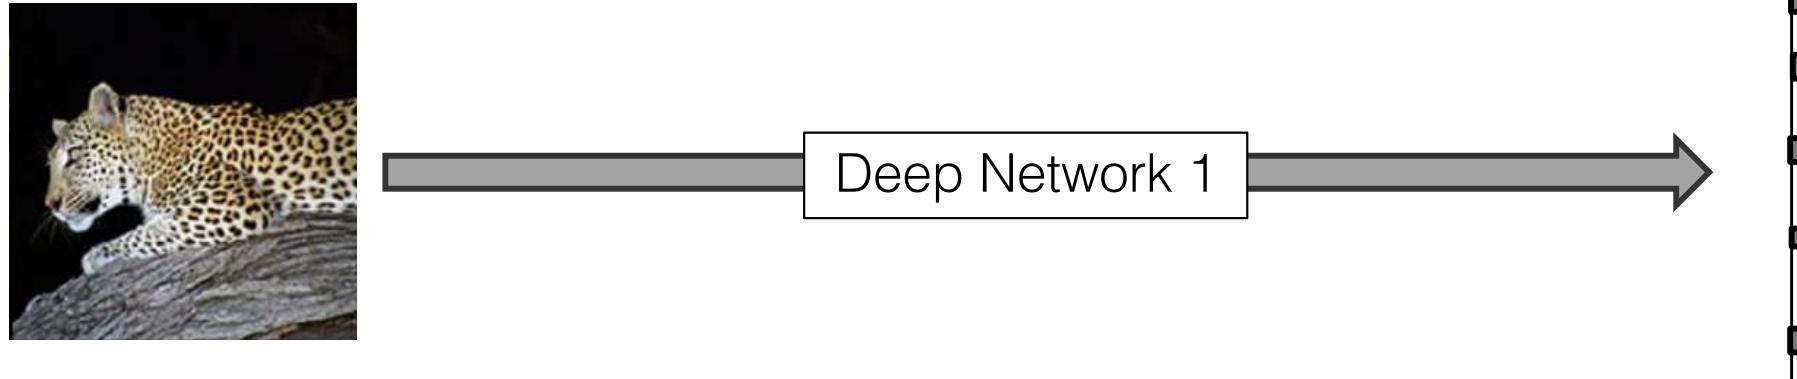
- we have few training data on our problem, but
- we have a lot of training data for a similar problem.

# transfer learning / domain transfer

A simple method for transfer learning:

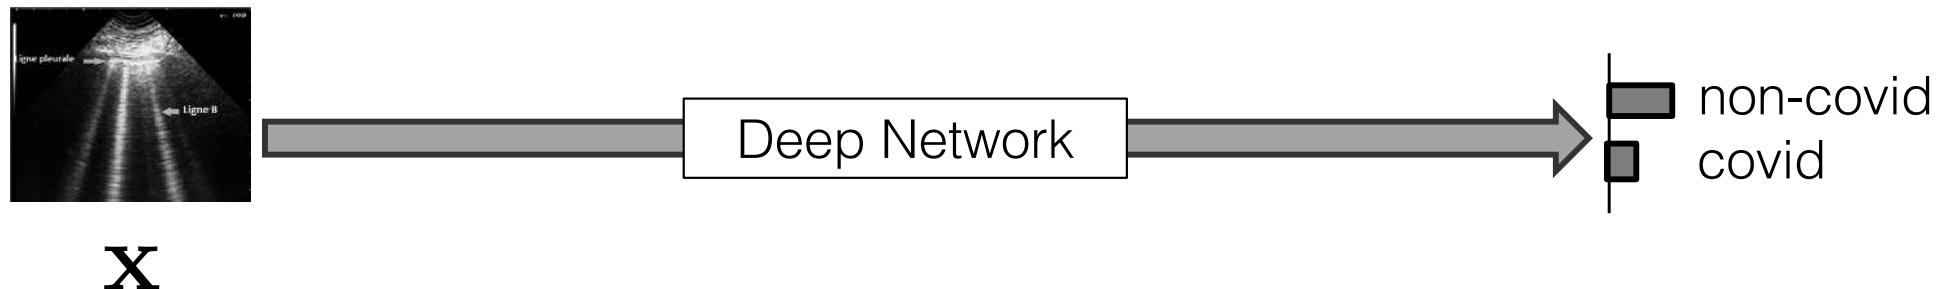


1. Training a deep network on a problem where a large amount of training data is available:

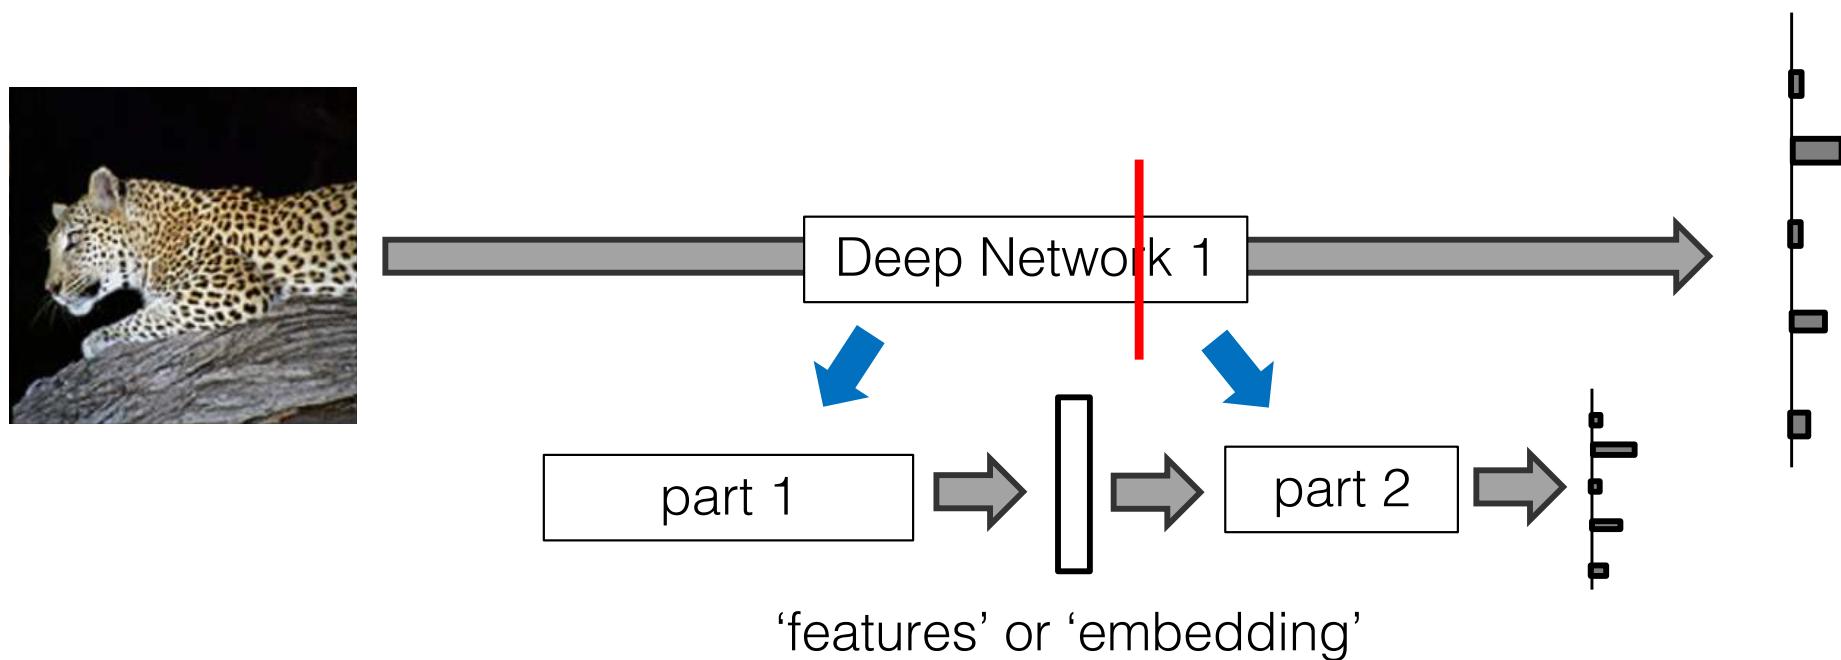


# transfer learning / domain transfer

A simple method for transfer learning:

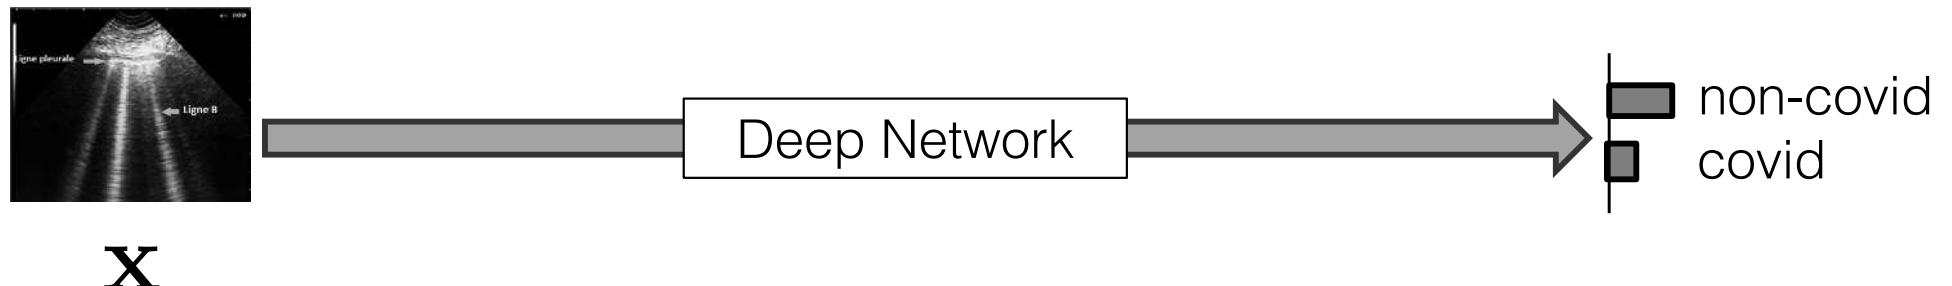


2. Cut this network into two parts (after training):

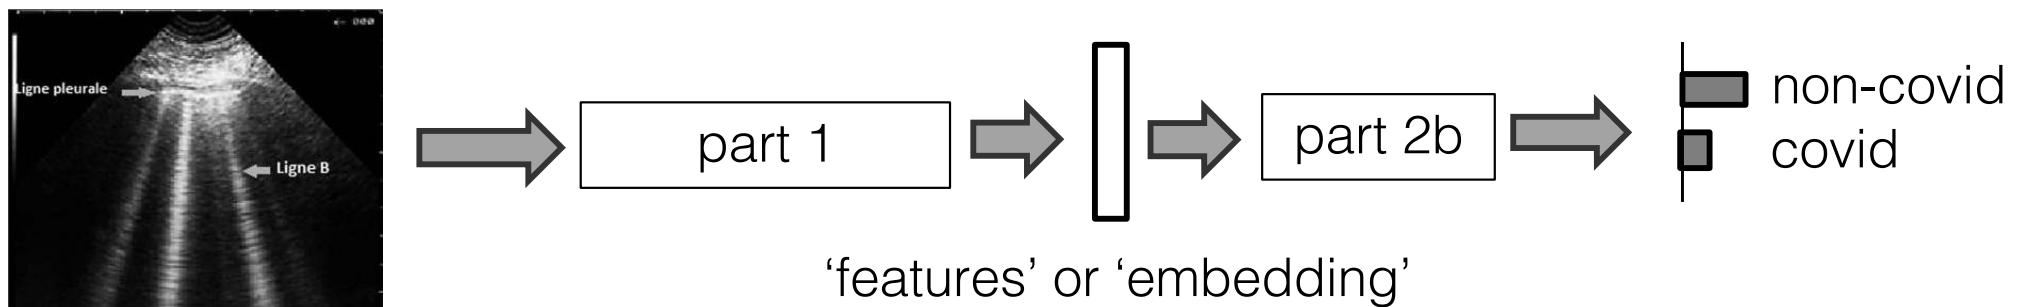


# transfer learning / domain transfer

A simple method for transfer learning:

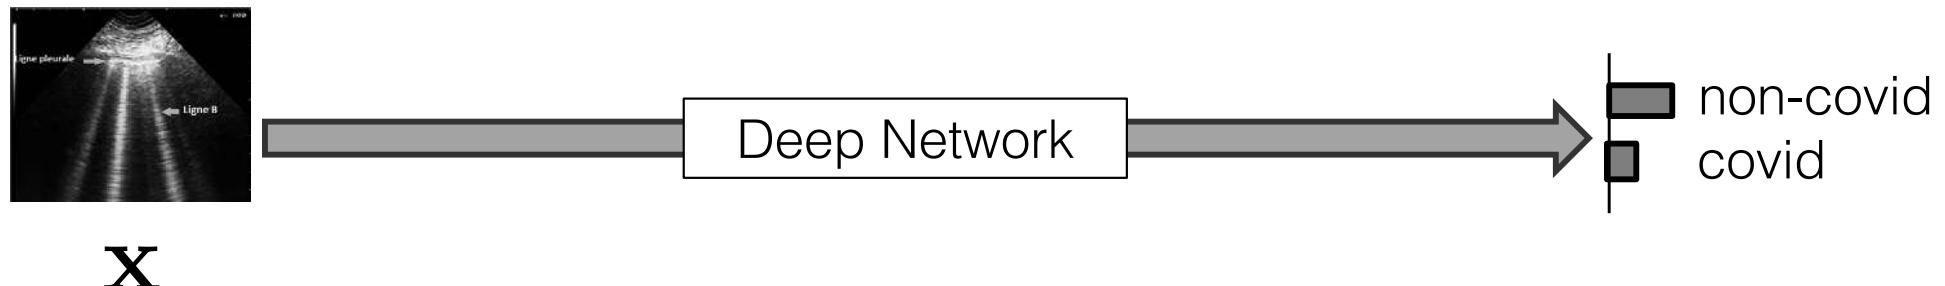


3. Keep the parameters of Part 1, initialize randomly Part 2b with the new number of classes



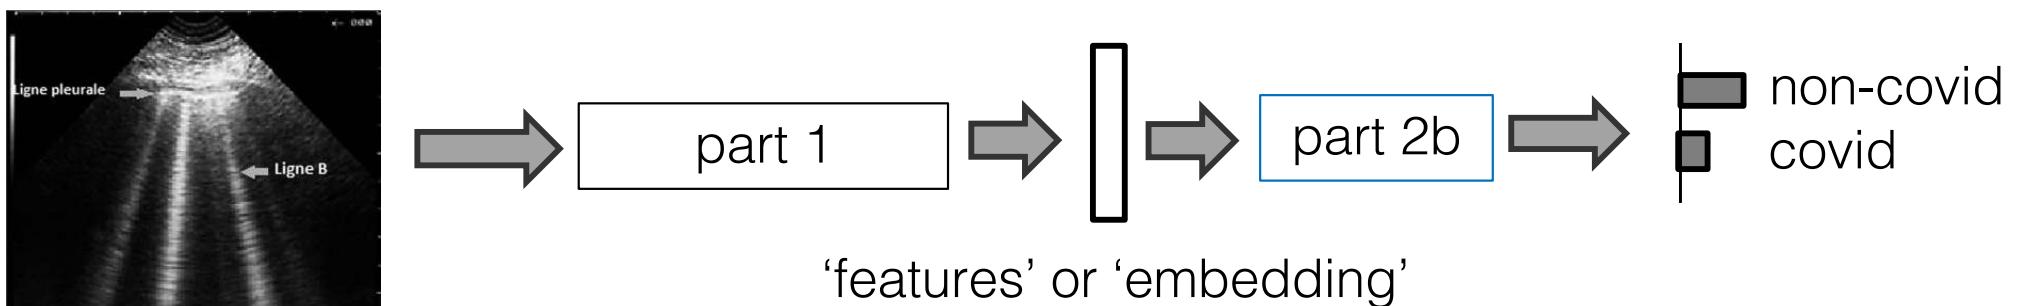
# transfer learning / domain transfer

A simple method for transfer learning:



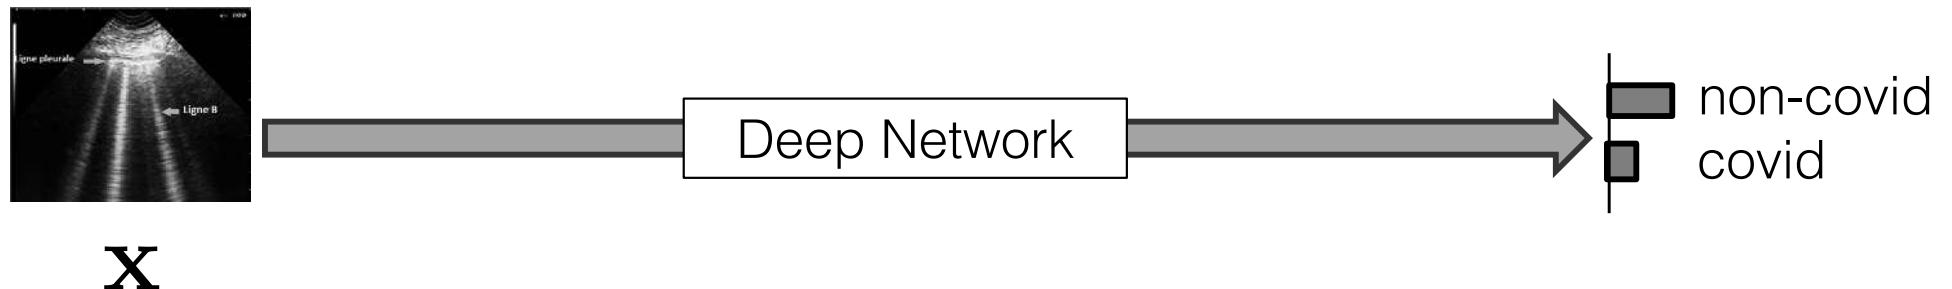
4. Keep the parameters of Part 1, optimize only the parameters of Part 2b on the available data

Alternatively, we can ‘fine-tune’ the parameters of Part 1.

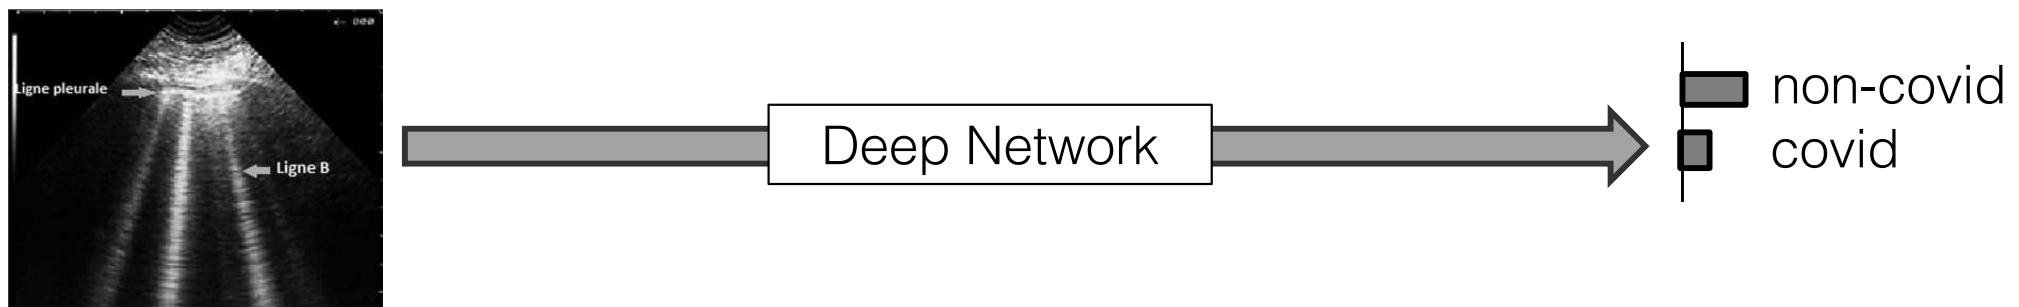


# transfer learning / domain transfer

A simple method for transfer learning:



Part 1 and Part 2b form a deep network:

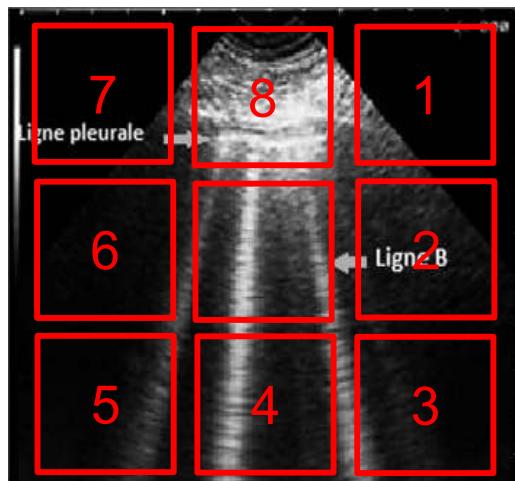


# self-learning

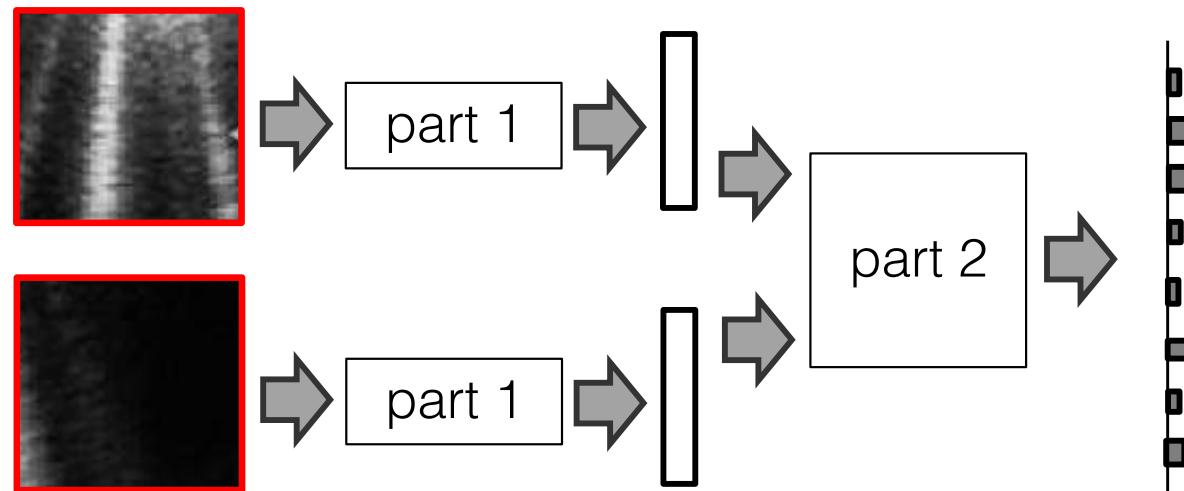
Self-learning for learning features.

Related to transfer learning but the first problem is ‘artificial’.

For example:



Given the center image, and one the 8 images, predict from where is taken this second image:



# TEXT PROCESSING

# TRANSLATION

Probleme kann man  
niemals mit deselben  
Denkweise lösen,  
durch die sie  
entstanden sind.



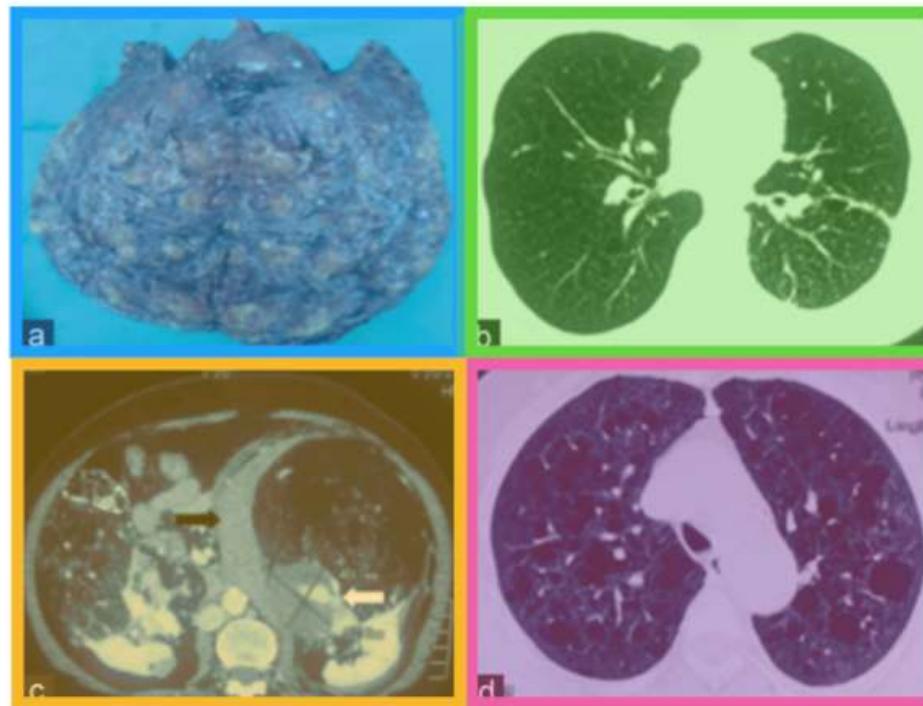
Problems can never  
be solved with the  
same way of  
thinking that  
caused them.

# IMAGE CAPTIONING



Deep Network →

The man at bat readies to swing at the pitch while the umpire looks on.



**Figure 1:** (a) Right renal angiomyolipoma (gross specimen postexcision). (b) High-resolution computed tomography chest images of Case 1 showing multiple variable sized cysts uniformly scattered in both lungs. (c) Computed tomography abdomen showing bilateral renal angiomyolipomas with fat densities, tortuous vessels, and pseudoaneurysm (white arrow). There is also the presence of perinephric hematoma (black arrow). (d) High-resolution computed tomography image of Case 2 showing bilateral lung cysts.

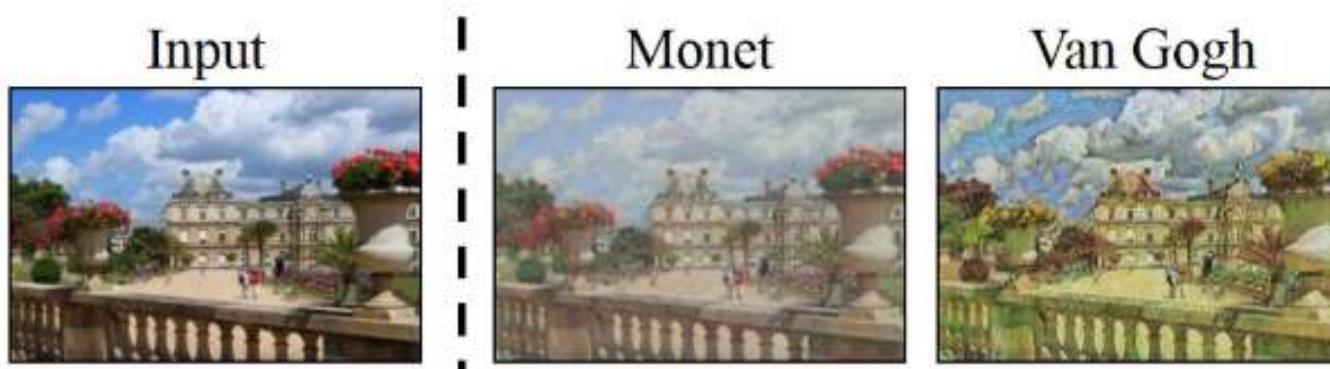
# GENERATIVE ADVERSARIAL NETWORKS (GANS)

# some applications of gans

Generating new images/data:



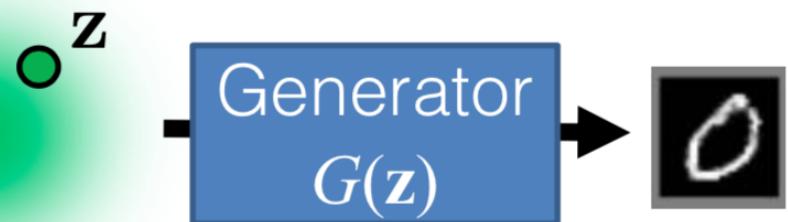
Style-transfer and Deep Fakes:



Generating potential drugs

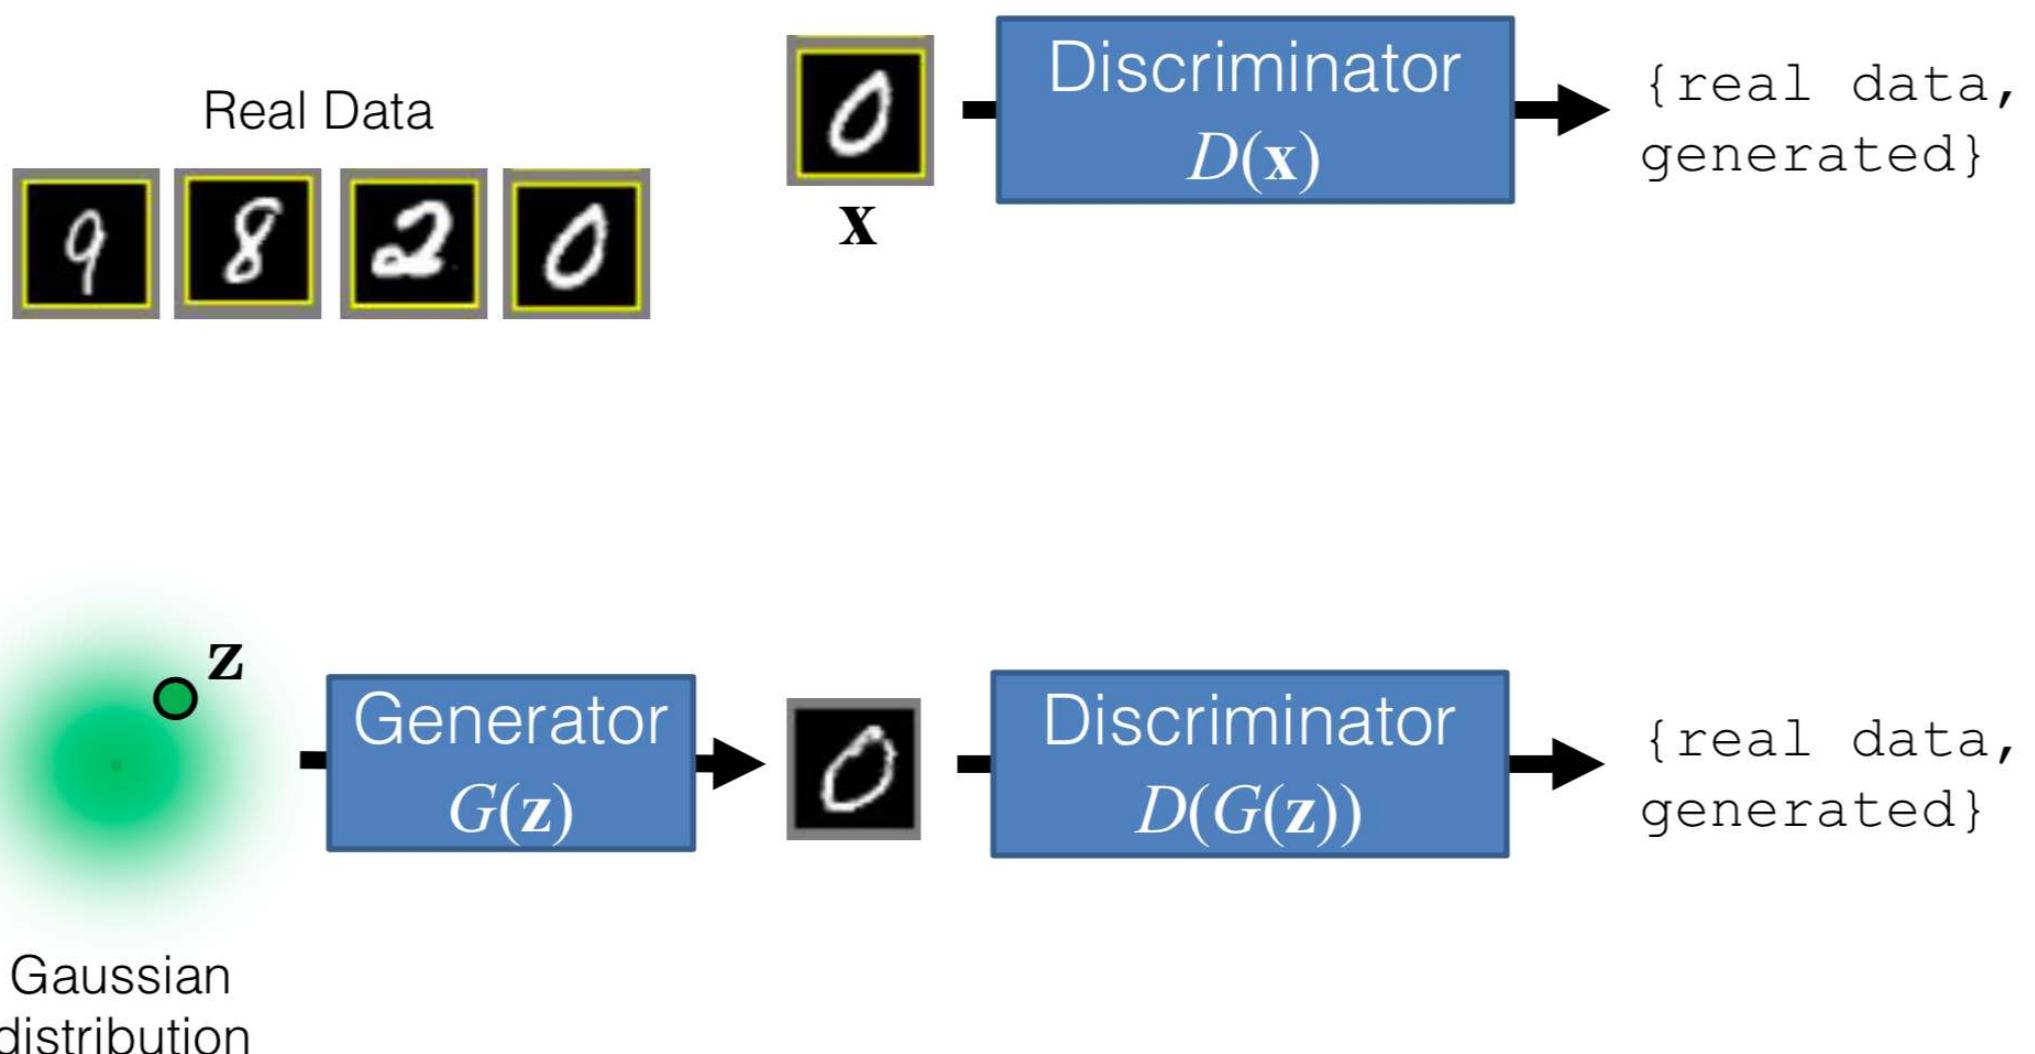
# GANs

We would like to train a network  $G$  to generate images of digits from noise vectors  $\mathbf{z}$ :



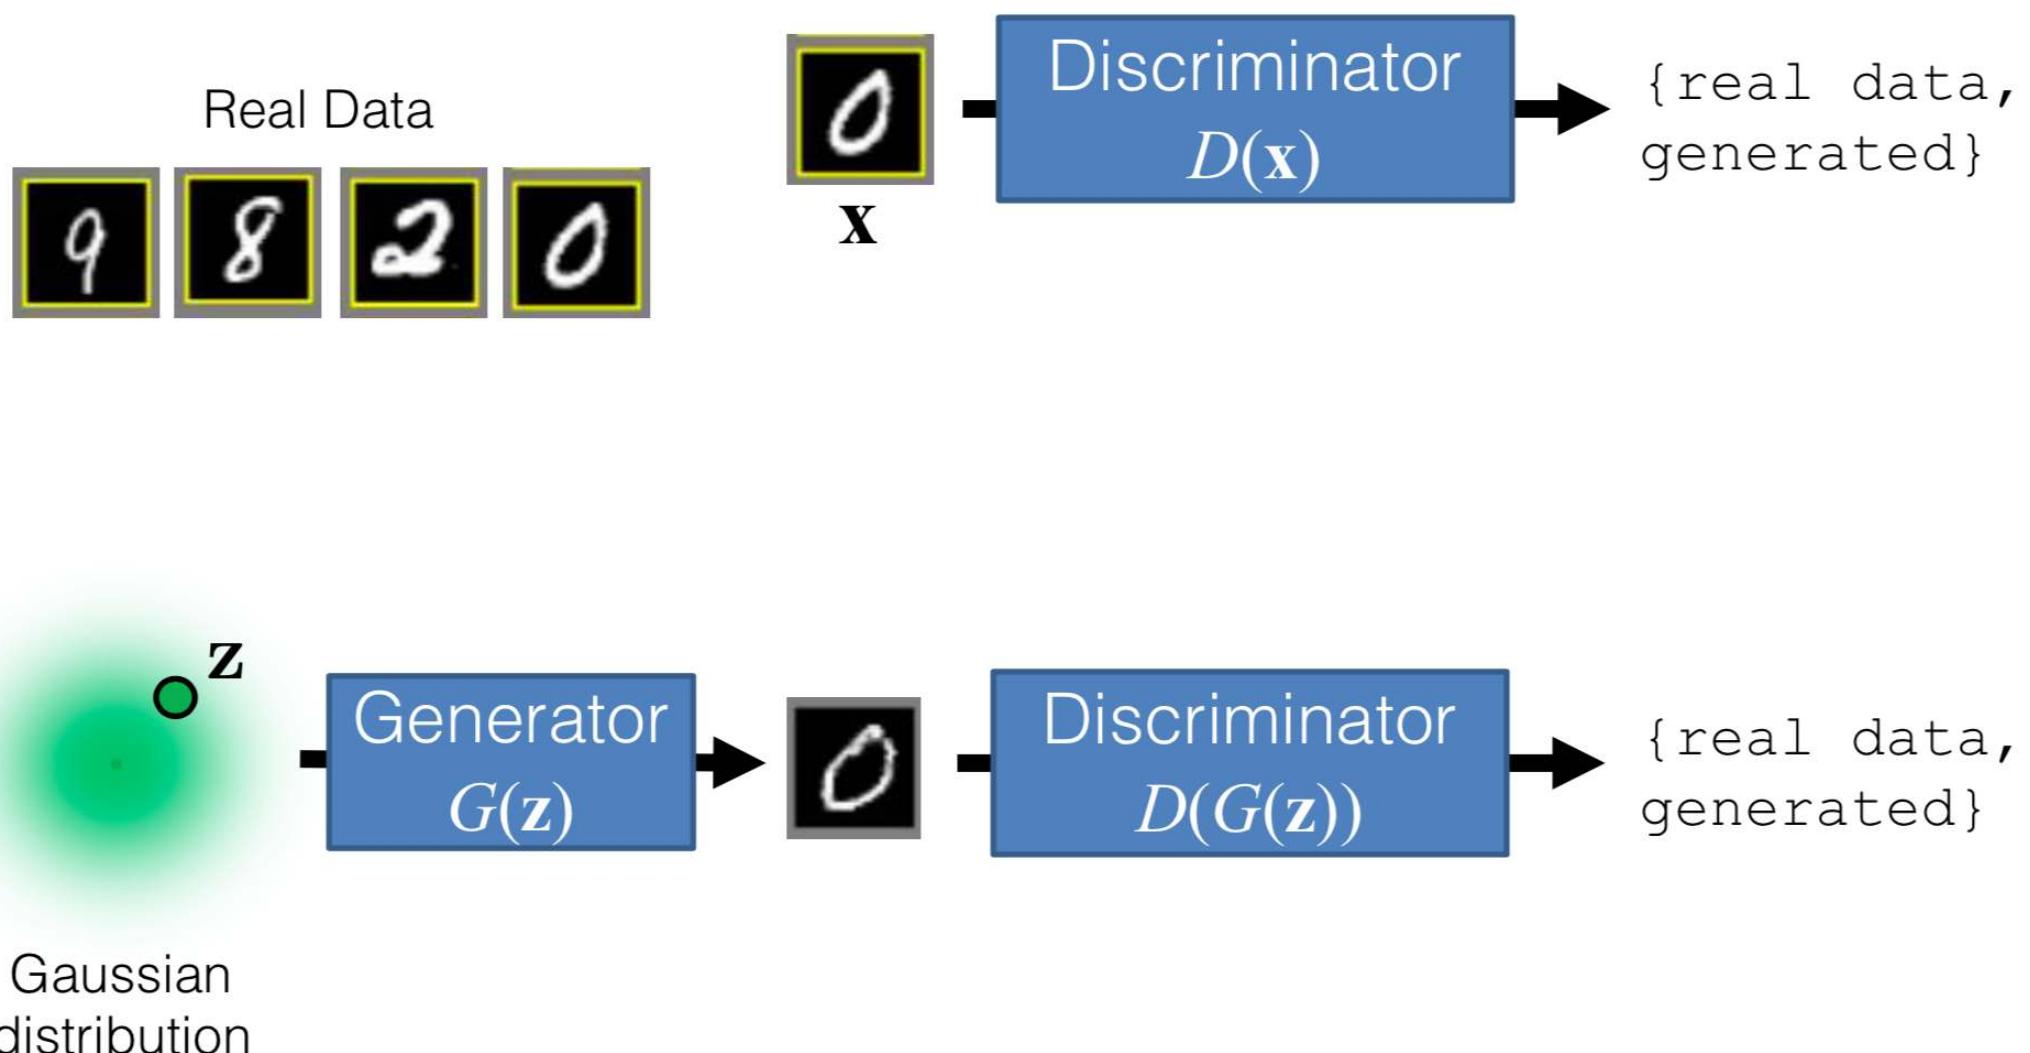
Gaussian  
distribution

# GANs



Gaussian  
distribution

# GANs

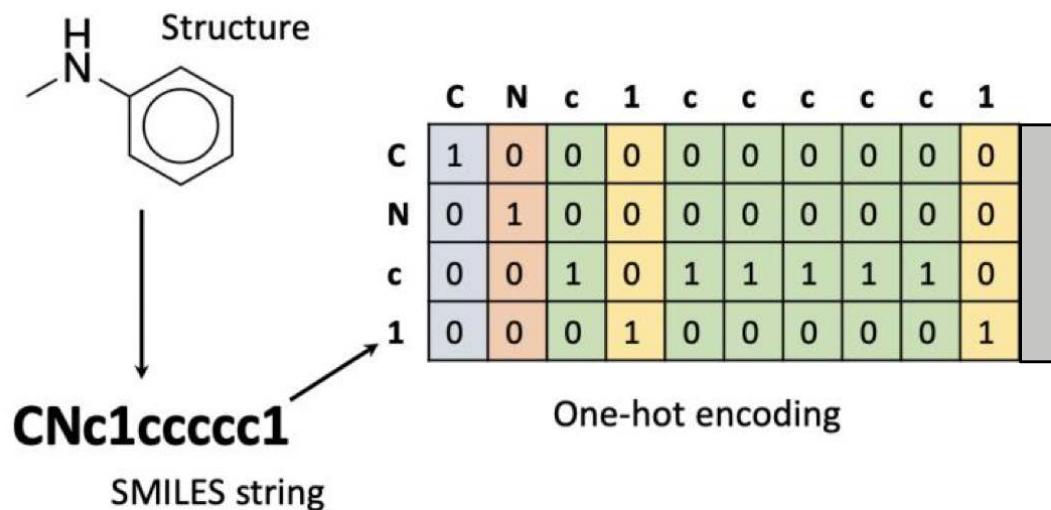


$$\min_G \max_D V(D, G) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})} [\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p_{\mathbf{z}}(\mathbf{z})} [\log(1 - D(G(\mathbf{z})))]$$

# generating molecules

Generative chemistry: drug discovery with deep learning generative models. Yuemin Bian and Xiang-Qun Xie. arXiv 2020.

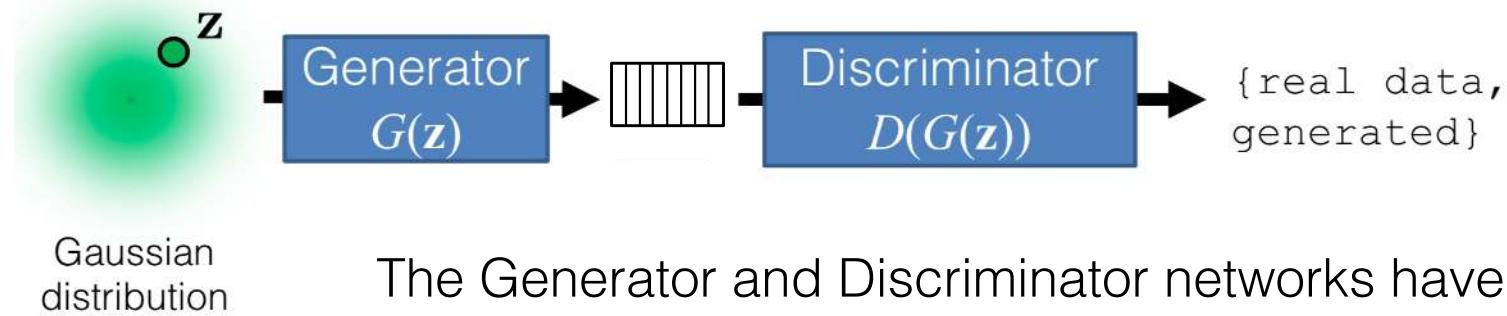
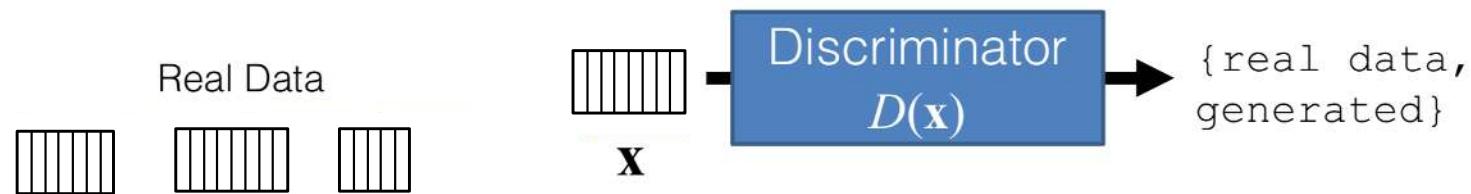
- 1) We need a representation for molecules:



this is a text-like representation.

# generating molecules

2) → we can train a GAN to learn generating new representations of molecules



The Generator and Discriminator networks have the same form as networks for text generation and text analysis.

# perspectives

- Black box models and explicability;
- Learning with less training data;
- Good practices for ‘AI engineering’ as for ‘software engineering’.